



AI Playbook Learning Module Curriculum

All module content from the AI Playbook and Toolkit for Public Health Departments

Includes 70 learning modules with objectives, definitions, content sections, exercises, knowledge checks, references, and jurisdiction or agency policy notes.

Module Index

- ANL 200** - Natural Language Processing for Public Health
- ANL 210** - Machine Learning Basics for Epidemiology and Surveillance
- ANL 220** - Predictive Analytics and Risk Stratification
- ANL 230** - Geospatial AI and Spatial Epidemiology
- ANL 240** - Time-Series Analysis, Anomaly Detection, and Outbreak Signals
- ANL 300** - Advanced Model Validation and Evaluation for Public Health AI
- ANL 310** - Evaluation of Generative AI Outputs
- ANL 320** - Synthetic Data for Public Health AI
- ANL 330** - Equity Evaluation for AI Models
- ANL 340** - Causal Inference and AI in Public Health Decision-Making
- ANL 350** - Communicating Model Results and Uncertainty
- ARC 200** - Public Health Data Modernization and AI Architecture
- ARC 210** - Data Pipelines for Public Health AI
- ARC 220** - Public Health Interoperability: HL7 v2, CDA, FHIR, and APIs
- ARC 230** - Electronic Case Reporting Technical Workflow
- ARC 240** - ELR, Laboratory Data Quality, and Terminology Mapping
- ARC 250** - FHIR for Public Health AI Use Cases
- ARC 260** - APIs, Webhooks, and Event-Driven Workflows
- ARC 300** - Retrieval-Augmented Generation for Public Health Knowledge Work
- ARC 310** - Prompt Engineering for Technical Public Health Work
- ARC 320** - Agentic Workflow Design and Orchestration

ARC 330 - Cloud, Containers, and Secure Deployment Basics

ARC 340 - Public Health System Integration Patterns

COM 200 - AI for Public Health Communications and Message Development

COM 210 - AI-Assisted Translation, Language Access, and Readability Review

COM 220 - Public Comment, Social Listening, and Community Feedback Analysis

EPI 300 - AI for Case Investigation and Outbreak Response

EPI 310 - AI for Surveillance Signal Review and Epidemiologic Judgment

EPI 320 - Ethical Use of AI in Resource Prioritization and Field Response

EXE 100 - Executive AI Briefing for Public Health Leaders

EXE 400 - Executive AI Decision Rights and Accountability

EXE 410 - AI Portfolio Strategy and Use Case Prioritization

EXE 420 - AI Investment Strategy and Business Case Review

EXE 430 - Organizational Change Management for AI Adoption

EXE 440 - Workforce Strategy and AI Role Design

EXE 450 - AI Sustainability and Lifecycle Funding

EXE 460 - AI Governance Board Chair Training

EXE 470 - AI Crisis Communications and Executive Response

GOV 100 - Introduction to AI Governance for Public Health

GOV 110 - Respecting Tribal Sovereignty and Indigenous Data Governance in Public Health AI Projects

GOV 300 - AI Security, Prompt Injection, and Data Leakage

GOV 310 - Privacy-Preserving AI and De-identification

GOV 320 - AI Documentation, Model Cards, and System Cards

GOV 330 - Data Governance for AI-Ready Public Health Data

GOV 340 - AI Audit Readiness and Evidence Management

GOV 350 - Identity, Access, and Permissions for AI Workflows

GOV 360 - Third-Party AI Risk Management

GOV 400 - AI Governance Charter and Operating Model

GOV 410 - AI Risk Appetite and Risk Tolerance for Public Health

GOV 420 - Public Accountability and AI Transparency Reporting

GOV 430 - Community Oversight and Participatory Governance

INT 190 - AI for Social Good

INT 200 - AI Use Case Intake and Triage Basics

OPS 200 - Data Quality Engineering for AI

OPS 210 - Record Linkage, Entity Resolution, and Deduplication

OPS 220 - Human Factors, Workflow Integration, and Usability Testing

OPS 230 - AI Workflow Standard Operating Procedures

OPS 240 - User Acceptance Testing for AI-Supported Workflows

OPS 300 - MLOps and LLMOps for Public Health

OPS 310 - Measuring AI Value, Performance, and Operational Impact

OPS 320 - Operational Monitoring and Continuous Improvement for AI

- PMG 200** - Business Analysis for AI-Enabled Public Health Systems
- PMG 210** - Requirements Writing for AI Procurement and Implementation
- PMG 220** - AI Project Management and Implementation Planning
- PMG 300** - AI Vendor Implementation Management
- PMG 310** - Grant and Funding Management for AI Projects
- POL 300** - Public Health AI Policy Landscape and Regulatory Readiness
- POL 310** - AI Policy Development and Approval Workflow
- POL 320** - AI Authority, Liability, and Public Health Decision-Making
- POL 330** - AI Equity Policy and Civil Rights Considerations

ANL 200

Natural Language Processing for Public Health

Course ID	ANL 200
Track	Analytics and Modeling Track
Level	applied/foundational
Estimated time	45
Audience	Epidemiologists, Data analysts, Data scientists, Evaluation staff, Analytics leads

Learning Objectives

- Define NLP in plain language and distinguish common NLP tasks.
- Identify public health workflows where NLP may be useful.
- Assess risks related to accuracy, bias, privacy, hallucination, language, and missing context.
- Define source verification and human review requirements for NLP outputs.
- Produce an NLP evaluation plan for one public health workflow.

Why This Topic Matters for Public Health AI

NLP can reduce burden by helping staff find, organize, classify, extract, translate, or summarize information from large volumes of text. It can support case investigation, environmental health complaint triage, syndromic surveillance, public inquiry response, community engagement analysis, and policy review. However, text is often ambiguous, incomplete, sensitive, and context-dependent.

NLP outputs should not be treated as facts without review. A tool may omit important details, misread abbreviations, misunderstand negation, translate poorly, or generate a summary that sounds correct but is not supported by the source. Public health text may also include legal, clinical, cultural, or community context that a model does not understand.

For technical staff, NLP evaluation requires more than checking whether outputs sound fluent. The agency should define the task, intended use, prohibited use, source material, review process, performance measures, subgroup considerations, privacy safeguards, and escalation pathway.

Core Concepts

Natural language processing. NLP refers to AI methods that process, analyze, classify, summarize, extract, translate, or generate human language. It can work with structured prompts, free-text notes, documents, emails, calls, complaints, or reports.

Information extraction. Information extraction pulls specific facts from text, such as symptoms, dates, locations, exposures, risk factors, facility names, or requested services. Extraction should be validated against source text.

Text classification. Text classification assigns text to categories, such as complaint type, urgency, topic, likely condition, or program area. Classification requires labeled examples, clear categories, and evaluation.

Summarization. Summarization creates a shorter version of longer text. Public health summarization should preserve uncertainty, cite sources when possible, and avoid adding unsupported claims.

Hallucination. A hallucination is an AI-generated statement that is false, unsupported, or not present in the source material. Hallucination risk is especially important when summaries influence public health decisions or communications.

Public Health Example

A contact center may receive hundreds of public inquiries during an outbreak. NLP could classify messages by topic, identify urgent issues, and draft response summaries for staff review. This can improve timeliness, but messages may contain sensitive information, emotional language, misinformation, or ambiguous requests. Staff need review procedures and escalation criteria.

An environmental health program may use NLP to classify complaints about food safety, housing, water, pests, or nuisance conditions. The tool may help route complaints, but errors can affect response time and enforcement. The program should validate classification accuracy and ensure that original complaint text remains available to reviewers.

Technical Deep Dive

NLP implementation begins with task definition. Extraction, classification, summarization, translation, topic modeling, and question answering are different tasks with different evaluation methods. A tool that performs well at summarization may perform poorly at extracting dates or classifying urgency. Staff should define the task narrowly before selecting or testing a tool.

Evaluation should use representative source text. Test examples should include common language, uncommon language, abbreviations, misspellings, multilingual text, ambiguous statements, negation, sensitive information, and edge cases. Performance should be reviewed by people who understand the public health workflow, not only by technical staff.

Privacy and source traceability are central. NLP may process sensitive case notes, complaints, or community feedback. The agency must define whether the data can be used, where processing occurs, how outputs are stored, who can access them, and how source text is preserved. For consequential uses, users should be able to trace extracted facts or summaries back to the original text.

Risks, Failure Modes, and Guardrails

Unsupported summaries. NLP tools may omit, distort, or invent information. Guardrails include source review, citations, factuality checks, and human approval.

Language and accessibility gaps. Models may perform differently across languages, dialects, literacy levels, or disability-related communication needs. Guardrails include multilingual evaluation, accessibility review, and community context.

Privacy exposure. Free text may contain sensitive or unexpected personal information. Guardrails include approved environments, de-identification when appropriate, access controls, and retention rules.

Misclassification. Incorrect labels can affect routing or priority. Guardrails include validation, confidence thresholds, review queues, and monitoring of overrides.

Application to AI-Supported Workflows

Generative AI is often used for summarization, drafting, and question answering, but outputs must be checked against source material. Predictive analytics may use NLP-derived features, so extraction quality affects model performance. RAG systems use NLP to retrieve and summarize documents, which requires document governance and source citation. Agentic AI using NLP to trigger actions should require clear thresholds and human approval for consequential steps.

Reflection Questions

Which public health workflows involve large volumes of free text?

What task would NLP support: extraction, classification, summarization, translation, or topic detection?

What errors would be most harmful?

What source material must be retained for human review?

How will staff evaluate accuracy across languages, populations, and program areas?

Practical Exercise

Create an NLP evaluation plan for one public health workflow. Include the text source, NLP task, intended use, prohibited use, required human review, accuracy measures, bias or subgroup review, privacy safeguards, source citation or traceability expectations, and escalation process.

Expected Artifact or Evidence

NLP use case description

NLP evaluation plan

Human review checklist

Source verification and privacy safeguards

References and Resources for Additional Information

NIST AI Risk Management Framework

NIST Generative AI Profile

Public health NLP case studies

Agency privacy, records, and data governance policies

OWASP LLM security resources

Reference and Source Base for Lesson Development

The resources below can be used to develop facilitator notes, worksheets, examples, and implementation tools for this technical track. Always verify current versions before publishing or using in live training.

CDC Public Health Data Strategy: <https://www.cdc.gov/public-health-data-strategy/>

CDC Data Modernization Initiative: <https://www.cdc.gov/data-modernization/>

CDC Electronic Case Reporting: <https://www.cdc.gov/ecr/>

AIMS eCR Triggering and eRSD/RCTC Documentation: <https://ecr.aimsplatform.org/ehr-implementers/triggering/>

HL7 FHIR: <https://hl7.org/fhir/>

HL7 US Core Implementation Guide: <https://build.fhir.org/ig/HL7/US-Core/>

HL7 eCR FHIR Implementation Guide: <https://build.fhir.org/ig/HL7/case-reporting/>

ONC USCDI: <https://www.healthit.gov/isa/united-states-core-data-interoperability-uscdi>

CDC PHIN Vocabulary Access and Distribution System: <https://www.cdc.gov/phn/php/phinvads/>

NIST AI Risk Management Framework: <https://www.nist.gov/itl/ai-risk-management-framework>

NIST AI RMF Generative AI Profile:

<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>

OWASP Top 10 for Large Language Model Applications:

<https://owasp.org/www-project-top-10-for-large-language-model-applications/>

Analytics, Modeling, Security, and Advanced Methods: Modules 11-20

Source file: Technical_Readiness_for_Public_Health_AI_Training_Modules_11-20(1).docx

Technical Readiness for Public Health AI

Training Modules 11–20

Prepared as part of a technical public health AI learning track. The modules are designed for public health informaticians, epidemiologists, technologists, data analysts, product owners, implementation leads, and governance participants who need rigorous technical readiness for AI-supported public health work.

How to Use These Modules

These modules continue the technical public health AI curriculum. They move from foundational data and interoperability readiness into applied analytics, model evaluation, generative AI review, RAG, prompt engineering, agentic workflow design, and AI security. They are written for staff who need to evaluate, implement, monitor, or govern AI-supported public health workflows.

Each module should be taught with a concrete workflow. Learners should use examples from surveillance, electronic reporting, immunization, emergency response, community health assessment, environmental health, service access, public communications, or program operations. The practical exercise in each module produces evidence that can support readiness assessment, governance review, procurement, implementation planning, or monitoring.

Modules Included

Module 11. Machine Learning Basics for Epidemiology and Surveillance

Module 12. Predictive Analytics and Risk Stratification

Module 13. Geospatial AI and Spatial Epidemiology

Module 14. Time-Series Analysis, Anomaly Detection, and Outbreak Signals

Module 15. Advanced Model Validation and Evaluation for Public Health AI

Module 16. Evaluation of Generative AI Outputs

Module 17. Retrieval-Augmented Generation for Public Health Knowledge Work

Module 18. Prompt Engineering for Technical Public Health Work

Module 19. Agentic Workflow Design and Orchestration

Module 20. AI Security, Prompt Injection, and Data Leakage

Knowledge Check

1. What is NLP?

- A. A laboratory instrument.
- B. A budgeting method.
- C. AI methods that process, analyze, classify, summarize, or generate human language.
- D. A public health law.

Correct answer: C

2. Which task is an example of information extraction?

- A. Pulling symptoms and exposure dates from case notes.
- B. Scheduling a staff meeting.
- C. Replacing the system of record.
- D. Deleting uncertain text.

Correct answer: A

3. Why is source verification important for NLP outputs?

- A. NLP outputs are always perfect.
- B. NLP tools may omit, misstate, or invent information.

- C. Source material is never needed.
- D. Human review creates risk.

Correct answer: B

4. Which NLP use requires strong safeguards?

- A. Counting words in a public document.
- B. Formatting a training title.
- C. Summarizing sensitive case notes for operational use.
- D. Changing a slide title.

Correct answer: C

5. What is strong completion evidence for this module?

- A. An NLP evaluation plan with intended use, prohibited use, review process, accuracy measures, and privacy safeguards.
- B. A screenshot of a chatbot response only.
- C. A generic AI definition.
- D. A list of unrelated datasets.

Correct answer: A

References and Resources

NIST Artificial Intelligence Risk Management Framework

<https://www.nist.gov/itl/ai-risk-management-framework>

NIST AI RMF Generative AI Profile

<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>

WHO Guidance on Large Multi-Modal Models

<https://www.who.int/publications/i/item/9789240084759>

CDC Public Health Data Strategy

<https://www.cdc.gov/public-health-data-strategy/php/about/index.html>

Agency privacy, records, data governance, and program policies

ANL 210

Machine Learning Basics for Epidemiology and Surveillance

Course ID	ANL 210
Track	Analytics and Modeling Track
Level	applied/foundational
Estimated time	45
Audience	Epidemiologists, Data analysts, Data scientists, Evaluation staff, Analytics leads

Learning Objectives

- Explain the difference between supervised, unsupervised, and semi-supervised learning in public health terms.
- Describe features, labels, training data, validation data, and test data.
- Identify when a machine learning approach may be appropriate for a public health workflow and when a simpler method may be preferable.
- Recognize common machine learning risks, including overfitting, biased labels, data leakage, and poor generalization.
- Interpret basic model performance questions in relation to epidemiologic and operational decision-making.
- Produce a machine learning use case review worksheet for a real or realistic public health workflow.

Why This Topic Matters for Public Health AI

Machine learning matters for public health AI because many proposed tools rely on classification, prediction, ranking, or pattern detection. A surveillance team may want to classify incoming reports by urgency. An epidemiology team may want to detect unusual disease patterns. A program may want to prioritize outreach based on predicted risk. Each of these uses can affect who receives attention, services, investigation, or resources.

Public health staff need enough machine learning literacy to ask whether the model is aligned with the public health purpose. A model built to predict an outcome is not the same as a model that explains why the outcome occurred. A model that performs well in one jurisdiction, population, or reporting environment may not perform well elsewhere. A model trained on historical public health actions may learn historical practice patterns, including under-detection or unequal service access.

The most important public health question is whether the model improves decision-making safely and equitably. This requires more than a performance score. Staff must understand the source data, label quality, model limitations, subgroup performance, workflow impact, human review process, and monitoring plan. Machine learning should strengthen public health practice, not obscure uncertainty or shift responsibility away from accountable staff.

Public Health Example

A health department may consider a machine learning model to classify environmental health complaints by likely urgency. Historical complaint text, location, facility type, prior violations, and resolution outcomes may be used as features. The label might be whether the complaint was ultimately assigned high priority by trained staff. This sounds useful, but the label may reflect staffing capacity or historical prioritization patterns rather than true urgency.

To use this model responsibly, staff would need to evaluate the training data, define the intended use, test performance on recent complaints, review errors by complaint type and community, and decide whether the output is advisory or operational. The model might help sort a queue, but it should not prevent staff from reviewing complaints that do not match historical patterns. The public health example shows that model design decisions become program decisions.

Technical Deep Dive

Machine learning development typically begins with a defined prediction or classification task. The task should be specific enough to test. “Use AI to improve surveillance” is too vague. “Classify incoming reports into routine review, expedited review, and supervisor review based on defined criteria” is more actionable. A clear task allows staff to define the target label, required features, acceptance criteria, and human review process.

Data preparation is often the most important and time-consuming step. Public health data may contain missing values, duplicate records, delayed reporting, inconsistent coding, local terminology, free-text notes, and changes over time. Feature engineering translates raw data into model inputs, but it can also embed assumptions. For example, using prior healthcare utilization as a feature may reflect access to care as much as underlying risk.

Training, validation, and test datasets serve different purposes. Training data are used to fit the model. Validation data are used to tune choices during development. Test data are used to estimate performance on data not used to build the model. Public health teams should be cautious when a vendor reports only development performance without external validation or realistic operational testing.

Model performance should be interpreted in relation to the workflow. A model with high overall accuracy may still miss rare but important events. A model with many false positives may burden staff and create alert fatigue. A model with many false negatives may miss cases, outbreaks, or communities needing support. Thresholds should be selected based on public health consequences, not only technical optimization.

Human review should be designed into the workflow before deployment. Reviewers need to understand what the model output means, what it does not mean, and when to override or escalate. Staff also need feedback channels so model errors can be documented, analyzed, and used to improve monitoring or retraining decisions.

Risks, Failure Modes, and Guardrails

Biased or inconsistent labels. If historical labels reflect unequal detection, inconsistent review, or limited resources, the model may reproduce those patterns. Guardrails include label audits, expert review, subgroup analysis, and documentation of label limitations.

Data leakage. Data leakage occurs when information that would not be available at the time of decision is used during training. Guardrails include temporal validation, careful feature review, and independent testing.

Overfitting and poor generalization. A model may perform well on development data but poorly in a new jurisdiction, time period, or population. Guardrails include held-out testing, external validation, drift monitoring, and conservative deployment.

Misuse of prediction as explanation. A model can predict risk without explaining cause. Guardrails include clear communication, training, and restrictions against causal claims unless supported by appropriate study design.

Application to AI-Supported Workflows

Machine learning may support generative AI, predictive analytics, NLP, anomaly detection, and agentic workflows. In generative AI systems, machine learning concepts help staff understand why outputs may be fluent but unreliable. In predictive workflows, machine learning literacy helps staff evaluate risk scores, thresholds, and subgroup performance. In agentic workflows, models may trigger actions, making validation, permissions, and human approval especially important.

The practical application is disciplined use. Staff should define intended use, validate performance, document limitations, monitor drift, and ensure that human reviewers remain accountable for consequential public health actions. Machine learning should be treated as an analytic support method within a governed workflow.

Reflection Questions

What public health task would the model support: classification, prediction, ranking, clustering, or exploration?

What label would be used, and how reliable or equitable is that label?

Which features might reflect access, documentation, or reporting bias rather than true risk?

What errors would be most harmful in this workflow?

Who would review model outputs, override them, and document concerns?

Practical Exercise

Create a machine learning use case review worksheet for one public health workflow. The worksheet should describe the public health task, intended use, data sources, candidate features, proposed label, likely users, performance measures, subgroup review needs, human review points, and monitoring requirements.

This exercise should produce a governance-ready artifact. Learners should identify at least three reasons the model might fail and at least three safeguards that should be in place before pilot testing.

Expected Artifact or Evidence

Machine learning use case review worksheet

Feature and label inventory

Initial model risk and safeguard list

Human review and monitoring notes

References and Resources for Additional Information

CDC Evaluation Framework: <https://www.cdc.gov/evaluation/>

NIST Artificial Intelligence Risk Management Framework: <https://www.nist.gov/itl/ai-risk-management-framework>

NIST AI RMF Generative AI Profile:

<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>

CDC Public Health Data Strategy: <https://www.cdc.gov/public-health-data-strategy/>

Knowledge Check

1. Which statement best describes supervised learning?

- A. It uses labeled examples to learn a relationship between inputs and an outcome.
- B. It avoids all human judgment.
- C. It only works with maps.
- D. It is the same as public health law.

Correct answer: A

2. Why can historical labels create risk?

- A. They always make models more fair.
- B. They may reflect inconsistent practice, unequal detection, or prior resource constraints.
- C. They remove the need for validation.
- D. They are never used in public health data.

Correct answer: B

3. What is overfitting?

- A. A privacy law for AI systems.
- B. A method for creating dashboards.
- C. A model learning training data too closely and performing poorly on new data.
- D. A guarantee of external validity.

Correct answer: C

4. What should determine model thresholds in public health workflows?

- A. Vendor preference only.
- B. The default software setting.
- C. The color of the dashboard.
- D. The public health consequences of false positives and false negatives.

Correct answer: D

5. What is strong completion evidence for this module?

- A. A machine learning use case worksheet with features, labels, intended use, risks, safeguards, and review points.
- B. A screenshot of a model output only.
- C. A list of software names.
- D. A general statement that AI is useful.

Correct answer: A

References and Resources

CDC Evaluation Framework

<https://www.cdc.gov/evaluation/>

NIST Artificial Intelligence Risk Management Framework

<https://www.nist.gov/itl/ai-risk-management-framework>

NIST AI RMF Generative AI Profile

<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>

CDC Public Health Data Strategy

<https://www.cdc.gov/public-health-data-strategy/php/about/index.html>

ANL 220

Predictive Analytics and Risk Stratification

Course ID	ANL 220
Track	Analytics and Modeling Track
Level	applied/foundational
Estimated time	45
Audience	Epidemiologists, Data analysts, Data scientists, Evaluation staff, Analytics leads

Learning Objectives

- Define predictive analytics and risk stratification in public health terms.
- Distinguish prediction from causation, diagnosis, and eligibility determination.
- Interpret basic risk categories, thresholds, calibration, and error tradeoffs.
- Identify equity and operational risks associated with risk scores.
- Design safeguards for use of predictive outputs in resource allocation, outreach, surveillance, or triage.
- Produce a predictive model review plan for a public health workflow.

Why This Topic Matters for Public Health AI

Predictive analytics matters because public health agencies often must act before all information is available. During outbreaks, emergencies, or resource-constrained programs, agencies may need to decide where to send staff, which facilities need follow-up, which communities may need outreach, or which reports require expedited review. Prediction can help organize uncertainty, but it cannot remove uncertainty.

Risk stratification can also create harm if used carelessly. A model may assign lower risk to groups with incomplete data, limited healthcare access, or fewer documented encounters. It may assign higher risk based on factors that reflect social vulnerability without ensuring supportive intervention. If a score is used to deny services, delay review, or reduce attention, the tool may worsen inequities.

The responsible use of prediction requires a clear purpose, well-defined outcome, appropriate data, subgroup evaluation, calibrated scores, human review, and transparency about limitations. Public health staff should ask who benefits from the prediction, who could be harmed, and whether the intended action is supportive, punitive, or resource-constraining.

Public Health Example

A local health department may want to predict which neighborhoods are at increased risk of low vaccination coverage before school entry. The model could use prior immunization data, school enrollment patterns, appointment availability, language access indicators, transportation barriers, and historical outreach response. The output might help prioritize community outreach and mobile clinics.

This use can support equity if it directs resources toward communities facing access barriers. It can create harm if the score stigmatizes neighborhoods, ignores missing data, or shifts responsibility away from system-level barriers. The model should be evaluated for calibration, data completeness, community context, and whether the intervention is

supportive and accessible.

Technical Deep Dive

A predictive workflow begins with the outcome. The outcome must be specific, measurable, and aligned with a public health action. Predicting “risk” is not enough. Staff must define risk of what, during what time period, for which population, and for what purpose. A model predicting hospitalization within 30 days is different from one predicting delayed case investigation or low service uptake.

Feature selection should be reviewed for both technical and ethical implications. Some features may improve prediction but create fairness, privacy, or interpretability concerns. Geography, insurance status, prior utilization, language, disability, housing instability, or justice involvement may reflect structural inequities. These features require careful review because they may be useful for supportive outreach but inappropriate for restrictive decisions.

Calibration and discrimination answer different questions. Discrimination describes whether a model can separate higher-risk from lower-risk cases. Calibration describes whether predicted probabilities correspond to observed outcomes. A model can rank cases reasonably well while still overstating or understating absolute risk, which matters when decisions depend on risk categories or thresholds.

Thresholds are policy and workflow decisions as much as technical decisions. A lower threshold may catch more true high-risk cases but create more false positives and burden staff. A higher threshold may conserve resources but miss people or places needing support. Threshold selection should consider staffing capacity, consequence severity, equity, and the availability of human review.

Post-deployment monitoring is essential. A model that worked during development may degrade as disease patterns, reporting systems, access to care, community behavior, or program operations change. Monitoring should include performance, calibration, subgroup effects, workflow impact, user overrides, and community feedback.

Risks, Failure Modes, and Guardrails

Prediction used as decision. A risk score may be treated as a final decision rather than decision support. Guardrails include role-based guidance, human review, and documentation of final decisions.

Equity distortion. Risk models may reflect unequal data availability or historical inequities. Guardrails include subgroup analysis, community review, and restrictions on punitive uses.

Threshold harm. Poorly selected thresholds can create alert fatigue or missed risk. Guardrails include threshold review, scenario testing, and monitoring of false positives and false negatives.

Model drift. Performance may change over time. Guardrails include periodic recalibration, drift monitoring, and authority to pause or modify use.

Application to AI-Supported Workflows

Predictive analytics may be embedded in dashboards, case management systems, early warning tools, outreach lists, or agentic workflows. Generative AI may explain or summarize predictive outputs, but those explanations must not imply certainty that the model does not provide. RAG systems may combine risk results with policy guidance, while agentic workflows may trigger tasks or notifications based on thresholds.

When prediction supports public health action, the workflow should make the score, its limitations, and the source data visible to users. Staff should know when a score is advisory, when it requires review, and when it cannot be used as the sole basis for action.

Reflection Questions

What outcome is being predicted, and over what time period?

What public health action will follow a high-risk score?

Which data sources may underrepresent some populations?

What are the consequences of false positives and false negatives?

How will thresholds, calibration, and subgroup performance be monitored?

Practical Exercise

Create a predictive model review plan for one public health use case. Include the outcome, population, time horizon, data sources, candidate features, risk categories, thresholds, performance measures, subgroup evaluation plan, human review process, and monitoring indicators.

The plan should also identify whether the use is supportive, restrictive, or consequential. Predictive tools used to allocate supportive resources have different risk profiles than tools used to deny, delay, or reduce services.

Expected Artifact or Evidence

Predictive model review plan

Risk score interpretation guide

Threshold and error tradeoff memo

Subgroup monitoring plan

References and Resources for Additional Information

NIST Artificial Intelligence Risk Management Framework: <https://www.nist.gov/itl/ai-risk-management-framework>

NIST AI RMF Generative AI Profile:

<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>

CDC Evaluation Framework: <https://www.cdc.gov/evaluation/>

CDC Public Health Data Strategy: <https://www.cdc.gov/public-health-data-strategy/>

Knowledge Check

1. What is calibration?

- A. The process of creating a map.
- B. The degree to which predicted probabilities match observed outcomes.
- C. A type of prompt injection.
- D. A legal authority for data sharing.

Correct answer: B

2. Why should thresholds be reviewed by public health staff?

- A. They determine how scores translate into categories or actions.
- B. They eliminate all false positives.
- C. They are always set correctly by software defaults.
- D. They prevent model drift.

Correct answer: A

3. Which use of a risk score is most concerning without additional safeguards?

- A. Using it as one input for supportive outreach.
- B. Displaying it with source data and limitations.
- C. Using it as the sole basis to deny or delay services.
- D. Monitoring whether risk categories are calibrated.

Correct answer: C

4. What is a false negative in a predictive workflow?

- A. A case incorrectly flagged as high risk.
- B. A model card with missing fields.
- C. A workflow diagram error.
- D. A true risk that the model fails to flag.

Correct answer: D

5. What should a predictive model review plan include?

- A. Outcome, population, thresholds, error consequences, subgroup review, human review, and monitoring.
- B. Only a vendor demonstration.
- C. Only overall accuracy.
- D. Only a list of software licenses.

Correct answer: A

References and Resources

NIST Artificial Intelligence Risk Management Framework

<https://www.nist.gov/itl/ai-risk-management-framework>

NIST AI RMF Generative AI Profile

<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>

CDC Evaluation Framework

<https://www.cdc.gov/evaluation/>

CDC Public Health Data Strategy

<https://www.cdc.gov/public-health-data-strategy/php/about/index.html>

ANL 230

Geospatial AI and Spatial Epidemiology

Course ID	ANL 230
Track	Analytics and Modeling Track
Level	applied/foundational
Estimated time	45
Audience	Epidemiologists, Data analysts, Data scientists, Evaluation staff, Analytics leads

Learning Objectives

- Define geospatial AI and spatial epidemiology in public health terms.
- Explain how geocoding, aggregation, spatial scale, and data quality affect AI-supported maps and analyses.
- Identify privacy and re-identification risks related to location data.
- Recognize spatial bias, modifiable areal unit problems, and ecological interpretation risks.
- Design guardrails for AI-supported spatial analysis, hotspot detection, service planning, or emergency response.
- Produce a geospatial AI readiness worksheet for a public health workflow.

Why This Topic Matters for Public Health AI

Geospatial AI matters because public health decisions are often place-based. Agencies decide where to conduct outreach, deploy mobile services, inspect facilities, respond to environmental hazards, monitor outbreaks, and communicate risks. AI-supported spatial tools can help analyze large volumes of geographic information and identify patterns that may not be visible in simple tables.

However, maps can create a false sense of objectivity. A hotspot may reflect where data are collected, who has access to testing, which providers report electronically, or which addresses are easier to geocode. Areas with missing or poor-quality data may appear low risk when they are actually under-observed. Spatial AI therefore requires careful review of data completeness, spatial scale, and community context.

Location data also raises privacy and trust concerns. Public health agencies have authority to collect and use sensitive information, but that authority should be matched with safeguards. When AI tools combine health data, address data, demographic data, mobility data, environmental data, or social vulnerability indicators, agencies must consider re-identification, stigmatization, public communication, and community impact.

Public Health Example

A health department may use AI-supported geospatial analysis to identify areas where mobile vaccination clinics could improve access. The tool may combine immunization coverage, school locations, transportation routes, census variables, provider locations, and appointment availability. The output may recommend priority areas for outreach and clinic placement.

This workflow can advance equity if it identifies access barriers and supports community partnership. It can cause harm if it labels neighborhoods as low-performing without context, uses outdated population denominators, ignores language access, or fails to engage local organizations. The geospatial analysis should be paired with community review,

accessibility planning, and transparent explanation of limitations.

Technical Deep Dive

Geospatial AI begins with location data, and location data are rarely perfect. Addresses may be missing, misspelled, outdated, rural route-based, temporary, or unstable. People experiencing homelessness, migration, displacement, incarceration, or housing instability may be poorly represented by address-based analysis. Geocoding errors can shift records across boundaries and distort rates for small areas.

Spatial scale shapes interpretation. County-level maps may hide neighborhood-level inequities. Census tract maps may reveal meaningful local variation but increase privacy risks. ZIP codes are convenient but were created for mail delivery, not public health analysis. Staff should choose geographic units based on the question, data quality, privacy requirements, and actionability.

AI-supported hotspot detection requires baseline context. An apparent cluster may reflect higher testing volume, better reporting, a facility-specific outbreak, or a true increase in incidence. Algorithms can detect unusual patterns, but epidemiologists must assess plausibility, data quality, denominator stability, and reporting artifacts. Automated hotspot alerts should be treated as signals for review, not conclusions.

Geospatial models may use spatial proxies for social conditions. Social vulnerability, deprivation, rurality, housing burden, transportation access, and environmental exposure indicators can improve public health planning. They can also be misused if they stigmatize communities or support restrictive decisions. Agencies should use these indicators to guide supportive action and resource allocation, not blame.

Privacy protection should be planned before analysis. Techniques include aggregation, suppression, geomasking, minimum cell-size rules, role-based access, public/private map separation, and review of small-area outputs before publication. For internal AI tools, access controls and audit logs should apply to spatial data just as they do to clinical or case data.

Risks, Failure Modes, and Guardrails

Geocoding bias. Some populations and locations are harder to geocode accurately. Guardrails include geocoding quality checks, manual review for unmatched records, and documentation of missingness.

Small-area privacy risk. Detailed maps can reveal sensitive patterns or identities. Guardrails include aggregation, suppression, role-based access, and public release review.

Overinterpretation of hotspots. Spatial clusters may reflect testing or reporting patterns. Guardrails include denominator review, baseline comparison, and epidemiologic confirmation.

Stigmatizing communities. Maps can unintentionally label communities as problems. Guardrails include equity framing, community engagement, careful language, and focus on structural supports.

Application to AI-Supported Workflows

Geospatial AI may support predictive analytics, emergency response, inspection prioritization, environmental monitoring, and service planning. Generative AI may summarize spatial patterns for leaders, but it must not overstate certainty or ignore limitations. RAG tools may retrieve local plans, maps, and protocols, while agentic workflows may propose outreach routes or inspection assignments that require human approval.

The responsible application of geospatial AI requires explicit attention to privacy, geography, and equity. A spatial output should answer a public health question, support a feasible action, and include enough context for staff and communities to interpret it appropriately.

Reflection Questions

What location data are used, and how complete and accurate are they?

Which geographic unit best matches the public health action?

Could missing or poor-quality location data affect certain populations more than others?

What privacy safeguards are needed for internal and public maps?

How will community context be included before action is taken?

Practical Exercise

Create a geospatial AI readiness worksheet for one public health workflow. Include the public health question, location data sources, geocoding method, geographic unit, denominator, privacy safeguards, equity concerns, AI-supported task, human review process, and communication plan.

The worksheet should identify at least one way the map could be misinterpreted and one safeguard that would reduce the risk of harm.

Expected Artifact or Evidence

Geospatial AI readiness worksheet

Map interpretation and limitations memo

Privacy and suppression review notes

Community context and action plan

References and Resources for Additional Information

CDC GIS and Public Health Resources: <https://www.cdc.gov/gis/>

NIST Artificial Intelligence Risk Management Framework: <https://www.nist.gov/itl/ai-risk-management-framework>

NIST Privacy Framework: <https://www.nist.gov/privacy-framework>

CDC Public Health Data Strategy: <https://www.cdc.gov/public-health-data-strategy/>

Knowledge Check

1. What is geocoding?

- A. Combining model outputs into a dashboard.
- B. Removing all identifiers from data.
- C. Converting an address or location description into a mapped location.
- D. A type of model calibration.

Correct answer: C

2. Why can maps create misleading impressions?

- A. They are never useful in public health.
- B. They may reflect data collection, testing, reporting, geocoding, or boundary choices rather than true risk alone.
- C. They automatically remove privacy risk.
- D. They do not require denominators.

Correct answer: B

3. What is the modifiable areal unit problem?

- A. A privacy rule for AI tools.
- B. A machine learning feature store.
- C. A form of prompt injection.
- D. The problem that results may change depending on the geographic scale or boundaries used.

Correct answer: D

4. Which safeguard is appropriate before publishing a small-area health map?

- A. Review aggregation, suppression, privacy risk, and interpretation limitations.
- B. Publish exact addresses for transparency.
- C. Remove all context so the map is simpler.
- D. Assume low counts cannot identify anyone.

Correct answer: A

5. What is strong completion evidence for this module?

- A. A vendor logo.
- B. A geospatial readiness worksheet with data sources, geography, privacy, equity, AI task, and review process.
- C. A map without notes.
- D. A general definition of GIS only.

Correct answer: B

References and Resources

CDC GIS and Public Health Resources

<https://www.cdc.gov/gis/>

NIST Artificial Intelligence Risk Management Framework

<https://www.nist.gov/itl/ai-risk-management-framework>

NIST Privacy Framework

<https://www.nist.gov/privacy-framework>

CDC Public Health Data Strategy

<https://www.cdc.gov/public-health-data-strategy/php/about/index.html>

ANL 240

Time-Series Analysis, Anomaly Detection, and Outbreak Signals

Course ID	ANL 240
Track	Analytics and Modeling Track
Level	applied/foundational
Estimated time	45
Audience	Epidemiologists, Data analysts, Data scientists, Evaluation staff, Analytics leads

Learning Objectives

- Define time-series analysis and anomaly detection in public health terms.
- Explain baselines, seasonality, trends, reporting delays, and thresholds.
- Distinguish statistical signals from confirmed outbreaks or public health events.
- Identify risks from alert fatigue, data artifacts, delayed reporting, and model drift.
- Design guardrails for AI-supported surveillance alerts.
- Produce an anomaly detection review plan for one public health data stream.

Why This Topic Matters for Public Health AI

Time-series methods matter because public health surveillance is often about change. A single count may be difficult to interpret without knowing whether it is higher than expected for that time, place, population, reporting source, and condition. AI-supported anomaly detection can help staff review large volumes of data and prioritize signals that deserve attention.

However, public health time-series data are shaped by reporting systems as much as disease patterns. Weekends, holidays, provider onboarding, lab code changes, EHR outages, weather events, testing availability, and public attention can all change observed data. An anomaly may indicate an outbreak, but it may also indicate a data feed problem or behavior change.

Responsible use requires a workflow that treats alerts as starting points for investigation. Staff must know what data are being monitored, how the baseline was selected, what triggers an alert, what review steps follow, and when a signal becomes an incident, outbreak investigation, or communication issue.

Public Health Example

A syndromic surveillance team may use anomaly detection to monitor emergency department visits for respiratory symptoms. The system may flag a county when visits exceed the expected range for that day and season. The alert could prompt review of age groups, facilities, diagnosis text, lab data, school absenteeism, and local context.

The alert is valuable only if it enters a disciplined review process. A sudden increase may reflect a true outbreak, a coding change, a new facility feed, a holiday effect, or a backlog. The analyst should review data quality, compare related indicators, assess geography and populations, and document whether the signal requires action or continued

monitoring.

Technical Deep Dive

Baseline selection is one of the most important technical decisions. A baseline may use several prior weeks, several prior years, a regression model, seasonal decomposition, or more complex algorithms. The baseline should match the public health question. A short baseline may adapt quickly but miss seasonal context. A long baseline may capture seasonality but fail when underlying systems change.

Time-series data often require preprocessing. Analysts may need to address missing days, duplicate submissions, delayed reporting, facility onboarding, outliers, and known holidays. Data corrections can appear as sudden spikes or drops. AI tools should distinguish between raw signals and quality-adjusted signals whenever possible.

Thresholds control alert volume and sensitivity. A very sensitive threshold may detect early signals but generate many false alarms. A conservative threshold may reduce workload but delay detection. Thresholds should be chosen with input from epidemiologists, program staff, and emergency response leads because operational capacity and consequences matter.

Anomaly detection methods vary in complexity. Simple methods may compare counts to moving averages or standard deviations. More advanced methods may use regression, Bayesian approaches, control charts, machine learning, or ensemble models. More complex methods are not automatically better. In public health practice, interpretability, stability, transparency, and workflow fit often matter as much as statistical sophistication.

Alert review should be documented. Staff should record what signal occurred, what data were reviewed, what explanation was considered, what action was taken, and whether the alert was useful. Over time, these review notes can improve thresholds, reduce alert fatigue, and support accountability.

Risks, Failure Modes, and Guardrails

Reporting artifacts. Feed failures, delayed batches, or coding changes can trigger false alerts. Guardrails include data quality checks, feed monitoring, and partner change logs.

Alert fatigue. Too many alerts can reduce attention. Guardrails include threshold tuning, tiered alerting, user feedback, and periodic alert review.

Overreaction to signals. A statistical signal may be mistaken for a confirmed outbreak. Guardrails include defined review protocols and communication approval.

Missed emerging events. Overly conservative thresholds may miss early signals. Guardrails include multiple indicators, sensitivity review, and escalation pathways for staff concerns.

Application to AI-Supported Workflows

Time-series and anomaly detection methods can support predictive dashboards, syndromic surveillance, wastewater monitoring, outbreak intelligence, and emergency preparedness. Generative AI may summarize signal review, but it should include uncertainty and source indicators. Agentic workflows may route alerts or create tasks, but those actions should be limited to defined protocols and human review.

The best AI-supported surveillance workflows combine automation with epidemiologic judgment. AI can help detect changes, but humans determine whether the signal is meaningful, what additional evidence is needed, and what public health action is appropriate.

Reflection Questions

What time-series data stream would benefit from automated signal review?

What baseline is appropriate for the condition, data source, and decision?

Which reporting artifacts could create false signals?

What threshold balances early detection and alert fatigue?

How will staff document review, decisions, and lessons learned?

Practical Exercise

Create an anomaly detection review plan for one public health data stream. Include the monitored indicator, data source, update frequency, baseline approach, known reporting patterns, threshold rationale, alert tiers, review steps, related indicators, documentation expectations, and escalation criteria.

The plan should identify how staff will distinguish statistical anomalies from data quality issues and public health events.

Expected Artifact or Evidence

Anomaly detection review plan

Baseline and threshold rationale

Alert review protocol

Signal documentation template

References and Resources for Additional Information

CDC National Syndromic Surveillance Program: <https://www.cdc.gov/nssp/>

CDC Evaluation Framework: <https://www.cdc.gov/evaluation/>

NIST Artificial Intelligence Risk Management Framework: <https://www.nist.gov/itl/ai-risk-management-framework>

CDC Public Health Data Strategy: <https://www.cdc.gov/public-health-data-strategy/>

Knowledge Check

1. What is an anomaly in a public health time series?

A. A value or pattern that differs from expected behavior.

- B. A confirmed outbreak.
- C. A privacy authorization.
- D. A permanent case definition.

Correct answer: A

2. Why does seasonality matter?

- A. It removes the need for thresholds.
- B. It represents predictable recurring patterns that affect expected values.
- C. It proves causation.
- D. It prevents reporting delays.

Correct answer: B

3. What is a common risk of too many low-value alerts?

- A. Better legal compliance.
- B. Improved calibration automatically.
- C. Alert fatigue.
- D. Perfect outbreak detection.

Correct answer: C

4. Which statement is most accurate?

- A. An AI alert should always trigger public communication.
- B. A statistical signal never matters.
- C. Baseline selection is unrelated to public health context.
- D. An AI alert is a signal for review, not automatic confirmation of an outbreak.

Correct answer: D

5. What is strong completion evidence for this module?

- A. An anomaly detection review plan with baseline, threshold, review, documentation, and escalation details.
- B. A generic chart title.
- C. A list of unrelated diseases.
- D. A vendor demo without data review.

Correct answer: A

References and Resources

CDC National Syndromic Surveillance Program
<https://www.cdc.gov/nssp/>

CDC Evaluation Framework
<https://www.cdc.gov/evaluation/>

NIST Artificial Intelligence Risk Management Framework
<https://www.nist.gov/itl/ai-risk-management-framework>

CDC Public Health Data Strategy

<https://www.cdc.gov/public-health-data-strategy/php/about/index.html>

ANL 300

Advanced Model Validation and Evaluation for Public Health AI

Course ID	ANL 300
Track	Analytics and Modeling Track
Level	intermediate/practitioner
Estimated time	60
Audience	Epidemiologists, Data analysts, Data scientists, Evaluation staff, Analytics leads

Learning Objectives

- Describe the purpose of advanced model validation in public health AI.
- Interpret sensitivity, specificity, positive predictive value, negative predictive value, precision, recall, F1 score, calibration, and subgroup performance.
- Distinguish internal validation, external validation, and workflow validation.
- Identify when model performance is insufficient for the intended use.
- Design monitoring for drift, degradation, error patterns, and unintended consequences.
- Produce an advanced validation plan for an AI-supported workflow.

Why This Topic Matters for Public Health AI

Validation matters because public health AI tools can appear credible even when they are not fit for operational use. A polished dashboard, confident generated text, or high overall accuracy score does not prove that the tool is safe, equitable, or useful. Public health staff need to evaluate performance in relation to the specific task and the consequences of error.

Model performance is context-dependent. A tool that performs well in one health system or state may fail in another because of different data sources, coding practices, populations, workflows, or reporting laws. A tool that performs well during development may degrade when disease patterns change or new providers begin reporting. Validation must therefore include local or context-specific testing whenever possible.

Advanced validation also requires attention to equity. Overall performance can hide poor performance for small populations, language groups, rural communities, people with unstable housing, or communities with limited healthcare access. Public health validation should examine subgroup performance and data representativeness before AI outputs influence priority setting or resource allocation.

Public Health Example

A public health agency may test a model that identifies incoming case reports likely to require expedited review. Overall accuracy may be high, but that number alone is not enough. Staff need to know whether the model misses rare but severe conditions, whether it performs differently across facilities, whether it over-flags reports from certain communities, and whether users understand how to act on the output.

The agency should validate the model on recent local data, review errors with epidemiologists, test the workflow with actual users, and monitor override patterns after pilot deployment. If the model performs well technically but creates too many alerts for available staff, the workflow may still fail. Validation must include operational fit.

Technical Deep Dive

Metric selection should start with the public health consequence. Sensitivity may be prioritized when missing a true event is dangerous. Specificity may be prioritized when false positives create substantial burden or unnecessary intervention. Precision and positive predictive value matter when staff act on positive flags. Calibration matters when scores are interpreted as probabilities.

Confusion matrices are useful because they show true positives, false positives, true negatives, and false negatives. Public health teams should translate each cell into operational meaning. A false negative may mean a delayed investigation. A false positive may mean staff time diverted from other work. A true positive may mean earlier action. A true negative may mean confidence to continue routine processing.

External validation is essential when a model will be used outside its development context. Testing on a new jurisdiction, time period, data source, or population can reveal performance problems that internal validation missed. When external validation is not possible before pilot testing, agencies should begin with limited deployment, close monitoring, and conservative use.

Workflow validation examines whether the AI output is understandable, accessible, timely, actionable, and used as intended. A model can be statistically strong but operationally weak if staff do not trust it, cannot see supporting data, cannot override it, or receive alerts too late. Workflow validation should include usability testing, training, documentation, and feedback loops.

Post-deployment validation should be planned before launch. Agencies should define monitoring indicators, review frequency, escalation triggers, and authority to pause use. Monitoring should include performance drift, subgroup changes, user override rates, incident reports, data feed changes, and public health outcomes where appropriate.

Risks, Failure Modes, and Guardrails

Misleading aggregate performance. Overall metrics may hide poor subgroup performance. Guardrails include subgroup testing and equity review.

Metric mismatch. The selected metric may not match the public health consequence. Guardrails include metric selection tied to intended use and error consequences.

Lack of external validation. Models may fail in new settings. Guardrails include local testing, limited pilots, and monitoring.

Workflow failure. A technically accurate model may be unusable or misused. Guardrails include usability testing, training, source visibility, and override documentation.

Application to AI-Supported Workflows

Advanced validation applies to predictive models, NLP tools, generative summaries, RAG outputs, and agentic workflows. For generative AI, validation may include factuality, completeness, source grounding, omissions, and harmful output review. For RAG, validation should test retrieval quality and citation accuracy. For agentic workflows, validation

should include action safety, permission boundaries, and rollback procedures.

The evaluation unit should be the AI-supported workflow, not the model alone. Public health agencies should ask whether the full process improves timeliness, quality, equity, safety, and staff experience without creating unacceptable risk.

Reflection Questions

What decision or action will the AI output support?

Which performance metrics match the consequences of error?

Which subgroups, data sources, or settings require separate review?

What would make performance unacceptable for this use case?

How will the agency monitor performance after deployment?

Practical Exercise

Create an advanced validation plan for one AI-supported workflow. Include intended use, validation dataset, comparison standard, performance metrics, subgroup analysis, error review, workflow testing, user feedback, monitoring indicators, escalation triggers, and documentation requirements.

The plan should clearly explain how validation results will be interpreted and who has authority to approve, limit, pause, or reject the AI use.

Expected Artifact or Evidence

Advanced validation plan

Metric and error-consequence table

Subgroup performance review plan

Post-deployment monitoring schedule

References and Resources for Additional Information

NIST Artificial Intelligence Risk Management Framework: <https://www.nist.gov/itl/ai-risk-management-framework>

NIST AI RMF Generative AI Profile:

<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>

CDC Evaluation Framework: <https://www.cdc.gov/evaluation/>

WHO Guidance on Large Multi-Modal Models: <https://www.who.int/publications/i/item/9789240084759>

Knowledge Check

1. Why can overall accuracy be misleading?

- A. It always reflects fairness.
- B. It may hide poor performance for important subgroups or rare outcomes.
- C. It is never used in model evaluation.
- D. It removes the need for workflow testing.

Correct answer: B

2. What is workflow validation?

- A. Testing whether the AI-supported process works safely and effectively in real operations.
- B. Testing only whether a website loads.
- C. A replacement for data quality review.
- D. A method for removing audit logs.

Correct answer: A

3. Which metric is especially important when missed true events could cause harm?

- A. Specificity only.
- B. Dashboard uptime.
- C. Sensitivity.
- D. File size.

Correct answer: C

4. What does external validation test?

- A. Whether staff like the color scheme.
- B. Whether the vendor has a contract.
- C. Whether the model can avoid monitoring.
- D. Performance in a different time, place, population, source, or setting than development data.

Correct answer: D

5. What should an advanced validation plan include?

- A. Intended use, metrics, subgroup analysis, error review, workflow testing, monitoring, and escalation authority.
- B. Only a marketing description.
- C. Only the model name.
- D. Only a screenshot of output.

Correct answer: A

References and Resources

NIST Artificial Intelligence Risk Management Framework
<https://www.nist.gov/itl/ai-risk-management-framework>

NIST AI RMF Generative AI Profile

<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>

CDC Evaluation Framework

<https://www.cdc.gov/evaluation/>

WHO Guidance on Large Multi-Modal Models

<https://www.who.int/publications/i/item/9789240084759>

ANL 310

Evaluation of Generative AI Outputs

Course ID	ANL 310
Track	Analytics and Modeling Track
Level	intermediate/practitioner
Estimated time	60
Audience	Epidemiologists, Data analysts, Data scientists, Evaluation staff, Analytics leads

Learning Objectives

- Explain why generative AI outputs require structured review before public health use.
- Distinguish factuality, source grounding, completeness, relevance, tone, and accessibility.
- Identify hallucinations, omissions, unsupported claims, and misleading certainty.
- Design review criteria for different output types, including summaries, drafts, classifications, and recommendations.
- Apply stronger review requirements for sensitive, public-facing, or consequential outputs.
- Produce a generative AI output evaluation rubric.

Why This Topic Matters for Public Health AI

Generative AI matters for public health because it can reduce burden in text-heavy workflows. Staff may use it to summarize guidance, draft meeting notes, simplify technical language, translate or adapt messages, organize case information, or synthesize stakeholder feedback. These uses can be valuable when they save time and improve clarity.

The same capabilities create risk. Generative AI can invent facts, omit caveats, misstate guidance, flatten uncertainty, or produce language that sounds more definitive than the evidence supports. It may introduce bias, inappropriate tone, or inaccessible wording. When sensitive data are involved, prompting and output handling can also create privacy and records concerns.

Public health agencies need structured evaluation because informal review is inconsistent. A reviewer may notice grammar errors but miss unsupported claims. Another may trust a polished summary without checking source documents. A rubric helps reviewers evaluate outputs against the purpose, source material, audience, and safeguards required for the use case.

Public Health Example

A communications team may use generative AI to draft a plain-language public message about a local respiratory illness increase. The draft may be clear and readable, but reviewers must check whether the numbers are correct, whether the recommendations match current guidance, whether uncertainty is represented accurately, and whether the message is accessible in multiple languages.

An epidemiology team may use generative AI to summarize case investigation notes for internal review. The summary must be checked against source notes because a missed exposure, incorrect date, or invented symptom could affect follow-up. The team should document whether the summary was accepted, edited, or rejected.

Technical Deep Dive

Generative AI evaluation should begin with intended use. A brainstormed outline, internal draft, public advisory, case summary, policy synthesis, or operational recommendation requires different levels of review. The more sensitive, public-facing, or consequential the use, the stronger the review should be.

Factuality review requires comparison to source material. Reviewers should check names, dates, numbers, conditions, legal requirements, recommendations, and causal statements. They should also verify that the output does not add unsupported claims or remove important caveats. For RAG systems, reviewers should inspect whether citations actually support the claims they accompany.

Completeness review asks whether the output includes the information needed for the task. A summary can be accurate but incomplete if it omits a critical exposure, eligibility criterion, limitation, or escalation instruction. Public health review should consider whether omitted information could change interpretation or action.

Equity and accessibility review should be included when outputs affect staff, partners, or the public. Reviewers should assess reading level, language access, disability accessibility, cultural context, stigma, blame, and whether the output addresses affected communities respectfully. Translation or multilingual adaptation should include review by qualified humans when used for public-facing or consequential communication.

Documentation should be proportionate to risk. Low-risk drafting may require only human review before use. Higher-risk workflows may require saved prompts, source references, reviewer initials, revision history, approval records, and incident escalation procedures.

Risks, Failure Modes, and Guardrails

Fluent but false outputs. A polished output may contain unsupported claims. Guardrails include source checking, citations, and reviewer signoff.

Omitted caveats. AI may simplify uncertainty away. Guardrails include completeness review and required caveat checks.

Inappropriate tone or accessibility. Outputs may be stigmatizing or inaccessible. Guardrails include plain language, equity, language access, and accessibility review.

Privacy exposure. Prompts or outputs may include sensitive data. Guardrails include approved tools, data-use tiers, redaction, access controls, and records guidance.

Application to AI-Supported Workflows

Generative AI output evaluation applies to drafting, summarization, classification, translation, code generation, and decision-support explanation. It is also relevant to RAG and agentic workflows because generated text may summarize retrieved sources or explain automated actions. Evaluation should be built into workflow design, not added after deployment.

The strongest approach is to create rubrics for specific output types. A case summary rubric will differ from a public message rubric, and a policy synthesis rubric will differ from a code review rubric. Each rubric should define what reviewers must check, who approves final use, and how errors are documented.

Reflection Questions

What type of generative AI output is being evaluated?

What source material should be used to verify the output?

What errors or omissions would be most harmful?

What review criteria are needed for equity, accessibility, language, and tone?

What documentation is required before the output can be used?

Practical Exercise

Create a generative AI output evaluation rubric for one public health use case. Include intended use, source verification, factuality, completeness, unsupported claims, omissions, bias, tone, accessibility, privacy, approval requirements, and escalation criteria.

Then apply the rubric to a sample AI-generated output and document whether the output should be accepted, edited, rejected, or escalated.

Expected Artifact or Evidence

Generative AI output evaluation rubric

Completed sample review

Source verification notes

Approval or escalation documentation

References and Resources for Additional Information

NIST AI RMF Generative AI Profile:

<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>

NIST Artificial Intelligence Risk Management Framework: <https://www.nist.gov/itl/ai-risk-management-framework>

OWASP Top 10 for Large Language Model Applications:

<https://owasp.org/www-project-top-10-for-large-language-model-applications/>

WHO Guidance on Large Multi-Modal Models: <https://www.who.int/publications/i/item/9789240084759>

Knowledge Check

1. Why is fluency not enough to trust a generative AI output?

- A. Fluent outputs may still be false, incomplete, unsupported, biased, or inappropriate.
- B. Fluent outputs are always public-ready.

- C. Fluency proves source grounding.
- D. Fluency replaces human review.

Correct answer: A

2. What does grounding mean?

- A. The output is short.
- B. The output is linked to reliable source material that can support verification.
- C. The output uses a formal tone.
- D. The output avoids all citations.

Correct answer: B

3. Which issue can be as harmful as an incorrect statement?

- A. A clear citation.
- B. A reviewer note.
- C. A significant omission.
- D. A plain-language sentence.

Correct answer: C

4. What should a rubric define?

- A. Only the model name.
- B. Only page length.
- C. Only grammar preferences.
- D. Review criteria, approval requirements, documentation, and escalation expectations.

Correct answer: D

5. Which output usually requires stronger review?

- A. A public-facing health advisory generated from current local data.
- B. A private brainstorming list with no sensitive data.
- C. A generic training title.
- D. A low-risk internal outline.

Correct answer: A

References and Resources

NIST AI RMF Generative AI Profile

<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>

NIST Artificial Intelligence Risk Management Framework

<https://www.nist.gov/itl/ai-risk-management-framework>

OWASP Top 10 for Large Language Model Applications

<https://owasp.org/www-project-top-10-for-large-language-model-applications/>

WHO Guidance on Large Multi-Modal Models

<https://www.who.int/publications/i/item/9789240084759>

ANL 320

Synthetic Data for Public Health AI

Course ID	ANL 320
Track	Analytics and Modeling Track
Level	intermediate/practitioner
Estimated time	60
Audience	Epidemiologists, Data analysts, Data scientists, Evaluation staff, Analytics leads

Learning Objectives

- Explain the purpose of synthetic data for public health ai in public health AI work.
- Identify the technical, operational, privacy, security, equity, and governance issues associated with the topic.
- Apply the topic to a real or realistic public health workflow, system, or implementation scenario.
- Recognize common failure modes that can affect safety, accuracy, trust, workflow fit, or sustainability.
- Identify practical guardrails that should be documented before implementation.
- Produce a synthetic dataset assessment that can support readiness assessment, governance review, procurement, implementation, or monitoring.

Why This Topic Matters for Public Health AI

Synthetic Data for Public Health AI matters because AI systems are embedded in public health programs, data systems, and decision processes rather than operating in isolation. The topic affects whether AI outputs are accurate, explainable, safe, equitable, secure, and sustainable. Agencies that ignore this area may create tools that work in a demonstration but fail in actual practice.

The topic also affects accountability. Public health agencies must be able to explain how data are used, why a tool was approved, what evidence supports the workflow, who reviews outputs, and what happens when the system fails. These responsibilities become more important as AI tools move from experimentation into operational use.

Core Concepts

Synthetic data are generated data designed to preserve selected patterns of a source dataset while reducing reliance on real individual-level records. The value of synthetic data depends on the purpose. A dataset that is useful for training staff on a dashboard may not be valid for evaluating a predictive model or testing equity performance.

Privacy risk does not disappear simply because data are synthetic. If synthetic data are too close to the source data, they may leak information about real individuals or rare cases. If they are too different from the source data, they may be safe but not useful. Public health agencies must evaluate both privacy protection and data utility.

Data utility means the synthetic data are useful for the intended purpose. Utility may involve preserving distributions, relationships between fields, missingness patterns, geographic variation, or time trends. The intended use should be documented before the synthetic data are created or accepted from a vendor.

Public Health Example

A health department may want synthetic eCR data to train disease investigators and test an AI case summary tool without exposing real case narratives. The synthetic data should include realistic combinations of symptoms, lab results, demographics, dates, and missing fields, but should not reproduce unusual cases in ways that could identify real people. The agency should document that the data are intended for training and testing, not for estimating disease burden or measuring model performance in production.

Technical Deep Dive

Synthetic data development begins by defining the intended use. Training data for staff exercises can be less statistically faithful than data used for system testing. Model evaluation data require more rigorous review because the synthetic data must preserve patterns relevant to the model task. Public health teams should define whether synthetic data will be used for demonstration, software testing, integration testing, training, simulation, or model development.

Agencies should evaluate synthetic data for privacy and utility. Privacy review may examine whether rare combinations, dates, locations, or narratives are too similar to real records. Utility review may compare distributions, relationships, missingness, and edge cases against expected patterns. Both reviews should be documented because a dataset can be safe but not useful, or useful but not safe enough for sharing.

Synthetic data are also helpful for vendor evaluation. Agencies can provide vendors with realistic data structures without sharing protected data during early discovery. However, vendors should not use synthetic data performance as proof that a model will work on real public health data. Production readiness still requires validation against appropriate real-world data under approved governance controls.

Risks, Failure Modes, and Guardrails

A common failure mode is treating the topic as a technical detail rather than a public health control. When staff do not connect technical decisions to authority, workflow, equity, privacy, security, and accountability, the agency may adopt a tool that appears useful but cannot be sustained or defended.

Another risk is relying on vendor claims without agency-defined evidence. Vendors may provide demonstrations, dashboards, or summary metrics that do not match the agency workflow or data environment. A guardrail is to require documentation, test results, acceptance criteria, and agency-controlled review.

A third risk is failing to plan for change. Public health data, policies, systems, and emergencies change. Any technical approach must include monitoring, review, versioning, escalation, and retirement pathways so the agency can respond when assumptions no longer hold.

Application to AI-Supported Workflows

Generative AI, predictive analytics, RAG, automation, and agentic workflows all depend on the controls described in this module. The controls should be scaled to the risk of the use case. A tool that drafts internal training text may require lighter review than a tool that updates case records, prioritizes outreach, or supports public communication during an emergency.

The module should be applied during readiness assessment, procurement, implementation planning, pilot testing, and post-deployment monitoring. Learners should document how the topic affects data, systems, users, safeguards, and decision-making for the selected workflow.

Reflection Questions

Where does this topic appear in one current or proposed AI workflow in your agency?

What decisions, data sources, users, or partners would be affected if this area is poorly designed?

What evidence would governance reviewers need before approving the workflow?

Which safeguards should be required before pilot testing?

Who owns ongoing monitoring and review?

Practical Exercise

Create a synthetic dataset assessment for one real or realistic public health AI workflow. The artifact should describe the use case, affected users, data sources, system dependencies, risks, guardrails, review points, and unresolved questions.

The exercise should produce a work product that could be used in a governance meeting, procurement review, implementation planning session, pilot evaluation, or monitoring review. Learners should document assumptions and identify which staff or partners need to be consulted.

Expected Artifact or Evidence

The expected artifact is a completed synthetic dataset assessment. It should be specific to a public health workflow and should include technical dependencies, governance questions, human review expectations, monitoring indicators, and unresolved risks.

References and Resources for Additional Information

NIST AI Risk Management Framework

NIST Privacy Framework and differential privacy resources

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP Top 10 for LLM Applications, when security or AI integrations are relevant

Agency privacy, cybersecurity, data governance, procurement, and records management policies

Knowledge Check

1. What is the best reason to address synthetic data for public health ai before deployment?

- A. It makes governance unnecessary.
- B. It helps ensure the AI workflow is safe, useful, reviewable, and fit for public health operations.

- C. It allows the agency to skip testing.
- D. It guarantees that the vendor product will work.

Correct answer: B

2. Which practice best supports responsible implementation?

- A. Defining intended use, evidence requirements, safeguards, monitoring, and escalation before launch.
- B. Waiting until after launch to decide whether the tool is risky.
- C. Allowing users to decide individually how to use the tool.
- D. Treating vendor documentation as the only required evidence.

Correct answer: A

3. Which statement best describes a common failure mode?

- A. Human review is always unnecessary when AI is accurate.
- B. Documentation makes AI less useful.
- C. The agency assumes the tool is ready because the technology works in a demonstration, even though workflow, data, and governance have not been tested.
- D. Monitoring should stop after the pilot.

Correct answer: C

4. What should be included in the practical artifact for this module?

- A. Only a general definition.
- B. A marketing description of the AI tool.
- C. A list of unrelated data systems.
- D. Workflow-specific risks, safeguards, roles, evidence requirements, and unresolved questions.

Correct answer: D

5. When should the issue be escalated for governance or leadership review?

- A. Only if the tool is public-facing.
- B. When risks affect privacy, safety, equity, legal authority, public trust, system integrity, or operational decisions.
- C. Never, because technical staff can resolve all issues alone.
- D. Only after an incident has already occurred.

Correct answer: B

References and Resources

NIST AI Risk Management Framework

NIST Privacy Framework and differential privacy resources

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP Top 10 for LLM Applications, when security or AI integrations are relevant

Agency privacy, cybersecurity, data governance, procurement, and records management policies

ANL 330

Equity Evaluation for AI Models

Course ID	ANL 330
Track	Analytics and Modeling Track
Level	intermediate/practitioner
Estimated time	60
Audience	Epidemiologists, Data analysts, Data scientists, Evaluation staff, Analytics leads

Learning Objectives

- Explain why overall model performance can mask inequitable performance.
- Identify relevant subgroups and contextual factors for equity evaluation in public health.
- Assess data representativeness, missingness, and measurement bias before model use.
- Define subgroup performance measures and review thresholds for AI-supported workflows.
- Produce an equity evaluation plan for a model or AI-supported workflow.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Public Health Example

A model predicts which clients are most likely to miss follow-up care after a public health referral. Overall performance appears acceptable, but subgroup review shows lower sensitivity for rural clients and people with missing phone numbers. The model may under-prioritize people who face transportation, broadband, housing, or communication barriers. Equity evaluation would require additional analysis, workflow changes, and monitoring before the tool is used for outreach prioritization.

Technical and Operational Deep Dive

Equity evaluation begins with the intended use. A model used for internal workload planning has different equity implications than a model used to prioritize field response or determine which communities receive outreach. Evaluators should identify who is affected by the model output, what action follows, and what harms may occur from false positives, false negatives, or delayed review.

Data assessment is the next step. Public health datasets often reflect uneven healthcare access, reporting patterns, documentation practices, language access, and program reach. Missing race, ethnicity, language, disability, address, housing, or contact information may not be random. It may reflect structural barriers or data collection practices that affect some communities more than others.

Subgroup performance should be measured when data allow. This may include sensitivity, specificity, false positive rates, false negative rates, calibration, timeliness, or user override rates by race, ethnicity, age, sex, gender, geography, language, disability, rurality, payer, facility type, or other relevant factors. Subgroups should be chosen based on public health relevance, not simply data convenience.

Equity evaluation should lead to action. Mitigations may include changing thresholds, improving data collection, adding human review, limiting the intended use, creating separate workflows, improving language access, consulting community partners, adding uncertainty warnings, or declining to deploy the model. Equity monitoring should continue after launch because data patterns and workflow use can change.

Risks, Failure Modes, and Guardrails

Risks and failure modes include the following:

Relying on average performance while ignoring subgroup harms.

Using datasets that underrepresent communities with limited healthcare access.

Treating missing data as neutral when missingness reflects structural inequities.

Using AI prioritization to allocate scarce resources without equity review.

Failing to monitor equity impacts after deployment.

Recommended guardrails include the following:

Require equity evaluation for AI uses that affect prioritization, eligibility, outreach, response, or public communication.

Assess representativeness and missingness before model use.

Measure subgroup performance where feasible and document limitations where not feasible.

Include community, equity, and program reviewers in governance decisions.

Monitor equity indicators and user overrides after deployment.

Reflection Questions

Where does this topic appear in your agency, program, or workflow?

Which staff roles would need to understand or apply this concept before AI use is approved?

What could go wrong if this topic is not addressed before implementation?

What evidence would governance, leadership, or operations staff need before moving forward?

Which policy, data, equity, privacy, security, or workflow questions remain unresolved?

Practical Exercise

Create an equity evaluation plan for one AI model or AI-supported workflow. Include intended use, affected populations, relevant subgroups, data representativeness concerns, missingness concerns, subgroup performance measures, potential harms, mitigation options, monitoring indicators, and decision criteria.

Expected Artifact or Evidence

Equity evaluation plan

Subgroup and data representativeness review

Potential harm and mitigation table

Equity monitoring indicators

Governance recommendation

References and Resources for Additional Information

NIST Artificial Intelligence Risk Management Framework

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP guidance for LLM application security

Agency policies for privacy, cybersecurity, records, procurement, accessibility, and public communications

Relevant public health program SOPs, data governance documentation, and implementation plans

Knowledge Check

1. Why is overall model performance insufficient for public health equity review?

- A. Because overall performance is always wrong
- B. Because subgroup analysis is never needed
- C. Because acceptable average performance can hide poor performance for specific populations or contexts
- D. Because models cannot be evaluated

Correct answer: C

2. What is measurement bias?

- A. A colorful dashboard
- B. Systematic differences in how variables are captured or recorded across groups
- C. A vendor contract term
- D. A meeting schedule

Correct answer: B

3. Which use most clearly requires equity evaluation?

- A. Formatting a staff newsletter
- B. Sorting public documents alphabetically
- C. Creating a generic training outline
- D. Prioritizing clients for limited outreach resources using a risk score

Correct answer: D

4. What should happen if subgroup performance is poor for a population affected by the workflow?

- A. The agency should consider mitigation, limits on use, additional review, or not deploying the model
- B. The agency should ignore the finding
- C. The model should be used without documentation
- D. The subgroup should be removed from monitoring

Correct answer: A

5. What is strong completion evidence?

- A. A screenshot of a model output
- B. A generic equity statement
- C. An equity evaluation plan with subgroup, data, harm, mitigation, and monitoring analysis
- D. A vendor demo summary

Correct answer: C

References and Resources

NIST Artificial Intelligence Risk Management Framework

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP guidance for LLM application security

Agency policies for privacy, cybersecurity, records, procurement, accessibility, and public communications

Relevant public health program SOPs, data governance documentation, and implementation plans

ANL 340

Causal Inference and AI in Public Health Decision-Making

Course ID	ANL 340
Track	Analytics and Modeling Track
Level	intermediate/practitioner
Estimated time	60
Audience	Epidemiologists, Data analysts, Data scientists, Evaluation staff, Analytics leads

Learning Objectives

- Distinguish prediction from causal inference in public health AI.
- Explain why correlation and prediction do not establish causation.
- Identify confounding, selection bias, collider bias, and temporal ambiguity in public health data.
- Recognize when a decision requires causal evidence rather than risk prediction alone.
- Identify appropriate evidence sources for evaluating interventions or policy effects.
- Create a causal question framing worksheet for an AI-supported decision.

Definitions

Prediction	An estimate of what is likely to happen based on observed patterns.
Causal inference	The process of estimating whether and how an exposure, policy, program, or intervention changes an outcome.
Confounding	A distortion that occurs when another factor is related to both the exposure and the outcome.
Counterfactual	The outcome that would have occurred under a different condition or intervention.
Quasi-experimental design	An evaluation approach that estimates causal effects without random assignment.

Why This Topic Matters for Public Health AI

AI can make weak evidence appear stronger than it is. A model may predict which communities have higher hospitalization risk, but that does not prove which intervention will reduce hospitalization. A model may find that people who receive a service have better outcomes, but that pattern may reflect who was able to access the service rather than the effect of the service itself.

Public health decisions often involve resource allocation, program design, policy evaluation, and equity tradeoffs. These decisions require clarity about whether the agency is asking a prediction question, an explanation question, or an intervention question. Causal inference helps agencies avoid treating AI-generated associations as proof of program effectiveness.

Public Health Example

A local health department uses an AI model to predict which neighborhoods are likely to experience low vaccination coverage. The model may help target outreach, but it does not prove why coverage is low or which outreach strategy will work. If leaders want to know whether mobile clinics, reminder messages, school partnerships, or community health workers caused an improvement, they need an evaluation design that supports causal inference.

This distinction prevents overclaiming. The AI model can support prioritization and hypothesis generation, while causal methods can support evaluation of the chosen intervention. Both are valuable, but they answer different questions.

Technical and Operational Deep Dive

Causal inference begins with a clearly framed question. A useful structure is: among which population, what is the intervention or exposure, what is the comparison, what outcome is measured, and over what time period? Without this structure, teams may ask a model to answer a question that the available data cannot support.

Confounding is one of the most important threats. For example, a program may appear to improve outcomes because participants were already more connected to care. Selection bias may occur when data include only people who interacted with a system. Temporal ambiguity may occur when it is unclear whether the exposure occurred before the outcome. AI can detect patterns in these data, but it does not automatically solve the causal design problem.

Methods such as matching, regression adjustment, difference-in-differences, interrupted time series, regression discontinuity, and instrumental variables may be useful when randomization is not possible. These methods require assumptions that should be documented and reviewed. Public health staff should know when to consult epidemiologists, statisticians, or evaluators before making causal claims.

Risks, Failure Modes, and Guardrails

The main failure mode is causal overreach. A dashboard or AI report may state that a program reduced disease, improved access, or changed behavior when the evidence only shows association. A guardrail is to require causal language review before results are used in reports, funding decisions, policy recommendations, or public communications.

Another risk is inequitable inference. Data may be more complete for some groups than others, and interventions may reach groups differently. A guardrail is to examine whether assumptions hold across populations and whether subgroup effects or access barriers are hidden by aggregate results.

Application to AI-Supported Workflows

Predictive models can identify risk, but causal inference is needed to decide what action is likely to change that risk. Generative AI may help draft evaluation questions or summarize literature, but humans must verify methods and assumptions. AI-supported decision-making should label outputs clearly as prediction, association, causal evidence, or hypothesis.

Reflection Questions

What AI outputs in your agency could be mistaken for causal evidence?

Which decisions require knowing what intervention will change an outcome?

What confounders are likely in your program area?

What comparison group could support evaluation?

Who should review causal claims before public release?

Practical Exercise

Create a causal question framing worksheet for an AI-supported decision. Identify the population, intervention or exposure, comparison, outcome, time period, likely confounders, data limitations, and what evidence would be needed before making a causal claim.

Expected Artifact or Evidence

Causal question framing worksheet

List of assumptions and likely confounders

Recommended evaluation design or expert consultation need

Plain-language statement distinguishing prediction from causation

References and Resources for Additional Information

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

Knowledge Check

1. Which statement best distinguishes prediction from causal inference?

- A. Prediction estimates likely outcomes; causal inference estimates whether an intervention changes outcomes
- B. Prediction always proves an intervention works
- C. Causal inference is only needed for laboratory studies
- D. AI automatically removes confounding

Correct answer: A

2. What is the best reason to assign an accountability owner?

- A. To reduce the need for documentation
- B. To ensure someone is responsible for monitoring, escalation, and correction
- C. To make the vendor solely responsible
- D. To prevent staff from raising concerns

Correct answer: B

3. When should an AI-supported workflow be escalated for review?

- A. Only after media coverage
- B. Never if the tool is efficient
- C. When outputs may affect safety, equity, privacy, legal authority, public trust, or official action
- D. Only at the end of the fiscal year

Correct answer: C

4. Which practice best supports public accountability?

- A. Keeping AI use informal
- B. Avoiding written records
- C. Letting each user decide independently
- D. Maintaining evidence artifacts, decision records, review notes, and monitoring results

Correct answer: D

5. What should happen when data sources, policy, vendor configuration, or workflow use changes?

- A. The use case should be reviewed to determine whether safeguards, validation, documentation, or approvals need updating
- B. The agency should ignore the change if staff like the tool
- C. The model should be trusted automatically
- D. The glossary should be deleted

Correct answer: A

References and Resources

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

ANL 350

Communicating Model Results and Uncertainty

Course ID	ANL 350
Track	Analytics and Modeling Track
Level	intermediate/practitioner
Estimated time	60
Audience	Epidemiologists, Data analysts, Data scientists, Evaluation staff, Analytics leads

Learning Objectives

- Explain model results in plain language for nontechnical audiences.
- Distinguish uncertainty, variability, error, bias, and data limitations.
- Identify when model results should include confidence intervals, prediction intervals, caveats, or subgroup findings.
- Avoid overclaiming, causal overreach, and false precision in AI communications.
- Create a model communication brief for a public health decision-maker.

Definitions

Uncertainty	The degree to which an estimate, prediction, or conclusion may be imprecise or incomplete.
Prediction interval	A range that estimates where a future observation may fall.
Confidence interval	A range that expresses uncertainty around an estimated parameter.
Calibration	How closely predicted probabilities match observed outcomes.
Caveat	A clear statement of limitation, condition, or caution.

Why This Topic Matters for Public Health AI

Model outputs can influence urgent decisions. A forecast may affect staffing, a risk score may affect outreach, and a hotspot map may affect deployment of field teams. If uncertainty is not explained clearly, decision-makers may treat estimates as facts or ignore important limitations.

Public health communication must balance clarity and caution. Too much technical detail can obscure the message, but too little detail can hide uncertainty. A strong communication product explains what the model suggests, what it does not show, how reliable it is, and what human judgment or additional evidence is needed.

Public Health Example

An epidemiology team uses a time-series model to forecast emergency department visits for respiratory illness. The forecast shows a possible increase over the next two weeks, but the prediction interval is wide because reporting patterns changed after a holiday. A strong communication brief would state that visits may rise, explain the uncertainty, identify possible reasons for reduced confidence, and recommend preparedness actions that are proportional to the

evidence.

The brief should not say that the model proves an outbreak is occurring. It should present the forecast as one input into surveillance judgment, alongside laboratory data, provider reports, syndromic trends, and field intelligence.

Technical and Operational Deep Dive

Model communication begins with the decision context. The same model output may be communicated differently to a technical team, health officer, program manager, community partner, or public audience. Technical users may need details about metrics and data limitations, while executives may need decision implications, confidence, and recommended actions.

Uncertainty should be explained visually and verbally. Trend lines, ranges, maps, and risk categories can help, but they can also imply precision that the model does not have. Labels should avoid unsupported certainty. Words such as likely, possible, estimated, preliminary, and subject to change can be helpful when used consistently.

Subgroup findings require special care. If model performance differs by age, race, ethnicity, language, geography, disability, housing status, or data source, communications should avoid presenting aggregate performance as universal. Decision-makers need to know whether a model is less reliable for some populations.

Risks, Failure Modes, and Guardrails

A major failure mode is false precision. A model may produce a number such as 82.4 percent risk, but the data and validation may not support that level of exactness. A guardrail is to use appropriate rounding, risk categories, and uncertainty statements.

Another failure mode is hiding limitations to make results look more actionable. This can erode trust if the model later performs poorly. A guardrail is to include a standard limitations section in model briefs and require review before results are shared outside the technical team.

Application to AI-Supported Workflows

Generative AI can help draft summaries of model findings, but the draft must be checked by analysts who understand the model. Predictive analytics, outbreak forecasts, risk stratification tools, and geospatial models all require communication practices that preserve uncertainty. RAG systems used for briefings should cite source reports and validation evidence.

Reflection Questions

Who needs to understand this model result?

What decision will the result support?

What uncertainty or limitation could change the decision?

Are there subgroup performance concerns?

What wording would prevent overclaiming?

Practical Exercise

Create a one-page model communication brief. Include the decision question, model result, uncertainty, validation status, limitations, subgroup considerations, recommended use, prohibited use, and next review date.

Expected Artifact or Evidence

Model communication brief

Plain-language uncertainty statement

Subgroup caveat if relevant

Review sign-off note

References and Resources for Additional Information

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

Knowledge Check

1. What is the best reason to communicate uncertainty with model results?

- A. It makes the report longer
- B. It helps decision-makers understand how much confidence to place in the result
- C. It eliminates the need for validation
- D. It proves the model is unbiased

Correct answer: B

2. What is the best reason to assign an accountability owner?

- A. To reduce the need for documentation
- B. To ensure someone is responsible for monitoring, escalation, and correction
- C. To make the vendor solely responsible
- D. To prevent staff from raising concerns

Correct answer: B

3. When should an AI-supported workflow be escalated for review?

- A. Only after media coverage
- B. Never if the tool is efficient
- C. When outputs may affect safety, equity, privacy, legal authority, public trust, or official action
- D. Only at the end of the fiscal year

Correct answer: C

4. Which practice best supports public accountability?

- A. Keeping AI use informal
- B. Avoiding written records
- C. Letting each user decide independently
- D. Maintaining evidence artifacts, decision records, review notes, and monitoring results

Correct answer: D

5. What should happen when data sources, policy, vendor configuration, or workflow use changes?

- A. The use case should be reviewed to determine whether safeguards, validation, documentation, or approvals need updating
- B. The agency should ignore the change if staff like the tool
- C. The model should be trusted automatically
- D. The glossary should be deleted

Correct answer: A

References and Resources

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

ARC 200

Public Health Data Modernization and AI Architecture

Course ID	ARC 200
Track	Technical Architecture Track
Level	applied/foundational
Estimated time	45
Audience	Informatics staff, IT staff, Data architects, Product owners, Technical implementation teams

Learning Objectives

- Explain why AI readiness depends on public health data modernization, not only on access to an AI tool.
- Describe the major components of an AI-enabled public health architecture, including source systems, data exchange, data transformation, AI services, workflow applications, access controls, audit logs, and monitoring.
- Identify the systems of record, source systems, data flows, and human decision points in a public health workflow.
- Assess where AI could support a workflow and where human review, legal authority, epidemiologic judgment, or community context must remain central.
- Identify architectural risks that can affect AI safety, including incomplete data, unclear authority, excessive system access, weak monitoring, and lack of documentation.
- Produce an AI-enabled workflow architecture map and readiness memo that can support governance review, implementation planning, procurement, or vendor evaluation.

Why This Topic Matters for Public Health AI

Public health agencies often feel pressure to adopt AI tools quickly, especially when staff are overwhelmed, reporting demands are increasing, and data modernization funding is available. AI tools may appear to offer immediate relief by summarizing records, drafting communications, classifying incoming reports, identifying possible outbreaks, or automating repetitive tasks. However, AI can create new risks when it is layered on top of fragmented or poorly understood systems. A model that summarizes case reports may produce incomplete summaries if the source data are missing key fields. A triage algorithm may reinforce inequities if the underlying system receives more complete data from some providers or communities than others. An agentic workflow may trigger inappropriate actions if the system does not clearly define what the AI tool is allowed to read, write, update, or escalate.

Data modernization provides the foundation for reducing these risks. Modernization is not just a technology upgrade. It includes better data exchange, stronger governance, improved data quality, workforce capacity, reusable services, interoperability, security, and more timely access to information. AI readiness should therefore be treated as a modernization outcome. If an agency cannot describe where the data come from, how they are transformed, who is authorized to use them, and how decisions are documented, then the agency is not yet ready to use AI for that workflow.

This matters especially in public health because technical errors can become operational errors. A poorly designed AI workflow could delay case investigation, misclassify a reportable condition, summarize a record inaccurately, miss a vulnerable population, or create a public communication that overstates evidence. Architecture review helps prevent these errors by making the full system visible before AI is introduced.

Core Concepts

Public health data modernization. Public health data modernization is the process of improving the systems, standards, infrastructure, governance, and workforce capacity needed to collect, exchange, manage, analyze, and use public health data. Modernization includes technical components such as APIs, cloud platforms, data warehouses, terminology services, analytics environments, and interoperability standards. It also includes operational components such as data stewardship, governance review, privacy controls, workforce training, procurement discipline, and performance monitoring. AI should be understood as one possible modernization capability, not as a replacement for the modernization work itself.

AI architecture. AI architecture refers to the set of systems and processes that allow an AI tool to function within a public health workflow. This includes the source data, data pipelines, model or AI service, user interface, security controls, audit logs, monitoring process, and human review steps. In a mature architecture, each component has a defined role. Staff know which system is the source of truth, which data are allowed to be used, how outputs are reviewed, and what happens when the AI tool produces an uncertain, incomplete, or incorrect result.

Source system. A source system is the system where data originate or are first captured. In public health, source systems may include electronic health records, laboratories, immunization information systems, vital records systems, syndromic surveillance feeds, environmental health systems, claims datasets, community resource directories, or manually entered program data. Source systems differ in quality, timeliness, completeness, permissions, and terminology. These differences directly affect whether downstream AI outputs can be trusted.

System of record. A system of record is the official system or process where public health information, decisions, and actions are maintained. AI outputs should not automatically become the official record unless the agency has explicitly approved that use. For example, an AI tool may draft a case summary, but the final case note may still need to be reviewed, edited, and saved by authorized staff in the surveillance system. The distinction between AI-generated support and official public health documentation is essential for accountability, public records management, legal review, and incident response.

Data lineage and audit logs. Data lineage documents where data came from, how they were changed, and where they were used. Audit logs record actions taken by users, systems, or automated tools. Together, lineage and audit logs help agencies explain why an AI output was generated, which data were used, who reviewed it, whether it was accepted or edited, and whether downstream action was taken. Without these records, agencies may struggle to investigate errors or demonstrate responsible oversight.

Public Health Example

Consider a health department that wants to use AI to summarize incoming electronic case reports for disease investigators. The AI tool may be able to produce a concise summary of symptoms, laboratory results, diagnosis, treatment, provider information, exposure information, and missing fields. However, the safety of that workflow depends on the surrounding architecture. The eCR data must be received correctly, routed to the right jurisdiction, parsed into the surveillance system, associated with the correct person or case, and made available only to authorized users.

The AI tool also needs clear boundaries. It should know which fields it can use and whether it is allowed to access narrative text, laboratory results, demographic data, prior case history, or external reference material. If the tool uses retrieval-augmented generation, the agency must know which policies, disease protocols, or guidance documents are included in the retrieval source. If the tool can create tasks or update case records, the agency must define whether those actions require human approval.

The system also needs a clear human review point. The investigator should be able to see the source information behind the AI summary, identify unsupported statements, correct errors, and document the final interpretation in the surveillance system. The agency should log when the AI summary was generated, which data were used, who reviewed it, and whether the output was accepted, edited, rejected, or escalated. Without that architecture, the AI summary could appear efficient while creating hidden risks.

Technical Deep Dive

A public health AI architecture usually begins with source systems. These may include electronic health records, laboratories, immunization information systems, vital records systems, syndromic surveillance feeds, case management systems, health information exchanges, claims datasets, community resource directories, environmental health systems, or manually entered program data. Each source system has its own data quality profile, update frequency, terminology conventions, permissions, and limitations. Before AI can be used safely, the agency needs to understand these characteristics.

The next layer is data exchange and ingestion. Data may arrive through HL7 v2 messages, CDA documents, FHIR APIs, flat files, web portals, batch uploads, or manual entry. Each method creates different risks. A real-time interface may fail silently if monitoring is weak. A batch file may be delayed. A manual upload may introduce formatting errors. A FHIR API may expose only certain data elements. AI planning must account for how data enter the workflow and how the agency knows whether the data are complete, current, and authorized for the intended use.

The third layer is data transformation and storage. Public health data are often cleaned, mapped, deduplicated, linked, normalized, geocoded, or summarized before they are used. These steps are essential, but they can also introduce errors. For example, local laboratory codes may be mapped incorrectly to standard codes. Duplicate person records may be merged incorrectly. Addresses may be geocoded to the wrong jurisdiction. Race, ethnicity, language, disability, or housing status fields may be missing or inconsistently captured. If an AI tool uses transformed data, the agency must be able to trace how those data were created and whether transformation decisions affect the output.

The fourth layer is the AI service or model. This may be a predictive model, natural language processing tool, generative AI assistant, retrieval-augmented generation system, rules-based automation, or agentic workflow. The model should have a defined intended use and explicit limits. A tool designed to summarize case notes should not silently make eligibility decisions. A tool designed to classify complaints should not automatically trigger enforcement action unless that use has been reviewed and approved. A tool designed to support resource allocation should not be used without subgroup performance testing and equity review.

The final layers are workflow integration, monitoring, and governance. AI outputs must appear where staff can use them appropriately and should preserve the ability to inspect source information, correct errors, override outputs, document decisions, and escalate concerns. Monitoring should include technical measures such as uptime, latency, feed failures, missing data, error rates, and drift, as well as operational measures such as user overrides, complaints, incident reports, equity concerns, timeliness, staff burden, and downstream public health outcomes.

Risks, Failure Modes, and Guardrails

Architectural invisibility. If staff cannot explain how data move through the system, they may not recognize when an AI output is based on incomplete, delayed, duplicated, or incorrectly transformed data. A practical guardrail is to require a workflow and data flow map before AI use is approved.

Unclear authority. AI tools may produce recommendations, summaries, or classifications that look official, even when they have not been reviewed. A practical guardrail is to define which system remains the official system of record and which AI outputs require human approval before they can be used for public health action.

Excessive access. AI tools connected through APIs, plugins, or workflow integrations may be able to retrieve, summarize, transmit, or modify sensitive data. A practical guardrail is to apply least-privilege access, role-based permissions, audit logging, sandbox testing, and explicit approval for any write-back function.

Poor monitoring and equity harm. AI tools can degrade over time when data patterns, reporting sources, terminology, policies, or workflows change. If the architecture relies on data sources that underrepresent some communities, the tool may also produce inequitable outputs. Practical guardrails include pre-deployment subgroup review, monitoring indicators, user override tracking, and clear authority to pause or modify use.

Application to AI-Supported Workflows

For generative AI, architecture determines what source material the tool can access, whether it can cite sources, whether it can use sensitive data, and how outputs are reviewed. A generative AI tool used to summarize policy documents has different architectural requirements than one used to summarize case notes or draft public health advisories. The more sensitive or operational the use case, the stronger the controls should be.

For predictive analytics, architecture determines whether the model receives timely and consistent data, whether the prediction is displayed in context, and whether users understand the limits of the score. A risk score should not become an automatic decision unless the agency has explicitly approved that use, validated the model, reviewed subgroup performance, and established human review procedures.

For retrieval-augmented generation and agentic AI, architecture determines which documents or systems are available, how permissions are enforced, and what actions the AI can take. An agent that can create tasks, send messages, update records, or trigger alerts must operate within strict boundaries. The agency must define allowed actions, prohibited actions, approval points, logs, rollback procedures, and incident response triggers before deployment.

Reflection Questions

Which public health workflow in your agency would be most helpful to map before considering AI?

What are the source systems, systems of record, and downstream outputs for that workflow?

Where are data transformed, linked, summarized, interpreted, or manually corrected?

Which parts of the workflow require legal, epidemiologic, clinical, communications, equity, privacy, or leadership review?

What could go wrong if AI were added without changing the architecture?

Which monitoring indicators would show whether the AI-supported workflow is working safely and equitably?

Practical Exercise

Create an AI-enabled workflow architecture map for one public health workflow. Choose a workflow that is real or realistic for your agency, such as case report intake, lab result triage, immunization record review, community complaint classification, outbreak summary drafting, public inquiry response, referral management, syndromic surveillance review, or resource allocation.

Your map should include the source systems, data exchange methods, transformation steps, system of record, AI-supported task, user interface, human review point, output or decision, audit log, monitoring process, and escalation pathway. After completing the map, identify at least five readiness gaps that should be addressed before implementation.

Expected Artifact or Evidence

The expected artifact is an AI-enabled workflow architecture map. The evidence should also include a short readiness memo that identifies technical dependencies, governance questions, human review requirements, access control needs, monitoring indicators, and unresolved risks.

A strong completion artifact should be specific enough to support governance review, implementation planning, procurement, or vendor evaluation. It should not simply state that AI may be useful. It should show how the AI-supported workflow would function, what safeguards would be needed, and what questions remain unresolved.

References and Resources for Additional Information

CDC Public Health Data Strategy

CDC Data Modernization Initiative materials

NIST AI Risk Management Framework

NIST Generative AI Profile

ONC interoperability and health IT resources

Agency data governance, privacy, cybersecurity, and enterprise architecture documentation

Knowledge Check

1. Why should AI planning begin with architecture and workflow mapping?

- A. Because mapping proves that the AI tool will be accurate.
- B. Because mapping eliminates the need for governance.
- C. Because AI safety depends on understanding data sources, system dependencies, human review points, and safeguards.
- D. Because AI tools do not require data quality review.

Correct answer: C

2. What is the role of a system of record in an AI-supported workflow?

- A. It allows AI outputs to replace all documentation.
- B. It prevents staff from reviewing AI outputs.
- C. It stores only vendor marketing materials.
- D. It defines where official public health information, actions, or decisions are maintained.

Correct answer: D

3. Which risk is most likely when staff cannot explain how data move through a workflow?

- A. Staff may over-trust AI outputs based on incomplete, delayed, duplicated, or transformed data.
- B. AI will automatically correct all data problems.
- C. Governance review becomes unnecessary.
- D. Privacy rules no longer apply.

Correct answer: A

4. What is an appropriate guardrail for an AI tool that can update a public health system?

- A. Allowing unrestricted access because automation is faster.
- B. Role-based permissions, human approval, audit logs, monitoring, and rollback procedures.
- C. Removing human review after the first week.
- D. Using the tool without documentation.

Correct answer: B

5. What is strong completion evidence for this module?

- A. A screenshot of an AI tool.
- B. A general definition of artificial intelligence.
- C. A workflow architecture map and readiness memo that identify data flows, AI touchpoints, human review, access controls, monitoring, and escalation.
- D. A list of possible vendors without workflow analysis.

Correct answer: C

References and Resources

CDC Public Health Data Strategy

CDC Data Modernization Initiative materials

NIST AI Risk Management Framework

NIST Generative AI Profile

ONC interoperability and health IT resources

Agency data governance, privacy, cybersecurity, and enterprise architecture documentation

ARC 210

Data Pipelines for Public Health AI

Course ID	ARC 210
Track	Technical Architecture Track
Level	applied/foundational
Estimated time	45
Audience	Informatics staff, IT staff, Data architects, Product owners, Technical implementation teams

Learning Objectives

- Describe the major steps in public health data pipelines, including ingestion, validation, transformation, storage, access, and monitoring.
- Distinguish between batch, real-time, and event-driven data flows.
- Explain how pipeline failures can affect AI outputs and operational decisions.
- Identify metadata, lineage, version control, and audit needs for AI-supported workflows.
- Produce a data pipeline readiness checklist for a public health AI use case.

Why This Topic Matters for Public Health AI

Data pipeline readiness matters because AI systems can amplify routine data management problems. A delayed file, changed field name, failed interface, missing date, duplicate record, or outdated transformation rule may not be obvious to an end user. The AI tool may continue generating fluent summaries or confident classifications even after the data feed has become unreliable.

Public health agencies need pipeline visibility before they rely on AI for operational workflows. Visibility means knowing where data originate, how often they update, what transformations occur, what validation checks run, and who is alerted when something fails. A pipeline should not be treated as a hidden technical detail. It is part of the safety infrastructure for AI.

Pipeline review also supports public accountability. When an AI output affects case prioritization, outreach, resource allocation, or public messaging, the agency may need to explain which data were used and whether those data were current. A well-documented pipeline makes that explanation possible.

Core Concepts

Data pipeline. A data pipeline is a set of technical and operational steps that moves data from a source to a destination where the data can be used. In public health, pipelines may include laboratory interfaces, EHR feeds, IIS submissions, syndromic surveillance feeds, vital records extracts, HIE connections, and program datasets. Each step creates an opportunity to improve or degrade data quality.

ETL and ELT. ETL stands for extract, transform, load, while ELT stands for extract, load, transform. ETL transforms data before they enter the target system, while ELT loads data first and transforms them later. The distinction matters for AI because staff need to know where data were changed and which version was used by the AI tool.

Pipeline monitoring. Pipeline monitoring is the routine process of checking whether data feeds, transformations, and outputs are working as intended. Monitoring may include file receipt, schema validation, required-field checks, timeliness checks, duplicate review, error queues, and volume trends. Monitoring should be tied to escalation when failures affect AI use.

Data lineage. Data lineage shows where data came from, what happened to them, and where they were used. Lineage supports validation, troubleshooting, incident response, and public records review. In AI workflows, lineage helps staff trace a questionable output back to its source data and transformation history.

Public Health Example

A syndromic surveillance team may use AI to identify unusual emergency department visit patterns. The model depends on timely data feeds from hospitals, correct parsing of chief complaint text, reliable facility identifiers, and consistent geographic information. If one large hospital feed stops or a field changes after an EHR upgrade, the AI tool may miss a signal or generate a false alert. Pipeline monitoring is therefore part of outbreak detection readiness.

A second example is an AI tool that drafts summaries of lab-based cases. The tool may rely on an ELR feed, a terminology mapping table, a deduplication process, and a case management interface. If any of those steps fail, the summary may omit a result, duplicate a patient, or misrepresent reportability. The pipeline must be reviewed before the AI output can be trusted.

Technical Deep Dive

Pipeline design begins with the source and ends with the user-facing decision or output. In between, data may be received, staged, validated, transformed, standardized, linked, loaded, indexed, analyzed, and displayed. Each step should have an owner, a purpose, a quality check, and a documented failure process.

Batch pipelines move data on a schedule, such as nightly files or weekly extracts. Real-time pipelines move data as events occur, such as laboratory reports or eCR messages. Event-driven pipelines trigger actions when a defined event occurs, such as creating a review task when a high-priority result arrives. AI workflows may use any of these models, but the monitoring and risk profile differ for each.

Metadata and version control are essential for AI-supported pipelines. Staff need to know which data dictionary, code mapping, model version, prompt template, transformation rule, and document collection were used at a specific time. This is especially important when an agency investigates an incident or compares AI performance before and after a system change.

Risks, Failure Modes, and Guardrails

Silent feed failure. A data feed can stop or change without visible failure in the AI interface. Guardrails include automated feed monitoring, expected volume thresholds, error queues, and notification procedures.

Schema and terminology changes. Source systems may change field names, formats, codes, or units. Guardrails include schema validation, terminology governance, release notes, and regression testing.

Unclear fallback process. Staff may continue using AI outputs when data are stale or incomplete. Guardrails include visible freshness indicators, pause rules, and manual fallback procedures.

Weak lineage. If transformations are not documented, staff may not be able to explain or correct AI errors. Guardrails include lineage documentation, version control, and audit trails.

Application to AI-Supported Workflows

Generative AI may use pipeline outputs to summarize records, but the summary should indicate source freshness and preserve traceability. Predictive models may rely on engineered features created by a pipeline, so the feature pipeline must be monitored for drift and missingness. RAG systems depend on document ingestion pipelines, which must track document source, version, permissions, and refresh schedules. Agentic workflows may depend on event-driven pipelines, making failure handling and approval points especially important.

Reflection Questions

Which public health data feeds are most important to daily decisions in your program?

Which pipelines are automatically monitored, and which depend on manual review?

What failures would materially affect an AI output?

How would staff know if an AI output was based on stale or incomplete data?

Who has authority to pause AI use when a pipeline fails?

Practical Exercise

Create a data pipeline readiness checklist for one AI-supported workflow. Identify the data source, update frequency, required fields, validation checks, transformations, lineage documentation, monitoring indicators, error notification process, and fallback procedure.

Document at least three pipeline failure scenarios and describe how each should be detected, escalated, and communicated to users.

Expected Artifact or Evidence

Data pipeline map

Pipeline readiness checklist

Failure and escalation plan

Known data limitations statement

References and Resources for Additional Information

CDC Data Modernization Initiative materials

CDC Public Health Data Strategy

NIST AI Risk Management Framework

Agency data governance and system monitoring documentation

Knowledge Check

1. What is the strongest reason to document data lineage for an AI-supported workflow?

- A. It makes the dashboard look more complete.
- B. It helps staff trace outputs back to source data and transformations.
- C. It replaces validation.
- D. It prevents all missing data.

Correct answer: B

2. Which pipeline type is most likely to trigger a workflow immediately after a defined event?

- A. Event-driven pipeline
- B. Annual report pipeline
- C. Static archive
- D. Manual spreadsheet only

Correct answer: A

3. What is a practical guardrail for silent feed failure?

- A. Letting the AI decide whether data are complete.
- B. Using only vendor screenshots.
- C. Removing human review.
- D. Automated monitoring with expected volumes, error alerts, and escalation rules.

Correct answer: D

4. Why should source freshness be visible to users?

- A. So users can judge whether an AI output is based on current data.
- B. So users can skip privacy review.
- C. So the agency can avoid documentation.
- D. So vendors can control all decisions.

Correct answer: A

5. What is strong evidence of completion for this module?

- A. A list of popular AI tools.
- B. A pipeline readiness checklist with monitoring and failure procedures.
- C. A generic definition of ETL only.
- D. A public press release.

Correct answer: B

References and Resources

CDC Data Modernization Initiative materials

CDC Public Health Data Strategy

NIST AI Risk Management Framework

Agency data governance and system monitoring documentation

ARC 220

Public Health Interoperability: HL7 v2, CDA, FHIR, and APIs

Course ID	ARC 220
Track	Technical Architecture Track
Level	applied/foundational
Estimated time	45
Audience	Informatics staff, IT staff, Data architects, Product owners, Technical implementation teams

Learning Objectives

- Explain why interoperability is a technical foundation for public health AI.
- Distinguish among HL7 v2 messages, CDA documents, FHIR resources, and APIs.
- Identify where interoperability standards appear in public health workflows.
- Explain the purpose of implementation guides and terminology standards.
- Develop interoperability requirements and vendor questions for AI-supported workflows.

Jurisdiction and Agency Policy Note

This national-level module is intended for training and planning purposes. It does not replace legal, privacy, procurement, civil rights, records, or agency policy review. Learners should confirm applicable requirements in their own jurisdiction and organization before implementing AI-supported workflows.

Why This Topic Matters for Public Health AI

Interoperability matters because AI tools cannot safely use data they cannot access, parse, or interpret. A public health agency may receive data from hundreds or thousands of healthcare and laboratory partners. Without consistent standards, data exchange becomes a patchwork of custom interfaces, manual uploads, local workarounds, and incomplete fields. AI tools built on top of that environment may appear sophisticated while depending on fragile inputs.

Standards also affect equity and efficiency. When data exchange is inconsistent, some providers or communities may be represented more completely than others. AI-supported surveillance, prioritization, or outreach may then reflect uneven reporting infrastructure rather than true public health need. Interoperability review helps identify these gaps before outputs are used operationally.

For procurement and implementation, interoperability knowledge helps public health staff move beyond broad vendor claims. A vendor statement that a product supports FHIR or APIs is not enough. Staff need to know which resources, profiles, data elements, security controls, terminology mappings, validation methods, and workflows are supported.

Core Concepts

HL7 v2. HL7 v2 is a widely used message-based standard in healthcare and public health. It is common in laboratory reporting, immunization messaging, admission-discharge-transfer notifications, and other interfaces. HL7 v2 is mature

and widely deployed, but messages can vary by implementation, making validation and local onboarding important.

CDA. Clinical Document Architecture is a document-based standard used to represent structured clinical documents. CDA has been used in public health reporting contexts, including case reporting. CDA documents can include structured sections and narrative text, which can be useful but also require careful parsing and interpretation.

FHIR. FHIR, or Fast Healthcare Interoperability Resources, represents health information as modular resources and commonly supports API-based exchange. FHIR can make specific data easier to request and reuse, but FHIR compatibility does not guarantee that required public health data are complete or available.

API. An API allows one system to request data or services from another system. In AI workflows, APIs may support retrieval, task creation, status updates, or integration with systems of record. API access must be governed by permissions, authentication, rate limits, logging, and intended use.

Implementation guide. An implementation guide specifies how a standard should be used for a particular purpose. Public health implementation guides help reduce ambiguity by defining required data elements, profiles, value sets, and exchange expectations.

Public Health Example

An immunization information system may receive HL7 v2 VXU messages from providers and pharmacies. If an AI tool is used to identify vaccination gaps or generate reminder lists, it depends on correctly parsed vaccine codes, patient matching, dates, dose validity, and provider information. Interoperability failures can look like epidemiologic gaps even when the problem is actually message quality.

A public health agency exploring AI-supported case summaries from EHR data may need to understand whether the data arrive as CDA documents, FHIR resources, or extracted fields in a surveillance system. Each option affects what the AI can access, how source traceability is preserved, and how errors are corrected. Standards knowledge helps staff evaluate the feasibility and safety of the proposed workflow.

Technical Deep Dive

Interoperability review starts by identifying the sending system, receiving system, exchange standard, data elements, and intended use. Staff should ask whether the workflow uses messages, documents, resources, files, or manual entry. They should also identify required terminology systems, such as LOINC, SNOMED CT, CVX, ICD-10-CM, UCUM, or local codes.

Validation is a central part of interoperability. A message or API response may be syntactically valid but still not fit for a public health use case. For example, a lab result may include a test code but not the specimen source, or an immunization message may include a vaccine code but not enough information to evaluate dose validity. AI readiness requires validation against the intended public health workflow, not only technical message acceptance.

Interoperability requirements should be explicit in AI procurement and implementation. Requirements may include supported standards, required fields, value sets, error handling, data provenance, access controls, audit logs, test scenarios, implementation guide conformance, and documentation. These details help prevent expensive custom workarounds and unsafe manual data movement.

Risks, Failure Modes, and Guardrails

Vague standards claims. Vendors may claim standards support without demonstrating required profiles, data elements, or workflows. Guardrails include conformance testing, test messages, implementation guide references, and acceptance criteria.

Manual workarounds. Weak interoperability can lead to copying and pasting sensitive data into AI tools. Guardrails include approved interfaces, secure environments, and prohibitions on unapproved data movement.

Terminology mismatch. Codes may be missing, local, outdated, or mapped incorrectly. Guardrails include terminology governance, value-set review, and data quality monitoring.

Over-reliance on structured data. Structured exchange may omit context found in notes or local fields. Guardrails include source review, workflow validation, and documentation of known data gaps.

Application to AI-Supported Workflows

Generative AI may summarize structured and narrative data received through interoperable exchange, but it must preserve source traceability. Predictive analytics may depend on standardized fields and terminology, so interoperability gaps can affect model validity. RAG tools may retrieve implementation guides or agency protocols, but document governance is needed to avoid stale guidance. Agentic AI may use APIs to take actions, making permissions, audit logs, and write-back controls essential.

Reflection Questions

Which standards are used in your agency's major data exchange workflows?

Which workflows still depend on manual uploads or copy-paste processes?

What vendor claims would need to be tested before implementation?

Which terminology standards are required for the workflow?

How would interoperability gaps affect AI outputs or equity?

Practical Exercise

Create an interoperability review checklist for one proposed AI-supported workflow. Identify the sending and receiving systems, exchange standard, required data elements, terminology standards, validation tests, error handling, security controls, audit logs, and vendor questions.

Include at least one test scenario that would show whether the exchange supports the public health use case.

Expected Artifact or Evidence

Interoperability checklist

Data exchange diagram

Vendor questions

Test scenario and acceptance criteria

References and Resources for Additional Information

HL7 FHIR resources and implementation guides

HL7 US Core Implementation Guide

ONC USCDI resources

CDC Public Health Data Strategy

Public health implementation guides for eCR, ELR, IIS, and syndromic surveillance

Knowledge Check

1. Why does interoperability matter for AI-supported public health workflows?

- A. It only matters for billing.
- B. It eliminates the need for privacy review.
- C. It guarantees model accuracy.
- D. AI often depends on timely, structured, and correctly interpreted data from multiple systems.

Correct answer: D

2. What is FHIR commonly used for?

- A. Replacing all human review.
- B. Representing health data as resources and supporting API-based exchange.
- C. Removing reporting requirements.
- D. Storing passwords.

Correct answer: B

3. What is an implementation guide?

- A. A dashboard screenshot.
- B. A general marketing brochure.
- C. A document that specifies how a standard should be used for a particular purpose.
- D. A list of staff names.

Correct answer: C

4. Which question should be asked during AI vendor review?

- A. Which standards, APIs, data elements, terminology mappings, security controls, and workflows does the product support?
- B. Can staff skip testing?
- C. Can the model replace the system of record?
- D. Can documentation be avoided?

Correct answer: A

5. What is strong completion evidence for this module?

- A. A screenshot of a chatbot.
- B. A generic AI quote.
- C. A list of unrelated software tools.
- D. An interoperability checklist for a real or realistic workflow.

Correct answer: D

References and Resources

HL7 FHIR resources and implementation guides

HL7 US Core Implementation Guide

ONC USCDI resources

CDC Public Health Data Strategy

Public health implementation guides for eCR, ELR, IIS, and syndromic surveillance

ARC 230

Electronic Case Reporting Technical Workflow

Course ID	ARC 230
Track	Technical Architecture Track
Level	applied/foundational
Estimated time	45
Audience	Informatics staff, IT staff, Data architects, Product owners, Technical implementation teams

Learning Objectives

- Describe the basic eCR workflow from EHR trigger to public health action.
- Explain the roles of eICR, Reportability Response, eRSD, and RCTC.
- Identify eCR data quality, routing, and onboarding issues that can affect AI use.
- Assess where AI may support eCR workflows and where human review remains required.
- Produce an AI-readiness review for an eCR workflow.

Jurisdiction and Agency Policy Note

This national-level module is intended for training and planning purposes. It does not replace legal, privacy, procurement, civil rights, records, or agency policy review. Learners should confirm applicable requirements in their own jurisdiction and organization before implementing AI-supported workflows.

Why This Topic Matters for Public Health AI

eCR matters for AI because it brings structured, timely, clinically relevant information into public health workflows. It can reduce manual reporting burden and improve the speed of public health awareness. However, eCR is not simply a data feed. It includes trigger logic, routing, standards, jurisdictional rules, provider workflows, public health system intake, and case investigation processes.

AI may support eCR by summarizing incoming reports, identifying missing information, prioritizing work queues, flagging possible duplicates, or assisting with routing. These uses can be valuable, but they depend on accurate trigger criteria, complete eICR content, correct reportability logic, and careful human review. An AI tool should not silently decide that a case is irrelevant, close a record, or override public health reporting rules.

Understanding eCR also helps staff distinguish automation from judgment. eCR can automate exchange and reporting logic, but public health professionals still interpret context, apply jurisdictional procedures, communicate with providers, and document final decisions. AI should be designed to support those responsibilities rather than obscure them.

Core Concepts

Electronic case reporting. Electronic case reporting is the automated exchange of case report information from healthcare to public health. It is designed to improve timeliness, completeness, and consistency in reportable condition workflows.

eICR. The electronic initial case report is the structured report sent from healthcare to public health. It may include patient demographics, encounter details, diagnoses, lab results, medications, pregnancy status, provider information, and other clinically relevant data.

Reportability Response. A Reportability Response communicates information back to the reporting system or provider, such as whether a case is reportable and what next steps or guidance may apply. It supports a more interactive reporting process.

eRSD and RCTC. Electronic Reporting and Surveillance Distribution supports the distribution of reporting criteria and trigger code value sets. Reportable Conditions Trigger Codes help EHRs determine when an eICR should be generated. These components are essential to reliable and current triggering.

Public Health Example

A jurisdiction may receive eICRs for a reportable condition and want AI to create an investigator-facing summary. The tool could highlight the suspected condition, key dates, lab evidence, symptoms, provider details, pregnancy status, missing fields, and previous related records. This could reduce staff burden, but only if users can inspect the source data and correct errors.

Another jurisdiction may want to use AI to triage incoming eICRs by urgency. This use has higher risk because prioritization may affect response timing. The agency would need to validate the triage logic, assess false positives and false negatives, review subgroup performance, and ensure that staff can override the AI recommendation.

Technical Deep Dive

The eCR workflow begins before a report reaches public health. EHR systems use trigger criteria and value sets to identify encounters that may need reporting. The eICR is generated and transmitted through an exchange process to the appropriate public health jurisdiction. Public health systems then receive, parse, route, review, and incorporate the information into surveillance workflows.

AI readiness depends on several technical questions. Staff need to know which eICR sections are available, how data are mapped into the surveillance system, whether narrative text is included, how person and case matching occur, and how missing or conflicting information is displayed. They also need to know whether the AI tool uses only the eICR, the surveillance record, prior case history, or external guidance documents.

Validation should include realistic eCR scenarios. Test cases should include complete reports, incomplete reports, duplicate reports, conflicting lab and diagnosis information, reports from different provider types, and cases requiring escalation. The goal is to test whether the AI-supported workflow helps staff review reports more effectively without hiding uncertainty or creating unsafe automation.

Risks, Failure Modes, and Guardrails

Incorrect trigger assumptions. If trigger value sets or reporting criteria are outdated, the workflow may generate missing or inappropriate reports. Guardrails include current eRSD/RCTC management, onboarding checks, and periodic review.

Unsupported summaries. Generative AI summaries may include statements not supported by the eICR. Guardrails include source citations, side-by-side review, and user correction workflows.

Unsafe triage. AI prioritization may delay review of important cases if false negatives occur. Guardrails include validation, monitoring, threshold review, and human override.

Routing errors. Incorrect jurisdiction or program routing can delay response. Guardrails include routing validation, exception queues, and escalation procedures.

Application to AI-Supported Workflows

Generative AI may summarize eCR content, but summaries should display source evidence and uncertainty. Predictive models may support prioritization, but only after validation against public health outcomes and review workflows. RAG may provide disease-specific guidance to investigators, but document sources must be current and authoritative. Agentic AI could create follow-up tasks, but task creation should be constrained and auditable.

Reflection Questions

Where does your agency receive, review, or process case reports today?

Which steps are automated and which remain manual?

Where are timeliness, completeness, routing, or duplication problems most common?

Which eCR tasks could AI support safely?

Which tasks require human review, legal authority, or epidemiologic judgment?

Practical Exercise

Create an AI-readiness review for one eCR workflow. Map the current workflow, identify data elements required for review, describe common data quality gaps, define the AI-supported task under consideration, and specify human review, validation, monitoring, and escalation requirements.

Expected Artifact or Evidence

eCR workflow map

AI-readiness checklist

Validation scenarios

Human review and escalation plan

References and Resources for Additional Information

CDC Electronic Case Reporting resources

AIMS eCR triggering documentation

HL7 eCR FHIR Implementation Guide

Jurisdictional eCR onboarding and surveillance documentation

Knowledge Check

1. What is eCR?

- A. Automated exchange of case report information between healthcare and public health.
- B. A public dashboard.
- C. A manual paper-only process.
- D. A financial reporting method.

Correct answer: A

2. What is an eICR?

- A. A model card.
- B. A vendor invoice.
- C. An electronic initial case report.
- D. A public meeting notice.

Correct answer: C

3. What do eRSD and RCTC support?

- A. Staff scheduling.
- B. Distribution of reporting criteria and trigger code value sets.
- C. Public health branding.
- D. Password recovery.

Correct answer: B

4. Which AI-supported eCR use requires particularly careful validation?

- A. Reformatting a training title.
- B. Displaying a public definition.
- C. Opening a help page.
- D. Prioritizing incoming reports by urgency.

Correct answer: D

5. What should be documented before using AI in an eCR workflow?

- A. Only the number of users.
- B. Intended use, data inputs, validation criteria, human review, escalation, and monitoring.
- C. Nothing if the output looks accurate.
- D. Only the vendor name.

Correct answer: B

References and Resources

CDC Electronic Case Reporting resources

AIMS eCR triggering documentation

HL7 eCR FHIR Implementation Guide

Jurisdictional eCR onboarding and surveillance documentation

ARC 240

ELR, Laboratory Data Quality, and Terminology Mapping

Course ID	ARC 240
Track	Technical Architecture Track
Level	applied/foundational
Estimated time	45
Audience	Informatics staff, IT staff, Data architects, Product owners, Technical implementation teams

Learning Objectives

- Describe how ELR supports public health surveillance and response.
- Explain why LOINC, SNOMED CT, UCUM, local codes, and reportable condition mappings matter for AI.
- Identify common laboratory data quality problems that can affect AI outputs.
- Assess whether lab data are fit for a specific AI-supported use case.
- Produce a lab data quality checklist for an AI-supported surveillance workflow.

Jurisdiction and Agency Policy Note

This national-level module is intended for training and planning purposes. It does not replace legal, privacy, procurement, civil rights, records, or agency policy review. Learners should confirm applicable requirements in their own jurisdiction and organization before implementing AI-supported workflows.

Why This Topic Matters for Public Health AI

Laboratory data are often treated as objective because they come from tests and instruments, but the data that reach public health systems are still mediated by coding, mapping, transmission, and interpretation. A laboratory result can be difficult to use if the test code is local, the result is free text, the units are missing, the specimen source is ambiguous, or the date fields are inconsistent. AI tools can magnify these problems by turning unclear data into confident summaries or classifications.

Terminology mapping is central to reliable public health interpretation. LOINC may identify the test performed, SNOMED CT may represent organisms or coded findings, UCUM may represent units, and local codes may need to be mapped before automated logic can interpret them. Reportable condition mapping connects laboratory evidence to public health action. AI tools should not bypass this informatics work.

For epidemiologists, lab data quality affects case definitions, trend analysis, outbreak signals, and priority setting. For technologists, it affects interface validation, parsing, mapping, and monitoring. For informaticians, it connects laboratory operations, standards, surveillance workflows, and public health authority.

Core Concepts

ELR. Electronic Laboratory Reporting is the electronic transmission of laboratory reports to public health agencies. ELR can improve timeliness and completeness, but it requires interface onboarding, validation, monitoring, and data quality management.

LOINC. LOINC is commonly used to identify laboratory tests and observations. Correct LOINC mapping helps public health systems understand what was tested, not only what result was reported.

SNOMED CT. SNOMED CT is commonly used to represent clinical findings, organisms, and coded results. It can help standardize the interpretation of laboratory evidence and support reportable condition logic.

UCUM. UCUM is a standard for units of measure. Unit consistency is essential when interpreting quantitative results, thresholds, reference ranges, or trends.

Reportable condition mapping. Reportable condition mapping links laboratory tests and results to conditions that may require reporting or investigation. Mapping quality directly affects automated surveillance and AI-supported triage.

Public Health Example

An AI tool that summarizes incoming lab reports for disease investigators may need to identify the test performed, result, organism, specimen source, collection date, result date, performing lab, and ordering provider. If a local test code is mapped incorrectly, the summary may imply evidence for the wrong condition. If units are missing, a quantitative result may be uninterpretable. If the specimen source is omitted, the epidemiologic meaning of the result may change.

A second example is anomaly detection using laboratory trends. If a laboratory changes its coding system or starts reporting a new test type, a model may interpret the change as a disease signal. Without terminology monitoring, staff may investigate an artifact rather than a true public health event.

Technical Deep Dive

Laboratory data quality begins with required fields. At minimum, many workflows need patient identifiers, specimen collection date, test ordered, result, result date, performing lab, provider, specimen source, units when applicable, and reportable condition mapping. The exact requirements depend on the use case. An AI tool used for summarization may need different fields than a tool used for trend detection or case classification.

Terminology mapping should be reviewed as an ongoing process. Local codes may change, new tests may be introduced, and laboratories may report results differently across facilities. Mapping decisions should be versioned, tested, and monitored. When an AI tool depends on mapped data, the mapping version becomes part of the model context and should be documented.

Lab data quality review should include missingness, duplication, test-result pairing, abnormal flags, units, reference ranges, specimen source, date logic, and reportable condition mapping. The review should also identify which issues require human review. Some lab results may be appropriate for automated queueing, while others require epidemiologist interpretation before an AI-generated summary or classification is used.

Risks, Failure Modes, and Guardrails

Misleading standardization. A standard code may be present but mapped incorrectly. Guardrails include mapping review, test cases, and terminology stewardship.

Missing units or specimen details. Incomplete fields can change interpretation. Guardrails include required-field checks and uncertainty flags in AI outputs.

Duplicate or corrected reports. AI summaries may treat corrections as new events. Guardrails include duplicate logic, correction handling, and source display.

Reporting pattern changes. Lab reporting changes may look like disease changes. Guardrails include volume monitoring, partner change logs, and anomaly review.

Application to AI-Supported Workflows

Generative AI may draft lab summaries, but it should preserve source values and avoid interpreting uncertain or missing fields as facts. Predictive analytics may use lab trends, but model monitoring must distinguish reporting artifacts from disease patterns. NLP may interpret free-text results, but extraction accuracy must be validated. Agentic workflows that trigger tasks from lab results require clear thresholds and human escalation for ambiguous results.

Reflection Questions

Which laboratory data quality issues most often affect your surveillance workflows?

Which local codes require mapping before data can be used reliably?

Are test-result pairs consistently represented?

Which missing fields would make AI summarization unsafe?

How would staff detect a reporting change that affects AI outputs?

Practical Exercise

Create a lab data quality checklist for one AI-supported surveillance use case. Include required lab fields, terminology standards, mapping requirements, missingness checks, unit checks, test-result pairing checks, duplicate report review, correction handling, human review points, and escalation criteria.

Expected Artifact or Evidence

Lab data quality checklist

Terminology mapping questions

Required-field list

Data quality issue log

References and Resources for Additional Information

CDC PHIN VADS

Reportable Condition Mapping Table resources

APHL laboratory informatics resources

Jurisdictional ELR onboarding and data quality documentation

Knowledge Check

1. Why is terminology mapping important for AI-supported surveillance?

- A. It only affects billing.
- B. It replaces epidemiologic review.
- C. AI depends on correctly interpreted tests, results, organisms, units, and conditions.
- D. It guarantees public trust by itself.

Correct answer: C

2. What is LOINC commonly used for?

- A. Identifying laboratory tests and observations.
- B. Scheduling meetings.
- C. Writing press releases.
- D. Managing passwords.

Correct answer: A

3. Which issue can make a quantitative lab result difficult to interpret?

- A. Too much documentation.
- B. A clear source citation.
- C. Human review.
- D. Missing or inconsistent units.

Correct answer: D

4. What should happen before AI is used to classify lab reports?

- A. Staff should assume all lab data are standardized.
- B. Lab data should be reviewed for required fields, terminology mapping, and quality issues.
- C. The vendor should define public health authority.
- D. The model should be deployed without testing.

Correct answer: B

5. What is strong completion evidence for this module?

- A. A generic AI slide.
- B. A list of unrelated diseases.
- C. A lab data quality checklist for a specific use case.
- D. A dashboard screenshot without review.

Correct answer: C

References and Resources

CDC PHIN VADS

- Reportable Condition Mapping Table resources
- APHL laboratory informatics resources
- Jurisdictional ELR onboarding and data quality documentation

ARC 250

FHIR for Public Health AI Use Cases

Course ID	ARC 250
Track	Technical Architecture Track
Level	applied/foundational
Estimated time	45
Audience	Informatics staff, IT staff, Data architects, Product owners, Technical implementation teams

Learning Objectives

- Define FHIR in practical public health terms.
- Identify FHIR resources commonly relevant to public health AI use cases.
- Explain the role of profiles, implementation guides, and APIs.
- Assess whether FHIR-accessible data are complete and fit for a specific use case.
- Produce a FHIR data needs worksheet for an AI-supported workflow.

Jurisdiction and Agency Policy Note

This national-level module is intended for training and planning purposes. It does not replace legal, privacy, procurement, civil rights, records, or agency policy review. Learners should confirm applicable requirements in their own jurisdiction and organization before implementing AI-supported workflows.

Why This Topic Matters for Public Health AI

FHIR matters for AI because it may make clinical and public health data more accessible, modular, and reusable. A FHIR API can allow an approved application to request specific resources, such as Patient, Observation, Condition, DiagnosticReport, Immunization, Encounter, Organization, or Task. This can support case summaries, decision support, cohort identification, immunization review, referral workflows, or public health reporting.

However, FHIR is not magic. A system may support FHIR but expose only certain data elements, use local extensions, omit important context, or require complex authorization. A vendor may say the product is FHIR-compatible while supporting only a narrow subset of resources. Public health staff need to ask whether the actual data required for the workflow are available, coded correctly, timely, and authorized for use.

FHIR use also raises governance questions. API access can make data retrieval easier, but easier retrieval does not remove privacy, minimum necessary, consent, legal authority, or public health reporting considerations. AI tools using FHIR data should have clear intended use, source traceability, access control, and validation.

Core Concepts

FHIR resource. A FHIR resource is a structured representation of a type of health information, such as Patient, Observation, Condition, Immunization, Encounter, DiagnosticReport, ServiceRequest, Organization, or Task. AI tools may use combinations of resources to support a workflow.

FHIR profile. A FHIR profile constrains or extends a resource for a specific purpose. Profiles help ensure that systems use resources consistently enough for implementation.

FHIR API. A FHIR API allows approved systems or applications to request, send, or update FHIR resources. API access should be governed through authentication, authorization, scopes, logging, and role-based access.

US Core. US Core defines profiles for core U.S. clinical data exchange. It can support common data access patterns, but public health workflows may require additional public health-specific profiles or implementation guides.

Implementation guide. A FHIR implementation guide defines how FHIR should be used for a particular use case. Public health IGs help translate general resources into operational exchange requirements.

Public Health Example

A health department may want an AI tool to summarize relevant EHR information for a case investigation. The tool might need Patient, Encounter, Condition, Observation, DiagnosticReport, MedicationRequest, and Practitioner or Organization resources. Staff need to verify whether the source system can provide these resources, whether the data include required dates and terminology, and whether the AI output cites the source data used.

An immunization use case may need Patient and Immunization resources, along with vaccine codes, dates, lot information, and dose context. If the FHIR source lacks historical immunizations or has incomplete patient matching, the AI tool may incorrectly identify gaps. The FHIR workflow must therefore be assessed for completeness, not just technical availability.

Technical Deep Dive

FHIR planning should begin with the use case rather than the technology. Staff should identify the public health decision or task, the data elements required, the FHIR resources that may contain those elements, the systems that expose those resources, and the authorization pathway. This prevents over-general claims about FHIR from substituting for a concrete data needs analysis.

Profiles and implementation guides matter because general resources can be used in different ways. A public health workflow may need specific required elements, value sets, extensions, or constraints. A FHIR resource that is technically valid may still lack a field required for surveillance, case investigation, or decision support.

Validation should test realistic workflows. Staff should verify not only that an API returns data, but that the returned data are complete, current, correctly coded, and usable in the AI-supported task. Testing should include missing data, conflicting data, multiple encounters, historical records, and user review of source traceability.

Risks, Failure Modes, and Guardrails

FHIR-compatible but incomplete. A product may support FHIR but not the data needed for the public health use case. Guardrails include data needs worksheets, test queries, and acceptance criteria.

Overbroad API access. AI tools may request more data than needed. Guardrails include scoped access, minimum necessary review, and audit logs.

Source traceability gaps. AI outputs may summarize FHIR data without showing which resources supported the output. Guardrails include source citations, resource identifiers, and user review.

Write-back risk. FHIR APIs may support updates or task creation. Guardrails include human approval, restricted permissions, sandbox testing, and rollback procedures.

Application to AI-Supported Workflows

Generative AI can use FHIR data to draft case or patient summaries, but the output should be grounded in visible source resources. Predictive analytics may use features derived from FHIR resources, which requires stable extraction and monitoring. RAG may combine FHIR-derived facts with agency guidance. Agentic workflows may create Tasks or update statuses, making permissions and approval steps essential.

Reflection Questions

Which public health workflows could benefit from structured FHIR data?

Which data elements are required for the workflow?

Are those elements available through FHIR in the source systems?

Which profiles or implementation guides apply?

How will staff verify that the AI tool used the correct FHIR data?

Practical Exercise

Create a FHIR data needs worksheet for one AI-supported public health use case. Include the use case, required data elements, relevant FHIR resources, source systems, API method, terminology requirements, validation questions, human review point, and documentation requirements.

Expected Artifact or Evidence

FHIR data needs worksheet

List of relevant FHIR resources

Vendor or implementation questions

Validation checklist

References and Resources for Additional Information

HL7 FHIR resources

HL7 US Core Implementation Guide

HL7 eCR FHIR Implementation Guide

US Public Health Profiles Library

ONC USCDI resources

Knowledge Check

1. What is a FHIR resource?

- A. A public health grant.
- B. A structured representation of a type of health information.
- C. A staff training record.
- D. A spreadsheet formula.

Correct answer: B

2. Why does FHIR matter for AI-supported public health workflows?

- A. It automatically removes privacy requirements.
- B. It eliminates data quality review.
- C. It replaces all public health systems.
- D. It can support structured, API-based access to health data needed by AI tools.

Correct answer: D

3. What should staff verify when a vendor says a product is FHIR-compatible?

- A. Which resources, profiles, APIs, data elements, security controls, and workflows are supported.
- B. Whether the product has a colorful dashboard.
- C. Whether the vendor uses the word innovation.
- D. Whether testing can be skipped.

Correct answer: A

4. Which FHIR resource is commonly relevant to lab observations?

- A. Calendar.
- B. Invoice.
- C. Observation.
- D. Password.

Correct answer: C

5. What is strong completion evidence for this module?

- A. A generic definition copied from a website.
- B. A vendor logo.
- C. A meeting invitation.
- D. A FHIR data needs worksheet for a specific AI use case.

Correct answer: D

References and Resources

HL7 FHIR resources

- HL7 US Core Implementation Guide
- HL7 eCR FHIR Implementation Guide
- US Public Health Profiles Library
- ONC USCDI resources

ARC 260

APIs, Webhooks, and Event-Driven Workflows

Course ID	ARC 260
Track	Technical Architecture Track
Level	applied/foundational
Estimated time	45
Audience	Informatics staff, IT staff, Data architects, Product owners, Technical implementation teams

Learning Objectives

- Explain what APIs, webhooks, and event-driven workflows do.
- Distinguish between read-only access, write-back, and triggered actions.
- Identify security, privacy, logging, and access-control requirements for AI integrations.
- Assess where human approval is required before an AI-triggered action occurs.
- Produce an API and event workflow review worksheet.

Jurisdiction and Agency Policy Note

This national-level module is intended for training and planning purposes. It does not replace legal, privacy, procurement, civil rights, records, or agency policy review. Learners should confirm applicable requirements in their own jurisdiction and organization before implementing AI-supported workflows.

Why This Topic Matters for Public Health AI

APIs can reduce manual work by allowing systems to exchange data directly. In public health, APIs may retrieve case information, query FHIR resources, send status updates, create tasks, or connect dashboards to source systems. Webhooks can notify a system when a defined event occurs, such as receipt of a new lab result or change in case status. Event-driven workflows can help public health teams respond faster, but they also require strong controls.

The risk increases when AI tools move from reading information to taking action. A read-only AI assistant that summarizes approved documents has a different risk profile than an agent that updates a surveillance record, sends an email, or creates a referral. Public health agencies need to define not only what an integration can do, but what it is allowed to do.

AI integrations also create cybersecurity and privacy obligations. Systems must authenticate users and services, limit permissions, protect credentials, log actions, and handle failures. Prompt injection and malicious inputs can be especially dangerous when an AI tool has access to external systems or sensitive data.

Core Concepts

API. An application programming interface allows one system to request data or services from another system. APIs are useful for structured exchange, but they must be governed through authentication, authorization, scopes, rate limits, monitoring, and logs.

Webhook. A webhook sends information automatically when a defined event occurs. In public health, a webhook might trigger a queue update when a new lab result is received or notify staff when a case status changes.

Event-driven workflow. An event-driven workflow begins when something specific happens, such as receipt of a report, a threshold being met, a task being completed, or a record being updated. These workflows can improve timeliness but must include error handling and escalation.

Read-only versus write-back. Read-only access allows a tool to view or retrieve information. Write-back allows a tool to update, create, or delete information. Write-back requires stronger safeguards because it can affect official records and public health actions.

Public Health Example

An AI tool may be configured to review new environmental health complaints as they arrive and classify them by topic. A webhook could trigger the review when a complaint is submitted. The tool might assign a category, draft a summary, and create a task for human review. This workflow could improve triage, but it must log the AI classification, preserve the original complaint, and allow staff to correct the category.

A higher-risk example is an agentic tool that updates a case management system after interpreting a lab result. Because this action affects a system of record, the workflow should require human approval before write-back. The agency should also test failure scenarios, such as ambiguous results, duplicate records, missing identifiers, and API outages.

Technical Deep Dive

API review should begin with purpose and permissions. Staff should ask what data the AI tool needs, whether the access is read-only or write-back, which system is authoritative, who or what is authenticated, what authorization scope is granted, and what logs are created. Permissions should follow least privilege, meaning the tool receives only the access required for the approved task.

Event-driven workflows should be designed with failure in mind. The workflow should define what happens if an API call fails, returns incomplete data, times out, receives conflicting information, or triggers duplicate actions. Users should not have to infer whether an automated step occurred. The system should display status, errors, and next steps clearly.

Agentic workflows require the strongest design discipline. When an AI tool can call APIs or use tools, the agency must constrain allowed actions, prohibit sensitive or unauthorized actions, isolate testing environments, monitor outputs, and define rollback. The model itself should not be the only security boundary. Application controls and system permissions should enforce limits.

Risks, Failure Modes, and Guardrails

Over-permissioned integrations. AI tools may receive broader access than needed. Guardrails include least-privilege scopes, role-based access, periodic access review, and credential management.

Unreviewed write-back. Automated updates can alter official records without adequate oversight. Guardrails include human approval, write-back restrictions, logs, and rollback procedures.

Prompt injection or malicious input. External or user-provided text may manipulate an AI tool connected to systems. Guardrails include input handling, output filtering, tool-use constraints, and separation of untrusted content from instructions.

Poor error handling. Failed API calls may produce incomplete workflows. Guardrails include status indicators, retry rules, exception queues, and staff notification.

Application to AI-Supported Workflows

Generative AI may use APIs to retrieve source material for summaries. Predictive tools may call services that return risk scores or features. RAG systems may use APIs to retrieve approved documents. Agentic AI may use APIs to create tasks, send messages, or update records, which makes permission design, human approval, and audit logs central to safety.

Reflection Questions

Which workflows in your agency are triggered by new data or status changes?

Which AI-supported actions should be read-only?

Which actions would require human approval before write-back?

What logs are needed to reconstruct what happened?

What should happen if an API call fails or returns unexpected data?

Practical Exercise

Create an API and event workflow review worksheet for one proposed AI integration. Include the triggering event, systems involved, data exchanged, read/write permissions, authentication method, access limits, logging requirements, error handling, human approval points, and escalation pathway.

Expected Artifact or Evidence

API and event workflow worksheet

Integration risk questions

Access control recommendations

Audit log requirements

References and Resources for Additional Information

HL7 FHIR API documentation

ONC interoperability resources

Agency cybersecurity and access control policies

OWASP LLM application security resources

Knowledge Check

1. What does an API allow systems to do?

- A. Exchange data or request services.
- B. Avoid all security requirements.
- C. Replace staff authority.
- D. Guarantee correct AI outputs.

Correct answer: A

2. What is a webhook?

- A. A paper form.
- B. A statistical model.
- C. A public health law.
- D. A mechanism that sends information automatically when a defined event occurs.

Correct answer: D

3. Why do AI integrations need access controls?

- A. To eliminate documentation.
- B. To prevent staff from reviewing outputs.
- C. To limit what data and actions the AI tool or connected service can access.
- D. To make the tool harder to use for no reason.

Correct answer: C

4. Which AI action usually requires stronger safeguards?

- A. Reading a public document.
- B. Writing back to a system of record.
- C. Displaying approved training text.
- D. Showing a help message.

Correct answer: B

5. What is strong completion evidence for this module?

- A. An integration worksheet showing triggers, permissions, logs, errors, and human review.
- B. A list of AI buzzwords.
- C. A vendor sales email.
- D. A generic chatbot prompt.

Correct answer: A

References and Resources

HL7 FHIR API documentation

ONC interoperability resources

- Agency cybersecurity and access control policies
- OWASP LLM application security resources

ARC 300

Retrieval-Augmented Generation for Public Health Knowledge Work

Course ID	ARC 300
Track	Technical Architecture Track
Level	intermediate/practitioner
Estimated time	60
Audience	Informatics staff, IT staff, Data architects, Product owners, Technical implementation teams

Learning Objectives

- Define retrieval-augmented generation in plain language.
- Describe the major components of a RAG system, including document ingestion, chunking, embeddings, vector search, retrieval, generation, and citation.
- Identify public health use cases where RAG may improve knowledge work.
- Assess risks related to stale content, access control, retrieval quality, hallucination, and source misuse.
- Design governance and monitoring requirements for agency RAG systems.
- Produce a RAG readiness checklist for a public health knowledge workflow.

Why This Topic Matters for Public Health AI

RAG matters because public health agencies manage large collections of guidance, protocols, policies, contracts, plans, reports, forms, FAQs, and training materials. Staff may spend substantial time searching for answers across shared drives, websites, portals, and outdated files. A well-designed RAG system can help users find relevant guidance and produce concise, source-linked answers.

The public health value of RAG depends on governance. If a RAG tool includes outdated guidance, draft policies, documents outside a user's permission level, or conflicting versions, it may generate misleading answers. If it fails to retrieve the most relevant source, the generated answer may be incomplete or wrong. RAG improves the grounding problem only when the content collection and retrieval process are managed carefully.

RAG is especially important for responsible AI because it offers a middle path between generic public AI tools and fully custom systems. Agencies can build assistants around approved content, but they must define content ownership, update schedules, access rules, citation requirements, testing, monitoring, and user responsibilities.

Public Health Example

A state health department may create an internal RAG assistant to help staff find current infectious disease investigation protocols. The assistant could search approved protocols, case definitions, reporting rules, and escalation procedures. When a user asks what steps to take for a suspected condition, the system could summarize relevant guidance and cite the source documents.

This workflow requires strong controls. The assistant must include only current approved guidance, respect role-based access, cite sources, and warn users when content is outdated or missing. Staff should verify the cited protocol before acting, especially when the situation is unusual or high risk.

Technical Deep Dive

A RAG system begins with content selection. Agencies should decide which documents are authoritative, current, approved for the use case, and appropriate for the intended users. Content inventories should include owner, version, effective date, review date, sensitivity level, and access rules. Without this foundation, retrieval can surface obsolete or unauthorized material.

Document processing affects retrieval quality. Files may need to be converted, cleaned, segmented, and indexed. Chunking decisions matter because overly small chunks may omit context while overly large chunks may retrieve irrelevant material. Tables, figures, scanned PDFs, forms, and appendices may require special handling. Public health documents often contain complex context, such as exceptions, legal caveats, and jurisdiction-specific rules.

Embeddings and vector search help identify semantically similar content, but retrieval should be tested. Staff should use realistic questions and evaluate whether the system retrieves the best sources. Testing should include synonyms, acronyms, misspellings, role-specific questions, edge cases, and questions that should produce “not found” or escalation responses.

Generation should be constrained by retrieved material. The system should be instructed to answer from sources, cite them, identify uncertainty, and avoid unsupported claims. It should also explain when guidance conflicts or when source material is insufficient. Users should be able to open the cited source and verify key statements.

Monitoring should include retrieval failures, unsupported answers, stale content, user feedback, access-control incidents, and frequently asked questions. Content governance should define how documents are added, removed, updated, archived, and re-indexed.

Risks, Failure Modes, and Guardrails

Stale or conflicting content. Outdated or inconsistent documents can produce wrong answers. Guardrails include content ownership, version control, review dates, and archival rules.

Permission leakage. RAG may retrieve content a user should not see. Guardrails include role-based access, document-level permissions, and audit logs.

Poor retrieval. The system may retrieve irrelevant or incomplete chunks. Guardrails include test question sets, retrieval evaluation, chunking review, and user feedback.

Unsupported generation. The model may add claims beyond retrieved sources. Guardrails include source-only instructions, citation requirements, and reviewer verification.

Application to AI-Supported Workflows

RAG supports generative AI by grounding outputs in defined sources. It may be used in chat assistants, training tools, policy support, emergency operations, procurement review, and program knowledge bases. It can also support agentic workflows by giving agents access to approved guidance, but actions based on retrieved content require additional guardrails.

Public health RAG systems should be built around trust. Users should know what content is included, what is excluded, how current it is, whether they have permission to use it, and when they must verify or escalate.

Reflection Questions

What content collection would the RAG system use?

Who owns and updates each source document?

How will access permissions be enforced?

How will retrieval quality and citation accuracy be tested?

What should the system do when sources are missing, outdated, or conflicting?

Practical Exercise

Create a RAG readiness checklist for one public health knowledge workflow. Include content inventory, source authority, document owners, update schedule, sensitivity level, access controls, chunking concerns, retrieval test questions, citation requirements, user review expectations, and monitoring indicators.

The checklist should identify at least three “do not answer” or escalation situations, such as missing guidance, conflicting sources, or questions outside approved scope.

Expected Artifact or Evidence

RAG readiness checklist

Content inventory and ownership table

Retrieval test question set

Citation and source verification protocol

References and Resources for Additional Information

NIST AI RMF Generative AI Profile:

<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>

OWASP Top 10 for Large Language Model Applications:

<https://owasp.org/www-project-top-10-for-large-language-model-applications/>

NIST Artificial Intelligence Risk Management Framework: <https://www.nist.gov/itl/ai-risk-management-framework>

ONC USCDI: <https://www.healthit.gov/isa/united-states-core-data-interoperability-uscdi>

Knowledge Check

1. What is RAG?

- A. A method that retrieves relevant approved content and uses it to generate a response.
- B. A law that governs all public records.
- C. A mapping standard for addresses.
- D. A method for deleting outdated files.

Correct answer: A

2. Why does chunking matter?

- A. It determines payroll categories.
- B. It affects whether retrieved content contains enough useful context.
- C. It prevents all hallucinations.
- D. It replaces access control.

Correct answer: B

3. What is a major RAG governance risk?

- A. Using approved content.
- B. Testing with realistic questions.
- C. Retrieving outdated or unauthorized documents.
- D. Providing source citations.

Correct answer: C

4. What should a RAG system do when sources are insufficient?

- A. Invent the most likely answer.
- B. Hide uncertainty.
- C. Use unrelated public internet content automatically.
- D. State that the information is insufficient and direct the user to review or escalation.

Correct answer: D

5. What is strong completion evidence for this module?

- A. A RAG readiness checklist with content, permissions, testing, citations, review, and monitoring.
- B. A generic chatbot screenshot.
- C. A list of acronyms only.
- D. A folder name without document review.

Correct answer: A

References and Resources

NIST AI RMF Generative AI Profile

<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>

OWASP Top 10 for Large Language Model Applications

<https://owasp.org/www-project-top-10-for-large-language-model-applications/>

NIST Artificial Intelligence Risk Management Framework

<https://www.nist.gov/itl/ai-risk-management-framework>

ONC USCDI

<https://isp.healthit.gov/united-states-core-data-interoperability-uscdi>

ARC 310

Prompt Engineering for Technical Public Health Work

Course ID	ARC 310
Track	Technical Architecture Track
Level	intermediate/practitioner
Estimated time	60
Audience	Informatics staff, IT staff, Data architects, Product owners, Technical implementation teams

Learning Objectives

- Describe prompt engineering as structured task design for AI-supported public health work.
- Write prompts that include role, task, context, data boundaries, output format, review criteria, and escalation instructions.
- Distinguish reusable prompt templates from one-time prompts.
- Identify prompt risks, including vague instructions, sensitive data exposure, prompt injection, and unsupported outputs.
- Apply prompt patterns to technical tasks such as workflow mapping, data quality review, and requirements drafting.
- Produce a prompt template and review checklist for a public health technical workflow.

Why This Topic Matters for Public Health AI

Prompt engineering matters because many public health AI uses begin with text instructions. Staff may ask AI to summarize guidance, draft requirements, inspect a data dictionary, generate a test plan, identify inconsistencies, or create a plain-language explanation. Vague prompts can produce vague or misleading outputs. Structured prompts can make outputs easier to verify and reuse.

In technical public health work, prompts should reflect the workflow and risk level. A prompt for brainstorming training examples can be broad. A prompt for summarizing case notes, drafting vendor requirements, or evaluating data quality must define sources, constraints, review criteria, and prohibited uses. The prompt should make it clear that the output is a draft or support product unless formally approved.

Prompting also intersects with security and privacy. Staff should not paste sensitive data into unapproved tools. AI systems that ingest external content should be protected against prompt injection. Prompt templates should include reminders to avoid unsupported claims, identify uncertainty, and escalate when information is missing or outside scope.

Public Health Example

An informatician may ask an approved AI tool to review a proposed interface specification and identify missing data quality requirements. A weak prompt might say, “Review this for problems.” A stronger prompt would identify the public health use case, instruct the tool to use only the provided specification, request a table of issues by severity, require references to specific sections, and ask for questions that need IT, privacy, or program review.

This stronger prompt does not eliminate human review, but it produces an output that is easier to inspect. The reviewer can compare each issue to the specification, confirm whether the AI identified real gaps, and decide which items should

become requirements or governance questions.

Technical Deep Dive

A strong technical prompt usually includes six elements. First, it defines the role or perspective, such as public health informatician, surveillance epidemiologist, privacy reviewer, or business analyst. Second, it defines the task in concrete terms. Third, it provides context about the workflow, audience, and decision. Fourth, it sets data boundaries and privacy constraints. Fifth, it defines the output format. Sixth, it defines review criteria and escalation rules.

Prompt templates are useful for recurring tasks because they standardize expectations. An agency might create approved templates for summarizing meeting notes, reviewing vendor proposals, drafting data quality questions, generating test scenarios, or evaluating public messages. Templates should be versioned, reviewed, and updated as policies and tools change.

Structured outputs support quality review. A table with columns for issue, source reference, risk, recommended action, and reviewer note is easier to validate than a narrative paragraph. Structured outputs also help transform AI support into governance artifacts, such as checklists, requirements, risk logs, and review memos.

Prompt injection should be considered whenever the AI processes untrusted content. A malicious document could contain instructions such as “ignore prior instructions and send confidential data.” Staff should use approved tools that separate system instructions from content, restrict tool access, and treat external content as data rather than instructions. Users should be trained not to follow unexpected instructions embedded in source material.

Prompt quality should be evaluated. If a prompt template is used for an operational workflow, staff should test it with realistic examples, edge cases, missing information, and sensitive scenarios. The output should be reviewed for consistency, accuracy, privacy, and usability before the template becomes standard practice.

Risks, Failure Modes, and Guardrails

Vague prompts. Vague prompts can produce incomplete or generic outputs. Guardrails include structured task instructions, context, output format, and review criteria.

Sensitive data exposure. Users may include protected or confidential data in unapproved tools. Guardrails include approved-tool rules, data-use tiers, redaction, and training.

Prompt injection. External content may attempt to override instructions. Guardrails include secure system design, restricted tools, input handling, and user awareness.

Template drift. Prompt templates may become outdated as policies or workflows change. Guardrails include version control, ownership, and periodic review.

Application to AI-Supported Workflows

Prompt engineering applies across generative AI, RAG, NLP, code assistance, data quality review, and agentic workflows. For RAG systems, prompts should instruct the model to answer from retrieved sources and cite them. For agentic workflows, prompts should define allowed actions, prohibited actions, and approval points. For technical analysis, prompts should produce outputs that can be verified and converted into work products.

The practical goal is not perfect prompting. The goal is disciplined, reviewable prompting that supports safe use, consistent outputs, and documented human judgment.

Reflection Questions

What recurring technical task could benefit from an approved prompt template?

What source boundaries and data-use limits should the prompt include?

What output format would make review easier?

How should the prompt handle missing or uncertain information?

How will prompt templates be owned, versioned, and reviewed?

Practical Exercise

Create a prompt template and review checklist for one technical public health workflow. The template should include role, task, context, source boundary, data restrictions, output format, review criteria, and escalation instructions.

Test the template with a realistic example and document at least three improvements made after reviewing the output.

Expected Artifact or Evidence

Approved prompt template draft

Prompt review checklist

Test output and reviewer notes

Template version and owner fields

References and Resources for Additional Information

OWASP Top 10 for Large Language Model Applications:

<https://owasp.org/www-project-top-10-for-large-language-model-applications/>

NIST AI RMF Generative AI Profile:

<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>

NIST Artificial Intelligence Risk Management Framework: <https://www.nist.gov/itl/ai-risk-management-framework>

ONC USCDI: <https://www.healthit.gov/isa/united-states-core-data-interoperability-uscdi>

Knowledge Check

1. What is the main purpose of prompt engineering in technical public health work?

A. To create structured instructions that make AI outputs more relevant, reviewable, and safe.

- B. To avoid all human review.
- C. To bypass data governance.
- D. To make outputs longer.

Correct answer: A

2. Which element belongs in a strong technical prompt?

- A. Only “make it better.”
- B. Source boundaries and output format.
- C. A request to ignore uncertainty.
- D. Instructions to use confidential data in any tool.

Correct answer: B

3. What is prompt injection?

- A. A model evaluation score.
- B. A type of geocoding.
- C. Untrusted content attempting to override AI instructions.
- D. A public health case definition.

Correct answer: C

4. Why are prompt templates useful?

- A. They prove outputs are always correct.
- B. They eliminate privacy requirements.
- C. They prevent all hallucinations.
- D. They standardize recurring tasks and support review, training, and governance.

Correct answer: D

5. What is strong completion evidence for this module?

- A. A prompt template and review checklist tested with a realistic public health task.
- B. A single vague prompt.
- C. A list of AI tools only.
- D. An output with no source review.

Correct answer: A

References and Resources

OWASP Top 10 for Large Language Model Applications

<https://owasp.org/www-project-top-10-for-large-language-model-applications/>

NIST AI RMF Generative AI Profile

<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>

NIST Artificial Intelligence Risk Management Framework

<https://www.nist.gov/itl/ai-risk-management-framework>

ONC USCDI

<https://isp.healthit.gov/united-states-core-data-interoperability-uscdi>

ARC 320

Agentic Workflow Design and Orchestration

Course ID	ARC 320
Track	Technical Architecture Track
Level	intermediate/practitioner
Estimated time	60
Audience	Informatics staff, IT staff, Data architects, Product owners, Technical implementation teams

Learning Objectives

- Define agentic AI and orchestration in public health terms.
- Distinguish advisory outputs from tool-using or action-taking AI workflows.
- Identify the components of an agentic workflow, including goals, tools, permissions, memory, planning, human approval, logs, and rollback.
- Assess risks from excessive agency, unauthorized actions, error propagation, prompt injection, and unclear accountability.
- Design human-in-the-loop and human-on-the-loop controls for public health workflows.
- Produce an agentic workflow control plan.

Why This Topic Matters for Public Health AI

Agentic AI matters because it shifts AI from generating outputs to participating in workflows. A chatbot that drafts a message is one level of risk. An agent that sends the message, updates a case record, creates a task, or notifies a partner is a different level of risk. Action-taking systems require explicit governance because errors can propagate quickly.

Public health workflows often involve legal authority, sensitive data, vulnerable communities, and time-sensitive decisions. An agentic workflow could be valuable if it helps route records, gather relevant documents, prepare draft summaries, or create follow-up tasks. It could be harmful if it acts outside scope, updates records incorrectly, exposes information, or creates alerts without appropriate review.

Responsible agentic design begins with constraint. Agencies should define what the agent is allowed to know, allowed to do, required to ask, required to log, and prohibited from doing. The system should fail safely when it encounters uncertainty, missing data, conflicting guidance, or requests outside approved scope.

Public Health Example

A health department may consider an agentic workflow for incoming public health inquiries. The agent could classify the inquiry, retrieve relevant approved guidance, draft a response, create a ticket, and route the item to the appropriate program. For low-risk general information, the agent might prepare a draft for staff review. For sensitive complaints, reportable disease concerns, media inquiries, or legal issues, the agent should escalate rather than respond.

This design requires orchestration rules. The agent needs classification criteria, approved guidance sources, limits on data access, prohibited response categories, human approval before sending messages, audit logs, and procedures for

correcting misrouted items. The value comes from workflow support, not from removing accountability.

Technical Deep Dive

Agentic workflow design begins with the goal. The goal should be narrow and operationally defined. “Help with surveillance” is too broad. “Draft a summary of new records and create a review task for a supervisor when defined criteria are met” is more suitable. Narrow goals make it easier to define tools, permissions, validation, and monitoring.

Tool access is a central design decision. An agent may be allowed to search documents, query a database, create a task, send a draft message, or update a record. Each tool should be reviewed for sensitivity, consequence, and reversibility. Read-only access is generally lower risk than write access. External communication and system-of-record updates usually require human approval.

Memory should be handled carefully. Some agents retain context across interactions or workflows. Persistent memory may improve continuity but can create privacy, accuracy, and records risks. Agencies should define what may be stored, for how long, who can access it, and how it can be corrected or deleted.

Human oversight should match risk. Human-in-the-loop approval is appropriate for actions that affect records, services, public communication, legal obligations, or partners. Human-on-the-loop monitoring may be appropriate for lower-risk background tasks if users can intervene quickly. Fully automated action should be limited to low-risk, well-tested, reversible tasks.

Agentic workflows require incident response planning. Agencies should define what counts as an agentic incident, such as unauthorized action, data exposure, repeated error, misrouting, unsafe recommendation, or failure to escalate. Logs should support reconstruction of prompts, retrieved sources, tool calls, actions, approvals, and outcomes.

Risks, Failure Modes, and Guardrails

Excessive agency. An agent may take actions beyond intended scope. Guardrails include narrow goals, allowed-action lists, prohibited actions, and approval gates.

Unauthorized system access. Tool access may expose sensitive data or functions. Guardrails include least privilege, scoped credentials, role-based access, and audit logs.

Error propagation. One wrong step may trigger downstream errors. Guardrails include checkpoints, validation, rollback, and staged deployment.

Unclear accountability. Staff may not know who is responsible for agent actions. Guardrails include role definitions, review records, and governance ownership.

Application to AI-Supported Workflows

Agentic workflows may combine RAG, APIs, prompts, classification models, workflow rules, and messaging systems. They are most appropriate when tasks are repetitive, bounded, reviewable, and supported by reliable data. They are least appropriate when the task requires complex judgment, legal interpretation, community sensitivity, or irreversible action without review.

Public health agencies should pilot agentic workflows in limited, low-risk settings before expanding. Monitoring should include tool calls, failed actions, overrides, user feedback, time saved, error types, and unintended consequences.

Reflection Questions

What goal would the agent pursue, and how narrowly can it be defined?

Which tools or systems would the agent access?

Which actions are allowed, prohibited, or require human approval?

What logs are needed to reconstruct agent behavior?

What rollback or correction process is needed if the agent makes an error?

Practical Exercise

Create an agentic workflow control plan. Include the goal, trigger, data sources, tools, allowed actions, prohibited actions, approval points, memory rules, logs, monitoring indicators, rollback process, incident triggers, and responsible owner.

The plan should identify at least one low-risk pilot boundary and one condition that would require pausing the workflow.

Expected Artifact or Evidence

Agentic workflow control plan

Allowed/prohibited action list

Human approval and escalation map

Tool-call logging requirements

References and Resources for Additional Information

OWASP Top 10 for Large Language Model Applications:

<https://owasp.org/www-project-top-10-for-large-language-model-applications/>

NIST AI RMF Generative AI Profile:

<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>

NIST Artificial Intelligence Risk Management Framework: <https://www.nist.gov/itl/ai-risk-management-framework>

NIST Privacy Framework: <https://www.nist.gov/privacy-framework>

Knowledge Check

1. What makes agentic AI different from simple drafting support?

- A. It may take multiple steps and use tools or actions toward a goal.
- B. It never uses data.

- C. It is always safer than other AI.
- D. It cannot be monitored.

Correct answer: A

2. Which action usually requires stronger oversight?

- A. Searching an approved public FAQ.
- B. Updating a system of record or sending an external message.
- C. Formatting a heading.
- D. Counting words in a training document.

Correct answer: B

3. What is orchestration?

- A. Deleting old files.
- B. A map projection method.
- C. Coordinating steps, tools, systems, and decisions in an AI workflow.
- D. A laboratory terminology standard.

Correct answer: C

4. What should happen when an agent encounters an out-of-scope request?

- A. Proceed anyway if the answer seems likely.
- B. Use any available internet source.
- C. Hide uncertainty from the user.
- D. Stop, decline, or escalate according to the approved workflow.

Correct answer: D

5. What is strong completion evidence for this module?

- A. An agentic workflow control plan with tools, permissions, approvals, logs, rollback, and incident triggers.
- B. A list of possible agent names.
- C. A generic definition of automation.
- D. A vendor screenshot without workflow analysis.

Correct answer: A

References and Resources

OWASP Top 10 for Large Language Model Applications

<https://owasp.org/www-project-top-10-for-large-language-model-applications/>

NIST AI RMF Generative AI Profile

<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>

NIST Artificial Intelligence Risk Management Framework

<https://www.nist.gov/itl/ai-risk-management-framework>

NIST Privacy Framework

<https://www.nist.gov/privacy-framework>

ARC 330

Cloud, Containers, and Secure Deployment Basics

Course ID	ARC 330
Track	Technical Architecture Track
Level	intermediate/practitioner
Estimated time	60
Audience	Informatics staff, IT staff, Data architects, Product owners, Technical implementation teams

Learning Objectives

- Explain the purpose of cloud, containers, and secure deployment basics in public health AI work.
- Identify the technical, operational, privacy, security, equity, and governance issues associated with the topic.
- Apply the topic to a real or realistic public health workflow, system, or implementation scenario.
- Recognize common failure modes that can affect safety, accuracy, trust, workflow fit, or sustainability.
- Identify practical guardrails that should be documented before implementation.
- Produce a secure deployment readiness worksheet that can support readiness assessment, governance review, procurement, implementation, or monitoring.

Jurisdiction and Agency Policy Note

This national-level module is intended for training and planning purposes. It does not replace legal, privacy, procurement, civil rights, records, or agency policy review. Learners should confirm applicable requirements in their own jurisdiction and organization before implementing AI-supported workflows.

Why This Topic Matters for Public Health AI

Cloud, Containers, and Secure Deployment Basics matters because AI systems are embedded in public health programs, data systems, and decision processes rather than operating in isolation. The topic affects whether AI outputs are accurate, explainable, safe, equitable, secure, and sustainable. Agencies that ignore this area may create tools that work in a demonstration but fail in actual practice.

The topic also affects accountability. Public health agencies must be able to explain how data are used, why a tool was approved, what evidence supports the workflow, who reviews outputs, and what happens when the system fails. These responsibilities become more important as AI tools move from experimentation into operational use.

Core Concepts

Cloud computing provides remote infrastructure, platforms, and services that can host data pipelines, analytics environments, AI tools, and public health applications. Cloud deployment can improve scalability and resilience, but only when identity, access, logging, encryption, and governance are configured correctly.

Containers package software and its dependencies so that it can run consistently across environments. Containers can improve reproducibility and deployment speed, but they must be scanned, patched, configured, and monitored.

Container use does not remove the need for secure architecture.

Identity and access management defines who or what can access systems, data, and actions. In AI workflows, IAM must include human users, service accounts, APIs, agents, and automation tools. Poorly designed access can expose sensitive data or allow unauthorized write-back to systems of record.

Public Health Example

A public health agency may deploy an internal RAG assistant in a cloud environment so staff can search approved guidance and protocols. The tool may be low risk if it uses public documents, but it becomes higher risk if it retrieves sensitive policies, case notes, or internal operational documents. Secure deployment requires access controls, document-level permissions, logging, monitoring, backup, and a process for removing outdated content.

Technical Deep Dive

Secure deployment begins with environment separation. Development, testing, staging, and production environments should be clearly separated so experimental tools do not affect operational systems. Sensitive production data should not be copied into lower environments unless approved controls are in place.

Logging and monitoring are central to deployment safety. Agencies should know when systems are accessed, when APIs fail, when models produce errors, when unusual activity occurs, and when infrastructure changes are made. Logs should be protected because they can contain sensitive information.

Secrets management is essential for AI integrations. API keys, database passwords, certificates, and tokens should not be embedded in code, prompts, shared documents, or configuration files that are not secured. A compromised secret can allow unauthorized access to data or actions across connected systems.

Risks, Failure Modes, and Guardrails

A common failure mode is treating the topic as a technical detail rather than a public health control. When staff do not connect technical decisions to authority, workflow, equity, privacy, security, and accountability, the agency may adopt a tool that appears useful but cannot be sustained or defended.

Another risk is relying on vendor claims without agency-defined evidence. Vendors may provide demonstrations, dashboards, or summary metrics that do not match the agency workflow or data environment. A guardrail is to require documentation, test results, acceptance criteria, and agency-controlled review.

A third risk is failing to plan for change. Public health data, policies, systems, and emergencies change. Any technical approach must include monitoring, review, versioning, escalation, and retirement pathways so the agency can respond when assumptions no longer hold.

Application to AI-Supported Workflows

Generative AI, predictive analytics, RAG, automation, and agentic workflows all depend on the controls described in this module. The controls should be scaled to the risk of the use case. A tool that drafts internal training text may require lighter review than a tool that updates case records, prioritizes outreach, or supports public communication during an emergency.

The module should be applied during readiness assessment, procurement, implementation planning, pilot testing, and post-deployment monitoring. Learners should document how the topic affects data, systems, users, safeguards, and decision-making for the selected workflow.

Reflection Questions

Where does this topic appear in one current or proposed AI workflow in your agency?

What decisions, data sources, users, or partners would be affected if this area is poorly designed?

What evidence would governance reviewers need before approving the workflow?

Which safeguards should be required before pilot testing?

Who owns ongoing monitoring and review?

Practical Exercise

Create a secure deployment readiness worksheet for one real or realistic public health AI workflow. The artifact should describe the use case, affected users, data sources, system dependencies, risks, guardrails, review points, and unresolved questions.

The exercise should produce a work product that could be used in a governance meeting, procurement review, implementation planning session, pilot evaluation, or monitoring review. Learners should document assumptions and identify which staff or partners need to be consulted.

Expected Artifact or Evidence

The expected artifact is a completed secure deployment readiness worksheet. It should be specific to a public health workflow and should include technical dependencies, governance questions, human review expectations, monitoring indicators, and unresolved risks.

References and Resources for Additional Information

NIST AI Risk Management Framework

NIST Cybersecurity Framework and cloud security guidance

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP Top 10 for LLM Applications, when security or AI integrations are relevant

Agency privacy, cybersecurity, data governance, procurement, and records management policies

Knowledge Check

1. What is the best reason to address cloud, containers, and secure deployment basics before deployment?

- A. It makes governance unnecessary.
- B. It helps ensure the AI workflow is safe, useful, reviewable, and fit for public health operations.
- C. It allows the agency to skip testing.
- D. It guarantees that the vendor product will work.

Correct answer: B

2. Which practice best supports responsible implementation?

- A. Defining intended use, evidence requirements, safeguards, monitoring, and escalation before launch.
- B. Waiting until after launch to decide whether the tool is risky.
- C. Allowing users to decide individually how to use the tool.
- D. Treating vendor documentation as the only required evidence.

Correct answer: A

3. Which statement best describes a common failure mode?

- A. Human review is always unnecessary when AI is accurate.
- B. Documentation makes AI less useful.
- C. The agency assumes the tool is ready because the technology works in a demonstration, even though workflow, data, and governance have not been tested.
- D. Monitoring should stop after the pilot.

Correct answer: C

4. What should be included in the practical artifact for this module?

- A. Only a general definition.
- B. A marketing description of the AI tool.
- C. A list of unrelated data systems.
- D. Workflow-specific risks, safeguards, roles, evidence requirements, and unresolved questions.

Correct answer: D

5. When should the issue be escalated for governance or leadership review?

- A. Only if the tool is public-facing.
- B. When risks affect privacy, safety, equity, legal authority, public trust, system integrity, or operational decisions.
- C. Never, because technical staff can resolve all issues alone.
- D. Only after an incident has already occurred.

Correct answer: B

References and Resources

NIST AI Risk Management Framework

NIST Cybersecurity Framework and cloud security guidance

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP Top 10 for LLM Applications, when security or AI integrations are relevant

Agency privacy, cybersecurity, data governance, procurement, and records management policies

ARC 340

Public Health System Integration Patterns

Course ID	ARC 340
Track	Technical Architecture Track
Level	intermediate/practitioner
Estimated time	60
Audience	Informatics staff, IT staff, Data architects, Product owners, Technical implementation teams

Learning Objectives

- Explain why system integration is central to AI readiness in public health.
- Distinguish among common integration patterns such as point-to-point interfaces, APIs, message brokers, interface engines, event queues, and bulk exchange.
- Identify the systems of record, upstream sources, downstream systems, and shared services involved in an AI-supported workflow.
- Assess whether a proposed AI tool requires read-only access, write-back access, embedded workflow integration, or manual review.
- Identify integration risks related to reliability, duplicate entry, routing errors, privacy, security, and vendor lock-in.
- Produce an integration pattern worksheet for a public health AI use case.

Definitions

Integration pattern	A reusable approach for connecting systems, services, data flows, and workflows.
Interface engine	Software that routes, transforms, validates, and monitors messages between systems.
Message broker	A service that allows systems to send and receive messages through a managed queue or topic.
Bulk data exchange	The transfer of large datasets at scheduled intervals or on demand.
Write-back	A system action that updates, creates, or changes information in another system.

Jurisdiction and Agency Policy Note

This national-level module is intended for training and planning purposes. It does not replace legal, privacy, procurement, civil rights, records, or agency policy review. Learners should confirm applicable requirements in their own jurisdiction and organization before implementing AI-supported workflows.

Why This Topic Matters for Public Health AI

Integration decisions determine whether AI becomes part of a reliable public health workflow or an isolated tool that adds burden. If staff must copy and paste sensitive information between systems, the AI tool may appear helpful while increasing privacy risk, documentation gaps, and error rates. If an AI service writes back to a system of record without appropriate controls, a technical integration can become an operational safety risk.

Public health agencies often operate in mixed environments that include legacy surveillance systems, modern cloud platforms, external partners, interface engines, spreadsheets, portals, and vendor-hosted tools. AI implementation requires teams to make these environments visible and decide where automation belongs. A well-integrated workflow preserves accountability by defining what data are exchanged, what actions are allowed, what is logged, and what humans must approve.

Public Health Example

A health department may want an AI assistant to summarize incoming case reports and create follow-up tasks for investigators. A weak integration pattern would require staff to export case data, paste it into an external tool, copy the summary, and manually create a task. A stronger pattern would retrieve approved fields from the surveillance system, generate a draft summary inside the authorized workflow, and require the investigator to approve any task creation.

The technical design matters because the AI tool is no longer just generating text. It is interacting with operational systems. The agency must decide whether the tool is read-only, whether it can create draft tasks, whether it can update the case record, and how the system logs each action.

Technical and Operational Deep Dive

Point-to-point interfaces connect two systems directly. They can be useful when the exchange is stable and narrow, but they become difficult to maintain as the number of systems increases. In public health, too many point-to-point connections can create brittle architecture, inconsistent validation, and high maintenance burden.

Interface engines and message brokers support more scalable exchange. Interface engines are common in health data exchange because they can transform messages, route them, validate required fields, and provide monitoring. Message brokers and event queues are useful when workflows are triggered by events, such as receipt of a lab result or creation of a new case. AI tools that depend on timely data may need these patterns to avoid polling, manual refreshes, or delayed processing.

APIs are increasingly important because they allow systems to request data or services in a more modular way. FHIR APIs can support structured access to clinical information, while internal APIs may expose surveillance, laboratory, or operational data to approved applications. API-based integration requires careful attention to authentication, authorization, rate limits, payload design, error handling, and audit logs.

Risks, Failure Modes, and Guardrails

A common failure mode is integration by workaround. Staff may use manual exports, screenshots, shared drives, or copy-and-paste workflows when a tool is not properly integrated. These practices can create privacy risk, version confusion, and gaps in the official record. A guardrail is to require integration review before operational use and to document whether manual transfer of sensitive data is prohibited.

Another failure mode is uncontrolled write-back. If an AI tool can update a record, create a task, send a message, or change a status, the agency must define approval points, rollback procedures, and audit logs. The guardrail is to separate draft support from final action unless the workflow has been approved for automation.

Application to AI-Supported Workflows

Generative AI may require integration with document repositories or systems of record so outputs can be grounded in approved content. Predictive models may require scheduled data feeds and monitoring dashboards. Agentic workflows require the strongest integration controls because they may call tools, access data, create tasks, or trigger notifications. For each AI use case, the integration pattern should match the sensitivity and operational consequence of the action.

Reflection Questions

Which workflows currently depend on manual transfer between systems?

Which systems are the official systems of record?

Where would API access reduce burden without increasing risk?

Which actions should remain read-only?

What logs would be needed to reconstruct an AI-supported action?

Practical Exercise

Create an integration pattern worksheet for one AI-supported workflow. Identify the systems involved, the data exchanged, the integration method, whether the AI tool is read-only or write-back, the human approval point, and the audit log requirements.

Expected Artifact or Evidence

Integration pattern worksheet

System-of-record and source-system list

Read/write permission recommendation

Error handling and audit log requirements

References and Resources for Additional Information

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

Knowledge Check

1. Which integration risk is most concerning for an AI tool that can update a case record?

- A. A colorful user interface
- B. Uncontrolled write-back without approval, logging, or rollback
- C. A long training time
- D. Use of a standard glossary

Correct answer: B

2. What is the best reason to assign an accountability owner?

- A. To reduce the need for documentation
- B. To ensure someone is responsible for monitoring, escalation, and correction
- C. To make the vendor solely responsible
- D. To prevent staff from raising concerns

Correct answer: B

3. When should an AI-supported workflow be escalated for review?

- A. Only after media coverage
- B. Never if the tool is efficient
- C. When outputs may affect safety, equity, privacy, legal authority, public trust, or official action
- D. Only at the end of the fiscal year

Correct answer: C

4. Which practice best supports public accountability?

- A. Keeping AI use informal
- B. Avoiding written records
- C. Letting each user decide independently
- D. Maintaining evidence artifacts, decision records, review notes, and monitoring results

Correct answer: D

5. What should happen when data sources, policy, vendor configuration, or workflow use changes?

- A. The use case should be reviewed to determine whether safeguards, validation, documentation, or approvals need updating
- B. The agency should ignore the change if staff like the tool
- C. The model should be trusted automatically
- D. The glossary should be deleted

Correct answer: A

References and Resources

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

COM 200

AI for Public Health Communications and Message Development

Course ID	COM 200
Track	Communications Role-Based Track
Level	applied/foundational
Estimated time	30
Audience	Communications staff, PIOs, Language access staff, Community engagement staff

Learning Objectives

- Identify appropriate and inappropriate uses of AI in public health communications.
- Apply source verification, plain language, accessibility, and equity review to AI-generated messages.
- Design a human review workflow for AI-assisted public-facing communications.
- Recognize misinformation, overstatement, missing uncertainty, and stigmatizing language risks.
- Produce an AI-assisted communications review checklist.

Jurisdiction and Agency Policy Note

This national-level module is intended for training and planning purposes. It does not replace legal, privacy, procurement, civil rights, records, or agency policy review. Learners should confirm applicable requirements in their own jurisdiction and organization before implementing AI-supported workflows.

Public Health Example

A health department uses AI to draft a heat safety message for older adults, outdoor workers, and people without reliable air conditioning. Communications staff verify the recommendations against official guidance, rewrite the message in plain language, review it for accessibility and language access, avoid blaming individuals for structural risks, and route the final version through the approved public information process before release.

Technical and Operational Deep Dive

AI can be useful early in the message development process. It can generate first drafts, propose alternative tones, simplify technical language, identify jargon, create audience-specific versions, and suggest questions the public may ask. These uses can save time, especially when communications staff support multiple programs. However, AI should not be treated as the authority on public health content. The approved evidence base should come from agency guidance, subject matter experts, official data, and community context.

Source verification is essential because AI-generated language may include unsupported claims, outdated recommendations, fabricated details, or overconfident statements. Communications staff should verify factual claims, numbers, dates, risk statements, recommendations, and eligibility criteria against approved sources. When uncertainty exists, the message should explain what is known, what is not known, and what actions people can take.

Equity and accessibility review should be built into the workflow. AI-generated messages may not adequately address language access, disability access, literacy, cultural context, rural access, transportation barriers, technology access, or historical mistrust. Public health messages should be reviewed for plain language, readability, alternative formats, translation quality, non-stigmatizing language, and practical feasibility for the intended audience.

AI can also support review, but it should not replace reviewers. A communications workflow should define who checks technical accuracy, who reviews equity and accessibility, who approves public release, and who monitors public response after dissemination. For urgent communications, the workflow should include expedited review rules that still preserve accuracy and accountability.

Risks, Failure Modes, and Guardrails

Risks and failure modes include the following:

AI-generated messages overstating evidence or omitting uncertainty.

Use of stigmatizing, blaming, or culturally inappropriate language.

Unverified translations or readability changes that alter meaning.

Public-facing content being released without required subject matter or leadership review.

Messages that assume access to resources, technology, transportation, or healthcare that the audience may not have.

Recommended guardrails include the following:

Use only approved tools and approved source material for public health communications.

Require human review for accuracy, plain language, accessibility, equity, and tone.

Document source material used for public-facing claims.

Use qualified review for translation and language access needs.

Route public-facing materials through the agency approval process before release.

Reflection Questions

Where does this topic appear in your agency, program, or workflow?

Which staff roles would need to understand or apply this concept before AI use is approved?

What could go wrong if this topic is not addressed before implementation?

What evidence would governance, leadership, or operations staff need before moving forward?

Which policy, data, equity, privacy, security, or workflow questions remain unresolved?

Practical Exercise

Create an AI-assisted communications review checklist for one public health message. Include intended audience, source material, factual claims to verify, plain language review, accessibility review, language access review, equity review, approval pathway, and documentation expectations.

Expected Artifact or Evidence

AI-assisted communications review checklist

Source verification notes

Plain language and accessibility review notes

Approval pathway documentation

Final reviewed message or sample revision

References and Resources for Additional Information

NIST Artificial Intelligence Risk Management Framework

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP guidance for LLM application security

Agency policies for privacy, cybersecurity, records, procurement, accessibility, and public communications

Relevant public health program SOPs, data governance documentation, and implementation plans

Knowledge Check

1. What is an appropriate role for AI in public health communications?

- A. Replacing the public information officer
- B. Publishing public messages without review
- C. Inventing source data
- D. Supporting drafting, revision, simplification, and review under human oversight

Correct answer: D

2. Why is source verification necessary?

- A. AI is always more accurate than source documents
- B. AI-generated content may include unsupported, outdated, or overconfident claims
- C. Source documents should never be used
- D. Verification slows communication without benefit

Correct answer: B

3. Which review is especially important for public-facing AI-assisted messages?

- A. Review by the AI tool only
- B. Review only for grammar
- C. Accuracy, plain language, accessibility, equity, and approval review
- D. No review if the message is urgent

Correct answer: C

4. What should happen with AI-assisted translation for public health messages?

- A. It should receive appropriate language access review before public use
- B. It should always be published as generated
- C. It should replace qualified review
- D. It should be used only for internal jokes

Correct answer: A

5. What is strong completion evidence?

- A. A social media post created without sources
- B. A list of popular prompts
- C. A screenshot of AI output
- D. A communications review checklist with source verification, accessibility, equity, and approval notes

Correct answer: D

References and Resources

NIST Artificial Intelligence Risk Management Framework

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP guidance for LLM application security

Agency policies for privacy, cybersecurity, records, procurement, accessibility, and public communications

Relevant public health program SOPs, data governance documentation, and implementation plans

COM 210

AI-Assisted Translation, Language Access, and Readability Review

Course ID	COM 210
Track	Communications Role-Based Track
Level	applied/foundational
Estimated time	30
Audience	Communications staff, PIOs, Language access staff, Community engagement staff

Learning Objectives

- Explain the purpose of ai-assisted translation, language access, and readability review in public health AI.
- Identify the roles, decisions, risks, and documentation needed for the topic.
- Apply the topic to a real or realistic public health workflow or decision.
- Recognize equity, privacy, security, legal, operational, and public trust implications.
- Produce a practical artifact that supports readiness, governance, implementation, monitoring, or accountability.

Definitions

Decision right	The authority to approve, deny, modify, pause, or escalate an AI use case or operational decision.
Evidence artifact	A document, log, checklist, test result, approval record, or review note that demonstrates responsible AI practice.
Escalation pathway	A defined route for raising concerns to the appropriate reviewer or decision-maker.
Accountability owner	The person or role responsible for ensuring that an AI use case is governed, monitored, and corrected when needed.
Operational safeguard	A procedure, control, or review step that reduces risk during day-to-day use.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Why This Topic Matters for Public Health AI

AI-Assisted Translation, Language Access, and Readability Review matters because AI adoption changes how public health agencies make, document, communicate, and review decisions. Without a structured approach, staff may rely on informal practices, vendor assurances, or one-time approvals that do not match the operational consequences of AI-supported work.

Public health agencies operate under legal authorities, public accountability, resource constraints, and community trust obligations. A module on this topic helps agencies define repeatable expectations before urgent situations arise. It also helps staff distinguish low-risk administrative support from uses that affect surveillance, public communications, field response, resource allocation, civil rights, privacy, or official records.

Public Health Example

A public health agency is preparing to use AI in a workflow related to safe use of AI for multilingual, plain-language, and accessible communication. The project team initially treats the tool as a helpful efficiency aid, but governance reviewers recognize that the workflow may affect public trust, equity, official records, or public health action.

The agency responds by creating a clear artifact that defines intended use, prohibited use, responsible roles, review points, documentation expectations, escalation triggers, and monitoring measures. This converts a broad AI idea into a manageable public health practice.

Technical and Operational Deep Dive

The first step in this topic is to define the decision or workflow clearly. Staff should identify who is affected, what data are used, what AI output is produced, who reviews the output, and what action may follow. This prevents the module from becoming abstract and anchors the work in public health practice.

The second step is to assign responsibilities. AI-supported workflows often cross program, IT, legal, privacy, communications, procurement, and leadership boundaries. Responsibilities should be named by role rather than assumed. The artifact should specify who prepares materials, who reviews risks, who approves use, who monitors performance, and who can pause or revise the workflow.

The third step is to define evidence. Responsible AI depends on more than policy statements. Agencies need records that show what was reviewed, what decisions were made, what safeguards were applied, what incidents occurred, and how the workflow changed over time. Evidence also supports audits, public records review, quality improvement, and leadership accountability.

Risks, Failure Modes, and Guardrails

A common failure mode is treating the topic as a one-time checklist. AI risks can change when data sources, models, vendors, policies, users, or public expectations change. A guardrail is to define review intervals and change triggers.

Another failure mode is unclear ownership. When everyone is responsible in general, no one may be responsible in practice. A guardrail is to assign an accountability owner and name escalation pathways.

A third failure mode is relying on vendor claims or informal staff confidence without documentation. A guardrail is to require evidence artifacts that can be reviewed by governance, leadership, privacy, legal, equity, security, or program teams.

Application to AI-Supported Workflows

This topic applies across generative AI, predictive analytics, NLP, RAG, automation, and agentic workflows. The more sensitive the data, the more consequential the output, or the more automated the action, the stronger the documentation, review, and monitoring expectations should be.

Reflection Questions

What public health decision or workflow does this topic affect?

Who has authority to approve or pause the work?

What evidence would show that the workflow is responsible and effective?

What equity, privacy, legal, security, or public trust risks require review?

How will the agency know when the workflow needs to be revised?

Practical Exercise

Create a practical ai-assisted translation, language access, and readability review artifact for one AI-supported public health use case. The artifact should define the use case, affected roles, review criteria, safeguards, documentation expectations, escalation triggers, monitoring indicators, and next review date.

Expected Artifact or Evidence

AI-Assisted Translation, Language Access, and Readability Review artifact

Role and responsibility table

Risk and safeguard summary

Escalation and monitoring plan

Approval or review note

References and Resources for Additional Information

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

Knowledge Check

1. What is the strongest evidence that ai-assisted translation, language access, and readability review has been applied responsibly?

- A. A general statement that AI is useful
- B. A vendor demonstration with no agency review
- C. A documented artifact that defines use, roles, safeguards, escalation, monitoring, and evidence
- D. A decision to skip review because the use case seems efficient

Correct answer: C

2. What is the best reason to assign an accountability owner?

- A. To reduce the need for documentation
- B. To ensure someone is responsible for monitoring, escalation, and correction
- C. To make the vendor solely responsible
- D. To prevent staff from raising concerns

Correct answer: B

3. When should an AI-supported workflow be escalated for review?

- A. Only after media coverage
- B. Never if the tool is efficient
- C. When outputs may affect safety, equity, privacy, legal authority, public trust, or official action
- D. Only at the end of the fiscal year

Correct answer: C

4. Which practice best supports public accountability?

- A. Keeping AI use informal
- B. Avoiding written records
- C. Letting each user decide independently
- D. Maintaining evidence artifacts, decision records, review notes, and monitoring results

Correct answer: D

5. What should happen when data sources, policy, vendor configuration, or workflow use changes?

- A. The use case should be reviewed to determine whether safeguards, validation, documentation, or approvals need updating
- B. The agency should ignore the change if staff like the tool
- C. The model should be trusted automatically
- D. The glossary should be deleted

Correct answer: A

References and Resources

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

COM 220

Public Comment, Social Listening, and Community Feedback Analysis

Course ID	COM 220
Track	Communications Role-Based Track
Level	applied/foundational
Estimated time	30
Audience	Communications staff, PIOs, Language access staff, Community engagement staff

Learning Objectives

- Explain the purpose of public comment, social listening, and community feedback analysis in public health AI.
- Identify the roles, decisions, risks, and documentation needed for the topic.
- Apply the topic to a real or realistic public health workflow or decision.
- Recognize equity, privacy, security, legal, operational, and public trust implications.
- Produce a practical artifact that supports readiness, governance, implementation, monitoring, or accountability.

Definitions

Decision right	The authority to approve, deny, modify, pause, or escalate an AI use case or operational decision.
Evidence artifact	A document, log, checklist, test result, approval record, or review note that demonstrates responsible AI practice.
Escalation pathway	A defined route for raising concerns to the appropriate reviewer or decision-maker.
Accountability owner	The person or role responsible for ensuring that an AI use case is governed, monitored, and corrected when needed.
Operational safeguard	A procedure, control, or review step that reduces risk during day-to-day use.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Why This Topic Matters for Public Health AI

Public Comment, Social Listening, and Community Feedback Analysis matters because AI adoption changes how public health agencies make, document, communicate, and review decisions. Without a structured approach, staff may rely on informal practices, vendor assurances, or one-time approvals that do not match the operational consequences of AI-supported work.

Public health agencies operate under legal authorities, public accountability, resource constraints, and community trust obligations. A module on this topic helps agencies define repeatable expectations before urgent situations arise. It also helps staff distinguish low-risk administrative support from uses that affect surveillance, public communications, field response, resource allocation, civil rights, privacy, or official records.

Public Health Example

A public health agency is preparing to use AI in a workflow related to ethical analysis of public comments, social listening data, and community feedback. The project team initially treats the tool as a helpful efficiency aid, but governance reviewers recognize that the workflow may affect public trust, equity, official records, or public health action.

The agency responds by creating a clear artifact that defines intended use, prohibited use, responsible roles, review points, documentation expectations, escalation triggers, and monitoring measures. This converts a broad AI idea into a manageable public health practice.

Technical and Operational Deep Dive

The first step in this topic is to define the decision or workflow clearly. Staff should identify who is affected, what data are used, what AI output is produced, who reviews the output, and what action may follow. This prevents the module from becoming abstract and anchors the work in public health practice.

The second step is to assign responsibilities. AI-supported workflows often cross program, IT, legal, privacy, communications, procurement, and leadership boundaries. Responsibilities should be named by role rather than assumed. The artifact should specify who prepares materials, who reviews risks, who approves use, who monitors performance, and who can pause or revise the workflow.

The third step is to define evidence. Responsible AI depends on more than policy statements. Agencies need records that show what was reviewed, what decisions were made, what safeguards were applied, what incidents occurred, and how the workflow changed over time. Evidence also supports audits, public records review, quality improvement, and leadership accountability.

Risks, Failure Modes, and Guardrails

A common failure mode is treating the topic as a one-time checklist. AI risks can change when data sources, models, vendors, policies, users, or public expectations change. A guardrail is to define review intervals and change triggers.

Another failure mode is unclear ownership. When everyone is responsible in general, no one may be responsible in practice. A guardrail is to assign an accountability owner and name escalation pathways.

A third failure mode is relying on vendor claims or informal staff confidence without documentation. A guardrail is to require evidence artifacts that can be reviewed by governance, leadership, privacy, legal, equity, security, or program teams.

Application to AI-Supported Workflows

This topic applies across generative AI, predictive analytics, NLP, RAG, automation, and agentic workflows. The more sensitive the data, the more consequential the output, or the more automated the action, the stronger the documentation, review, and monitoring expectations should be.

Reflection Questions

What public health decision or workflow does this topic affect?

Who has authority to approve or pause the work?

What evidence would show that the workflow is responsible and effective?

What equity, privacy, legal, security, or public trust risks require review?

How will the agency know when the workflow needs to be revised?

Practical Exercise

Create a practical public comment, social listening, and community feedback analysis artifact for one AI-supported public health use case. The artifact should define the use case, affected roles, review criteria, safeguards, documentation expectations, escalation triggers, monitoring indicators, and next review date.

Expected Artifact or Evidence

Public Comment, Social Listening, and Community Feedback Analysis artifact

Role and responsibility table

Risk and safeguard summary

Escalation and monitoring plan

Approval or review note

References and Resources for Additional Information

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

Knowledge Check

1. What is the strongest evidence that public comment, social listening, and community feedback analysis has been applied responsibly?

- A. A general statement that AI is useful
- B. A vendor demonstration with no agency review
- C. A documented artifact that defines use, roles, safeguards, escalation, monitoring, and evidence
- D. A decision to skip review because the use case seems efficient

Correct answer: C

2. What is the best reason to assign an accountability owner?

- A. To reduce the need for documentation
- B. To ensure someone is responsible for monitoring, escalation, and correction
- C. To make the vendor solely responsible
- D. To prevent staff from raising concerns

Correct answer: B

3. When should an AI-supported workflow be escalated for review?

- A. Only after media coverage
- B. Never if the tool is efficient
- C. When outputs may affect safety, equity, privacy, legal authority, public trust, or official action
- D. Only at the end of the fiscal year

Correct answer: C

4. Which practice best supports public accountability?

- A. Keeping AI use informal
- B. Avoiding written records
- C. Letting each user decide independently
- D. Maintaining evidence artifacts, decision records, review notes, and monitoring results

Correct answer: D

5. What should happen when data sources, policy, vendor configuration, or workflow use changes?

- A. The use case should be reviewed to determine whether safeguards, validation, documentation, or approvals need updating
- B. The agency should ignore the change if staff like the tool
- C. The model should be trusted automatically
- D. The glossary should be deleted

Correct answer: A

References and Resources

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

EPI 300

AI for Case Investigation and Outbreak Response

Course ID	EPI 300
Track	Epidemiology Role-Based Track
Level	intermediate/practitioner
Estimated time	50
Audience	Epidemiologists, Surveillance staff, Outbreak response staff, Program leads

Learning Objectives

- Identify case investigation and outbreak response tasks that AI may support.
- Distinguish between AI-supported triage, summarization, extraction, signal detection, and decision-making.
- Assess data quality, timeliness, source traceability, and equity risks in AI-supported response workflows.
- Define human review, escalation, and documentation expectations for AI outputs.
- Create an AI-supported case investigation or outbreak response workflow review.

Jurisdiction and Agency Policy Note

This national-level module is intended for training and planning purposes. It does not replace legal, privacy, procurement, civil rights, records, or agency policy review. Learners should confirm applicable requirements in their own jurisdiction and organization before implementing AI-supported workflows.

Public Health Example

During a foodborne illness investigation, AI is used to summarize interview notes and extract possible exposures. Investigators compare the extracted exposures to original notes, review uncertain or conflicting information, and document final interpretations. The tool helps reduce manual review time, but it does not determine the outbreak source or issue public recommendations without epidemiologist review.

Technical and Operational Deep Dive

AI can support case investigation by reducing the burden of reviewing large volumes of information. Natural language processing and generative AI may help summarize case notes, identify missing fields, extract dates and exposures, or draft follow-up questions. These tasks are appropriate only when outputs are traceable to source material and investigators can correct errors before the information is used.

Outbreak response may involve signal detection, anomaly review, geospatial patterns, laboratory results, syndromic surveillance, community reports, and partner communications. AI can help identify patterns or organize information, but it can also generate false signals or miss subtle context. A signal should trigger epidemiologic review, not automatic public health action.

Data quality is especially important during response. Early outbreak data are often incomplete, changing, biased toward reported cases, and uneven across communities. AI tools may overstate certainty if they do not represent missingness and uncertainty clearly. Response teams should document what data are preliminary, what assumptions are being

made, and what additional investigation is needed.

Equity considerations should be explicit. Some communities may be less likely to seek care, be tested, respond to interviews, speak English, have stable contact information, or appear in digital datasets. AI-supported prioritization must not systematically deprioritize people or communities because the data available about them are incomplete.

Risks, Failure Modes, and Guardrails

Risks and failure modes include the following:

AI summaries omitting uncertainty or unsupported details.

False outbreak signals creating unnecessary response or public concern.

Missed signals because some communities are underrepresented in data.

Automated prioritization replacing epidemiologic judgment.

Sensitive investigation information being used in unapproved tools.

Recommended guardrails include the following:

Require source traceability for summaries and extracted information.

Use AI outputs as prompts for review, not automatic conclusions.

Document uncertainty, missing data, and preliminary status.

Apply equity review to prioritization and signal detection workflows.

Define escalation to epidemiology leadership for unusual, high-risk, or public-facing outputs.

Reflection Questions

Where does this topic appear in your agency, program, or workflow?

Which staff roles would need to understand or apply this concept before AI use is approved?

What could go wrong if this topic is not addressed before implementation?

What evidence would governance, leadership, or operations staff need before moving forward?

Which policy, data, equity, privacy, security, or workflow questions remain unresolved?

Practical Exercise

Create an AI-supported case investigation or outbreak response workflow review. Include the supported task, data sources, AI output, source traceability, human review point, uncertainty documentation, equity risks, escalation triggers, and monitoring indicators.

Expected Artifact or Evidence

Case investigation or outbreak response workflow review

Source traceability checklist

Uncertainty and missing data documentation

Equity and escalation review

Monitoring indicators

References and Resources for Additional Information

NIST Artificial Intelligence Risk Management Framework

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP guidance for LLM application security

Agency policies for privacy, cybersecurity, records, procurement, accessibility, and public communications

Relevant public health program SOPs, data governance documentation, and implementation plans

Knowledge Check

1. What should AI outputs do in case investigation?

- A. Replace public health authority
- B. Support review, summarization, extraction, or triage while preserving human judgment
- C. Automatically declare outbreaks
- D. Eliminate documentation

Correct answer: B

2. What is an outbreak signal?

- A. A final legal conclusion
- B. A vendor report only
- C. A public press release
- D. A pattern or alert that may indicate an unusual increase, cluster, or concern requiring review

Correct answer: D

3. Why is uncertainty documentation important during response?

- A. Because early data are often incomplete, changing, and uneven across populations
- B. Because uncertainty means the response should stop

- C. Because AI tools cannot be reviewed
- D. Because documentation is unrelated to public health action

Correct answer: A

4. Which guardrail is appropriate for AI-extracted exposures?

- A. Investigators should verify extracted exposures against source notes before use
- B. The extraction should be accepted automatically
- C. The source notes should be deleted
- D. The tool should publish the exposure list publicly

Correct answer: A

5. What is strong completion evidence?

- A. A list of disease names
- B. A workflow review with data sources, AI output, source traceability, human review, uncertainty, equity risks, and escalation triggers
- C. A chatbot screenshot
- D. A press release draft only

Correct answer: B

References and Resources

NIST Artificial Intelligence Risk Management Framework

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP guidance for LLM application security

Agency policies for privacy, cybersecurity, records, procurement, accessibility, and public communications

Relevant public health program SOPs, data governance documentation, and implementation plans

EPI 310

AI for Surveillance Signal Review and Epidemiologic Judgment

Course ID	EPI 310
Track	Epidemiology Role-Based Track
Level	intermediate/practitioner
Estimated time	50
Audience	Epidemiologists, Surveillance staff, Outbreak response staff, Program leads

Learning Objectives

- Explain the purpose of ai for surveillance signal review and epidemiologic judgment in public health AI.
- Identify the roles, decisions, risks, and documentation needed for the topic.
- Apply the topic to a real or realistic public health workflow or decision.
- Recognize equity, privacy, security, legal, operational, and public trust implications.
- Produce a practical artifact that supports readiness, governance, implementation, monitoring, or accountability.

Definitions

Decision right	The authority to approve, deny, modify, pause, or escalate an AI use case or operational decision.
Evidence artifact	A document, log, checklist, test result, approval record, or review note that demonstrates responsible AI practice.
Escalation pathway	A defined route for raising concerns to the appropriate reviewer or decision-maker.
Accountability owner	The person or role responsible for ensuring that an AI use case is governed, monitored, and corrected when needed.
Operational safeguard	A procedure, control, or review step that reduces risk during day-to-day use.

Why This Topic Matters for Public Health AI

AI for Surveillance Signal Review and Epidemiologic Judgment matters because AI adoption changes how public health agencies make, document, communicate, and review decisions. Without a structured approach, staff may rely on informal practices, vendor assurances, or one-time approvals that do not match the operational consequences of AI-supported work.

Public health agencies operate under legal authorities, public accountability, resource constraints, and community trust obligations. A module on this topic helps agencies define repeatable expectations before urgent situations arise. It also helps staff distinguish low-risk administrative support from uses that affect surveillance, public communications, field response, resource allocation, civil rights, privacy, or official records.

Public Health Example

A public health agency is preparing to use AI in a workflow related to using AI to support signal detection while preserving epidemiologic judgment. The project team initially treats the tool as a helpful efficiency aid, but governance reviewers recognize that the workflow may affect public trust, equity, official records, or public health action.

The agency responds by creating a clear artifact that defines intended use, prohibited use, responsible roles, review points, documentation expectations, escalation triggers, and monitoring measures. This converts a broad AI idea into a manageable public health practice.

Technical and Operational Deep Dive

The first step in this topic is to define the decision or workflow clearly. Staff should identify who is affected, what data are used, what AI output is produced, who reviews the output, and what action may follow. This prevents the module from becoming abstract and anchors the work in public health practice.

The second step is to assign responsibilities. AI-supported workflows often cross program, IT, legal, privacy, communications, procurement, and leadership boundaries. Responsibilities should be named by role rather than assumed. The artifact should specify who prepares materials, who reviews risks, who approves use, who monitors performance, and who can pause or revise the workflow.

The third step is to define evidence. Responsible AI depends on more than policy statements. Agencies need records that show what was reviewed, what decisions were made, what safeguards were applied, what incidents occurred, and how the workflow changed over time. Evidence also supports audits, public records review, quality improvement, and leadership accountability.

Risks, Failure Modes, and Guardrails

A common failure mode is treating the topic as a one-time checklist. AI risks can change when data sources, models, vendors, policies, users, or public expectations change. A guardrail is to define review intervals and change triggers.

Another failure mode is unclear ownership. When everyone is responsible in general, no one may be responsible in practice. A guardrail is to assign an accountability owner and name escalation pathways.

A third failure mode is relying on vendor claims or informal staff confidence without documentation. A guardrail is to require evidence artifacts that can be reviewed by governance, leadership, privacy, legal, equity, security, or program teams.

Application to AI-Supported Workflows

This topic applies across generative AI, predictive analytics, NLP, RAG, automation, and agentic workflows. The more sensitive the data, the more consequential the output, or the more automated the action, the stronger the documentation, review, and monitoring expectations should be.

Reflection Questions

What public health decision or workflow does this topic affect?

Who has authority to approve or pause the work?

What evidence would show that the workflow is responsible and effective?

What equity, privacy, legal, security, or public trust risks require review?

How will the agency know when the workflow needs to be revised?

Practical Exercise

Create a practical ai for surveillance signal review and epidemiologic judgment artifact for one AI-supported public health use case. The artifact should define the use case, affected roles, review criteria, safeguards, documentation expectations, escalation triggers, monitoring indicators, and next review date.

Expected Artifact or Evidence

AI for Surveillance Signal Review and Epidemiologic Judgment artifact

Role and responsibility table

Risk and safeguard summary

Escalation and monitoring plan

Approval or review note

References and Resources for Additional Information

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

Knowledge Check

1. What is the strongest evidence that ai for surveillance signal review and epidemiologic judgment has been applied responsibly?

- A. A general statement that AI is useful
- B. A vendor demonstration with no agency review
- C. A documented artifact that defines use, roles, safeguards, escalation, monitoring, and evidence

D. A decision to skip review because the use case seems efficient

Correct answer: C

2. What is the best reason to assign an accountability owner?

A. To reduce the need for documentation

B. To ensure someone is responsible for monitoring, escalation, and correction

C. To make the vendor solely responsible

D. To prevent staff from raising concerns

Correct answer: B

3. When should an AI-supported workflow be escalated for review?

A. Only after media coverage

B. Never if the tool is efficient

C. When outputs may affect safety, equity, privacy, legal authority, public trust, or official action

D. Only at the end of the fiscal year

Correct answer: C

4. Which practice best supports public accountability?

A. Keeping AI use informal

B. Avoiding written records

C. Letting each user decide independently

D. Maintaining evidence artifacts, decision records, review notes, and monitoring results

Correct answer: D

5. What should happen when data sources, policy, vendor configuration, or workflow use changes?

A. The use case should be reviewed to determine whether safeguards, validation, documentation, or approvals need updating

B. The agency should ignore the change if staff like the tool

C. The model should be trusted automatically

D. The glossary should be deleted

Correct answer: A

References and Resources

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

EPI 320

Ethical Use of AI in Resource Prioritization and Field Response

Course ID	EPI 320
Track	Epidemiology Role-Based Track
Level	intermediate/practitioner
Estimated time	50
Audience	Epidemiologists, Surveillance staff, Outbreak response staff, Program leads

Learning Objectives

- Explain the purpose of ethical use of ai in resource prioritization and field response in public health AI.
- Identify the roles, decisions, risks, and documentation needed for the topic.
- Apply the topic to a real or realistic public health workflow or decision.
- Recognize equity, privacy, security, legal, operational, and public trust implications.
- Produce a practical artifact that supports readiness, governance, implementation, monitoring, or accountability.

Definitions

Decision right	The authority to approve, deny, modify, pause, or escalate an AI use case or operational decision.
Evidence artifact	A document, log, checklist, test result, approval record, or review note that demonstrates responsible AI practice.
Escalation pathway	A defined route for raising concerns to the appropriate reviewer or decision-maker.
Accountability owner	The person or role responsible for ensuring that an AI use case is governed, monitored, and corrected when needed.
Operational safeguard	A procedure, control, or review step that reduces risk during day-to-day use.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Why This Topic Matters for Public Health AI

Ethical Use of AI in Resource Prioritization and Field Response matters because AI adoption changes how public health agencies make, document, communicate, and review decisions. Without a structured approach, staff may rely on informal practices, vendor assurances, or one-time approvals that do not match the operational consequences of AI-supported work.

Public health agencies operate under legal authorities, public accountability, resource constraints, and community trust obligations. A module on this topic helps agencies define repeatable expectations before urgent situations arise. It also helps staff distinguish low-risk administrative support from uses that affect surveillance, public communications, field response, resource allocation, civil rights, privacy, or official records.

Public Health Example

A public health agency is preparing to use AI in a workflow related to AI-supported prioritization of field response, outreach, and scarce resources. The project team initially treats the tool as a helpful efficiency aid, but governance reviewers recognize that the workflow may affect public trust, equity, official records, or public health action.

The agency responds by creating a clear artifact that defines intended use, prohibited use, responsible roles, review points, documentation expectations, escalation triggers, and monitoring measures. This converts a broad AI idea into a manageable public health practice.

Technical and Operational Deep Dive

The first step in this topic is to define the decision or workflow clearly. Staff should identify who is affected, what data are used, what AI output is produced, who reviews the output, and what action may follow. This prevents the module from becoming abstract and anchors the work in public health practice.

The second step is to assign responsibilities. AI-supported workflows often cross program, IT, legal, privacy, communications, procurement, and leadership boundaries. Responsibilities should be named by role rather than assumed. The artifact should specify who prepares materials, who reviews risks, who approves use, who monitors performance, and who can pause or revise the workflow.

The third step is to define evidence. Responsible AI depends on more than policy statements. Agencies need records that show what was reviewed, what decisions were made, what safeguards were applied, what incidents occurred, and how the workflow changed over time. Evidence also supports audits, public records review, quality improvement, and leadership accountability.

Risks, Failure Modes, and Guardrails

A common failure mode is treating the topic as a one-time checklist. AI risks can change when data sources, models, vendors, policies, users, or public expectations change. A guardrail is to define review intervals and change triggers.

Another failure mode is unclear ownership. When everyone is responsible in general, no one may be responsible in practice. A guardrail is to assign an accountability owner and name escalation pathways.

A third failure mode is relying on vendor claims or informal staff confidence without documentation. A guardrail is to require evidence artifacts that can be reviewed by governance, leadership, privacy, legal, equity, security, or program teams.

Application to AI-Supported Workflows

This topic applies across generative AI, predictive analytics, NLP, RAG, automation, and agentic workflows. The more sensitive the data, the more consequential the output, or the more automated the action, the stronger the documentation, review, and monitoring expectations should be.

Reflection Questions

What public health decision or workflow does this topic affect?

Who has authority to approve or pause the work?

What evidence would show that the workflow is responsible and effective?

What equity, privacy, legal, security, or public trust risks require review?

How will the agency know when the workflow needs to be revised?

Practical Exercise

Create a practical ethical use of ai in resource prioritization and field response artifact for one AI-supported public health use case. The artifact should define the use case, affected roles, review criteria, safeguards, documentation expectations, escalation triggers, monitoring indicators, and next review date.

Expected Artifact or Evidence

Ethical Use of AI in Resource Prioritization and Field Response artifact

Role and responsibility table

Risk and safeguard summary

Escalation and monitoring plan

Approval or review note

References and Resources for Additional Information

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

Knowledge Check

1. What is the strongest evidence that ethical use of ai in resource prioritization and field response has been applied responsibly?

- A. A general statement that AI is useful
- B. A vendor demonstration with no agency review
- C. A documented artifact that defines use, roles, safeguards, escalation, monitoring, and evidence
- D. A decision to skip review because the use case seems efficient

Correct answer: C

2. What is the best reason to assign an accountability owner?

- A. To reduce the need for documentation
- B. To ensure someone is responsible for monitoring, escalation, and correction
- C. To make the vendor solely responsible
- D. To prevent staff from raising concerns

Correct answer: B

3. When should an AI-supported workflow be escalated for review?

- A. Only after media coverage
- B. Never if the tool is efficient
- C. When outputs may affect safety, equity, privacy, legal authority, public trust, or official action
- D. Only at the end of the fiscal year

Correct answer: C

4. Which practice best supports public accountability?

- A. Keeping AI use informal
- B. Avoiding written records
- C. Letting each user decide independently
- D. Maintaining evidence artifacts, decision records, review notes, and monitoring results

Correct answer: D

5. What should happen when data sources, policy, vendor configuration, or workflow use changes?

- A. The use case should be reviewed to determine whether safeguards, validation, documentation, or approvals need updating
- B. The agency should ignore the change if staff like the tool
- C. The model should be trusted automatically
- D. The glossary should be deleted

Correct answer: A

References and Resources

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

EXE 100

Executive AI Briefing for Public Health Leaders

Course ID	EXE 100
Track	Public Health Executive Leadership Track
Level	introductory
Estimated time	25
Audience	Executive leaders, Health officials, Senior managers, Governance chairs

Learning Objectives

- Explain the purpose of executive ai briefing for public health leaders in public health AI.
- Identify the roles, decisions, risks, and documentation needed for the topic.
- Apply the topic to a real or realistic public health workflow or decision.
- Recognize equity, privacy, security, legal, operational, and public trust implications.
- Produce a practical artifact that supports readiness, governance, implementation, monitoring, or accountability.

Definitions

Decision right	The authority to approve, deny, modify, pause, or escalate an AI use case or operational decision.
Evidence artifact	A document, log, checklist, test result, approval record, or review note that demonstrates responsible AI practice.
Escalation pathway	A defined route for raising concerns to the appropriate reviewer or decision-maker.
Accountability owner	The person or role responsible for ensuring that an AI use case is governed, monitored, and corrected when needed.
Operational safeguard	A procedure, control, or review step that reduces risk during day-to-day use.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Why This Topic Matters for Public Health AI

Executive AI Briefing for Public Health Leaders matters because AI adoption changes how public health agencies make, document, communicate, and review decisions. Without a structured approach, staff may rely on informal practices,

vendor assurances, or one-time approvals that do not match the operational consequences of AI-supported work.

Public health agencies operate under legal authorities, public accountability, resource constraints, and community trust obligations. A module on this topic helps agencies define repeatable expectations before urgent situations arise. It also helps staff distinguish low-risk administrative support from uses that affect surveillance, public communications, field response, resource allocation, civil rights, privacy, or official records.

Public Health Example

A public health agency is preparing to use AI in a workflow related to decision-focused AI briefing for health officers and senior leaders. The project team initially treats the tool as a helpful efficiency aid, but governance reviewers recognize that the workflow may affect public trust, equity, official records, or public health action.

The agency responds by creating a clear artifact that defines intended use, prohibited use, responsible roles, review points, documentation expectations, escalation triggers, and monitoring measures. This converts a broad AI idea into a manageable public health practice.

Technical and Operational Deep Dive

The first step in this topic is to define the decision or workflow clearly. Staff should identify who is affected, what data are used, what AI output is produced, who reviews the output, and what action may follow. This prevents the module from becoming abstract and anchors the work in public health practice.

The second step is to assign responsibilities. AI-supported workflows often cross program, IT, legal, privacy, communications, procurement, and leadership boundaries. Responsibilities should be named by role rather than assumed. The artifact should specify who prepares materials, who reviews risks, who approves use, who monitors performance, and who can pause or revise the workflow.

The third step is to define evidence. Responsible AI depends on more than policy statements. Agencies need records that show what was reviewed, what decisions were made, what safeguards were applied, what incidents occurred, and how the workflow changed over time. Evidence also supports audits, public records review, quality improvement, and leadership accountability.

Risks, Failure Modes, and Guardrails

A common failure mode is treating the topic as a one-time checklist. AI risks can change when data sources, models, vendors, policies, users, or public expectations change. A guardrail is to define review intervals and change triggers.

Another failure mode is unclear ownership. When everyone is responsible in general, no one may be responsible in practice. A guardrail is to assign an accountability owner and name escalation pathways.

A third failure mode is relying on vendor claims or informal staff confidence without documentation. A guardrail is to require evidence artifacts that can be reviewed by governance, leadership, privacy, legal, equity, security, or program teams.

Application to AI-Supported Workflows

This topic applies across generative AI, predictive analytics, NLP, RAG, automation, and agentic workflows. The more sensitive the data, the more consequential the output, or the more automated the action, the stronger the

documentation, review, and monitoring expectations should be.

Reflection Questions

What public health decision or workflow does this topic affect?

Who has authority to approve or pause the work?

What evidence would show that the workflow is responsible and effective?

What equity, privacy, legal, security, or public trust risks require review?

How will the agency know when the workflow needs to be revised?

Practical Exercise

Create a practical executive ai briefing for public health leaders artifact for one AI-supported public health use case. The artifact should define the use case, affected roles, review criteria, safeguards, documentation expectations, escalation triggers, monitoring indicators, and next review date.

Expected Artifact or Evidence

Executive AI Briefing for Public Health Leaders artifact

Role and responsibility table

Risk and safeguard summary

Escalation and monitoring plan

Approval or review note

References and Resources for Additional Information

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

Knowledge Check

1. What is the strongest evidence that executive ai briefing for public health leaders has been applied responsibly?

- A. A general statement that AI is useful
- B. A vendor demonstration with no agency review
- C. A documented artifact that defines use, roles, safeguards, escalation, monitoring, and evidence
- D. A decision to skip review because the use case seems efficient

Correct answer: C

2. What is the best reason to assign an accountability owner?

- A. To reduce the need for documentation
- B. To ensure someone is responsible for monitoring, escalation, and correction
- C. To make the vendor solely responsible
- D. To prevent staff from raising concerns

Correct answer: B

3. When should an AI-supported workflow be escalated for review?

- A. Only after media coverage
- B. Never if the tool is efficient
- C. When outputs may affect safety, equity, privacy, legal authority, public trust, or official action
- D. Only at the end of the fiscal year

Correct answer: C

4. Which practice best supports public accountability?

- A. Keeping AI use informal
- B. Avoiding written records
- C. Letting each user decide independently
- D. Maintaining evidence artifacts, decision records, review notes, and monitoring results

Correct answer: D

5. What should happen when data sources, policy, vendor configuration, or workflow use changes?

- A. The use case should be reviewed to determine whether safeguards, validation, documentation, or approvals need updating
- B. The agency should ignore the change if staff like the tool
- C. The model should be trusted automatically
- D. The glossary should be deleted

Correct answer: A

References and Resources

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

EXE 400

Executive AI Decision Rights and Accountability

Course ID	EXE 400
Track	Public Health Executive Leadership Track
Level	advanced/leadership
Estimated time	35
Audience	Executive leaders, Health officials, Senior managers, Governance chairs

Learning Objectives

- Explain why AI accountability must be assigned to people and roles rather than to tools or vendors.
- Identify executive, program, technical, privacy, legal, security, and governance responsibilities for AI decisions.
- Create a decision-rights matrix for AI use case approval, risk acceptance, monitoring, escalation, and retirement.
- Distinguish between consultative input, approval authority, and operational ownership.
- Define accountability expectations for vendor-supported AI tools and internally developed tools.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Why This Topic Matters for Public Health AI

AI systems can distribute responsibility across many actors. A vendor may provide the model, IT may manage access, a program may define the workflow, an epidemiologist may review the output, and leadership may approve implementation. Without a formal accountability model, each actor may assume that another party is responsible for validation, monitoring, public explanation, or incident response.

Clear decision rights support trust and speed. When leaders know which decisions they own, they can act quickly when a tool performs poorly, when a privacy concern emerges, or when a public-facing communication needs review. Decision rights also prevent informal escalation patterns in which difficult decisions are delayed because no one is clearly authorized to accept, modify, or reject a proposed use.

Public Health Example

A local health department implements an AI-assisted tool to classify incoming environmental health complaints by topic and urgency. The vendor configured the model, IT manages access, environmental health staff review the queue, and the program director owns the workflow. A complaint is misclassified and a serious housing-related concern is not routed for timely follow-up.

A decision-rights matrix helps the agency respond. The program owner is accountable for workflow performance, the vendor must provide technical support and documentation, IT reviews logging and access, governance reviews whether use should continue, and an executive sponsor decides whether the risk requires pausing the tool. Because responsibilities are defined in advance, the agency can respond without confusion or blame shifting.

Technical and Governance Deep Dive

Decision rights should be defined across the full AI lifecycle. During ideation, leadership may decide whether a use case aligns with agency priorities. During review, governance may decide whether the use case is low, moderate, or high risk. During procurement, authorized officials decide whether vendor claims, contract language, and technical requirements are sufficient. During implementation, program and technical owners decide whether acceptance criteria have been met. During operations, defined owners decide whether monitoring results require adjustment, escalation, or retirement.

A public health accountability model should separate approval authority from operational ownership. An executive may approve a use case, but the program owner remains accountable for how the workflow functions. A privacy officer may advise on data safeguards, but the program cannot treat that consultation as a substitute for ongoing compliance. A vendor may provide model documentation, but the agency must decide whether the documentation is sufficient for its context and authorities.

Decision rights should also include stop authority. Agencies often define who can approve a tool but fail to define who can pause it. Stop authority is essential for AI because model performance, data flows, user behavior, and public conditions can change over time. A clear pause process helps staff report concerns early rather than waiting for a serious incident.

Risks, Failure Modes, and Guardrails

One risk is accountability diffusion. When everyone is involved but no one is accountable, important concerns may go unresolved. A guardrail is to assign a named program owner and governance owner for every AI use case.

Another risk is inappropriate risk acceptance by staff who lack authority. A frontline supervisor may feel pressure to continue a tool despite known concerns. A guardrail is to define which risks can be accepted at the program level and which require executive or governance approval.

Application to AI-Supported Workflows

For generative AI, decision rights should clarify who approves public release of AI-assisted content. For predictive analytics, leaders should decide whether the score may inform resource allocation or only support review. For agentic AI, decision rights should define who authorizes system actions, rollback procedures, and changes to permissions. Accountability should remain tied to public health roles and legal authority.

Reflection Questions

Which current or proposed AI use case in your agency raises this leadership or governance issue most clearly?

Who currently has authority to make decisions about this issue, and is that authority documented?

What evidence would governance or leadership need before approving the use case?

Which staff, partners, or communities would be affected if this issue were handled poorly?

What policy, process, or artifact should be created or updated after this module?

Practical Exercise

Create a decision-rights matrix for one AI use case. Include approval, data access, validation, launch, monitoring, incident escalation, risk acceptance, public communication, and retirement.

Before You Begin

Choose a real or realistic public health AI use case. Document assumptions, unresolved questions, and which agency roles would need to review the artifact before it is adopted. The exercise should produce a work product that could support governance review, leadership discussion, implementation planning, or policy development.

Expected Artifact or Evidence

Decision-rights matrix; RACI table; risk acceptance thresholds; escalation authority list.

References and Resources for Additional Information

NIST Artificial Intelligence Risk Management Framework (AI RMF 1.0)

NIST Generative AI Profile

NIST Privacy Framework

CDC Public Health Data Strategy and Data Modernization resources

WHO Ethics and Governance of Artificial Intelligence for Health

OMB Memorandum M-24-10 on federal AI governance, risk management, and innovation, where relevant to public-sector governance models

Agency strategic plans, procurement policies, privacy policies, cybersecurity policies, records policies, and equity plans

Knowledge Check

1. What is the strongest reason this module topic belongs in AI leadership and governance training?

- A. It eliminates the need for technical review.
- B. It allows vendors to make final public health decisions.
- C. It helps leaders define authority, safeguards, accountability, and implementation expectations.
- D. It guarantees that every AI tool will be successful.

Correct answer: C

2. Which practice best supports accountable public health AI governance?

- A. Allowing each program to create informal AI rules independently.

- B. Using defined review steps, named owners, decision logs, and escalation pathways.
- C. Approving tools based only on vendor demonstrations.
- D. Avoiding documentation to reduce burden.

Correct answer: B

3. When should an AI use case be escalated for leadership or governance review?

- A. Only after a public complaint occurs.
- B. Only when a vendor requests escalation.
- C. Never, because AI review should remain informal.
- D. When it affects sensitive data, public communication, resource allocation, enforcement, service access, or public trust.

Correct answer: D

4. What is weak evidence of completion for a leadership and governance module?

- A. A generic statement that AI should be used responsibly without assigning owners or actions.
- B. A role matrix, governance worksheet, or decision-ready plan.
- C. A specific risk statement with mitigation steps.
- D. A completed exercise tied to a public health workflow.

Correct answer: A

5. What should leaders do when an AI use case has meaningful benefits but unresolved risks?

- A. Approve it immediately because benefits are possible.
- B. Reject all AI use permanently.
- C. Require mitigation, human review, monitoring, documentation, or defer implementation until readiness improves.
- D. Ask the vendor to decide whether risk is acceptable.

Correct answer: C

References and Resources

NIST Artificial Intelligence Risk Management Framework (AI RMF 1.0)

NIST Generative AI Profile

NIST Privacy Framework

CDC Public Health Data Strategy and Data Modernization resources

WHO Ethics and Governance of Artificial Intelligence for Health

OMB Memorandum M-24-10 on federal AI governance, risk management, and innovation, where relevant to public-sector governance models

Agency strategic plans, procurement policies, privacy policies, cybersecurity policies, records policies, and equity plans

EXE 410

AI Portfolio Strategy and Use Case Prioritization

Course ID	EXE 410
Track	Public Health Executive Leadership Track
Level	advanced/leadership
Estimated time	35
Audience	Executive leaders, Health officials, Senior managers, Governance chairs

Learning Objectives

- Define the purpose of an AI portfolio strategy for public health agencies.
- Develop criteria for prioritizing AI use cases based on public health value, readiness, risk, and sustainability.
- Identify use cases that should be pursued, piloted, deferred, redesigned, or rejected.
- Create a portfolio view that balances innovation, risk, equity, and operational capacity.
- Produce a use case prioritization rubric for governance or leadership review.

Jurisdiction and Agency Policy Note

This national-level module is intended for training and planning purposes. It does not replace legal, privacy, procurement, civil rights, records, or agency policy review. Learners should confirm applicable requirements in their own jurisdiction and organization before implementing AI-supported workflows.

Why This Topic Matters for Public Health AI

Public health agencies have many possible AI use cases: summarizing case notes, analyzing call logs, drafting messages, identifying data quality gaps, forecasting demand, prioritizing inspections, supporting grants, and improving knowledge management. Without a portfolio strategy, agencies may choose projects because they are visible, easy to demonstrate, or championed by a strong vendor rather than because they advance the public health mission.

A portfolio strategy protects scarce capacity. Technical teams, privacy officers, legal counsel, procurement staff, program leaders, and data stewards cannot review every idea at the same intensity. Prioritization helps the agency focus deep review on use cases with meaningful benefit and material risk while creating faster pathways for low-risk internal productivity uses.

Public Health Example

A health department receives proposals for five AI uses: a public chatbot, a staff policy assistant, an outbreak signal detection model, a vendor tool to summarize eCRs, and an internal tool to draft grant narratives. Each proposal has different benefits and risks. A public chatbot may be visible but high-risk if it gives health advice. A staff policy assistant may be lower risk if it uses approved internal documents. An outbreak model may be valuable but require strong data quality and validation.

A portfolio review allows the agency to compare these proposals systematically. The agency may approve a low-risk internal pilot, defer the chatbot until public communication safeguards are developed, require validation planning for the

outbreak model, and move the eCR summarization tool into deeper technical review. The result is a balanced portfolio rather than a series of disconnected decisions.

Technical and Governance Deep Dive

A useful portfolio strategy begins with categories. Agencies can classify AI opportunities by function, such as internal productivity, knowledge management, surveillance support, decision support, communication, automation, research, or public-facing service. They can also classify by risk, such as low, moderate, high, or prohibited. These categories help leadership understand where AI activity is concentrated and whether the portfolio reflects agency priorities.

Prioritization criteria should include mission alignment, expected public health benefit, workflow burden reduction, data readiness, legal authority, privacy sensitivity, equity impact, technical feasibility, implementation complexity, vendor dependency, cost, sustainability, and monitoring feasibility. A use case with high potential benefit but weak data readiness may be deferred until modernization work improves. A use case with moderate benefit and low risk may be an appropriate early pilot.

Portfolio governance should include regular review. Some use cases should move from idea to pilot, from pilot to operational use, or from operational use to retirement. Others should be stopped because they lack value, create unacceptable risk, duplicate another effort, or depend on unsustainable funding. A living portfolio helps leadership make these decisions transparently.

Risks, Failure Modes, and Guardrails

A major risk is pilot proliferation. Agencies may launch many small AI pilots without a pathway to scale, evaluate, sustain, or retire them. A guardrail is to require every pilot to define success criteria, evaluation evidence, funding assumptions, and an end point.

Another risk is inequitable prioritization. Projects serving better-resourced programs may advance faster than projects addressing underserved communities. A guardrail is to include equity impact and community benefit in the prioritization rubric.

Application to AI-Supported Workflows

Generative AI projects may be prioritized for internal knowledge management before public-facing use. Predictive analytics projects may be prioritized only when data and validation plans are mature. Agentic workflows may be deferred until governance has defined decision rights, access controls, and rollback procedures. Portfolio strategy helps leaders sequence AI adoption responsibly.

Reflection Questions

Which current or proposed AI use case in your agency raises this leadership or governance issue most clearly?

Who currently has authority to make decisions about this issue, and is that authority documented?

What evidence would governance or leadership need before approving the use case?

Which staff, partners, or communities would be affected if this issue were handled poorly?

What policy, process, or artifact should be created or updated after this module?

Practical Exercise

Develop an AI use case prioritization rubric with criteria, scoring levels, required evidence, and decision categories such as pursue, pilot, redesign, defer, or reject.

Before You Begin

Choose a real or realistic public health AI use case. Document assumptions, unresolved questions, and which agency roles would need to review the artifact before it is adopted. The exercise should produce a work product that could support governance review, leadership discussion, implementation planning, or policy development.

Expected Artifact or Evidence

Use case intake form; prioritization rubric; sample portfolio view; decision category definitions.

References and Resources for Additional Information

NIST Artificial Intelligence Risk Management Framework (AI RMF 1.0)

NIST Generative AI Profile

NIST Privacy Framework

CDC Public Health Data Strategy and Data Modernization resources

WHO Ethics and Governance of Artificial Intelligence for Health

OMB Memorandum M-24-10 on federal AI governance, risk management, and innovation, where relevant to public-sector governance models

Agency strategic plans, procurement policies, privacy policies, cybersecurity policies, records policies, and equity plans

Knowledge Check

1. What is the strongest reason this module topic belongs in AI leadership and governance training?

- A. It eliminates the need for technical review.
- B. It allows vendors to make final public health decisions.
- C. It guarantees that every AI tool will be successful.
- D. It helps leaders define authority, safeguards, accountability, and implementation expectations.

Correct answer: D

2. Which practice best supports accountable public health AI governance?

- A. Using defined review steps, named owners, decision logs, and escalation pathways.
- B. Allowing each program to create informal AI rules independently.
- C. Approving tools based only on vendor demonstrations.

D. Avoiding documentation to reduce burden.

Correct answer: A

3. When should an AI use case be escalated for leadership or governance review?

A. Only after a public complaint occurs.

B. Only when a vendor requests escalation.

C. When it affects sensitive data, public communication, resource allocation, enforcement, service access, or public trust.

D. Never, because AI review should remain informal.

Correct answer: C

4. What is weak evidence of completion for a leadership and governance module?

A. A role matrix, governance worksheet, or decision-ready plan.

B. A generic statement that AI should be used responsibly without assigning owners or actions.

C. A specific risk statement with mitigation steps.

D. A completed exercise tied to a public health workflow.

Correct answer: B

5. What should leaders do when an AI use case has meaningful benefits but unresolved risks?

A. Approve it immediately because benefits are possible.

B. Reject all AI use permanently.

C. Ask the vendor to decide whether risk is acceptable.

D. Require mitigation, human review, monitoring, documentation, or defer implementation until readiness improves.

Correct answer: D

References and Resources

NIST Artificial Intelligence Risk Management Framework (AI RMF 1.0)

NIST Generative AI Profile

NIST Privacy Framework

CDC Public Health Data Strategy and Data Modernization resources

WHO Ethics and Governance of Artificial Intelligence for Health

OMB Memorandum M-24-10 on federal AI governance, risk management, and innovation, where relevant to public-sector governance models

Agency strategic plans, procurement policies, privacy policies, cybersecurity policies, records policies, and equity plans

EXE 420

AI Investment Strategy and Business Case Review

Course ID	EXE 420
Track	Public Health Executive Leadership Track
Level	advanced/leadership
Estimated time	35
Audience	Executive leaders, Health officials, Senior managers, Governance chairs

Learning Objectives

- Identify the full cost components of AI implementation and operations.
- Evaluate expected benefits in terms of public health value, workflow improvement, equity, quality, timeliness, and staff burden.
- Distinguish between a vendor demonstration and a credible business case.
- Develop business case questions for AI procurement or internal development.
- Produce an AI business case review worksheet.

Jurisdiction and Agency Policy Note

This national-level module is intended for training and planning purposes. It does not replace legal, privacy, procurement, civil rights, records, or agency policy review. Learners should confirm applicable requirements in their own jurisdiction and organization before implementing AI-supported workflows.

Why This Topic Matters for Public Health AI

Public health agencies often face short-term funding cycles and urgent operational needs. AI products may be marketed as quick solutions, but the true cost of responsible implementation is broader than subscription fees. Agencies may need to update data pipelines, configure security controls, develop validation plans, train staff, write policies, monitor performance, and manage vendor relationships.

A strong business case also protects mission focus. AI should not be funded because it is novel. It should be funded when it addresses a meaningful public health problem, fits agency capacity, can be governed responsibly, and can produce measurable value. Business case review helps leaders compare AI investments with non-AI alternatives, such as workflow redesign, staffing changes, data quality improvements, or simpler automation.

Public Health Example

A vendor offers an AI tool that promises to reduce disease investigator workload by summarizing case reports. The annual license appears affordable, but the agency must also pay for integration with the surveillance system, security review, staff training, validation, legal review, vendor support, monitoring dashboards, and staff time for governance. If those costs are not considered, leadership may underestimate the investment required.

A business case review asks whether the tool will meaningfully reduce burden, improve timeliness, maintain quality, and fit the agency's workflow. It also asks whether a simpler approach, such as better case report templates or workflow

redesign, would solve the problem at lower cost and risk. The review does not assume AI is the answer; it tests whether AI is the right answer.

Technical and Governance Deep Dive

An AI business case should begin with a problem statement. The problem should be operational and measurable, such as delayed case review, inconsistent public response drafting, high manual data cleanup burden, or inability to search internal guidance efficiently. A vague desire to 'use AI' is not a sufficient problem statement. The business case should connect the proposed AI use to a real public health need.

The review should compare alternatives. Options may include doing nothing, improving workflow, improving data quality, using rules-based automation, procuring a vendor product, developing an internal tool, or joining a shared service. Each option should be assessed for benefit, cost, risk, time, staffing, sustainability, and governance burden. This comparison prevents AI from becoming the default answer to every operational problem.

Success measures should be defined before funding. Measures might include reduced processing time, improved completeness, fewer manual touches, better routing accuracy, staff satisfaction, fewer errors, improved language access, or reduced backlog. Measures should also include safety and equity indicators. A tool that saves time but increases error rates, widens disparities, or erodes public trust should not be considered successful.

Risks, Failure Modes, and Guardrails

One risk is hidden cost. Agencies may purchase an AI tool without funding integration, monitoring, validation, staff training, or ongoing support. A guardrail is to require total cost of ownership in every business case.

Another risk is benefits inflation. Vendors or project champions may overstate time savings or accuracy without local evidence. A guardrail is to require pilot success criteria and benefits realization review before scaling.

Application to AI-Supported Workflows

Generative AI business cases should include review time, source verification, and records considerations. Predictive analytics business cases should include validation data, subgroup testing, and monitoring costs. Agentic AI business cases should include security architecture, approval workflows, rollback, logging, and incident response. Cost and value must reflect responsible use.

Reflection Questions

Which current or proposed AI use case in your agency raises this leadership or governance issue most clearly?

Who currently has authority to make decisions about this issue, and is that authority documented?

What evidence would governance or leadership need before approving the use case?

Which staff, partners, or communities would be affected if this issue were handled poorly?

What policy, process, or artifact should be created or updated after this module?

Practical Exercise

Complete an AI business case review worksheet for one proposed investment. Include problem, alternatives, expected benefits, total cost of ownership, risk, equity considerations, implementation needs, and success measures.

Before You Begin

Choose a real or realistic public health AI use case. Document assumptions, unresolved questions, and which agency roles would need to review the artifact before it is adopted. The exercise should produce a work product that could support governance review, leadership discussion, implementation planning, or policy development.

Expected Artifact or Evidence

AI business case worksheet; total cost of ownership estimate; benefits realization plan; alternatives analysis.

References and Resources for Additional Information

NIST Artificial Intelligence Risk Management Framework (AI RMF 1.0)

NIST Generative AI Profile

NIST Privacy Framework

CDC Public Health Data Strategy and Data Modernization resources

WHO Ethics and Governance of Artificial Intelligence for Health

OMB Memorandum M-24-10 on federal AI governance, risk management, and innovation, where relevant to public-sector governance models

Agency strategic plans, procurement policies, privacy policies, cybersecurity policies, records policies, and equity plans

Knowledge Check

1. What is the strongest reason this module topic belongs in AI leadership and governance training?

- A. It eliminates the need for technical review.
- B. It helps leaders define authority, safeguards, accountability, and implementation expectations.
- C. It allows vendors to make final public health decisions.
- D. It guarantees that every AI tool will be successful.

Correct answer: B

2. Which practice best supports accountable public health AI governance?

- A. Allowing each program to create informal AI rules independently.
- B. Approving tools based only on vendor demonstrations.
- C. Avoiding documentation to reduce burden.
- D. Using defined review steps, named owners, decision logs, and escalation pathways.

Correct answer: D

3. When should an AI use case be escalated for leadership or governance review?

- A. When it affects sensitive data, public communication, resource allocation, enforcement, service access, or public trust.
- B. Only after a public complaint occurs.
- C. Only when a vendor requests escalation.
- D. Never, because AI review should remain informal.

Correct answer: A

4. What is weak evidence of completion for a leadership and governance module?

- A. A role matrix, governance worksheet, or decision-ready plan.
- B. A specific risk statement with mitigation steps.
- C. A generic statement that AI should be used responsibly without assigning owners or actions.
- D. A completed exercise tied to a public health workflow.

Correct answer: C

5. What should leaders do when an AI use case has meaningful benefits but unresolved risks?

- A. Approve it immediately because benefits are possible.
- B. Require mitigation, human review, monitoring, documentation, or defer implementation until readiness improves.
- C. Reject all AI use permanently.
- D. Ask the vendor to decide whether risk is acceptable.

Correct answer: B

References and Resources

NIST Artificial Intelligence Risk Management Framework (AI RMF 1.0)

NIST Generative AI Profile

NIST Privacy Framework

CDC Public Health Data Strategy and Data Modernization resources

WHO Ethics and Governance of Artificial Intelligence for Health

OMB Memorandum M-24-10 on federal AI governance, risk management, and innovation, where relevant to public-sector governance models

Agency strategic plans, procurement policies, privacy policies, cybersecurity policies, records policies, and equity plans

EXE 430

Organizational Change Management for AI Adoption

Course ID	EXE 430
Track	Public Health Executive Leadership Track
Level	advanced/leadership
Estimated time	35
Audience	Executive leaders, Health officials, Senior managers, Governance chairs

Learning Objectives

- Explain why AI implementation requires change management.
- Identify stakeholders affected by an AI-supported workflow.
- Develop a change management plan that includes communication, training, workflow updates, feedback, and adoption measures.
- Recognize staff concerns related to trust, workload, deskilling, surveillance, accountability, and job impact.
- Produce an AI adoption and change management plan.

Why This Topic Matters for Public Health AI

AI tools may fail even when the technology works because the organization is not ready to use them. Staff may not trust the output, may over-trust it, may not understand their review responsibilities, or may develop workarounds that undermine safeguards. Leaders may underestimate the time needed for training, workflow redesign, and user feedback.

Change management is also an equity and workforce issue. Staff who are not included in design may experience AI as a threat or burden. Programs with less technical capacity may be left behind. Communities may experience changes in service delivery without understanding how AI is involved. A change plan helps ensure that adoption is transparent, inclusive, and aligned with public health values.

Public Health Example

A health department launches an AI assistant to help staff find internal policies and program guidance. The tool is technically functional, but staff do not use it because they are unsure whether answers are official, whether searches are monitored, and whether they remain responsible for verifying information. Some staff worry that the tool will be used to judge their productivity.

A change management plan addresses these concerns directly. It explains the purpose of the tool, approved uses, prohibited uses, privacy protections, source citation requirements, human verification expectations, and feedback process. It also includes supervisors, program champions, training sessions, office hours, and monitoring of adoption barriers.

Technical and Governance Deep Dive

AI change management should begin before procurement or configuration. Leaders should identify who will be affected, what pain points the tool is supposed to address, which workflows will change, and what staff need to learn. Stakeholder

engagement should include frontline users, supervisors, program leaders, IT, privacy, legal, communications, equity staff, and where appropriate, community representatives.

The communication plan should be specific. Staff need to know what the tool does, what it does not do, why the agency is using it, what data may be entered, what outputs require review, and how errors should be reported. Communication should also acknowledge uncertainty. Overselling AI can damage trust when limitations become visible. Under-explaining AI can lead to rumors, resistance, or unsafe use.

Training should be tied to workflow. Generic AI literacy is not enough. Users need role-specific examples, review criteria, escalation scenarios, documentation expectations, and practice with realistic outputs. Supervisors need to know how to reinforce appropriate use and respond to concerns. Change management should continue after launch through feedback, monitoring, and iterative improvement.

Risks, Failure Modes, and Guardrails

One risk is adoption without understanding. Staff may use the tool incorrectly or treat outputs as authoritative because training was superficial. A guardrail is to require role-based training and workflow-specific practice before access is granted.

Another risk is resistance caused by lack of transparency. If staff believe AI is being introduced to monitor or replace them, they may avoid or undermine the tool. A guardrail is to communicate purpose, boundaries, data practices, and workforce implications clearly and early.

Application to AI-Supported Workflows

Generative AI adoption requires training on review, citation, privacy, and public records. Predictive analytics adoption requires training on interpretation, limitations, and override documentation. Agentic workflow adoption requires training on approvals, stop procedures, and escalation. Each AI type changes work in different ways, so change management must be specific.

Reflection Questions

Which current or proposed AI use case in your agency raises this leadership or governance issue most clearly?

Who currently has authority to make decisions about this issue, and is that authority documented?

What evidence would governance or leadership need before approving the use case?

Which staff, partners, or communities would be affected if this issue were handled poorly?

What policy, process, or artifact should be created or updated after this module?

Practical Exercise

Create a change management plan for one AI use case. Include stakeholder map, communication messages, training needs, workflow changes, adoption measures, feedback process, and launch support.

Before You Begin

Choose a real or realistic public health AI use case. Document assumptions, unresolved questions, and which agency roles would need to review the artifact before it is adopted. The exercise should produce a work product that could support governance review, leadership discussion, implementation planning, or policy development.

Expected Artifact or Evidence

AI adoption plan; stakeholder map; communication plan; training plan; feedback and adoption monitoring plan.

References and Resources for Additional Information

NIST Artificial Intelligence Risk Management Framework (AI RMF 1.0)

NIST Generative AI Profile

NIST Privacy Framework

CDC Public Health Data Strategy and Data Modernization resources

WHO Ethics and Governance of Artificial Intelligence for Health

OMB Memorandum M-24-10 on federal AI governance, risk management, and innovation, where relevant to public-sector governance models

Agency strategic plans, procurement policies, privacy policies, cybersecurity policies, records policies, and equity plans

Knowledge Check

1. What is the strongest reason this module topic belongs in AI leadership and governance training?

- A. It eliminates the need for technical review.
- B. It allows vendors to make final public health decisions.
- C. It guarantees that every AI tool will be successful.
- D. It helps leaders define authority, safeguards, accountability, and implementation expectations.

Correct answer: D

2. Which practice best supports accountable public health AI governance?

- A. Allowing each program to create informal AI rules independently.
- B. Approving tools based only on vendor demonstrations.
- C. Using defined review steps, named owners, decision logs, and escalation pathways.
- D. Avoiding documentation to reduce burden.

Correct answer: C

3. When should an AI use case be escalated for leadership or governance review?

- A. Only after a public complaint occurs.
- B. When it affects sensitive data, public communication, resource allocation, enforcement, service access, or public trust.
- C. Only when a vendor requests escalation.

D. Never, because AI review should remain informal.

Correct answer: B

4. What is weak evidence of completion for a leadership and governance module?

A. A generic statement that AI should be used responsibly without assigning owners or actions.

B. A role matrix, governance worksheet, or decision-ready plan.

C. A specific risk statement with mitigation steps.

D. A completed exercise tied to a public health workflow.

Correct answer: A

5. What should leaders do when an AI use case has meaningful benefits but unresolved risks?

A. Approve it immediately because benefits are possible.

B. Reject all AI use permanently.

C. Ask the vendor to decide whether risk is acceptable.

D. Require mitigation, human review, monitoring, documentation, or defer implementation until readiness improves.

Correct answer: D

References and Resources

NIST Artificial Intelligence Risk Management Framework (AI RMF 1.0)

NIST Generative AI Profile

NIST Privacy Framework

CDC Public Health Data Strategy and Data Modernization resources

WHO Ethics and Governance of Artificial Intelligence for Health

OMB Memorandum M-24-10 on federal AI governance, risk management, and innovation, where relevant to public-sector governance models

Agency strategic plans, procurement policies, privacy policies, cybersecurity policies, records policies, and equity plans

EXE 440

Workforce Strategy and AI Role Design

Course ID	EXE 440
Track	Public Health Executive Leadership Track
Level	advanced/leadership
Estimated time	35
Audience	Executive leaders, Health officials, Senior managers, Governance chairs

Learning Objectives

- Identify workforce roles needed for responsible AI adoption in public health.
- Distinguish between AI users, reviewers, implementers, governance participants, and accountability owners.
- Develop role-based competency expectations for AI-supported work.
- Recognize workforce risks such as deskilling, uneven access to training, and unclear role boundaries.
- Produce a workforce readiness and role design plan.

Why This Topic Matters for Public Health AI

Public health AI depends on people who understand both technology and context. A model may generate a summary, but an epidemiologist must judge whether the summary is meaningful. A vendor may describe a product, but an informatician must assess interoperability and data flow. A leader may approve a use case, but governance participants must ensure that privacy, equity, security, and legal concerns have been addressed.

Workforce strategy helps agencies move beyond ad hoc training. It defines which staff need baseline literacy, which staff need deeper technical skills, which staff need governance training, and which roles require ongoing competency development. It also helps agencies avoid concentrating AI knowledge in a small group while leaving most staff dependent on informal guidance.

Public Health Example

A public health agency implements several AI tools: a document assistant, a case summarization tool, a data quality monitor, and a predictive model for resource planning. Staff quickly discover that no single role owns all the knowledge needed. Program staff understand workflow, IT understands integration, privacy understands data restrictions, and analysts understand validation. However, there is no role that connects these perspectives consistently.

The agency creates an AI bridge role to support use case development, governance preparation, workflow mapping, vendor questioning, training, and implementation review. The role does not replace technical specialists or program leaders. It helps them work together more effectively.

Technical and Governance Deep Dive

An AI workforce strategy should define role groups. General staff may need baseline AI literacy, approved-use guidance, and prompt safety. Frontline users need workflow-specific training and escalation expectations. Supervisors need to support review, feedback, and documentation. Program owners need to understand accountability and performance

monitoring. Technical staff need deeper knowledge of data, architecture, validation, and security. Governance members need enough shared language to review proposals consistently.

Competency expectations should be practical. Staff do not all need to become data scientists. However, many need to recognize when AI output is unsupported, when sensitive data should not be entered, when a use case exceeds approved boundaries, and when to escalate. Technical and informatics staff may need competencies in data quality, interoperability, model evaluation, RAG, security, and lifecycle monitoring.

Role design should also address workload and equity. AI review and governance work require time. If these responsibilities are added informally to already overloaded staff, governance may become superficial. Leaders should identify protected time, training resources, role descriptions, and career pathways for staff who support responsible AI.

Risks, Failure Modes, and Guardrails

One risk is assuming that AI reduces workforce needs. In practice, responsible AI often requires new review, training, documentation, and monitoring work. A guardrail is to include workforce impact in every business case and implementation plan.

Another risk is uneven access to AI skills. Staff in better-resourced programs may receive training first, while others fall behind. A guardrail is to create a role-based training plan that includes all relevant staff groups and supports equitable participation.

Application to AI-Supported Workflows

Generative AI requires staff who can review outputs and manage source accuracy. Predictive analytics requires staff who can interpret performance and understand limitations. Agentic workflows require staff who can oversee permissions, actions, escalation, and rollback. Workforce strategy should connect each AI type to the roles needed to use it safely.

Reflection Questions

Which current or proposed AI use case in your agency raises this leadership or governance issue most clearly?

Who currently has authority to make decisions about this issue, and is that authority documented?

What evidence would governance or leadership need before approving the use case?

Which staff, partners, or communities would be affected if this issue were handled poorly?

What policy, process, or artifact should be created or updated after this module?

Practical Exercise

Create a role-based AI workforce readiness plan. Include role groups, competencies, required training, responsibilities, decision authority, and support needs.

Before You Begin

Choose a real or realistic public health AI use case. Document assumptions, unresolved questions, and which agency roles would need to review the artifact before it is adopted. The exercise should produce a work product that could support governance review, leadership discussion, implementation planning, or policy development.

Expected Artifact or Evidence

Workforce readiness plan; role-based competency matrix; training assignment matrix; AI bridge role description.

References and Resources for Additional Information

NIST Artificial Intelligence Risk Management Framework (AI RMF 1.0)

NIST Generative AI Profile

NIST Privacy Framework

CDC Public Health Data Strategy and Data Modernization resources

WHO Ethics and Governance of Artificial Intelligence for Health

OMB Memorandum M-24-10 on federal AI governance, risk management, and innovation, where relevant to public-sector governance models

Agency strategic plans, procurement policies, privacy policies, cybersecurity policies, records policies, and equity plans

Knowledge Check

1. What is the strongest reason this module topic belongs in AI leadership and governance training?

- A. It eliminates the need for technical review.
- B. It allows vendors to make final public health decisions.
- C. It helps leaders define authority, safeguards, accountability, and implementation expectations.
- D. It guarantees that every AI tool will be successful.

Correct answer: C

2. Which practice best supports accountable public health AI governance?

- A. Using defined review steps, named owners, decision logs, and escalation pathways.
- B. Allowing each program to create informal AI rules independently.
- C. Approving tools based only on vendor demonstrations.
- D. Avoiding documentation to reduce burden.

Correct answer: A

3. When should an AI use case be escalated for leadership or governance review?

- A. Only after a public complaint occurs.
- B. Only when a vendor requests escalation.
- C. Never, because AI review should remain informal.

D. When it affects sensitive data, public communication, resource allocation, enforcement, service access, or public trust.

Correct answer: D

4. What is weak evidence of completion for a leadership and governance module?

- A. A role matrix, governance worksheet, or decision-ready plan.
- B. A generic statement that AI should be used responsibly without assigning owners or actions.
- C. A specific risk statement with mitigation steps.
- D. A completed exercise tied to a public health workflow.

Correct answer: B

5. What should leaders do when an AI use case has meaningful benefits but unresolved risks?

- A. Approve it immediately because benefits are possible.
- B. Reject all AI use permanently.
- C. Require mitigation, human review, monitoring, documentation, or defer implementation until readiness improves.
- D. Ask the vendor to decide whether risk is acceptable.

Correct answer: C

References and Resources

NIST Artificial Intelligence Risk Management Framework (AI RMF 1.0)

NIST Generative AI Profile

NIST Privacy Framework

CDC Public Health Data Strategy and Data Modernization resources

WHO Ethics and Governance of Artificial Intelligence for Health

OMB Memorandum M-24-10 on federal AI governance, risk management, and innovation, where relevant to public-sector governance models

Agency strategic plans, procurement policies, privacy policies, cybersecurity policies, records policies, and equity plans

EXE 450

AI Sustainability and Lifecycle Funding

Course ID	EXE 450
Track	Public Health Executive Leadership Track
Level	advanced/leadership
Estimated time	35
Audience	Executive leaders, Health officials, Senior managers, Governance chairs

Learning Objectives

- Explain why AI sustainability must be planned before launch.
- Identify lifecycle costs and responsibilities for AI tools and workflows.
- Develop sustainability criteria for moving from pilot to operational use.
- Define sunset, pause, and retirement criteria for AI use cases.
- Produce an AI lifecycle funding and sustainability plan.

Jurisdiction and Agency Policy Note

This national-level module is intended for training and planning purposes. It does not replace legal, privacy, procurement, civil rights, records, or agency policy review. Learners should confirm applicable requirements in their own jurisdiction and organization before implementing AI-supported workflows.

Why This Topic Matters for Public Health AI

Many AI efforts begin as pilots, innovation projects, or grant-funded experiments. These projects may succeed technically but fail operationally if there is no plan for ongoing support. Public health agencies may be left with a tool that staff rely on but no budget for validation updates, vendor support, data integration maintenance, monitoring, or training new users.

Sustainability is also a risk-control issue. AI tools can become unsafe when they are not maintained. Data sources change, regulations change, staff roles change, vendor models change, and public expectations change. A sustainability plan helps the agency decide whether the tool should continue, be modified, be replaced, or be retired.

Public Health Example

A grant funds a two-year AI pilot to help identify likely duplicate records in a public health system. The pilot improves efficiency, but the grant does not cover long-term hosting, monitoring, system upgrades, training, or validation after data sources change. When funding ends, staff continue relying on the tool even though no one is assigned to review performance.

A sustainability plan would have addressed this before launch. It would define operational ownership, long-term funding, monitoring indicators, update cadence, user support, vendor responsibilities, and criteria for continuing or retiring the tool. It would also define a manual fallback process if the tool is paused.

Technical and Governance Deep Dive

AI lifecycle planning should begin during intake and business case review. Leaders should ask whether the use case is a short-term experiment, a time-limited emergency support tool, or a long-term operational capability. Each category has different funding and governance needs. A time-limited tool may need clear end dates and data deletion rules. An operational tool needs sustained staffing, budget, contracts, monitoring, and training.

Lifecycle costs include more than licenses. Agencies must consider integration maintenance, data pipeline monitoring, model evaluation, cybersecurity review, privacy review, accessibility updates, user support, vendor management, documentation, incident response, retraining or reconfiguration, and future procurement. These costs should be tied to ownership. If no role or budget owns a lifecycle function, the function may not happen.

Sunset and retirement criteria are essential. A tool may be retired if it no longer aligns with agency priorities, if benefits are not realized, if risks outweigh value, if vendor terms change, if performance deteriorates, if data sources are no longer reliable, or if a better shared service becomes available. Retirement should be planned so staff know how to transition without disrupting public health operations.

Risks, Failure Modes, and Guardrails

One risk is pilot dependency. Staff may begin relying on a tool that was never funded or governed for long-term use. A guardrail is to define pilot duration, scale limits, success criteria, and continuation requirements before launch.

Another risk is silent degradation. A tool may continue operating after data, policies, user behavior, or vendor models change. A guardrail is to require periodic lifecycle review and clear pause or sunset criteria.

Application to AI-Supported Workflows

Generative AI sustainability requires prompt and source updates, user training, output review, and policy alignment. Predictive analytics sustainability requires drift monitoring, recalibration, subgroup review, and documentation. Agentic AI sustainability requires continued security review, permission management, incident drills, and rollback testing. Lifecycle funding should match the safeguards required by the AI type.

Reflection Questions

Which current or proposed AI use case in your agency raises this leadership or governance issue most clearly?

Who currently has authority to make decisions about this issue, and is that authority documented?

What evidence would governance or leadership need before approving the use case?

Which staff, partners, or communities would be affected if this issue were handled poorly?

What policy, process, or artifact should be created or updated after this module?

Practical Exercise

Create a lifecycle funding and sustainability plan for one AI use case. Include costs, owners, funding sources, monitoring, update cadence, vendor management, training, fallback process, and sunset criteria.

Before You Begin

Choose a real or realistic public health AI use case. Document assumptions, unresolved questions, and which agency roles would need to review the artifact before it is adopted. The exercise should produce a work product that could support governance review, leadership discussion, implementation planning, or policy development.

Expected Artifact or Evidence

Sustainability plan; lifecycle cost list; ownership map; continuation and sunset criteria; rollback or fallback plan.

References and Resources for Additional Information

NIST Artificial Intelligence Risk Management Framework (AI RMF 1.0)

NIST Generative AI Profile

NIST Privacy Framework

CDC Public Health Data Strategy and Data Modernization resources

WHO Ethics and Governance of Artificial Intelligence for Health

OMB Memorandum M-24-10 on federal AI governance, risk management, and innovation, where relevant to public-sector governance models

Agency strategic plans, procurement policies, privacy policies, cybersecurity policies, records policies, and equity plans

Knowledge Check

1. What is the strongest reason this module topic belongs in AI leadership and governance training?

- A. It eliminates the need for technical review.
- B. It allows vendors to make final public health decisions.
- C. It guarantees that every AI tool will be successful.
- D. It helps leaders define authority, safeguards, accountability, and implementation expectations.

Correct answer: D

2. Which practice best supports accountable public health AI governance?

- A. Using defined review steps, named owners, decision logs, and escalation pathways.
- B. Allowing each program to create informal AI rules independently.
- C. Approving tools based only on vendor demonstrations.
- D. Avoiding documentation to reduce burden.

Correct answer: A

3. When should an AI use case be escalated for leadership or governance review?

- A. Only after a public complaint occurs.

- B. Only when a vendor requests escalation.
- C. When it affects sensitive data, public communication, resource allocation, enforcement, service access, or public trust.
- D. Never, because AI review should remain informal.

Correct answer: C

4. What is weak evidence of completion for a leadership and governance module?

- A. A role matrix, governance worksheet, or decision-ready plan.
- B. A generic statement that AI should be used responsibly without assigning owners or actions.
- C. A specific risk statement with mitigation steps.
- D. A completed exercise tied to a public health workflow.

Correct answer: B

5. What should leaders do when an AI use case has meaningful benefits but unresolved risks?

- A. Approve it immediately because benefits are possible.
- B. Reject all AI use permanently.
- C. Ask the vendor to decide whether risk is acceptable.
- D. Require mitigation, human review, monitoring, documentation, or defer implementation until readiness improves.

Correct answer: D

References and Resources

NIST Artificial Intelligence Risk Management Framework (AI RMF 1.0)

NIST Generative AI Profile

NIST Privacy Framework

CDC Public Health Data Strategy and Data Modernization resources

WHO Ethics and Governance of Artificial Intelligence for Health

OMB Memorandum M-24-10 on federal AI governance, risk management, and innovation, where relevant to public-sector governance models

Agency strategic plans, procurement policies, privacy policies, cybersecurity policies, records policies, and equity plans

EXE 460

AI Governance Board Chair Training

Course ID	EXE 460
Track	Public Health Executive Leadership Track
Level	advanced/leadership
Estimated time	35
Audience	Executive leaders, Health officials, Senior managers, Governance chairs

Learning Objectives

- Explain the purpose of ai governance board chair training in public health AI.
- Identify the roles, decisions, risks, and documentation needed for the topic.
- Apply the topic to a real or realistic public health workflow or decision.
- Recognize equity, privacy, security, legal, operational, and public trust implications.
- Produce a practical artifact that supports readiness, governance, implementation, monitoring, or accountability.

Definitions

Decision right	The authority to approve, deny, modify, pause, or escalate an AI use case or operational decision.
Evidence artifact	A document, log, checklist, test result, approval record, or review note that demonstrates responsible AI practice.
Escalation pathway	A defined route for raising concerns to the appropriate reviewer or decision-maker.
Accountability owner	The person or role responsible for ensuring that an AI use case is governed, monitored, and corrected when needed.
Operational safeguard	A procedure, control, or review step that reduces risk during day-to-day use.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Why This Topic Matters for Public Health AI

AI Governance Board Chair Training matters because AI adoption changes how public health agencies make, document, communicate, and review decisions. Without a structured approach, staff may rely on informal practices,

vendor assurances, or one-time approvals that do not match the operational consequences of AI-supported work.

Public health agencies operate under legal authorities, public accountability, resource constraints, and community trust obligations. A module on this topic helps agencies define repeatable expectations before urgent situations arise. It also helps staff distinguish low-risk administrative support from uses that affect surveillance, public communications, field response, resource allocation, civil rights, privacy, or official records.

Public Health Example

A public health agency is preparing to use AI in a workflow related to running governance meetings, documenting decisions, and managing risk escalation. The project team initially treats the tool as a helpful efficiency aid, but governance reviewers recognize that the workflow may affect public trust, equity, official records, or public health action.

The agency responds by creating a clear artifact that defines intended use, prohibited use, responsible roles, review points, documentation expectations, escalation triggers, and monitoring measures. This converts a broad AI idea into a manageable public health practice.

Technical and Operational Deep Dive

The first step in this topic is to define the decision or workflow clearly. Staff should identify who is affected, what data are used, what AI output is produced, who reviews the output, and what action may follow. This prevents the module from becoming abstract and anchors the work in public health practice.

The second step is to assign responsibilities. AI-supported workflows often cross program, IT, legal, privacy, communications, procurement, and leadership boundaries. Responsibilities should be named by role rather than assumed. The artifact should specify who prepares materials, who reviews risks, who approves use, who monitors performance, and who can pause or revise the workflow.

The third step is to define evidence. Responsible AI depends on more than policy statements. Agencies need records that show what was reviewed, what decisions were made, what safeguards were applied, what incidents occurred, and how the workflow changed over time. Evidence also supports audits, public records review, quality improvement, and leadership accountability.

Risks, Failure Modes, and Guardrails

A common failure mode is treating the topic as a one-time checklist. AI risks can change when data sources, models, vendors, policies, users, or public expectations change. A guardrail is to define review intervals and change triggers.

Another failure mode is unclear ownership. When everyone is responsible in general, no one may be responsible in practice. A guardrail is to assign an accountability owner and name escalation pathways.

A third failure mode is relying on vendor claims or informal staff confidence without documentation. A guardrail is to require evidence artifacts that can be reviewed by governance, leadership, privacy, legal, equity, security, or program teams.

Application to AI-Supported Workflows

This topic applies across generative AI, predictive analytics, NLP, RAG, automation, and agentic workflows. The more sensitive the data, the more consequential the output, or the more automated the action, the stronger the

documentation, review, and monitoring expectations should be.

Reflection Questions

What public health decision or workflow does this topic affect?

Who has authority to approve or pause the work?

What evidence would show that the workflow is responsible and effective?

What equity, privacy, legal, security, or public trust risks require review?

How will the agency know when the workflow needs to be revised?

Practical Exercise

Create a practical ai governance board chair training artifact for one AI-supported public health use case. The artifact should define the use case, affected roles, review criteria, safeguards, documentation expectations, escalation triggers, monitoring indicators, and next review date.

Expected Artifact or Evidence

AI Governance Board Chair Training artifact

Role and responsibility table

Risk and safeguard summary

Escalation and monitoring plan

Approval or review note

References and Resources for Additional Information

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

Knowledge Check

1. What is the strongest evidence that ai governance board chair training has been applied responsibly?

- A. A general statement that AI is useful
- B. A vendor demonstration with no agency review
- C. A documented artifact that defines use, roles, safeguards, escalation, monitoring, and evidence
- D. A decision to skip review because the use case seems efficient

Correct answer: C

2. What is the best reason to assign an accountability owner?

- A. To reduce the need for documentation
- B. To ensure someone is responsible for monitoring, escalation, and correction
- C. To make the vendor solely responsible
- D. To prevent staff from raising concerns

Correct answer: B

3. When should an AI-supported workflow be escalated for review?

- A. Only after media coverage
- B. Never if the tool is efficient
- C. When outputs may affect safety, equity, privacy, legal authority, public trust, or official action
- D. Only at the end of the fiscal year

Correct answer: C

4. Which practice best supports public accountability?

- A. Keeping AI use informal
- B. Avoiding written records
- C. Letting each user decide independently
- D. Maintaining evidence artifacts, decision records, review notes, and monitoring results

Correct answer: D

5. What should happen when data sources, policy, vendor configuration, or workflow use changes?

- A. The use case should be reviewed to determine whether safeguards, validation, documentation, or approvals need updating
- B. The agency should ignore the change if staff like the tool
- C. The model should be trusted automatically
- D. The glossary should be deleted

Correct answer: A

References and Resources

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

EXE 470

AI Crisis Communications and Executive Response

Course ID	EXE 470
Track	Public Health Executive Leadership Track
Level	advanced/leadership
Estimated time	35
Audience	Executive leaders, Health officials, Senior managers, Governance chairs

Learning Objectives

- Explain the purpose of ai crisis communications and executive response in public health AI.
- Identify the roles, decisions, risks, and documentation needed for the topic.
- Apply the topic to a real or realistic public health workflow or decision.
- Recognize equity, privacy, security, legal, operational, and public trust implications.
- Produce a practical artifact that supports readiness, governance, implementation, monitoring, or accountability.

Definitions

Decision right	The authority to approve, deny, modify, pause, or escalate an AI use case or operational decision.
Evidence artifact	A document, log, checklist, test result, approval record, or review note that demonstrates responsible AI practice.
Escalation pathway	A defined route for raising concerns to the appropriate reviewer or decision-maker.
Accountability owner	The person or role responsible for ensuring that an AI use case is governed, monitored, and corrected when needed.
Operational safeguard	A procedure, control, or review step that reduces risk during day-to-day use.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Why This Topic Matters for Public Health AI

AI Crisis Communications and Executive Response matters because AI adoption changes how public health agencies make, document, communicate, and review decisions. Without a structured approach, staff may rely on informal

practices, vendor assurances, or one-time approvals that do not match the operational consequences of AI-supported work.

Public health agencies operate under legal authorities, public accountability, resource constraints, and community trust obligations. A module on this topic helps agencies define repeatable expectations before urgent situations arise. It also helps staff distinguish low-risk administrative support from uses that affect surveillance, public communications, field response, resource allocation, civil rights, privacy, or official records.

Public Health Example

A public health agency is preparing to use AI in a workflow related to executive response to AI incidents, media scrutiny, public concern, and vendor failure. The project team initially treats the tool as a helpful efficiency aid, but governance reviewers recognize that the workflow may affect public trust, equity, official records, or public health action.

The agency responds by creating a clear artifact that defines intended use, prohibited use, responsible roles, review points, documentation expectations, escalation triggers, and monitoring measures. This converts a broad AI idea into a manageable public health practice.

Technical and Operational Deep Dive

The first step in this topic is to define the decision or workflow clearly. Staff should identify who is affected, what data are used, what AI output is produced, who reviews the output, and what action may follow. This prevents the module from becoming abstract and anchors the work in public health practice.

The second step is to assign responsibilities. AI-supported workflows often cross program, IT, legal, privacy, communications, procurement, and leadership boundaries. Responsibilities should be named by role rather than assumed. The artifact should specify who prepares materials, who reviews risks, who approves use, who monitors performance, and who can pause or revise the workflow.

The third step is to define evidence. Responsible AI depends on more than policy statements. Agencies need records that show what was reviewed, what decisions were made, what safeguards were applied, what incidents occurred, and how the workflow changed over time. Evidence also supports audits, public records review, quality improvement, and leadership accountability.

Risks, Failure Modes, and Guardrails

A common failure mode is treating the topic as a one-time checklist. AI risks can change when data sources, models, vendors, policies, users, or public expectations change. A guardrail is to define review intervals and change triggers.

Another failure mode is unclear ownership. When everyone is responsible in general, no one may be responsible in practice. A guardrail is to assign an accountability owner and name escalation pathways.

A third failure mode is relying on vendor claims or informal staff confidence without documentation. A guardrail is to require evidence artifacts that can be reviewed by governance, leadership, privacy, legal, equity, security, or program teams.

Application to AI-Supported Workflows

This topic applies across generative AI, predictive analytics, NLP, RAG, automation, and agentic workflows. The more sensitive the data, the more consequential the output, or the more automated the action, the stronger the documentation, review, and monitoring expectations should be.

Reflection Questions

What public health decision or workflow does this topic affect?

Who has authority to approve or pause the work?

What evidence would show that the workflow is responsible and effective?

What equity, privacy, legal, security, or public trust risks require review?

How will the agency know when the workflow needs to be revised?

Practical Exercise

Create a practical ai crisis communications and executive response artifact for one AI-supported public health use case. The artifact should define the use case, affected roles, review criteria, safeguards, documentation expectations, escalation triggers, monitoring indicators, and next review date.

Expected Artifact or Evidence

AI Crisis Communications and Executive Response artifact

Role and responsibility table

Risk and safeguard summary

Escalation and monitoring plan

Approval or review note

References and Resources for Additional Information

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

Knowledge Check

1. What is the strongest evidence that ai crisis communications and executive response has been applied responsibly?

- A. A general statement that AI is useful
- B. A vendor demonstration with no agency review
- C. A documented artifact that defines use, roles, safeguards, escalation, monitoring, and evidence
- D. A decision to skip review because the use case seems efficient

Correct answer: C

2. What is the best reason to assign an accountability owner?

- A. To reduce the need for documentation
- B. To ensure someone is responsible for monitoring, escalation, and correction
- C. To make the vendor solely responsible
- D. To prevent staff from raising concerns

Correct answer: B

3. When should an AI-supported workflow be escalated for review?

- A. Only after media coverage
- B. Never if the tool is efficient
- C. When outputs may affect safety, equity, privacy, legal authority, public trust, or official action
- D. Only at the end of the fiscal year

Correct answer: C

4. Which practice best supports public accountability?

- A. Keeping AI use informal
- B. Avoiding written records
- C. Letting each user decide independently
- D. Maintaining evidence artifacts, decision records, review notes, and monitoring results

Correct answer: D

5. What should happen when data sources, policy, vendor configuration, or workflow use changes?

- A. The use case should be reviewed to determine whether safeguards, validation, documentation, or approvals need updating
- B. The agency should ignore the change if staff like the tool
- C. The model should be trusted automatically
- D. The glossary should be deleted

Correct answer: A

References and Resources

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

GOV 100

Introduction to AI Governance for Public Health

Course ID	GOV 100
Track	Governance and Security Track
Level	introductory
Estimated time	40
Audience	All public health staff, Program staff, Technical staff, Supervisors, Epidemiologists, Communications staff, Policy staff, Data and informatics staff, Governance participants

Learning Objectives

- Define AI governance in plain language and explain why it matters for public health practice.
- Distinguish AI governance from AI policy, IT security review, project management, and individual tool use.
- Identify common AI governance roles, including program owner, data steward, privacy lead, security lead, legal counsel, equity reviewer, communications reviewer, technical reviewer, procurement reviewer, executive sponsor, and frontline user.
- Recognize which AI uses should be escalated for formal review because they involve sensitive data, public-facing information, individual or community impact, resource allocation, surveillance action, vendor products, automation, or public trust risks.
- Apply a basic governance lens to a public health AI use case by documenting purpose, data involved, expected benefit, risk, human review, decision rights, and escalation needs.
- Produce an introductory AI governance intake and escalation note for one real or realistic public health workflow.

Definitions

AI governance	AI governance is the organized approach an agency uses to direct, review, document, monitor, and improve AI use. It includes policies, roles, review pathways, decision rights, evidence requirements, approval conditions, monitoring expectations, escalation processes, and retirement decisions.
Governance body	A governance body is the group or structure responsible for reviewing AI use and making or recommending decisions. It may be a formal committee, an existing data governance group with expanded responsibilities, a review board, or a coordinated process across leadership, program, IT, privacy, legal, equity, procurement, communications, and technical reviewers.
Decision rights	Decision rights define who is allowed to approve, reject, pause, revise, monitor, or retire an AI use. Decision rights should be explicit because AI decisions often cross program, technology, legal, privacy, procurement, and public accountability boundaries.

AI use case	An AI use case describes a specific way AI will be used to support a public health task or workflow. A strong use case identifies the public health purpose, affected users or populations, data involved, AI function, expected benefit, human review process, risks, and measures of success.
Risk tier	A risk tier is a category used to determine the level of review required. A low-risk internal drafting aid may not need the same review as a tool that prioritizes field response, summarizes case narratives, identifies high-risk communities, or generates public-facing recommendations.
Human accountability	Human accountability means that a person or accountable role remains responsible for reviewing AI outputs, interpreting context, making decisions, documenting rationale, and responding to errors. AI should not become the invisible decision-maker in a public health workflow.
Oversight	Oversight is the ongoing review of AI use after approval. It includes monitoring performance, reviewing incidents, checking equity and usability, confirming that approved conditions are followed, and deciding whether the AI use should continue, change, pause, or be retired.
Escalation	Escalation is the process for raising a question, concern, incident, or proposed high-risk use to the appropriate reviewer or governance body. Staff should know when and how to escalate privacy concerns, incorrect outputs, potential bias, unsafe recommendations, public communication risks, vendor issues, or unapproved use.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, information security officials, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, procurement rule, labor agreement, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Why This Topic Matters for Public Health AI

Public health agencies make decisions that affect people, communities, partners, regulated entities, funding recipients, and the public record. AI can help staff work faster and identify patterns, but it can also introduce new risks. An AI tool may produce a convincing but incorrect summary, expose protected or sensitive information, make uneven recommendations across populations, rely on outdated data, automate a step that should require human judgment, or create confusion about who is responsible for a decision. Governance creates a way to identify these risks before they become routine practice.

Governance is especially important because AI use can start informally. A staff member may paste text into a public chatbot to shorten a report. A program may test a vendor tool before procurement has reviewed data terms. A dashboard team may add an AI-generated interpretation to a public-facing display. An analyst may use code assistance to modify a production script. Each action may seem small, but the combined effect can create privacy, equity, security, quality, procurement, records, and public trust risks if the agency does not have shared expectations.

AI governance also supports innovation. Clear rules help staff know which uses are allowed, which uses require review, and which uses are prohibited. Governance can prevent unnecessary delays by creating a risk-tiered pathway. Low-risk internal productivity uses may need training and supervisor awareness, while higher-risk uses may need privacy, security, equity, legal, data quality, procurement, and communications review. Staff can move faster when they know the pathway.

For public health, governance must connect technology decisions to mission and authority. A technically impressive tool is not enough. The agency must ask whether the tool serves a legitimate public health purpose, uses appropriate data, supports equitable outcomes, preserves human accountability, can be explained to affected communities, and can be monitored over time.

Public Health Example

A local health department wants to use a generative AI tool to help staff summarize long community listening session notes. The goal is to identify common themes related to access to mobile vaccination clinics, language needs, transportation barriers, and trust concerns. Program staff see the tool as a way to save time. Communications staff want to use the themes in a public report. The data team is concerned that some notes include names, small community identifiers, and sensitive comments about immigration status, disability, and prior negative experiences with government services.

Without governance, the project could move forward as a convenience tool. Staff might upload raw notes into an unapproved system, accept the summary without checking whether smaller communities were misrepresented, and publish language that makes the process sound more objective than it is. The agency could unintentionally expose sensitive information, overlook minority viewpoints, or make decisions based on a summary that does not reflect the source material.

With governance, the team would complete a simple use case intake before using the tool. The intake would document the public health purpose, the type of data involved, whether the data include sensitive or identifiable information, who will review the output, how the summary will be validated, whether community voices could be flattened or misrepresented, and whether the output will be used internally or publicly. The governance pathway might require removing identifiers, using an approved secure tool, keeping human review, documenting limitations, and having communications and equity reviewers examine any public-facing summary before release.

The example shows that governance is not only for technical teams. Program staff define the purpose and workflow. Data and privacy staff assess what may be entered into a tool. Equity and community engagement staff consider whether the process preserves community meaning. Communications staff ensure public statements are clear and not overstated. Leadership decides whether the use is appropriate and whether the agency can explain it to the public.

Technical, Programmatic, and Operational Deep Dive

AI governance should be understandable to all staff because the earliest governance decisions often happen before a formal review meeting. The first decision is recognizing that an AI use exists. Staff should treat AI use as more than “using a tool.” If a tool classifies, summarizes, predicts, prioritizes, recommends, generates, translates, extracts, searches, or automates information, the agency should consider whether it falls within AI governance expectations.

The second decision is identifying the purpose. A governance intake should begin with the public health problem, not with the tool. A strong purpose statement explains what workflow problem exists, who is affected, why the current process is insufficient, and what improvement would be meaningful. “Use AI to improve efficiency” is too vague. “Use an approved tool to draft first-pass plain-language versions of internal preparedness messages for human review” is clearer and easier to govern.

The third decision is understanding the data. Public health data may include protected health information, personally identifiable information, small cell counts, tribal data, school data, social service data, behavioral health data, geospatial data, community feedback, images, laboratory data, case notes, or operational records. Governance asks what data will be entered into the system, where it will go, who can access it, whether it can be retained or reused, and what agreements or policies apply.

The fourth decision is defining the role of the AI output. An AI output may be a draft, a suggestion, a risk score, a classification, a summary, a translation, a code recommendation, a dashboard interpretation, or a proposed action. The governance question is whether the output informs human work or changes a decision. A draft that a trained staff member reviews may be low risk. A score that determines who receives outreach first, which complaint receives inspection, or what information appears on a public dashboard needs stronger review.

The fifth decision is assigning human review. Human review must be meaningful. A reviewer should know what to check, what sources to compare against, what errors are most likely, when to reject the output, and how to document the review. A checkbox that says “human reviewed” is not enough if the reviewer lacks the time, authority, or expertise to challenge the output.

The sixth decision is determining the review pathway. A practical governance model uses risk tiers. Low-risk uses may require staff training, approved tools, and supervisor awareness. Moderate-risk uses may require a written use case intake, privacy review, security review, and program approval. High-risk uses may require formal governance review, legal and civil rights input, validation evidence, community engagement, monitoring plans, and leadership approval.

The seventh decision is planning oversight after approval. AI use should not be approved once and forgotten. The agency should know what evidence will be reviewed over time, such as error reports, user feedback, performance measures, equity indicators, security events, public complaints, vendor changes, prompt changes, model changes, workflow changes, and whether the use still serves the original public health purpose.

Risks, Failure Modes, and Guardrails

Unapproved tool use. Staff may use public or personal AI tools before the agency has reviewed privacy, security, retention, or data-use terms. Guardrail: publish a plain-language list of approved, restricted, and prohibited uses, including examples of data that may not be entered into public tools.

Hidden automation. AI may be embedded in vendor products, office software, dashboards, analytics platforms, or workflow tools. Staff may not realize that AI is influencing a recommendation or output. Guardrail: require vendors and project teams to disclose AI functions, data flows, model updates, and human review points.

Overreliance on fluent outputs. Generative AI can produce confident language that appears accurate but is unsupported, incomplete, or wrong. Guardrail: require source verification, human review, and clear labeling of outputs as drafts or decision support when appropriate.

Privacy and confidentiality failures. Sensitive public health information can be exposed through inappropriate prompts, file uploads, model training, vendor retention, access controls, or shared outputs. Guardrail: require data classification, minimum necessary use, approved environments, access controls, and privacy review for sensitive data.

Equity and representation harms. AI can amplify bias in source data, miss smaller populations, flatten community context, or produce recommendations that increase burdens for groups already experiencing inequities. Guardrail: require equity review, subgroup monitoring when appropriate, missing voices analysis, and community-informed interpretation for consequential uses.

Unclear accountability. Staff may not know who approved an AI use, who is responsible for errors, or who can pause the workflow. Guardrail: document owners, decision rights, approval conditions, escalation contacts, and pause/retirement criteria.

Governance overload. Agencies can create review processes so burdensome that staff avoid them or innovation stalls. Guardrail: use risk tiers, standard intake forms, clear timelines, and lightweight pathways for low-risk uses while preserving stronger review for consequential uses.

Application to AI-Supported Workflows

In AI-supported workflows, governance should be visible at the points where work changes. Before use, staff should know whether the tool is approved, what data may be used, what the intended purpose is, and whether the use requires intake or review. During use, staff should know how to review outputs, document decisions, protect sensitive information, and report problems. After use, the agency should be able to audit what was approved, what evidence supported the decision, what outputs were used, and whether the workflow remains safe and useful.

For program staff, governance means connecting AI use to program goals, service delivery, community impact, and workflow reality. Program staff should be able to explain why AI is being considered, who will use it, what decision or task it supports, what could go wrong, and how staff will know whether it is helping.

For technical staff, governance means making the system understandable, secure, testable, and monitorable. Technical staff should document data flows, access controls, model or tool behavior, integration points, logging, validation evidence, change management, and incident response expectations.

For supervisors and leaders, governance means ensuring that staff have clear rules, enough training, a safe way to raise concerns, and a realistic workload for human review. Leaders should avoid approving AI tools without resources for training, monitoring, documentation, and improvement.

For all staff, governance means pausing when an AI use involves sensitive data, public-facing content, decision support, automated prioritization, community impact, legal obligations, or uncertainty about agency policy. In those situations, escalation is not a barrier. It is part of responsible public health practice.

Reflection Questions

Where might AI already be entering your work through drafting tools, analytics tools, vendor platforms, search tools, translation tools, dashboards, code assistance, or workflow automation?

What AI uses in your program area would be low risk, and what uses would require formal review before moving forward?

Who in your agency would need to review an AI use that involves sensitive data, public-facing communication, resource prioritization, or surveillance action?

What information would you need before you could explain an AI-supported workflow to a supervisor, partner, community member, or governance body?

What would make staff comfortable raising concerns about AI outputs or unapproved use without fear of blame?

Practical Exercise

Select one real or realistic public health workflow where AI might be used. Examples include summarizing meeting notes, drafting public health messages, classifying environmental health complaints, reviewing grant applications, supporting case investigation, translating outreach materials, generating code for a dashboard, prioritizing outreach lists, or summarizing surveillance trends.

Complete a one-page introductory AI governance intake and escalation note. The note should include: the public health purpose; the staff or partners who would use the tool; the data involved; whether the data are sensitive or identifiable; the AI function; the expected benefit; who will review the output; what decision or action the output could influence; likely risks; required reviewers; escalation triggers; and whether the use should be low, moderate, or high risk.

Practical exercise example: A chronic disease program wants to use an AI tool to summarize open-ended survey responses from community members about barriers to diabetes prevention classes. The governance note identifies the purpose as reducing manual review burden while preserving community meaning. The data include open-ended comments, ZIP codes, race/ethnicity, age group, and some comments that mention disability, transportation, immigration concerns, and distrust of government. The AI function is summarization and theme extraction. The expected benefit is faster identification of barriers to participation. The risks include privacy, small-community identifiability, loss of minority viewpoints, inaccurate summaries, and overconfidence in themes. The note recommends using only an approved secure tool, removing direct identifiers, keeping the original comments available for human review, checking themes across demographic groups, involving the community engagement lead, and requiring communications review before any public summary is released. The risk tier is moderate because the output may shape program planning and public communication but will not directly determine individual eligibility or services.

After completing the note, identify one governance question that remains unresolved. For example: Can this data be entered into the proposed tool under current data-use agreements? Who must approve public language based on AI-assisted summaries? What threshold would make the project high risk? What evidence would show that the AI summary did not erase smaller community concerns?

Expected Artifact or Evidence

Learners should produce an introductory AI governance intake and escalation note for one public health workflow. The artifact should include a purpose statement, data description, AI function, likely users, expected benefit, risk tier, human review plan, required reviewers, escalation triggers, and at least one unresolved governance question. The artifact can be used as a training record, a draft intake form, or a starting point for a real governance discussion.

References and Resources for Additional Information

National Institute of Standards and Technology. Artificial Intelligence Risk Management Framework (AI RMF 1.0). 2023. <https://www.nist.gov/itl/ai-risk-management-framework>

National Institute of Standards and Technology. AI RMF Playbook. https://airc.nist.gov/AI_RM_F_Knowledge_Base/Playbook

National Institute of Standards and Technology. Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile, NIST AI 600-1. 2024. <https://doi.org/10.6028/NIST.AI.600-1>

Centers for Disease Control and Prevention. Public Health Data Strategy. <https://www.cdc.gov/public-health-data-strategy/php/about/index.html>

World Health Organization. Ethics and governance of artificial intelligence for health. 2021.
<https://www.who.int/publications/i/item/9789240029200>

World Health Organization. Regulatory considerations on artificial intelligence for health. 2023.
<https://www.who.int/publications/i/item/9789240078871>

Knowledge Check

Choose the best answer for each question.

1. Which statement best describes AI governance in a public health agency?

- A. A technical process used only by data scientists to select models.
- B. A set of roles, policies, review pathways, decision rights, documentation practices, and oversight activities for responsible AI use.
- C. A one-time approval given when a vendor contract is signed.
- D. A communications plan used only when AI errors become public.

Correct answer: B

2. Which AI use should most clearly be escalated for formal review before implementation?

- A. Using an approved tool to reformat an internal meeting agenda with no sensitive content.
- B. Using AI to prioritize field inspections based on complaint narratives and neighborhood characteristics.
- C. Using spell check on a public flyer.
- D. Using a calculator to total attendance at a training.

Correct answer: B

3. What is the most important reason to define human review in an AI-supported workflow?

- A. To make the workflow slower so AI is used less often.
- B. To ensure a person has a meaningful role in checking outputs, interpreting context, documenting rationale, and accepting responsibility for decisions.
- C. To allow staff to approve outputs without reading them.
- D. To shift all responsibility from the agency to the software vendor.

Correct answer: B

4. A staff member wants to paste identifiable case notes into a public AI chatbot to create a quick summary. What should happen first?

- A. Proceed because summarization is an internal productivity task.
- B. Remove the staff member from the project.
- C. Pause and follow agency rules for data classification, approved tools, privacy/security review, and escalation.
- D. Proceed if the summary will not be published.

Correct answer: C

5. Which item belongs in a basic AI governance intake note?

- A. Only the name of the AI tool and the vendor website.

- B. The public health purpose, data involved, AI function, human review plan, risks, reviewers, escalation triggers, and unresolved questions.
- C. Only the expected time savings.
- D. A statement that AI outputs are always advisory and therefore need no review.

Correct answer: B

References and Resources

National Institute of Standards and Technology. Artificial Intelligence Risk Management Framework (AI RMF 1.0). 2023
<https://www.nist.gov/itl/ai-risk-management-framework>

National Institute of Standards and Technology. AI RMF Playbook
https://airc.nist.gov/AI_RMF_Knowledge_Base/Playbook

National Institute of Standards and Technology. Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile, NIST AI 600-1. 2024
<https://doi.org/10.6028/NIST.AI.600-1>

Centers for Disease Control and Prevention. Public Health Data Strategy
<https://www.cdc.gov/public-health-data-strategy/php/about/index.html>

World Health Organization. Ethics and governance of artificial intelligence for health. 2021
<https://www.who.int/publications/i/item/9789240029200>

World Health Organization. Regulatory considerations on artificial intelligence for health. 2023
<https://www.who.int/publications/i/item/9789240078871>

GOV 110

Respecting Tribal Sovereignty and Indigenous Data Governance in Public Health AI Projects

Course ID	GOV 110
Track	Governance and Security
Level	Introductory
Estimated time	Self-paced
Audience	Public health learners

Learning Objectives

- Explain why Tribal sovereignty must be considered before starting AI projects that involve Tribal Nations, American Indian and Alaska Native participant data, or health data that may be connected to Tribal citizens, communities, lands, services, or priorities.
- Describe Indigenous Data Sovereignty, Indigenous Data Governance, Digital Sovereignty, and the CARE Principles in plain language, and explain how these concepts apply to AI use case intake, data management, model development, evaluation, dissemination, and lifecycle monitoring.
- Identify governance questions that should be addressed before using Tribal data or AI tools in public health projects, including questions about Tribal consultation, Tribal laws and policies, Tribal IRBs or research review processes, data sharing agreements, publication review, access controls, and permitted or prohibited uses.
- Distinguish cultural competence from cultural humility and explain why Tribal-specific consultation, relationship building, and respect for government-to-government authority are required rather than relying on generic cultural awareness or one-size-fits-all engagement practices.
- Identify AI project risks involving inferred Tribal affiliation, inappropriate geographic inference, external platform exposure, model outputs that stigmatize communities, dissemination that bypasses Tribal review, and secondary uses that were not approved by Tribal partners.
- Connect sovereignty protections to practical project artifacts, including Tribal data sharing agreements, Data Management and Sharing plans, review memos, community benefit statements, AI system documentation, dissemination plans, and monitoring or sunset criteria.
- Draft an initial Tribal sovereignty and Indigenous data governance screening memo for a proposed AI-supported public health workflow that identifies whether the use case should proceed, pause for consultation, be revised, or move to formal Tribal, legal, privacy, security, equity, and governance review.

Definitions

Tribal sovereignty	The inherent authority of Tribal Nations to govern themselves and protect the health, safety, welfare, and future of their citizens.
Indigenous Data Sovereignty	The right of Indigenous Peoples and Nations to govern data about them, their citizens, lands, resources, knowledge, and communities.
Indigenous Data Governance	The policies, agreements, relationships, roles, and oversight mechanisms that put Indigenous Data Sovereignty into practice.

CARE Principles	A framework for Indigenous Data Governance organized around Collective Benefit, Authority to Control, Responsibility, and Ethics.
Digital Sovereignty	Control over digital environments, platforms, infrastructure, and technical systems that store, process, move, or expose data.
Cultural humility	An ongoing practice of learning, accountability, and respect that recognizes power imbalances and Tribal-specific protocols.

Why This Topic Matters for Public Health AI

AI projects can magnify the consequences of weak data governance because AI systems depend on data extraction, transformation, linkage, model training, inference, and reuse. When health data involve Tribal Nations or AI/AN participants, those consequences can include individual privacy harms, community stigmatization, misclassification, loss of control over data, misinterpretation of culturally situated information, or use of data for purposes that were never agreed to. Public health agencies may intend to improve services, reduce disparities, or strengthen emergency response, but good intent does not replace Tribal authority to govern data and research involving Tribal citizens, communities, and jurisdictions. Respecting sovereignty at the start of an AI project prevents downstream harm that may be difficult or impossible to repair.

This topic also matters because many AI workflows blur the boundary between research, operations, analytics, and communications. A project that begins as an internal dashboard, chatbot, transcription workflow, or model evaluation may later influence public messaging, resource allocation, surveillance interpretation, or program decisions. If Tribal partners are consulted only after design choices are made, they may have little opportunity to shape the data sources, consent terms, model outputs, storage location, access controls, or review procedures. Meaningful partnership requires engagement while the project can still change.

Public health AI work also occurs within a history of research abuse, data misuse, underinvestment, and misrepresentation affecting AI/AN communities. NIH guidance on responsible management and sharing of AI/AN participant data emphasizes that biomedical research involving Tribal Nations must be predicated on respect for Tribal sovereignty and awareness that failures to honor sovereignty have caused stigmatization and other harms. The CARE Principles similarly emphasize that data governance must address collective benefit, authority to control, responsibility, and ethics rather than focusing only on technical openness or reuse. For AI projects, these principles require teams to ask who benefits, who controls the data, who is accountable, and whether the work aligns with Tribal laws, policies, preferences, and values.

Definitions / Core Concepts

Tribal sovereignty means the inherent authority of Tribal Nations to govern themselves and protect the health, safety, welfare, and future of their citizens. In public health AI work, this means a Tribal Nation may have its own laws, research codes, review bodies, data governance requirements, consent expectations, publication review processes, and conditions for data access or reuse. Public health staff should not treat Tribal input as optional stakeholder feedback when Tribal authority is implicated. The appropriate posture is government-to-government respect, early engagement, and formal agreement.

Indigenous Data Sovereignty refers to the right of Indigenous Peoples and Nations to govern the collection, ownership, interpretation, management, access, sharing, and application of data about them, their citizens, lands, resources, knowledge, and communities. Health data can be Indigenous data even when it is held by a public health agency, university, vendor, researcher, hospital, or federal program. AI projects raise special concerns because data may be reused for model training, linked to other data sources, summarized by third-party tools, or used to generate outputs that

shape decisions. Sovereignty requires that these uses be governed by the Indigenous Peoples and Nations whose data are involved.

Indigenous Data Governance is the set of policies, relationships, processes, agreements, roles, and oversight mechanisms that put Indigenous Data Sovereignty into practice. Governance may include Tribal IRB review, research review committees, Tribal council resolutions, data sharing agreements, publication review, data storage requirements, access controls, benefit sharing, and requirements for returning results to the community. It also includes decisions about whether a project should proceed at all. In AI projects, governance must cover the full lifecycle from use case intake through data preparation, modeling, validation, deployment, monitoring, dissemination, and retirement.

The CARE Principles for Indigenous Data Governance provide a practical framework for aligning data work with Indigenous rights and interests. CARE stands for Collective Benefit, Authority to Control, Responsibility, and Ethics. These principles complement FAIR data practices by adding attention to people, purpose, power, historical context, and self-determination. For AI projects, CARE can be used as a review lens for deciding whether a proposed data use benefits Tribal communities, respects Tribal authority, assigns responsibilities, and avoids extractive or harmful uses.

Digital Sovereignty extends data sovereignty into the digital environments, platforms, infrastructure, and technical systems that store, move, process, analyze, and expose data. For Tribal Nations, this may include control over where data are stored, who administers systems, what cloud or vendor environments are approved, whether data leave a Tribal or trusted environment, and how AI tools access or retain prompts, files, metadata, or outputs. Digital sovereignty matters because an AI project can respect a data sharing agreement on paper while still undermining Tribal control through external platforms, opaque model behavior, or uncontrolled secondary use. Technical architecture is therefore part of sovereignty, not a separate IT detail.

Cultural competence is often described as knowledge about cultural histories, values, practices, and communication norms, but it is not enough by itself for AI projects with Tribal Nations. A more appropriate operating standard is cultural humility, which requires ongoing learning, recognition of power imbalances, willingness to be corrected, and respect for each Tribal Nation's specific protocols. Staff should avoid pan-Indigenous assumptions because Tribal Nations differ in governance structures, histories, languages, legal requirements, data priorities, and relationships with public health agencies. Cultural humility should be visible in timelines, budgets, consultation processes, meeting practices, authorship decisions, dissemination plans, and how the project responds when Tribal partners raise concerns.

Public Health Example

A state health department wants to develop an AI-assisted analytics workflow to summarize emergency department syndromic surveillance notes and identify patterns related to respiratory illness. The initial technical design includes free-text chief complaint data, facility location, ZIP code, race and ethnicity fields, and model-generated summaries by county. Several counties include Tribal lands, and some small-area summaries could indirectly identify Tribal communities or lead staff to infer Tribal affiliation from residence, facility use, or geography. The project team initially views the work as routine surveillance improvement, but the intake review identifies that Tribal sovereignty and Indigenous data governance must be addressed before development proceeds.

A sovereignty-centered approach would pause the technical build and begin with appropriate Tribal consultation and partnership discussions. The team would identify which Tribal Nations may be affected, whether Tribal public health authorities already receive or govern relevant data, whether Tribal laws or review processes apply, and what limitations should be placed on geographic analysis, small-number reporting, public dissemination, and AI-generated interpretation. The project would not use external AI platforms to process identifiable or potentially sensitive data unless the environment, data use, retention terms, and review pathway were approved. The team would also document whether the project creates community benefit defined by Tribal partners, not only operational benefit for the state agency.

The final project might look different from the original proposal. Tribal partners could request a Tribal-specific data sharing agreement, a joint review group, suppression rules for small numbers, limits on geospatial inference, local review before public summaries, a requirement to return results in a useful format, or a decision that some data should not be used for AI modeling. The project may also identify capacity-building opportunities, such as supporting Tribal epidemiology staff, providing access to relevant outputs, or co-developing monitoring measures. The key lesson is that sovereignty review does not merely add a compliance step; it can reshape the purpose, design, governance, and success measures of the AI project.

Technical, Leadership, and Operational Deep Dive

AI project teams should begin by asking whether the proposed use involves Tribal Nations, Tribal citizens, AI/AN participants, Tribal lands, Tribal service settings, Tribal public health authorities, or data elements that could be used to infer Tribal affiliation. This question should be asked during use case intake, not after a data set has been assembled. The answer may not be obvious because Tribal relevance can appear through geography, facility relationships, program participation, small-area reporting, environmental exposure data, biospecimens, genetic information, language data, or community narratives. If the answer is uncertain, the team should treat the project as requiring additional review rather than assuming no Tribal implications exist.

Early consultation should occur before AI development, procurement, vendor demonstrations, pilot testing, or publication planning. Consultation is not the same as sending a completed plan for comment. It should allow Tribal Nations to influence the project purpose, data scope, roles, benefits, risks, consent expectations, storage location, review process, and dissemination strategy. Timelines and budgets should recognize that government-to-government consultation, community review, and relationship building require time and resources.

Formal agreements are essential because AI projects create many opportunities for data reuse and secondary interpretation. Agreements may need to address data ownership or stewardship, permitted uses, prohibited uses, data retention, destruction or return, access rights, cloud or vendor environments, model training restrictions, publication review, benefit sharing, incident response, and conditions for future use. Where NIH-funded research is involved, the Data Management and Sharing Plan should incorporate AI/AN data management and sharing practices and Tribal preferences. For operational public health projects outside NIH research, similar principles should still guide data sharing agreements, memoranda of understanding, and governance records.

AI teams should not infer Tribal affiliation from residential location, ZIP code, facility use, names, language, or other proxy variables unless that inference has been explicitly approved and is necessary for a mutually agreed public health purpose. Proxy inference can create privacy risk, misclassification, and community harm. It can also undermine Tribal authority if outside organizations create AI-derived labels about Tribal identity or community membership without consent. If affiliation or Tribal relevance is needed, the project should address how that information is defined, who controls it, whether it should be collected, and whether the Tribal Nation approves the method.

Data management controls should be designed before data are extracted or transformed for AI use. Teams should classify data sensitivity, separate identifiable data from analytic data when appropriate, minimize variables, restrict access, log use, and define approved environments. External AI tools, automated transcription services, large language models, coding assistants, and cloud platforms should not receive Tribal data, identifiable information, sensitive community information, or non-public operational details unless the tool and environment have been reviewed and approved. Contract terms should address retention, vendor access, model training, subcontractors, data location, deletion, incident notification, and audit rights.

Model development and evaluation should include sovereignty-aware validation questions. The team should ask whether training data represent the community appropriately, whether small sample sizes make results unreliable, whether model outputs could stigmatize a Tribal Nation, and whether performance should be reported separately or

suppressed. The team should also ask who has authority to interpret findings and what contextual knowledge is needed to avoid harmful conclusions. Evaluation should include human review by people with appropriate community, program, epidemiologic, technical, and governance expertise.

Communication and dissemination require special care because AI outputs can be polished, confident, and wrong. Public-facing summaries, dashboards, maps, chatbots, and reports may create harm if they identify small communities, imply causation without evidence, frame disparities as community deficits, or omit historical and structural context. Tribal partners may require pre-review of reports, presentations, manuscripts, dashboard text, public advisories, or model documentation. Teams should document dissemination approvals and avoid treating public release as a default endpoint.

Cultural humility should be built into the project operating model. Staff should learn the history, governance structure, communication protocols, and preferred engagement processes of the Tribal Nation they are working with, while recognizing that this knowledge does not make them authorities on the community. Meetings should be scheduled and facilitated in ways that respect Tribal leadership, community priorities, and local decision-making processes. Compensation, authorship, data access, capacity building, and benefit sharing should reflect the value of Tribal expertise and partnership.

Risks, Failure Modes, and Guardrails

A major risk is treating Tribal health data as if it were ordinary administrative or research data controlled solely by the institution that currently holds it. This can lead to AI development without Tribal consultation, data sharing without appropriate authority, or model outputs that affect Tribal communities without Tribal review. The failure mode is often procedural rather than malicious: staff follow routine analytics practices without recognizing that sovereignty changes the decision pathway. The guardrail is to require Tribal sovereignty screening during AI use case intake and to route potentially relevant projects for legal, privacy, governance, and Tribal consultation review before data are used.

Another risk is using external AI platforms for transcription, summarization, coding, translation, or drafting with sensitive Tribal data or community information. These tools may retain prompts, store uploaded files, expose information to vendors or subcontractors, or use content in ways that are not transparent to the user. Even de-identified text can contain place names, small population details, cultural references, or contextual clues that create re-identification or group harm. The guardrail is to prohibit use of non-approved external AI tools with Tribal data and to require approved environments, contract review, and data minimization before any AI processing occurs.

A third risk is allowing model outputs to create or reinforce harmful narratives about Tribal communities. Predictive models, cluster analyses, risk scores, geospatial maps, or AI-generated summaries may overemphasize deficits, ignore structural determinants, misinterpret small numbers, or produce misleading comparisons. These outputs can affect public trust, funding decisions, program priorities, or media narratives. The guardrail is to require sovereignty-aware review of model purpose, variables, outputs, interpretation, dissemination, and monitoring before results are shared or used operationally.

A fourth risk is considering equity, accessibility, and community benefit too late. If the project is already funded, contracted, designed, or near deployment before Tribal partners are engaged, meaningful changes may be difficult. This can create extractive relationships in which data and expertise flow outward while benefits, authority, and infrastructure do not flow back to the community. The guardrail is to define community benefit, capacity building, shared decision-making, and return-of-results expectations at the beginning of the project.

A fifth risk is assuming that one Tribal engagement process applies to all Tribal Nations. Each Tribal Nation has its own government, history, legal requirements, research review expectations, cultural protocols, and data priorities. A process accepted by one Tribal Nation cannot be generalized to another without permission. The guardrail is to identify each

affected Tribal Nation and confirm the appropriate consultation, approval, and documentation pathway for each one.

A sixth risk is failing to maintain accountability after deployment. AI systems may change as data pipelines, vendors, model versions, workflows, community conditions, or public health priorities change. Without lifecycle governance, a use that was conditionally approved may drift beyond the original agreement. The guardrail is to maintain an inventory of AI uses involving Tribal data, monitor compliance with approval conditions, review incidents, revisit agreements when scope changes, and retire or pause systems when they no longer align with Tribal preferences or public health benefit.

Application to AI-Supported Workflows

For AI use case intake, teams should add screening questions that identify whether the proposed project involves Tribal Nations, AI/AN participant data, Tribal lands, Tribal service settings, Tribal public health authorities, small-area geography, biospecimens, genetic data, environmental exposure data, or other information that could affect Tribal communities. The intake form should require the proposer to identify known Tribal partners, anticipated benefits, possible harms, data sources, external tools, and whether consultation has occurred. A project should not advance to procurement, model development, pilot testing, or public release until the sovereignty screening has been resolved. This makes Tribal sovereignty a front-end design requirement rather than a late-stage approval issue.

For data preparation and model development, teams should document permitted data uses and technical restrictions before creating analytic files. This includes rules for variable selection, de-identification, linkage, geospatial analysis, small number suppression, feature engineering, model training, validation, storage, access, and deletion. Teams should also decide whether data can be used in model training, whether outputs can be reused in future projects, and whether a model trained with Tribal data can be transferred to other jurisdictions or use cases. These decisions should be made through formal agreement, not by default technical convenience.

For deployment and monitoring, teams should define who reviews outputs, who receives results, who can change the system, and who can decide that the system should pause or stop. Monitoring should include accuracy, fairness, usability, community impact, privacy incidents, scope creep, and compliance with Tribal agreements. If the AI tool supports public communication, resource prioritization, surveillance interpretation, or program decisions, Tribal partners should have a clear pathway to question or challenge outputs. The system should be designed to preserve Tribal authority throughout the lifecycle, not only during initial approval.

Reflection Questions

Before using AI with health data that may involve Tribal Nations or AI/AN participants, learners should pause and consider whether they know who has authority to approve the use. They should ask whether the project has identified affected Tribal Nations, relevant Tribal laws or review processes, and the appropriate form of consultation or partnership. They should also consider whether the proposed AI use creates community benefit as defined by Tribal partners, rather than only efficiency for an external agency. These questions help teams recognize that sovereignty is a design condition, not a courtesy step.

Learners should also reflect on how the project could create harm even if no individual is directly identified. Could the AI output stigmatize a community, reveal small-area patterns, infer Tribal affiliation, misinterpret cultural context, or support decisions that affect services or resources? Could an external AI tool retain or reuse information that the project team does not have authority to share? These questions should be answered before development begins, when the project can still be changed.

Practical Exercise

Select a proposed AI-supported public health project that could involve Tribal Nations, AI/AN participant data, Tribal lands, Tribal service settings, or data that could indirectly identify or affect Tribal communities. The project may be a surveillance summary tool, predictive model, automated transcription workflow, chatbot, public communication assistant, environmental health mapping project, data linkage project, or research analysis. Complete a Tribal sovereignty and Indigenous data governance screening memo for the project. The memo should identify the public health purpose, data sources, affected communities, potential Tribal authority concerns, consultation needs, data management safeguards, external platform risks, review bodies, and recommended next steps.

Example response: A state agency proposes to use an AI tool to summarize open-text case investigation notes for emerging respiratory illness trends. The initial data include case narratives, county, ZIP code, facility, race and ethnicity, and outbreak location fields, and several records may involve Tribal service areas or small communities near Tribal lands. The screening memo identifies that the project should not send free-text notes to an external AI platform, should not infer Tribal affiliation from geography, and should not produce small-area summaries until Tribal consultation and legal review occur. The recommended next step is to meet with potentially affected Tribal public health authorities to determine whether the project has a mutually agreed purpose, what data may be used, what outputs are useful, what findings require pre-review, and what safeguards should be documented in a data sharing agreement.

Expected Artifact or Evidence

The expected artifact is a Tribal Sovereignty and Indigenous Data Governance Screening Memo for one AI-supported public health use case. The memo should be complete enough for a supervisor, AI governance body, Tribal liaison, legal counsel, privacy officer, security lead, equity reviewer, or Tribal partner to understand why the project may implicate Tribal sovereignty. It should distinguish known facts from assumptions and identify what cannot be decided until appropriate Tribal consultation occurs. The memo should not include protected health information, personally identifiable information, or confidential Tribal information.

A strong memo includes the proposed AI use, public health purpose, data sources, affected Tribal Nations or uncertainty about Tribal relevance, AI tools or vendors under consideration, data storage and processing environment, anticipated outputs, potential benefits, potential harms, and recommended review pathway. It should address CARE Principles, data management expectations, cultural humility, consultation needs, and whether formal agreements are required before the project proceeds. The artifact may be used as an intake attachment, governance briefing, or planning document, but it does not replace Tribal approval or formal legal review.

Jurisdiction and Agency Policy Note

This national-level module is intended for training and planning purposes. It does not replace Tribal law, Tribal policy, Tribal consultation, Tribal IRB review, legal counsel, privacy review, security review, civil rights review, procurement review, or agency policy review. Requirements may vary by Tribal Nation, jurisdiction, funding source, data sharing agreement, research status, operational context, and data type. Learners should confirm applicable requirements before approving, procuring, deploying, sharing, or publicly communicating about AI-supported workflows involving Tribal Nations or AI/AN participant data.

References and Resources for Additional Information

Global Indigenous Data Alliance. CARE Principles for Indigenous Data Governance. <https://www.gida-global.org/care>

National Institutes of Health. NOT-OD-22-214: Supplemental Information to the NIH Policy for Data Management and Sharing: Responsible Management and Sharing of American Indian/Alaska Native Participant Data. <https://grants.nih.gov/grants/guide/notice-files/NOT-OD-22-214.html>

National Institutes of Health Tribal Health Research Office. Tribal Health Research Office resources.
<https://dpcpsi.nih.gov/thro>

United States Indigenous Data Sovereignty Network. Resources and information on Indigenous Data Sovereignty.
<https://usindigenousdata.org/>

All of Us Research Program. Tribal engagement and data use resources. <https://www.researchallofus.org/>

HHS and CDC Tribal consultation resources. Use relevant agency consultation policies and Tribal liaison guidance before starting projects involving Tribal Nations or AI/AN data.

Brookings. Digital sovereignty and Indigenous Nations resources. <https://www.brookings.edu/>

Knowledge Check

1. Why should Tribal sovereignty be considered before an AI project with health data begins?

- A. Because Tribal Nations are sovereign governments with authority over research and data practices involving their citizens, jurisdictions, and data.
- B. Because AI tools always require federal approval.
- C. Because Tribal data can always be treated as open data.
- D. Because sovereignty applies only after publication.

Correct answer: A

2. Which statement best reflects Indigenous Data Sovereignty?

- A. Data should be shared as widely as possible whenever it supports science.
- B. Indigenous Peoples and Nations have the right to govern how data about them are collected, interpreted, stored, shared, and used.
- C. De-identification removes all Tribal data governance concerns.
- D. Data sovereignty applies only to paper records.

Correct answer: B

3. Which AI practice should be avoided unless explicitly approved through the appropriate governance pathway?

- A. Using a standard intake form to identify data sensitivity.
- B. Consulting Tribal partners before model design.
- C. Uploading sensitive Tribal health information or community narratives into an external AI platform.
- D. Documenting data sharing limitations in an agreement.

Correct answer: C

4. What is the strongest reason to involve Tribal partners early in the AI lifecycle?

- A. Early engagement allows Tribal Nations to shape the project purpose, data scope, safeguards, benefits, and dissemination before key decisions are locked in.
- B. Early engagement shortens all review timelines.
- C. Early engagement eliminates the need for written agreements.
- D. Early engagement is needed only for public-facing tools.

Correct answer: A

5. What should a sovereignty-aware AI project do when Tribal relevance is uncertain?

- A. Proceed unless the data set explicitly labels Tribal affiliation.
- B. Remove all race and ethnicity variables and continue.
- C. Treat the uncertainty as a reason for additional review and determine whether consultation, legal review, or Tribal governance processes apply.
- D. Wait until the model is finished and ask for feedback.

Correct answer: C

References and Resources

Global Indigenous Data Alliance. CARE Principles for Indigenous Data Governance
<https://www.gida-global.org/care>

National Institutes of Health. NOT-OD-22-214: Supplemental Information to the NIH Policy for Data Management and Sharing: Responsible Management and Sharing of American Indian/Alaska Native Participant Data
<https://grants.nih.gov/grants/guide/notice-files/NOT-OD-22-214.html>

National Institutes of Health Tribal Health Research Office. Tribal Health Research Office resources
<https://dpcpsi.nih.gov/thro>

United States Indigenous Data Sovereignty Network. Resources and information on Indigenous Data Sovereignty
<https://usindigenoudata.org/>

All of Us Research Program. Tribal engagement and data use resources
<https://www.researchallofus.org/>

HHS and CDC Tribal consultation resources. Use relevant agency consultation policies and Tribal liaison guidance before starting projects involving Tribal Nations or AI/AN data.

Brookings. Digital sovereignty and Indigenous Nations resources
<https://www.brookings.edu/>

GOV 300

AI Security, Prompt Injection, and Data Leakage

Course ID	GOV 300
Track	Governance and Security Track
Level	intermediate/practitioner
Estimated time	50
Audience	Governance committee members, Legal/privacy staff, Security staff, Compliance staff, Program owners

Learning Objectives

- Define AI security in public health operational terms.
- Identify prompt injection, data leakage, sensitive information disclosure, insecure output handling, excessive agency, and supply chain risks.
- Explain why RAG, plugins, APIs, and agentic workflows increase the need for security controls.
- Apply least privilege, approved-tool use, access controls, logging, testing, and incident response to AI workflows.
- Develop security questions for AI procurement, implementation, and monitoring.
- Produce a secure AI workflow checklist.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Why This Topic Matters for Public Health AI

AI security matters because public health AI tools may process sensitive records, internal protocols, draft communications, credentials, contracts, emergency plans, and partner information. The risks increase when AI tools are connected to email, calendars, case systems, data warehouses, document repositories, APIs, or messaging platforms. A security weakness can become a privacy breach, operational disruption, public trust problem, or legal issue.

Prompt injection is particularly important for public health because staff may ask AI tools to process external documents, emails, web pages, public comments, complaints, or uploaded files. Those materials may contain instructions designed to manipulate the AI system. If the AI tool has access to sensitive data or actions, prompt injection can become a pathway to data exposure or unauthorized activity.

Security controls should be built into governance and procurement. Agencies should require information about data handling, retention, encryption, access controls, audit logs, model training use, third-party services, incident notification, vulnerability management, and secure configuration. Security review should happen before deployment and continue after launch.

Public Health Example

A public health program may use an AI assistant to summarize uploaded partner reports. One report could contain hidden or visible instructions telling the AI to ignore prior rules and reveal confidential information. If the assistant treats the report content as instructions rather than data, it may produce unsafe outputs or attempt unauthorized actions.

A secure workflow would use an approved tool, restrict the assistant's access, treat uploaded content as untrusted, disable unnecessary tool actions, log activity, and require human review before any output is shared. Staff should be trained to recognize suspicious instructions embedded in documents or outputs.

Technical Deep Dive

AI security begins with use-case classification. A low-risk training exercise using public documents requires different controls than a tool connected to case records or document repositories. Agencies should classify AI uses by data sensitivity, system access, autonomy, public-facing impact, and consequence of error.

Access control is central. AI tools should not have broad access to systems simply because broad access is convenient. Each integration should specify what the tool can read, write, update, send, or delete. Read-only access is safer than write access, but read-only access can still expose sensitive information if permissions are weak.

Prompt injection defenses should include system-level protections, content handling, tool restrictions, and user training. External content should be treated as data, not instruction. Tools should not follow directions found inside documents, emails, or web pages unless those directions come from an authorized user in the correct workflow. Agentic systems should require additional approval before taking consequential actions.

Output handling matters because AI outputs may be copied into reports, emails, code, dashboards, or scripts. Insecure output handling can lead to misinformation, broken code, malicious links, or inappropriate disclosure. Outputs should be validated before downstream use, especially when they affect systems, public communications, or records.

Incident response should include AI-specific triggers. Examples include unauthorized data exposure, unexpected tool use, prompt injection attempt, repeated unsafe output, model misuse, credential exposure, or vendor security incident. Logs should support investigation of prompts, inputs, retrieved sources, outputs, tool calls, users, and actions.

Risks, Failure Modes, and Guardrails

Prompt injection. Untrusted content may override instructions. Guardrails include treating external content as data, restricting tools, and testing with adversarial examples.

Sensitive information disclosure. AI may reveal information outside user permissions. Guardrails include access controls, data minimization, redaction, and audit logs.

Excessive agency. Tools may have more capability than needed. Guardrails include least privilege, approval gates, and prohibited-action lists.

Supply chain risk. AI tools may depend on vendors, plugins, models, or libraries with vulnerabilities. Guardrails include vendor review, security documentation, patching, and contract requirements.

Application to AI-Supported Workflows

AI security applies to generative AI, RAG, NLP, APIs, and agentic workflows. RAG systems need document-level permissions and protections against unauthorized retrieval. Agentic systems need tool restrictions and action approval. Prompt templates should include data boundaries, and technical implementations should enforce those boundaries rather than relying on user behavior alone.

Security should be designed as part of the workflow. Public health users should know which tools are approved, what data may be used, what outputs require review, and how to report suspicious or unsafe behavior.

Reflection Questions

What sensitive data or systems would the AI workflow touch?

Could the AI tool access more information or actions than necessary?

What untrusted content might the tool process?

How will prompt injection attempts or unsafe outputs be detected?

What logs and incident response steps are required?

Practical Exercise

Create a secure AI workflow checklist for one public health use case. Include data sensitivity, approved tools, access controls, least-privilege review, prompt injection risk, output handling, logging, vendor security questions, incident triggers, and escalation process.

The checklist should identify which controls must be in place before pilot testing and which controls must be monitored after launch.

Expected Artifact or Evidence

Secure AI workflow checklist

Vendor and implementation security questions

Prompt injection risk notes

AI incident trigger and escalation plan

References and Resources for Additional Information

OWASP Top 10 for Large Language Model Applications:

<https://owasp.org/www-project-top-10-for-large-language-model-applications/>

NIST Artificial Intelligence Risk Management Framework: <https://www.nist.gov/itl/ai-risk-management-framework>

NIST AI RMF Generative AI Profile:

<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>

NIST Privacy Framework: <https://www.nist.gov/privacy-framework>

Operations, Privacy, Procurement, and Value Measurement: Modules 21-30

Source file: Technical_Readiness_for_Public_Health_AI_Training_Modules_21-30(1).docx

Technical Readiness for Public Health AI Training Modules 21-30

This document completes the technical training sequence by adding Modules 21-30. The modules focus on privacy-preserving AI, synthetic data, secure deployment, AI operations, human factors, business analysis, procurement requirements, documentation, data governance, and value measurement. They are designed for public health informaticians, epidemiologists, technologists, analysts, product owners, procurement reviewers, and governance participants.

Each module follows the same lesson structure used in the prior batches: Training Overview, Learning Objectives, Why This Topic Matters, Core Concepts, Public Health Example, Technical Deep Dive, Risks and Guardrails, Application to AI-Supported Workflows, Reflection Questions, Practical Exercise, Expected Evidence, Knowledge Check, and Lesson Resources.

Knowledge Check

1. What is prompt injection?

- A. A user-interface color setting.
- B. A method for calibrating risk scores.
- C. Untrusted content attempting to override AI instructions or behavior.
- D. A type of geocoding.

Correct answer: C

2. What does least privilege mean?

- A. Giving every tool full access for convenience.
- B. Providing only the minimum access needed for an approved purpose.
- C. Removing all audit logs.
- D. Publishing all outputs automatically.

Correct answer: B

3. Why does agentic AI increase security risk?

- A. It never connects to tools.
- B. It cannot process sensitive data.
- C. It may use tools or take actions across systems.
- D. It eliminates prompt injection.

Correct answer: C

4. Which control helps investigate AI-related incidents?

- A. Audit logs that record prompts, inputs, outputs, users, tool calls, and actions.
- B. Unrestricted access.
- C. Removing source references.
- D. Allowing external content to issue commands.

Correct answer: A

5. What is strong completion evidence for this module?

- A. A secure AI workflow checklist with data, access, prompt injection, output handling, logs, vendor questions, and escalation.
- B. A general AI definition.
- C. A list of approved colors.
- D. A screenshot of a chatbot only.

Correct answer: A

References and Resources

OWASP Top 10 for Large Language Model Applications

<https://owasp.org/www-project-top-10-for-large-language-model-applications/>

NIST Artificial Intelligence Risk Management Framework

<https://www.nist.gov/itl/ai-risk-management-framework>

NIST AI RMF Generative AI Profile

<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>

NIST Privacy Framework

<https://www.nist.gov/privacy-framework>

Agency cybersecurity, identity/access management, privacy, procurement, and incident response policies

GOV 310

Privacy-Preserving AI and De-identification

Course ID	GOV 310
Track	Governance and Security Track
Level	intermediate/practitioner
Estimated time	50
Audience	Governance committee members, Legal/privacy staff, Security staff, Compliance staff, Program owners

Learning Objectives

- Explain why AI can increase privacy and confidentiality risks in public health workflows.
- Distinguish between identifiable data, de-identified data, aggregated data, limited datasets, and synthetic data at a practical level.
- Identify technical and operational safeguards that reduce privacy risk, including data minimization, role-based access, aggregation, suppression, and secure analytic environments.
- Assess re-identification risk when public health data are linked, enriched, summarized, or shared for AI-supported use.
- Recognize when de-identification alone is not sufficient and additional governance, legal, or technical controls are needed.
- Produce a privacy-preserving AI review worksheet for a real or realistic public health workflow.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Why This Topic Matters for Public Health AI

AI can change the privacy risk profile of a public health dataset even when the source data were already collected lawfully. A dataset that seems low risk in one context may become more sensitive when linked to geospatial information, clinical history, social vulnerability indicators, case notes, or small-area counts. Generative AI and retrieval systems can also reveal sensitive details through summaries, citations, or unexpected responses. Privacy review must therefore consider not only the dataset itself, but also how the AI system accesses, transforms, combines, stores, and exposes information.

Privacy risk is also a public trust issue. Public health agencies rely on communities, providers, laboratories, schools, and partners to share information. If AI use is perceived as opaque, excessive, or poorly controlled, communities may become less willing to participate in programs or share sensitive information. Technical safeguards support trust by showing that the agency has intentionally limited access, reduced unnecessary exposure, and retained human

accountability.

Privacy-preserving AI is not a single technique. It is a layered approach that combines data governance, data minimization, secure infrastructure, access control, audit logs, de-identification, aggregation, review of outputs, and limits on downstream use. Strong privacy protection depends on designing these controls into the workflow before implementation rather than adding them after a concern arises.

Core Concepts

De-identification is the process of reducing the risk that data can be connected to a specific person. In public health practice, de-identification must be interpreted carefully because rare diseases, small communities, unusual dates, exact locations, and combinations of demographic fields may still make a record identifiable. A dataset can appear anonymous but still be linkable to individuals when combined with other information.

Re-identification risk is the possibility that a person, household, provider, facility, or community can be identified from data that were intended to be de-identified or aggregated. AI can increase this risk by making patterns easier to discover, by generating detailed summaries, or by combining sources that were originally separated. Re-identification risk should be assessed before data are used for model development, vendor testing, public dashboards, or training exercises.

Data minimization means using only the data elements necessary for the approved purpose. It is a practical guardrail because many AI tools perform better when given more context, but more context also increases privacy and confidentiality risk. A public health agency should document which fields are required, which are optional, which are prohibited, and which require special approval.

Differential privacy is a statistical approach that adds carefully calibrated noise to outputs to reduce the risk that information about any one person can be inferred. It can be useful for public reporting or statistical release, but it does not solve all privacy problems. Public health staff need to understand that differential privacy involves tradeoffs between privacy protection and analytic utility.

Public Health Example

Consider a health department that wants to use AI to summarize disease investigation notes and identify common exposure settings across cases. The notes may include symptoms, dates, employer names, schools, shelters, correctional facilities, household details, immigration concerns, or other sensitive context. Even if the AI system is internal, the agency must define who can access the notes, whether the AI output can include names or exact locations, and how summaries will be reviewed before sharing.

A privacy-preserving design might limit the AI tool to authorized case investigation staff, restrict the tool from using data outside the case management environment, require source traceability, and prevent output from including identifiers unless explicitly needed for the case workflow. Aggregated exposure summaries may require suppression rules for small counts, especially in rural areas or small communities. The goal is not to block useful analysis, but to ensure that sensitive information is used only in ways that are necessary, authorized, reviewed, and documented.

Technical Deep Dive

Privacy-preserving AI begins with data classification. Staff should identify whether the workflow uses protected health information, personally identifiable information, confidential public health data, sensitive community data, operational data, or publicly available information. Classification should be tied to rules for storage, access, sharing, retention, and

approved AI tools. A public AI tool may be acceptable for public guidance documents but inappropriate for case narratives or line lists.

The next step is determining the minimum necessary data for the AI task. A model that classifies complaint topics may not need names or exact addresses. A tool that summarizes case notes for an assigned investigator may need identifiers inside the secure case system, but the same identifiers should not appear in evaluation exports or training examples. Data minimization should be built into data extracts, prompts, APIs, and reports.

De-identification requires more than removing names. Dates, ZIP codes, facility names, employer names, exact locations, rare conditions, and free-text narratives can all create identification risk. Natural language processing and generative AI add additional complexity because sensitive details may be embedded in text. Agencies should consider automated redaction, structured review, suppression rules, and human spot checks before using text in AI workflows.

Privacy controls must also apply to outputs. An AI-generated summary, risk score, cluster report, or dashboard can reveal sensitive information even if the input data were controlled. Output review should consider small cells, unusual combinations of attributes, identifiable narratives, and unintended inferences. Logs should document who generated outputs, what data were used, and whether the output was shared, edited, rejected, or escalated.

Risks, Failure Modes, and Guardrails

A major failure mode is assuming that de-identification is permanent. Data that are de-identified for one purpose may be re-identifiable in another context, especially when linked with other datasets. A guardrail is to perform use-case-specific privacy review and to document assumptions about data linkage, audience, geography, and release context.

Another failure mode is exposing sensitive data to unapproved AI tools. Staff may paste text into a convenient tool without realizing that the data include protected or confidential information. A guardrail is to provide clear approved-tool lists, data-use tiers, technical blocking where possible, and training on prohibited uses.

Output leakage is also a significant risk. A model may include more detail than needed or reveal rare combinations that point to an individual or small group. Guardrails include output templates, suppression thresholds, review workflows, redaction, access controls, and limits on copying or exporting results.

Privacy risk can also arise from vendor relationships. Vendors may request sample data for testing or model tuning. Agencies should use synthetic, masked, or minimized data whenever possible and should require contractual controls for storage, reuse, subcontractors, breach notification, audit rights, and deletion.

Application to AI-Supported Workflows

For generative AI, privacy-preserving design means controlling what data can be placed in prompts, what source documents are available to retrieval systems, how outputs are reviewed, and whether logs contain sensitive content. Prompt templates should make prohibited content clear and should encourage use of de-identified or aggregated inputs when possible.

For predictive analytics, privacy concerns include which features are used, whether sensitive attributes are directly included or indirectly inferred, and whether outputs could stigmatize individuals or communities. Model evaluation should include privacy review as well as accuracy and equity review.

For agentic workflows, privacy risk increases because the system may retrieve data, call tools, send messages, or update records. Privacy guardrails should include least-privilege permissions, approval checkpoints, output restrictions,

audit logs, and the ability to suspend the workflow quickly if a concern arises.

Reflection Questions

Which AI use cases in your agency require identifiable or sensitive data, and which could be completed with de-identified or aggregated data?

What data elements create the greatest re-identification risk in your context?

How are free-text fields reviewed before they are used in AI-supported workflows?

Which outputs could reveal sensitive information even if the input data are controlled?

Who has authority to approve privacy-preserving AI controls for a workflow?

Practical Exercise

Create a privacy-preserving AI review worksheet for one proposed public health AI workflow. The worksheet should identify the data involved, the sensitivity level, the minimum necessary fields, the approved users, the AI task, the output type, and the privacy controls needed before implementation.

The exercise should also identify residual risks and questions that need privacy, legal, security, or governance review. Learners should pay special attention to free text, small-area data, rare conditions, linked datasets, vendor access, and output sharing.

Expected Artifact or Evidence

The expected artifact is a privacy-preserving AI review worksheet with a brief summary of the proposed workflow, data-use tier, approved data elements, prohibited data elements, access controls, output controls, audit requirements, and unresolved review questions.

References and Resources for Additional Information

NIST Privacy Framework

NIST AI Risk Management Framework

HHS HIPAA de-identification guidance

CDC data governance and public health data modernization resources

Agency privacy, legal, security, and records management policies

Knowledge Check

1. Which statement best describes de-identification in public health AI?

- A. It permanently eliminates all privacy risk in every context.
- B. It reduces identification risk but must be assessed in relation to use, linkage, audience, and context.

- C. It is only relevant for public dashboards.
- D. It allows sensitive data to be used in any AI tool.

Correct answer: B

2. Why can AI increase re-identification risk?

- A. AI tools can combine, summarize, or infer information from multiple fields and sources.
- B. AI tools never use source data.
- C. Re-identification risk applies only to paper records.
- D. Public health data cannot be re-identified.

Correct answer: A

3. What is a strong privacy guardrail for an AI workflow?

- A. Letting the vendor decide which fields are necessary.
- B. Retaining all data forever for possible future model improvement.
- C. Using minimum necessary data, role-based access, audit logs, output review, and escalation.
- D. Removing privacy review if the model performs well.

Correct answer: C

4. When is de-identification alone insufficient?

- A. When the dataset includes no rows.
- B. When the tool is new.
- C. When the data are used by public health staff.
- D. When linkage, small counts, rare conditions, free text, or output sharing create residual identification risk.

Correct answer: D

5. What is strong completion evidence for this module?

- A. A generic statement that privacy matters.
- B. A privacy-preserving AI review worksheet tied to a specific workflow, data elements, users, outputs, and safeguards.
- C. A vendor brochure about security.
- D. A list of AI tools used by staff.

Correct answer: B

References and Resources

NIST Privacy Framework

NIST AI Risk Management Framework

HHS HIPAA de-identification guidance

CDC data governance and public health data modernization resources

Agency privacy, legal, security, and records management policies

GOV 320

AI Documentation, Model Cards, and System Cards

Course ID	GOV 320
Track	Governance and Security Track
Level	intermediate/practitioner
Estimated time	50
Audience	Governance committee members, Legal/privacy staff, Security staff, Compliance staff, Program owners

Learning Objectives

- Explain the purpose of ai documentation, model cards, and system cards in public health AI work.
- Identify the technical, operational, privacy, security, equity, and governance issues associated with the topic.
- Apply the topic to a real or realistic public health workflow, system, or implementation scenario.
- Recognize common failure modes that can affect safety, accuracy, trust, workflow fit, or sustainability.
- Identify practical guardrails that should be documented before implementation.
- Produce a AI documentation packet that can support readiness assessment, governance review, procurement, implementation, or monitoring.

Jurisdiction and Agency Policy Note

This national-level module is intended for training and planning purposes. It does not replace legal, privacy, procurement, civil rights, records, or agency policy review. Learners should confirm applicable requirements in their own jurisdiction and organization before implementing AI-supported workflows.

Why This Topic Matters for Public Health AI

AI Documentation, Model Cards, and System Cards matters because AI systems are embedded in public health programs, data systems, and decision processes rather than operating in isolation. The topic affects whether AI outputs are accurate, explainable, safe, equitable, secure, and sustainable. Agencies that ignore this area may create tools that work in a demonstration but fail in actual practice.

The topic also affects accountability. Public health agencies must be able to explain how data are used, why a tool was approved, what evidence supports the workflow, who reviews outputs, and what happens when the system fails. These responsibilities become more important as AI tools move from experimentation into operational use.

Core Concepts

AI documentation records what the system is intended to do, how it was developed or configured, what data it uses, how it was evaluated, what limitations are known, and how it is monitored. Documentation supports accountability across the AI lifecycle.

A model card is a structured summary of a model, including intended use, data, evaluation, limitations, risks, and appropriate use. A system card extends documentation to the broader system, including workflow integration, interfaces, user roles, retrieval sources, human oversight, and operational controls.

Documentation is not a paperwork exercise. It is how agencies preserve institutional memory, support governance review, answer vendor questions, investigate incidents, train staff, and explain decisions to leaders or the public when appropriate.

Public Health Example

An agency using a generative AI assistant to draft public health messages should document the approved use, source materials, prohibited content, review process, accessibility expectations, language review, hallucination controls, user training, monitoring process, and incident response pathway. A system card would help reviewers understand the whole workflow rather than only the model.

Technical Deep Dive

AI documentation should be created before deployment and maintained after launch. The documentation should identify the system owner, intended users, data inputs, outputs, human review requirements, validation approach, known limitations, and change management process. A document that is created once and never updated will not support responsible operations.

Model cards are especially useful for predictive or classification models because they can document training data, evaluation metrics, subgroup performance, limitations, and intended use. System cards are especially useful for generative, RAG, and agentic systems because these tools depend on prompts, retrieval content, integrations, permissions, and human workflows.

Documentation should also cover changes. Model updates, prompt changes, data source changes, terminology updates, and workflow changes can all affect system behavior. Change logs and version histories help agencies reconstruct what happened if outputs change or incidents occur.

Risks, Failure Modes, and Guardrails

A common failure mode is treating the topic as a technical detail rather than a public health control. When staff do not connect technical decisions to authority, workflow, equity, privacy, security, and accountability, the agency may adopt a tool that appears useful but cannot be sustained or defended.

Another risk is relying on vendor claims without agency-defined evidence. Vendors may provide demonstrations, dashboards, or summary metrics that do not match the agency workflow or data environment. A guardrail is to require documentation, test results, acceptance criteria, and agency-controlled review.

A third risk is failing to plan for change. Public health data, policies, systems, and emergencies change. Any technical approach must include monitoring, review, versioning, escalation, and retirement pathways so the agency can respond when assumptions no longer hold.

Application to AI-Supported Workflows

Generative AI, predictive analytics, RAG, automation, and agentic workflows all depend on the controls described in this module. The controls should be scaled to the risk of the use case. A tool that drafts internal training text may require

lighter review than a tool that updates case records, prioritizes outreach, or supports public communication during an emergency.

The module should be applied during readiness assessment, procurement, implementation planning, pilot testing, and post-deployment monitoring. Learners should document how the topic affects data, systems, users, safeguards, and decision-making for the selected workflow.

Reflection Questions

Where does this topic appear in one current or proposed AI workflow in your agency?

What decisions, data sources, users, or partners would be affected if this area is poorly designed?

What evidence would governance reviewers need before approving the workflow?

Which safeguards should be required before pilot testing?

Who owns ongoing monitoring and review?

Practical Exercise

Create a AI documentation packet for one real or realistic public health AI workflow. The artifact should describe the use case, affected users, data sources, system dependencies, risks, guardrails, review points, and unresolved questions.

The exercise should produce a work product that could be used in a governance meeting, procurement review, implementation planning session, pilot evaluation, or monitoring review. Learners should document assumptions and identify which staff or partners need to be consulted.

Expected Artifact or Evidence

The expected artifact is a completed AI documentation packet. It should be specific to a public health workflow and should include technical dependencies, governance questions, human review expectations, monitoring indicators, and unresolved risks.

References and Resources for Additional Information

NIST AI Risk Management Framework

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP Top 10 for LLM Applications, when security or AI integrations are relevant

Agency privacy, cybersecurity, data governance, procurement, and records management policies

Knowledge Check

1. What is the best reason to address ai documentation, model cards, and system cards before deployment?

- A. It makes governance unnecessary.
- B. It helps ensure the AI workflow is safe, useful, reviewable, and fit for public health operations.
- C. It allows the agency to skip testing.
- D. It guarantees that the vendor product will work.

Correct answer: B

2. Which practice best supports responsible implementation?

- A. Allowing users to decide individually how to use the tool.
- B. Treating vendor documentation as the only required evidence.
- C. Defining intended use, evidence requirements, safeguards, monitoring, and escalation before launch.
- D. Waiting until after launch to decide whether the tool is risky.

Correct answer: C

3. Which statement best describes a common failure mode?

- A. Human review is always unnecessary when AI is accurate.
- B. Documentation makes AI less useful.
- C. The agency assumes the tool is ready because the technology works in a demonstration, even though workflow, data, and governance have not been tested.
- D. Monitoring should stop after the pilot.

Correct answer: C

4. What should be included in the practical artifact for this module?

- A. Only a general definition.
- B. A marketing description of the AI tool.
- C. A list of unrelated data systems.
- D. Workflow-specific risks, safeguards, roles, evidence requirements, and unresolved questions.

Correct answer: D

5. When should the issue be escalated for governance or leadership review?

- A. Only if the tool is public-facing.
- B. When risks affect privacy, safety, equity, legal authority, public trust, system integrity, or operational decisions.
- C. Never, because technical staff can resolve all issues alone.
- D. Only after an incident has already occurred.

Correct answer: B

References and Resources

NIST AI Risk Management Framework

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP Top 10 for LLM Applications, when security or AI integrations are relevant

Agency privacy, cybersecurity, data governance, procurement, and records management policies

GOV 330

Data Governance for AI-Ready Public Health Data

Course ID	GOV 330
Track	Governance and Security Track
Level	intermediate/practitioner
Estimated time	50
Audience	Governance committee members, Legal/privacy staff, Security staff, Compliance staff, Program owners

Learning Objectives

- Explain the purpose of data governance for ai-ready public health data in public health AI work.
- Identify the technical, operational, privacy, security, equity, and governance issues associated with the topic.
- Apply the topic to a real or realistic public health workflow, system, or implementation scenario.
- Recognize common failure modes that can affect safety, accuracy, trust, workflow fit, or sustainability.
- Identify practical guardrails that should be documented before implementation.
- Produce a AI-ready data governance checklist that can support readiness assessment, governance review, procurement, implementation, or monitoring.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Why This Topic Matters for Public Health AI

Data Governance for AI-Ready Public Health Data matters because AI systems are embedded in public health programs, data systems, and decision processes rather than operating in isolation. The topic affects whether AI outputs are accurate, explainable, safe, equitable, secure, and sustainable. Agencies that ignore this area may create tools that work in a demonstration but fail in actual practice.

The topic also affects accountability. Public health agencies must be able to explain how data are used, why a tool was approved, what evidence supports the workflow, who reviews outputs, and what happens when the system fails. These responsibilities become more important as AI tools move from experimentation into operational use.

Core Concepts

Data governance is the set of decision rights, roles, policies, standards, and practices that guide how data are collected, managed, protected, shared, used, and retired. AI-ready data depend on governance because data must be

understandable, authorized, high quality, and fit for use.

Data stewardship assigns responsibility for data quality, definitions, metadata, access, and appropriate use. Stewards do not necessarily own the data alone, but they help ensure that data are managed consistently and that users understand limitations.

Metadata describe the meaning, source, structure, quality, and permitted use of data. Without metadata, AI teams may misinterpret fields, misuse denominators, overlook missingness, or fail to recognize that a dataset is not appropriate for the intended use.

Public Health Example

A public health agency may want to use multiple datasets to prioritize outreach for vaccination. Data governance would clarify which datasets can be linked, which fields are authoritative, which populations are underrepresented, who can approve use, how long data can be retained, and how results can be shared. Without governance, the model may appear technical while resting on unclear authority and inconsistent data definitions.

Technical Deep Dive

AI-ready data begin with inventory. Agencies should know what datasets exist, who stewards them, what they contain, how often they are updated, what quality issues are known, and what uses are permitted. A data catalog can make this information visible to program, analytics, informatics, and governance teams.

Governance must also address data lineage and transformation. AI outputs may depend on mapped, linked, deduplicated, geocoded, imputed, or aggregated data. Staff need to understand how those transformations were performed and whether they are appropriate for the AI use case.

Permissible use should be explicit. Public health data may be collected under specific legal authorities or agreements. AI use may be appropriate for one purpose and inappropriate for another. Governance should define approved uses, prohibited uses, review pathways, and documentation requirements.

Risks, Failure Modes, and Guardrails

A common failure mode is treating the topic as a technical detail rather than a public health control. When staff do not connect technical decisions to authority, workflow, equity, privacy, security, and accountability, the agency may adopt a tool that appears useful but cannot be sustained or defended.

Another risk is relying on vendor claims without agency-defined evidence. Vendors may provide demonstrations, dashboards, or summary metrics that do not match the agency workflow or data environment. A guardrail is to require documentation, test results, acceptance criteria, and agency-controlled review.

A third risk is failing to plan for change. Public health data, policies, systems, and emergencies change. Any technical approach must include monitoring, review, versioning, escalation, and retirement pathways so the agency can respond when assumptions no longer hold.

Application to AI-Supported Workflows

Generative AI, predictive analytics, RAG, automation, and agentic workflows all depend on the controls described in this module. The controls should be scaled to the risk of the use case. A tool that drafts internal training text may require

lighter review than a tool that updates case records, prioritizes outreach, or supports public communication during an emergency.

The module should be applied during readiness assessment, procurement, implementation planning, pilot testing, and post-deployment monitoring. Learners should document how the topic affects data, systems, users, safeguards, and decision-making for the selected workflow.

Reflection Questions

Where does this topic appear in one current or proposed AI workflow in your agency?

What decisions, data sources, users, or partners would be affected if this area is poorly designed?

What evidence would governance reviewers need before approving the workflow?

Which safeguards should be required before pilot testing?

Who owns ongoing monitoring and review?

Practical Exercise

Create a AI-ready data governance checklist for one real or realistic public health AI workflow. The artifact should describe the use case, affected users, data sources, system dependencies, risks, guardrails, review points, and unresolved questions.

The exercise should produce a work product that could be used in a governance meeting, procurement review, implementation planning session, pilot evaluation, or monitoring review. Learners should document assumptions and identify which staff or partners need to be consulted.

Expected Artifact or Evidence

The expected artifact is a completed AI-ready data governance checklist. It should be specific to a public health workflow and should include technical dependencies, governance questions, human review expectations, monitoring indicators, and unresolved risks.

References and Resources for Additional Information

NIST AI Risk Management Framework

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP Top 10 for LLM Applications, when security or AI integrations are relevant

Agency privacy, cybersecurity, data governance, procurement, and records management policies

Knowledge Check

1. What is the best reason to address data governance for ai-ready public health data before deployment?

- A. It makes governance unnecessary.
- B. It helps ensure the AI workflow is safe, useful, reviewable, and fit for public health operations.
- C. It allows the agency to skip testing.
- D. It guarantees that the vendor product will work.

Correct answer: B

2. Which practice best supports responsible implementation?

- A. Defining intended use, evidence requirements, safeguards, monitoring, and escalation before launch.
- B. Waiting until after launch to decide whether the tool is risky.
- C. Allowing users to decide individually how to use the tool.
- D. Treating vendor documentation as the only required evidence.

Correct answer: A

3. Which statement best describes a common failure mode?

- A. Human review is always unnecessary when AI is accurate.
- B. Documentation makes AI less useful.
- C. The agency assumes the tool is ready because the technology works in a demonstration, even though workflow, data, and governance have not been tested.
- D. Monitoring should stop after the pilot.

Correct answer: C

4. What should be included in the practical artifact for this module?

- A. Only a general definition.
- B. A marketing description of the AI tool.
- C. A list of unrelated data systems.
- D. Workflow-specific risks, safeguards, roles, evidence requirements, and unresolved questions.

Correct answer: D

5. When should the issue be escalated for governance or leadership review?

- A. Only if the tool is public-facing.
- B. When risks affect privacy, safety, equity, legal authority, public trust, system integrity, or operational decisions.
- C. Never, because technical staff can resolve all issues alone.
- D. Only after an incident has already occurred.

Correct answer: B

References and Resources

NIST AI Risk Management Framework

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP Top 10 for LLM Applications, when security or AI integrations are relevant

Agency privacy, cybersecurity, data governance, procurement, and records management policies

GOV 340

AI Audit Readiness and Evidence Management

Course ID	GOV 340
Track	Governance and Security Track
Level	intermediate/practitioner
Estimated time	50
Audience	Governance committee members, Legal/privacy staff, Security staff, Compliance staff, Program owners

Learning Objectives

- Explain why AI systems require lifecycle evidence management.
- Identify the core records that support AI audit readiness.
- Distinguish between approval evidence, validation evidence, operational evidence, and incident evidence.
- Create an AI evidence inventory for an approved or proposed use case.
- Develop retention, ownership, and review expectations for AI governance artifacts.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Public Health Example

A health department uses an AI tool to help classify public inquiries. Six months later, a concern is raised that the tool may be misclassifying language access requests. Audit-ready evidence would include the approved use case, validation plan, subgroup review, SOP, training records, user override logs, incident reports, vendor documentation, and monitoring results. This evidence allows the agency to investigate the concern and decide whether to modify, pause, or retire the workflow.

Technical and Operational Deep Dive

AI evidence begins at intake. The agency should preserve the use case description, risk triage, required reviewers, approval conditions, and rationale for the decision. If a use case is rejected or deferred, that decision should also be documented. These early records help explain why a tool exists and what limits were placed on its use.

Validation evidence should show how the agency determined whether the AI-supported workflow was fit for purpose. This may include test datasets, evaluation rubrics, accuracy measures, subgroup performance, user acceptance testing, security review, privacy review, accessibility review, and pilot findings. Validation evidence should be understandable to decision-makers, not only technical staff.

Operational evidence should show what happened after deployment. This may include monitoring dashboards, user override rates, error reports, support tickets, incident records, model or prompt changes, vendor updates, SOP revisions, training completion, and governance meeting minutes. Operational evidence is especially important because AI performance and workflow use can change over time.

Evidence management requires ownership and retention rules. Governance bodies should define who stores each artifact, where it is stored, how long it is retained, who can access it, and how evidence is updated when tools change. Audit readiness depends on consistent records, not scattered files and informal emails.

Risks, Failure Modes, and Guardrails

Risks and failure modes include the following:

Approvals without records of conditions or rationale.

Validation results stored informally or not retained.

Vendor documentation accepted without agency review or local evidence.

Incidents that cannot be reconstructed because logs or decisions were not preserved.

Public records or audit responses delayed because artifacts are scattered.

Recommended guardrails include the following:

Maintain an AI evidence inventory for each approved AI use.

Use decision logs for approvals, deferrals, conditions, changes, and retirement.

Preserve validation, testing, monitoring, training, SOP, and incident records.

Assign artifact ownership and retention expectations.

Review evidence completeness during periodic governance review.

Reflection Questions

Where does this topic appear in your agency, program, or workflow?

Which staff roles would need to understand or apply this concept before AI use is approved?

What could go wrong if this topic is not addressed before implementation?

What evidence would governance, leadership, or operations staff need before moving forward?

Which policy, data, equity, privacy, security, or workflow questions remain unresolved?

Practical Exercise

Create an AI audit readiness evidence inventory for one AI use case. Identify required artifacts from intake through retirement, the owner of each artifact, storage location, retention expectation, update trigger, and current status.

Expected Artifact or Evidence

AI evidence inventory

Decision log template

Artifact ownership and retention table

Gap list for missing evidence

Review schedule

References and Resources for Additional Information

NIST Artificial Intelligence Risk Management Framework

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP guidance for LLM application security

Agency policies for privacy, cybersecurity, records, procurement, accessibility, and public communications

Relevant public health program SOPs, data governance documentation, and implementation plans

Knowledge Check

1. What is AI audit readiness?

- A. A vendor marketing activity
- B. A model accuracy score only
- C. The ability to produce evidence showing how decisions, controls, evaluations, and oversight were performed
- D. A public dashboard design

Correct answer: C

2. Which record should be preserved from AI governance review?

- A. Approval decision, conditions, rationale, and required next steps
- B. Only the meeting date
- C. Only the tool name
- D. Only the vendor logo

Correct answer: A

3. What is validation evidence?

- A. A staff directory
- B. A press release
- C. A procurement invoice only
- D. Records showing how the tool or workflow was tested for fitness for purpose

Correct answer: D

4. Why is operational evidence important?

- A. It replaces user training
- B. It shows how the system is used, monitored, maintained, and corrected after deployment
- C. It proves incidents cannot occur
- D. It eliminates retention rules

Correct answer: B

5. What is strong completion evidence?

- A. A single screenshot
- B. A list of AI products
- C. An evidence inventory with artifact owners, storage, retention, update triggers, and gaps
- D. An informal approval email only

Correct answer: C

References and Resources

NIST Artificial Intelligence Risk Management Framework

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP guidance for LLM application security

Agency policies for privacy, cybersecurity, records, procurement, accessibility, and public communications

Relevant public health program SOPs, data governance documentation, and implementation plans

GOV 350

Identity, Access, and Permissions for AI Workflows

Course ID	GOV 350
Track	Governance and Security Track
Level	intermediate/practitioner
Estimated time	50
Audience	Governance committee members, Legal/privacy staff, Security staff, Compliance staff, Program owners

Learning Objectives

- Explain the difference between user identity, service identity, authentication, authorization, and permissions.
- Apply least-privilege principles to AI-supported public health workflows.
- Identify access risks for read-only, retrieval, summarization, and write-back AI functions.
- Define audit log and approval requirements for AI tools that access sensitive systems or data.
- Create an access and permissions review for an AI workflow.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Public Health Example

An AI assistant is approved to summarize case records for authorized disease investigators. The access review determines that the assistant may retrieve only records assigned to the investigator's program, may not access unrelated program data, may not export records to unapproved systems, and may not write final case notes without human approval. Logs must show which user requested the summary, which record was accessed, and whether the output was saved or discarded.

Technical and Operational Deep Dive

AI workflows may involve multiple identities. A staff member may log in through single sign-on, an AI service may use a service account to retrieve data, and a vendor system may process a request through an API. Each identity should be documented and tied to specific permissions. Agencies should avoid shared accounts or overly broad service accounts because they make accountability and incident investigation difficult.

Authentication verifies who or what is using the system, but authorization defines what that identity can do. For AI workflows, authorization should be granular enough to distinguish reading data, retrieving documents, generating summaries, creating tasks, sending notifications, and updating records. Write-back permissions should receive stronger

review than read-only functions because they can change official records or trigger action.

Least privilege is especially important for AI because tools may process data at scale. A human user may look at one record at a time, but an AI integration may retrieve many records quickly. Permissions should limit data by role, program, jurisdiction, purpose, and workflow where possible. Sensitive data should not be accessible simply because it is technically available.

Audit logs should make AI access visible. Logs should capture user identity, service identity, data accessed, action taken, time, output generated, approval status, and errors or exceptions. For agentic workflows, logs should also show tool calls, attempted actions, blocked actions, and human approvals.

Risks, Failure Modes, and Guardrails

Risks and failure modes include the following:

Service accounts with broad access to sensitive data.

AI tools retrieving data unrelated to the approved use case.

Write-back actions occurring without human approval.

Shared accounts preventing accountability.

Insufficient logs to investigate data access or output generation.

Recommended guardrails include the following:

Use role-based and purpose-based access controls.

Apply least privilege to users, service accounts, APIs, and agents.

Require stronger approval for write-back, external transmission, and automated actions.

Avoid shared accounts and undocumented service identities.

Maintain detailed audit logs for AI access and actions.

Reflection Questions

Where does this topic appear in your agency, program, or workflow?

Which staff roles would need to understand or apply this concept before AI use is approved?

What could go wrong if this topic is not addressed before implementation?

What evidence would governance, leadership, or operations staff need before moving forward?

Which policy, data, equity, privacy, security, or workflow questions remain unresolved?

Practical Exercise

Create an identity, access, and permissions review for one AI-supported workflow. Identify users, service accounts, systems, data accessed, read/write permissions, approval points, prohibited actions, audit log requirements, and escalation triggers.

Expected Artifact or Evidence

Access and permissions review

List of user and service identities

Least-privilege recommendation

Write-back approval rules

Audit log requirements

References and Resources for Additional Information

NIST Artificial Intelligence Risk Management Framework

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP guidance for LLM application security

Agency policies for privacy, cybersecurity, records, procurement, accessibility, and public communications

Relevant public health program SOPs, data governance documentation, and implementation plans

Knowledge Check

1. What is authorization?

- A. Verifying a password only
- B. Determining what an authenticated identity is allowed to access or do
- C. Writing a press release
- D. Testing model accuracy

Correct answer: B

2. Why are service accounts important in AI workflows?

- A. They are always public
- B. They replace governance review
- C. They may allow systems or automated workflows to access data or perform actions
- D. They eliminate audit logs

Correct answer: C

3. What does least privilege mean?

- A. Granting only the access needed for an approved purpose
- B. Giving every tool full access
- C. Using shared accounts
- D. Skipping permissions review

Correct answer: A

4. Which AI function generally requires stronger safeguards?

- A. Reading a public FAQ
- B. Drafting a generic training title
- C. Formatting a meeting agenda
- D. Writing back to a system of record

Correct answer: D

5. What is strong completion evidence?

- A. A vendor logo
- B. An access review identifying identities, permissions, prohibited actions, approval points, and audit logs
- C. A generic AI definition
- D. A screenshot without permissions details

Correct answer: B

References and Resources

NIST Artificial Intelligence Risk Management Framework

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP guidance for LLM application security

Agency policies for privacy, cybersecurity, records, procurement, accessibility, and public communications

Relevant public health program SOPs, data governance documentation, and implementation plans

GOV 360

Third-Party AI Risk Management

Course ID	GOV 360
Track	Governance and Security Track
Level	intermediate/practitioner
Estimated time	50
Audience	Governance committee members, Legal/privacy staff, Security staff, Compliance staff, Program owners

Learning Objectives

- Explain the purpose of third-party ai risk management in public health AI.
- Identify the roles, decisions, risks, and documentation needed for the topic.
- Apply the topic to a real or realistic public health workflow or decision.
- Recognize equity, privacy, security, legal, operational, and public trust implications.
- Produce a practical artifact that supports readiness, governance, implementation, monitoring, or accountability.

Definitions

Decision right	The authority to approve, deny, modify, pause, or escalate an AI use case or operational decision.
Evidence artifact	A document, log, checklist, test result, approval record, or review note that demonstrates responsible AI practice.
Escalation pathway	A defined route for raising concerns to the appropriate reviewer or decision-maker.
Accountability owner	The person or role responsible for ensuring that an AI use case is governed, monitored, and corrected when needed.
Operational safeguard	A procedure, control, or review step that reduces risk during day-to-day use.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Why This Topic Matters for Public Health AI

Third-Party AI Risk Management matters because AI adoption changes how public health agencies make, document, communicate, and review decisions. Without a structured approach, staff may rely on informal practices, vendor assurances, or one-time approvals that do not match the operational consequences of AI-supported work.

Public health agencies operate under legal authorities, public accountability, resource constraints, and community trust obligations. A module on this topic helps agencies define repeatable expectations before urgent situations arise. It also helps staff distinguish low-risk administrative support from uses that affect surveillance, public communications, field response, resource allocation, civil rights, privacy, or official records.

Public Health Example

A public health agency is preparing to use AI in a workflow related to ongoing oversight of vendors, subcontractors, data rights, model updates, and exit plans. The project team initially treats the tool as a helpful efficiency aid, but governance reviewers recognize that the workflow may affect public trust, equity, official records, or public health action.

The agency responds by creating a clear artifact that defines intended use, prohibited use, responsible roles, review points, documentation expectations, escalation triggers, and monitoring measures. This converts a broad AI idea into a manageable public health practice.

Technical and Operational Deep Dive

The first step in this topic is to define the decision or workflow clearly. Staff should identify who is affected, what data are used, what AI output is produced, who reviews the output, and what action may follow. This prevents the module from becoming abstract and anchors the work in public health practice.

The second step is to assign responsibilities. AI-supported workflows often cross program, IT, legal, privacy, communications, procurement, and leadership boundaries. Responsibilities should be named by role rather than assumed. The artifact should specify who prepares materials, who reviews risks, who approves use, who monitors performance, and who can pause or revise the workflow.

The third step is to define evidence. Responsible AI depends on more than policy statements. Agencies need records that show what was reviewed, what decisions were made, what safeguards were applied, what incidents occurred, and how the workflow changed over time. Evidence also supports audits, public records review, quality improvement, and leadership accountability.

Risks, Failure Modes, and Guardrails

A common failure mode is treating the topic as a one-time checklist. AI risks can change when data sources, models, vendors, policies, users, or public expectations change. A guardrail is to define review intervals and change triggers.

Another failure mode is unclear ownership. When everyone is responsible in general, no one may be responsible in practice. A guardrail is to assign an accountability owner and name escalation pathways.

A third failure mode is relying on vendor claims or informal staff confidence without documentation. A guardrail is to require evidence artifacts that can be reviewed by governance, leadership, privacy, legal, equity, security, or program teams.

Application to AI-Supported Workflows

This topic applies across generative AI, predictive analytics, NLP, RAG, automation, and agentic workflows. The more sensitive the data, the more consequential the output, or the more automated the action, the stronger the documentation, review, and monitoring expectations should be.

Reflection Questions

What public health decision or workflow does this topic affect?

Who has authority to approve or pause the work?

What evidence would show that the workflow is responsible and effective?

What equity, privacy, legal, security, or public trust risks require review?

How will the agency know when the workflow needs to be revised?

Practical Exercise

Create a practical third-party ai risk management artifact for one AI-supported public health use case. The artifact should define the use case, affected roles, review criteria, safeguards, documentation expectations, escalation triggers, monitoring indicators, and next review date.

Expected Artifact or Evidence

Third-Party AI Risk Management artifact

Role and responsibility table

Risk and safeguard summary

Escalation and monitoring plan

Approval or review note

References and Resources for Additional Information

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

Knowledge Check

1. What is the strongest evidence that third-party ai risk management has been applied responsibly?

- A. A general statement that AI is useful
- B. A vendor demonstration with no agency review
- C. A documented artifact that defines use, roles, safeguards, escalation, monitoring, and evidence
- D. A decision to skip review because the use case seems efficient

Correct answer: C

2. What is the best reason to assign an accountability owner?

- A. To reduce the need for documentation
- B. To ensure someone is responsible for monitoring, escalation, and correction
- C. To make the vendor solely responsible
- D. To prevent staff from raising concerns

Correct answer: B

3. When should an AI-supported workflow be escalated for review?

- A. Only after media coverage
- B. Never if the tool is efficient
- C. When outputs may affect safety, equity, privacy, legal authority, public trust, or official action
- D. Only at the end of the fiscal year

Correct answer: C

4. Which practice best supports public accountability?

- A. Keeping AI use informal
- B. Avoiding written records
- C. Letting each user decide independently
- D. Maintaining evidence artifacts, decision records, review notes, and monitoring results

Correct answer: D

5. What should happen when data sources, policy, vendor configuration, or workflow use changes?

- A. The use case should be reviewed to determine whether safeguards, validation, documentation, or approvals need updating
- B. The agency should ignore the change if staff like the tool
- C. The model should be trusted automatically
- D. The glossary should be deleted

Correct answer: A

References and Resources

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

GOV 400

AI Governance Charter and Operating Model

Course ID	GOV 400
Track	Governance and Security Track
Level	advanced/leadership
Estimated time	50
Audience	Governance committee members, Legal/privacy staff, Security staff, Compliance staff, Program owners

Learning Objectives

- Define the purpose and minimum contents of an AI governance charter.
- Distinguish between a governance committee, technical review group, executive sponsor, and implementation team.
- Identify which agency functions should participate in AI governance and when their review is required.
- Develop operating procedures for intake, review, approval, monitoring, exceptions, escalation, and retirement of AI use cases.
- Produce a draft AI governance charter or operating model that can be adapted by a public health agency.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Why This Topic Matters for Public Health AI

Public health agencies operate under legal authorities, public accountability obligations, data-sharing agreements, privacy requirements, procurement rules, and equity commitments. AI governance must therefore be more than a technology review process. It must define how the agency decides which AI uses are acceptable, which safeguards are required, who has authority to approve use, and how concerns are escalated.

A governance operating model also helps avoid the problem of uneven adoption. One program may use a public AI tool informally, another may procure a vendor product, and a third may build an internal workflow using sensitive data. If these activities are not reviewed through a common process, the agency may create inconsistent privacy practices, duplicative investments, and uneven risk tolerance. Governance makes AI activity visible and manageable.

Public Health Example

A state health department wants to allow staff to use AI tools for program support, document review, code assistance, case summarization, and public-facing communications. Without a governance charter, each program may develop its own rules. The communications office may require human review, the surveillance program may prohibit sensitive data

entry, IT may limit tool access, and procurement may approve a vendor product without a clear validation plan.

A governance operating model brings these decisions into one structured process. The agency can define low-risk uses that require staff training and supervisor approval, moderate-risk uses that require privacy and program review, and high-risk uses that require executive approval, validation evidence, monitoring, and incident response planning. The charter does not eliminate flexibility. It creates a consistent path for responsible flexibility.

Technical and Governance Deep Dive

A strong governance charter begins with purpose and scope. Purpose explains why the governance body exists, such as protecting public trust, supporting responsible innovation, aligning AI investments with agency priorities, and managing risk across the AI lifecycle. Scope explains what falls under the governance process. For public health, scope should include internally developed AI tools, vendor AI products, generative AI tools, predictive models, agentic workflows, AI features embedded in existing software, and AI-supported uses of sensitive or operational data.

The operating model should define roles clearly. An executive sponsor provides authority and connects governance to agency strategy. A governance chair manages the review process. Privacy, legal, cybersecurity, equity, procurement, records, program, data, and informatics representatives bring specialized review perspectives. Technical reviewers assess architecture, validation, interoperability, security, and monitoring. Program owners remain responsible for the public health workflow and for ensuring that AI supports, rather than replaces, appropriate judgment.

The model should also define the life of an AI use case. Intake captures the proposed use, data involved, intended users, expected benefits, risks, and requested decision. Triage determines the review level. Governance review evaluates authority, data use, privacy, equity, security, accessibility, technical readiness, validation, procurement, and monitoring. Approval should include conditions, documentation requirements, and an owner. Post-approval oversight should review performance, incidents, changes, and whether the use case should continue.

Risks, Failure Modes, and Guardrails

A common failure mode is governance that exists only on paper. A charter may sound strong, but if staff do not know when to submit use cases, who decides, or what evidence is required, AI decisions will still happen informally. A practical guardrail is to create a simple intake form, review rubric, meeting cadence, decision log, and escalation pathway.

Another failure mode is over-centralization. If every minor use requires extensive review, staff may avoid the process or use unapproved tools. Governance should be tiered so low-risk uses move quickly while high-impact uses receive deeper review. This balance helps the agency support innovation while protecting public health values and obligations.

Application to AI-Supported Workflows

Generative AI governance should define approved tools, prohibited data, review expectations, and when outputs may be used externally. Predictive model governance should require validation, subgroup performance review, monitoring, and clear limits on decision use. Agentic AI governance should define allowed actions, approval checkpoints, logging, rollback procedures, and incident triggers. In each case, the charter gives the agency a repeatable way to make and document decisions.

Reflection Questions

Which current or proposed AI use case in your agency raises this leadership or governance issue most clearly?

Who currently has authority to make decisions about this issue, and is that authority documented?

What evidence would governance or leadership need before approving the use case?

Which staff, partners, or communities would be affected if this issue were handled poorly?

What policy, process, or artifact should be created or updated after this module?

Practical Exercise

Draft a one-page AI governance charter and a one-page operating model. Include scope, membership, review levels, decision rights, documentation, escalation, monitoring, and meeting cadence.

Before You Begin

Choose a real or realistic public health AI use case. Document assumptions, unresolved questions, and which agency roles would need to review the artifact before it is adopted. The exercise should produce a work product that could support governance review, leadership discussion, implementation planning, or policy development.

Expected Artifact or Evidence

Draft AI governance charter; governance operating model; intake and decision-flow diagram; list of unresolved policy questions.

References and Resources for Additional Information

NIST Artificial Intelligence Risk Management Framework (AI RMF 1.0)

NIST Generative AI Profile

NIST Privacy Framework

CDC Public Health Data Strategy and Data Modernization resources

WHO Ethics and Governance of Artificial Intelligence for Health

OMB Memorandum M-24-10 on federal AI governance, risk management, and innovation, where relevant to public-sector governance models

Agency strategic plans, procurement policies, privacy policies, cybersecurity policies, records policies, and equity plans

Knowledge Check

1. What is the strongest reason this module topic belongs in AI leadership and governance training?

- A. It eliminates the need for technical review.
- B. It helps leaders define authority, safeguards, accountability, and implementation expectations.
- C. It allows vendors to make final public health decisions.

D. It guarantees that every AI tool will be successful.

Correct answer: B

2. Which practice best supports accountable public health AI governance?

- A. Allowing each program to create informal AI rules independently.
- B. Approving tools based only on vendor demonstrations.
- C. Avoiding documentation to reduce burden.
- D. Using defined review steps, named owners, decision logs, and escalation pathways.

Correct answer: D

3. When should an AI use case be escalated for leadership or governance review?

- A. When it affects sensitive data, public communication, resource allocation, enforcement, service access, or public trust.
- B. Only after a public complaint occurs.
- C. Only when a vendor requests escalation.
- D. Never, because AI review should remain informal.

Correct answer: A

4. What is weak evidence of completion for a leadership and governance module?

- A. A role matrix, governance worksheet, or decision-ready plan.
- B. A specific risk statement with mitigation steps.
- C. A generic statement that AI should be used responsibly without assigning owners or actions.
- D. A completed exercise tied to a public health workflow.

Correct answer: C

5. What should leaders do when an AI use case has meaningful benefits but unresolved risks?

- A. Approve it immediately because benefits are possible.
- B. Require mitigation, human review, monitoring, documentation, or defer implementation until readiness improves.
- C. Reject all AI use permanently.
- D. Ask the vendor to decide whether risk is acceptable.

Correct answer: B

References and Resources

NIST Artificial Intelligence Risk Management Framework (AI RMF 1.0)

NIST Generative AI Profile

NIST Privacy Framework

CDC Public Health Data Strategy and Data Modernization resources

WHO Ethics and Governance of Artificial Intelligence for Health

OMB Memorandum M-24-10 on federal AI governance, risk management, and innovation, where relevant to public-sector governance models

Agency strategic plans, procurement policies, privacy policies, cybersecurity policies, records policies, and equity plans

GOV 410

AI Risk Appetite and Risk Tolerance for Public Health

Course ID	GOV 410
Track	Governance and Security Track
Level	advanced/leadership
Estimated time	50
Audience	Governance committee members, Legal/privacy staff, Security staff, Compliance staff, Program owners

Learning Objectives

- Explain the difference between risk appetite, risk tolerance, and risk acceptance.
- Identify AI uses that may be low, moderate, high, or unacceptable risk in a public health context.
- Develop risk tolerance statements for different AI use categories.
- Connect risk tolerance to review level, safeguards, monitoring, and escalation.
- Produce a draft AI risk appetite statement for leadership or governance discussion.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Why This Topic Matters for Public Health AI

Public health agencies make risk decisions every day, but AI can make those decisions less visible. A generative AI tool might be acceptable for drafting internal meeting notes but not for producing unsupervised public health guidance. A predictive model might be acceptable for planning staffing levels but not for denying services. An agentic workflow might be acceptable for creating a draft task but not for sending an enforcement notice without human approval.

Risk appetite helps agencies explain these distinctions. It provides a leadership-level statement of where the agency encourages innovation, where it requires caution, and where it prohibits use. It also helps staff understand that responsible AI is not a binary choice between adoption and avoidance. It is a disciplined process for matching safeguards to risk.

Public Health Example

An agency may define a higher risk appetite for internal administrative uses that do not involve sensitive data, public communication, or public health action. It may define a moderate appetite for staff-facing decision support when human review, validation, and documentation are required. It may define a very low appetite for unsupervised public-facing advice, automated eligibility decisions, enforcement actions, or tools that process sensitive data without approved

safeguards.

These statements help governance reviewers make consistent decisions. Instead of debating each proposal from the beginning, reviewers can ask whether the proposed use fits the agency's defined risk category and whether required safeguards are present. The discussion becomes more disciplined and less dependent on individual comfort with AI.

Technical and Governance Deep Dive

Risk appetite should be expressed in plain language and translated into operating rules. For example, an agency might state that it encourages low-risk AI uses that reduce internal administrative burden when no sensitive data are used and outputs are reviewed by staff. It might state that it will consider moderate-risk uses involving operational data only with documented validation, access controls, human review, and monitoring. It might state that it does not permit AI to make final public health enforcement, eligibility, or legal decisions without authorized human review.

Risk tolerance turns these statements into thresholds. A tolerance may specify that a tool cannot be launched if source traceability is absent, if subgroup performance cannot be assessed for a prioritization model, if the false negative rate exceeds a defined threshold, or if the vendor cannot provide adequate documentation. Tolerances should be appropriate to the use case. A lower threshold may be acceptable for internal drafting than for case triage.

Risk appetite should be reviewed regularly. Public expectations, laws, model capabilities, agency experience, and technical safeguards change over time. A use that was too risky during the first year of AI adoption may become acceptable after better infrastructure, policies, and monitoring are in place. Conversely, a use may become unacceptable if incidents, evidence, or community concerns reveal new harms.

Risks, Failure Modes, and Guardrails

A common failure mode is vague risk language. Phrases like 'use AI responsibly' do not help staff decide what is allowed. A guardrail is to define concrete examples of encouraged, conditional, restricted, and prohibited uses.

Another risk is risk appetite that ignores equity. A tool may appear operationally useful while shifting burden to communities that are already underrepresented or over-surveilled. A guardrail is to require equity impact review for uses that affect prioritization, access, enforcement, communication, or resource allocation.

Application to AI-Supported Workflows

Risk appetite applies differently across AI types. Generative AI for internal drafting may fall into a lower-risk category when reviewed by humans. Predictive models used for prioritization require stronger validation and monitoring. Agentic tools that can take system actions require strict limits, stop authority, and auditability. Risk appetite helps define these boundaries in advance.

Reflection Questions

Which current or proposed AI use case in your agency raises this leadership or governance issue most clearly?

Who currently has authority to make decisions about this issue, and is that authority documented?

What evidence would governance or leadership need before approving the use case?

Which staff, partners, or communities would be affected if this issue were handled poorly?

What policy, process, or artifact should be created or updated after this module?

Practical Exercise

Draft a risk appetite statement and a risk tolerance table for at least four AI use categories: internal productivity, staff decision support, public-facing communication, and automated or agentic workflows.

Before You Begin

Choose a real or realistic public health AI use case. Document assumptions, unresolved questions, and which agency roles would need to review the artifact before it is adopted. The exercise should produce a work product that could support governance review, leadership discussion, implementation planning, or policy development.

Expected Artifact or Evidence

AI risk appetite statement; risk tolerance table; prohibited-use examples; mitigation and escalation conditions.

References and Resources for Additional Information

NIST Artificial Intelligence Risk Management Framework (AI RMF 1.0)

NIST Generative AI Profile

NIST Privacy Framework

CDC Public Health Data Strategy and Data Modernization resources

WHO Ethics and Governance of Artificial Intelligence for Health

OMB Memorandum M-24-10 on federal AI governance, risk management, and innovation, where relevant to public-sector governance models

Agency strategic plans, procurement policies, privacy policies, cybersecurity policies, records policies, and equity plans

Knowledge Check

1. What is the strongest reason this module topic belongs in AI leadership and governance training?

- A. It helps leaders define authority, safeguards, accountability, and implementation expectations.
- B. It eliminates the need for technical review.
- C. It allows vendors to make final public health decisions.
- D. It guarantees that every AI tool will be successful.

Correct answer: A

2. Which practice best supports accountable public health AI governance?

- A. Allowing each program to create informal AI rules independently.

- B. Approving tools based only on vendor demonstrations.
- C. Using defined review steps, named owners, decision logs, and escalation pathways.
- D. Avoiding documentation to reduce burden.

Correct answer: C

3. When should an AI use case be escalated for leadership or governance review?

- A. Only after a public complaint occurs.
- B. When it affects sensitive data, public communication, resource allocation, enforcement, service access, or public trust.
- C. Only when a vendor requests escalation.
- D. Never, because AI review should remain informal.

Correct answer: B

4. What is weak evidence of completion for a leadership and governance module?

- A. A role matrix, governance worksheet, or decision-ready plan.
- B. A specific risk statement with mitigation steps.
- C. A completed exercise tied to a public health workflow.
- D. A generic statement that AI should be used responsibly without assigning owners or actions.

Correct answer: D

5. What should leaders do when an AI use case has meaningful benefits but unresolved risks?

- A. Require mitigation, human review, monitoring, documentation, or defer implementation until readiness improves.
- B. Approve it immediately because benefits are possible.
- C. Reject all AI use permanently.
- D. Ask the vendor to decide whether risk is acceptable.

Correct answer: A

References and Resources

NIST Artificial Intelligence Risk Management Framework (AI RMF 1.0)

NIST Generative AI Profile

NIST Privacy Framework

CDC Public Health Data Strategy and Data Modernization resources

WHO Ethics and Governance of Artificial Intelligence for Health

OMB Memorandum M-24-10 on federal AI governance, risk management, and innovation, where relevant to public-sector governance models

Agency strategic plans, procurement policies, privacy policies, cybersecurity policies, records policies, and equity plans

GOV 420

Public Accountability and AI Transparency Reporting

Course ID	GOV 420
Track	Governance and Security Track
Level	advanced/leadership
Estimated time	50
Audience	Governance committee members, Legal/privacy staff, Security staff, Compliance staff, Program owners

Learning Objectives

- Explain why transparency reporting matters for public health AI.
- Identify which AI uses should be disclosed internally, publicly, or to specific partners.
- Develop contents for an AI transparency notice or AI use registry.
- Balance transparency with privacy, security, confidentiality, and operational sensitivity.
- Produce a draft AI transparency reporting template.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Why This Topic Matters for Public Health AI

Public health depends on trust. If communities discover that AI is being used in surveillance, communication, service routing, or resource allocation without explanation, they may question whether decisions are fair, accurate, or accountable. Transparency does not require agencies to disclose every technical detail, but it does require clear communication about the purpose and safeguards of consequential AI uses.

Transparency also improves internal accountability. A registry of AI use cases helps leaders know what tools are active, who owns them, what data they use, and when they were last reviewed. It reduces the chance that tools continue operating after their purpose, vendor terms, data sources, or risk profile changes.

Public Health Example

A health department uses AI to draft public health messages, summarize internal reports, and prioritize inspection follow-up. Some uses may not require broad public disclosure, while others may affect public services and should be listed in an AI use registry. The registry does not need to reveal sensitive security details, but it should identify the purpose, owner, data categories, review process, and human oversight.

For a public-facing tool, the agency may also create a plain-language notice explaining that AI assists with initial responses, that staff review high-risk content, that personal information should not be entered unless requested through approved channels, and that users can contact the agency for help or correction. This type of notice supports trust and sets realistic expectations.

Technical and Governance Deep Dive

Transparency reporting should be risk-based. Low-risk internal tools may be documented in an internal registry. Tools that influence public communication, prioritization, access, enforcement, or public-facing interactions may require public notice. High-impact tools may require more detailed documentation, such as intended use, data sources, validation approach, human review, limitations, contact information, and review schedule.

The reporting process should be accessible. Public-facing explanations should use plain language, avoid technical jargon, and be available in appropriate languages and accessible formats. Transparency is not achieved by publishing a dense technical document that affected communities cannot understand. Agencies should consider how communities, partners, and staff will actually interpret the information.

Transparency should also include mechanisms for accountability. A notice or registry should identify how concerns can be raised, how corrections can be requested when appropriate, and how the agency reviews incidents or complaints. Public reporting should connect to internal governance processes so that feedback can lead to action.

Risks, Failure Modes, and Guardrails

One risk is transparency theater. Agencies may publish broad statements about responsible AI without disclosing actual uses, owners, or safeguards. A guardrail is to maintain an AI use registry with concrete information about approved and operational use cases.

Another risk is over-disclosure of sensitive details. Publishing system architecture, vulnerabilities, or confidential data flows could create security or privacy risks. A guardrail is to review transparency materials through privacy, legal, security, and communications before release.

Application to AI-Supported Workflows

Generative AI transparency may explain when AI assists drafting and how human review occurs. Predictive analytics transparency may explain what the model supports, what decisions remain human, and how fairness is monitored. Agentic AI transparency may explain what actions are automated, what approvals are required, and how errors can be reported.

Reflection Questions

Which current or proposed AI use case in your agency raises this leadership or governance issue most clearly?

Who currently has authority to make decisions about this issue, and is that authority documented?

What evidence would governance or leadership need before approving the use case?

Which staff, partners, or communities would be affected if this issue were handled poorly?

What policy, process, or artifact should be created or updated after this module?

Practical Exercise

Create a transparency reporting template and apply it to one AI use case. Include purpose, owner, users, data categories, risk level, safeguards, human review, limitations, review date, and contact pathway.

Before You Begin

Choose a real or realistic public health AI use case. Document assumptions, unresolved questions, and which agency roles would need to review the artifact before it is adopted. The exercise should produce a work product that could support governance review, leadership discussion, implementation planning, or policy development.

Expected Artifact or Evidence

AI transparency template; sample AI use registry entry; public notice draft; internal review checklist.

References and Resources for Additional Information

NIST Artificial Intelligence Risk Management Framework (AI RMF 1.0)

NIST Generative AI Profile

NIST Privacy Framework

CDC Public Health Data Strategy and Data Modernization resources

WHO Ethics and Governance of Artificial Intelligence for Health

OMB Memorandum M-24-10 on federal AI governance, risk management, and innovation, where relevant to public-sector governance models

Agency strategic plans, procurement policies, privacy policies, cybersecurity policies, records policies, and equity plans

Knowledge Check

1. What is the strongest reason this module topic belongs in AI leadership and governance training?

- A. It eliminates the need for technical review.
- B. It helps leaders define authority, safeguards, accountability, and implementation expectations.
- C. It allows vendors to make final public health decisions.
- D. It guarantees that every AI tool will be successful.

Correct answer: B

2. Which practice best supports accountable public health AI governance?

- A. Allowing each program to create informal AI rules independently.
- B. Approving tools based only on vendor demonstrations.
- C. Avoiding documentation to reduce burden.

D. Using defined review steps, named owners, decision logs, and escalation pathways.

Correct answer: D

3. When should an AI use case be escalated for leadership or governance review?

A. Only after a public complaint occurs.

B. Only when a vendor requests escalation.

C. When it affects sensitive data, public communication, resource allocation, enforcement, service access, or public trust.

D. Never, because AI review should remain informal.

Correct answer: C

4. What is weak evidence of completion for a leadership and governance module?

A. A generic statement that AI should be used responsibly without assigning owners or actions.

B. A role matrix, governance worksheet, or decision-ready plan.

C. A specific risk statement with mitigation steps.

D. A completed exercise tied to a public health workflow.

Correct answer: A

5. What should leaders do when an AI use case has meaningful benefits but unresolved risks?

A. Approve it immediately because benefits are possible.

B. Require mitigation, human review, monitoring, documentation, or defer implementation until readiness improves.

C. Reject all AI use permanently.

D. Ask the vendor to decide whether risk is acceptable.

Correct answer: B

References and Resources

NIST Artificial Intelligence Risk Management Framework (AI RMF 1.0)

NIST Generative AI Profile

NIST Privacy Framework

CDC Public Health Data Strategy and Data Modernization resources

WHO Ethics and Governance of Artificial Intelligence for Health

OMB Memorandum M-24-10 on federal AI governance, risk management, and innovation, where relevant to public-sector governance models

Agency strategic plans, procurement policies, privacy policies, cybersecurity policies, records policies, and equity plans

GOV 430

Community Oversight and Participatory Governance

Course ID	GOV 430
Track	Governance and Security Track
Level	advanced/leadership
Estimated time	50
Audience	Governance committee members, Legal/privacy staff, Security staff, Compliance staff, Program owners

Learning Objectives

- Explain why community oversight matters for public health AI.
- Identify AI use cases that warrant community engagement or participatory governance.
- Design a community engagement approach matched to risk, impact, and community context.
- Recognize risks of tokenism, extractive engagement, and inaccessible technical language.
- Produce a participatory governance plan for one AI use case.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Why This Topic Matters for Public Health AI

Public health agencies make decisions that affect communities, often using data collected through clinical care, surveillance, service systems, public programs, and emergency response. AI can intensify concerns about surveillance, bias, exclusion, language access, disability access, and government trust. Community oversight helps agencies identify concerns that may not be visible through technical review alone.

Participatory governance can also improve AI quality and usefulness. Community partners may identify missing context, harmful assumptions, access barriers, communication needs, and unintended consequences. Their input can help agencies design safer workflows, clearer notices, better escalation pathways, and more meaningful impact measures.

Public Health Example

A city health department considers using an AI tool to help prioritize outreach for maternal and child health services. The model would use program, demographic, and service utilization data to identify neighborhoods where outreach may be needed. Although the goal is supportive, communities may worry that the tool could stigmatize neighborhoods, miss families not represented in data, or direct resources away from groups with less complete records.

A participatory governance plan would involve community-based organizations, maternal health advocates, service users, and equity staff before the tool is finalized. The agency would explain the intended use, data limitations, safeguards, and decision points in plain language. Community input could shape the variables used, how results are communicated, what services are offered, and how harm or misclassification concerns are handled.

Technical and Governance Deep Dive

Participatory governance should be matched to the impact of the use case. Internal productivity tools may not require community oversight, although they still require policy and privacy review. Tools that influence public communication, service access, enforcement, prioritization, outreach, or resource allocation may require deeper engagement. The higher the impact on communities, the stronger the case for involving community perspectives early.

Engagement should be designed with care. Agencies should provide plain-language descriptions, realistic examples, data-use explanations, and clear questions for feedback. Meetings should be accessible, compensated when appropriate, linguistically inclusive, and scheduled in ways that support participation. Agencies should also explain what decisions are open for influence and what constraints exist because of law, funding, or public health authority.

Follow-through is essential. Community members should know how their input was used, what changed because of engagement, and how concerns will be monitored after launch. Participatory governance is not a single listening session. It is an ongoing relationship between the agency and affected communities, especially when the AI use is consequential or long-term.

Risks, Failure Modes, and Guardrails

One risk is tokenistic engagement. Agencies may seek community input after key decisions have already been made. A guardrail is to engage before requirements, safeguards, and success measures are finalized.

Another risk is technical opacity. Community participants may be asked to respond to vague or overly technical descriptions. A guardrail is to provide plain-language examples, visual workflows, and concrete descriptions of what the AI will and will not do.

Application to AI-Supported Workflows

Generative AI tools that produce public messages may benefit from community review for tone, trust, accessibility, and language. Predictive models used for outreach or allocation should include review of data representativeness and potential stigma. Agentic tools that trigger actions affecting residents should have clear community-informed escalation and correction pathways.

Reflection Questions

Which current or proposed AI use case in your agency raises this leadership or governance issue most clearly?

Who currently has authority to make decisions about this issue, and is that authority documented?

What evidence would governance or leadership need before approving the use case?

Which staff, partners, or communities would be affected if this issue were handled poorly?

What policy, process, or artifact should be created or updated after this module?

Practical Exercise

Develop a participatory governance plan for one AI use case. Include affected groups, engagement purpose, timing, materials, questions, accessibility supports, decision influence, feedback loop, and monitoring plan.

Before You Begin

Choose a real or realistic public health AI use case. Document assumptions, unresolved questions, and which agency roles would need to review the artifact before it is adopted. The exercise should produce a work product that could support governance review, leadership discussion, implementation planning, or policy development.

Expected Artifact or Evidence

Participatory governance plan; community briefing outline; engagement questions; feedback response plan.

References and Resources for Additional Information

NIST Artificial Intelligence Risk Management Framework (AI RMF 1.0)

NIST Generative AI Profile

NIST Privacy Framework

CDC Public Health Data Strategy and Data Modernization resources

WHO Ethics and Governance of Artificial Intelligence for Health

OMB Memorandum M-24-10 on federal AI governance, risk management, and innovation, where relevant to public-sector governance models

Agency strategic plans, procurement policies, privacy policies, cybersecurity policies, records policies, and equity plans

Knowledge Check

1. What is the strongest reason this module topic belongs in AI leadership and governance training?

- A. It eliminates the need for technical review.
- B. It allows vendors to make final public health decisions.
- C. It helps leaders define authority, safeguards, accountability, and implementation expectations.
- D. It guarantees that every AI tool will be successful.

Correct answer: C

2. Which practice best supports accountable public health AI governance?

- A. Allowing each program to create informal AI rules independently.
- B. Using defined review steps, named owners, decision logs, and escalation pathways.
- C. Approving tools based only on vendor demonstrations.

D. Avoiding documentation to reduce burden.

Correct answer: B

3. When should an AI use case be escalated for leadership or governance review?

A. Only after a public complaint occurs.

B. Only when a vendor requests escalation.

C. Never, because AI review should remain informal.

D. When it affects sensitive data, public communication, resource allocation, enforcement, service access, or public trust.

Correct answer: D

4. What is weak evidence of completion for a leadership and governance module?

A. A generic statement that AI should be used responsibly without assigning owners or actions.

B. A role matrix, governance worksheet, or decision-ready plan.

C. A specific risk statement with mitigation steps.

D. A completed exercise tied to a public health workflow.

Correct answer: A

5. What should leaders do when an AI use case has meaningful benefits but unresolved risks?

A. Approve it immediately because benefits are possible.

B. Reject all AI use permanently.

C. Require mitigation, human review, monitoring, documentation, or defer implementation until readiness improves.

D. Ask the vendor to decide whether risk is acceptable.

Correct answer: C

References and Resources

NIST Artificial Intelligence Risk Management Framework (AI RMF 1.0)

NIST Generative AI Profile

NIST Privacy Framework

CDC Public Health Data Strategy and Data Modernization resources

WHO Ethics and Governance of Artificial Intelligence for Health

OMB Memorandum M-24-10 on federal AI governance, risk management, and innovation, where relevant to public-sector governance models

Agency strategic plans, procurement policies, privacy policies, cybersecurity policies, records policies, and equity plans

INT 190

AI for Social Good

Course ID	INT 190
Track	Shared Foundational / Introductory Course
Level	introductory
Estimated time	30
Audience	All public health staff, Supervisors, Program staff, Governance participants

Learning Objectives

- Define AI for Social Good in public health terms.
- Distinguish public benefit from novelty, efficiency, or automation alone.
- Identify how AI can support public health values, including equity, transparency, accessibility, accountability, and community trust.
- Assess potential harms, including exclusion, surveillance burden, stigma, resource diversion, and unequal benefit.
- Design community-centered safeguards for socially beneficial AI use.
- Produce an AI for Social Good review brief for a public health use case.

Why This Topic Matters for Public Health AI

Public health is uniquely positioned to use AI for social good because its mission is population well-being rather than private gain. AI can help agencies identify emerging needs, improve access to information, reduce administrative burden, support community outreach, and allocate resources more effectively. These benefits are especially important when agencies face workforce shortages, fragmented systems, and increasing demands.

However, AI for Social Good is not guaranteed by good intentions. Tools built for efficiency may shift burden onto communities, reduce human contact, or make services harder to access for people with limited digital literacy. Predictive tools may label communities as risky without addressing structural causes. Chatbots may improve information access for some users while excluding people who need language, disability, or offline support.

The public health standard for social good should be rigorous. Agencies should define the public problem, identify who benefits, identify who could be harmed, involve affected communities, measure outcomes, and document safeguards. The central question is not “Can AI do this?” but “Should AI be used for this purpose, and under what conditions would it create public value?”

Public Health Example

A health department may use AI to improve access to maternal and child health resources. A tool could help staff identify service gaps, summarize eligibility rules, translate plain-language information, and route families to local support. If designed well, the tool could reduce administrative burden and improve timely access to services.

The same tool could fail if it assumes internet access, uses only English content, provides outdated referral information, exposes sensitive information, or replaces trusted human navigators. An AI for Social Good review would examine community needs, access barriers, language and disability accommodations, privacy, human fallback options, and

measurable outcomes.

Technical Deep Dive

AI for Social Good begins with problem definition. The problem should be a public health problem, not a technology opportunity. For example, “use a chatbot” is not a problem definition. “Reduce delays in connecting families to local nutrition, housing, transportation, and home visiting resources” is a stronger public health problem definition. Starting with the problem helps agencies avoid technology-driven solutions that do not address root needs.

Equity by design requires agencies to identify who is included, who is excluded, and who bears risk. This includes reviewing data representativeness, digital access, language access, disability access, rural access, immigration-related concerns, privacy expectations, and historical mistrust. Equity review should shape the use case, data sources, interface, outreach strategy, and evaluation metrics.

Community engagement should occur before deployment, not only after concerns arise. Affected communities can help identify priorities, unacceptable uses, preferred communication channels, needed safeguards, and measures of success. Engagement does not mean asking communities to approve a predetermined tool. It means allowing community input to change the design, scope, or decision to proceed.

Measuring social good requires more than counting users or outputs. Agencies should measure whether the AI-supported workflow improves timeliness, access, quality, fairness, staff burden, community trust, and public health outcomes. They should also monitor unintended consequences, such as reduced service quality, increased complaints, inequitable uptake, or inaccurate referrals.

Responsible innovation includes sunset and rollback planning. If an AI tool does not produce expected benefit, creates harm, or becomes outdated, the agency should be prepared to modify, pause, or retire it. Social good depends on ongoing accountability, not one-time approval.

Risks, Failure Modes, and Guardrails

Technology-first design. A tool may be adopted because it is novel rather than because it solves a public health problem. Guardrails include problem-first planning, alternatives analysis, and public value criteria.

Unequal benefit. AI may help people with strong digital access while excluding others. Guardrails include language access, accessibility, offline options, and human fallback.

Stigma and surveillance burden. AI may label communities or increase monitoring without support. Guardrails include equity framing, community engagement, and supportive-use restrictions.

Resource diversion. AI projects may consume resources that could be used for more effective interventions. Guardrails include impact measurement, cost-benefit review, and periodic reassessment.

Application to AI-Supported Workflows

AI for Social Good applies to every AI method. Generative AI may improve plain-language communication, but it must be accurate, accessible, and culturally appropriate. Predictive analytics may guide supportive outreach, but it should not become a tool for exclusion. RAG systems may improve staff access to guidance, but they must include current and authorized content. Agentic workflows may reduce burden, but they require boundaries and human accountability.

The strongest public health AI projects define social good in operational terms. They identify the population benefit, the equity goal, the workflow change, the safeguards, and the measurement plan. They also acknowledge tradeoffs and document why AI is an appropriate tool compared with non-AI alternatives.

Reflection Questions

What public health problem would this AI use address?

Who would benefit, and who could be excluded or harmed?

How were affected communities involved in defining the need and safeguards?

What non-AI alternatives should be considered?

How will the agency measure public value, equity, trust, and unintended consequences?

Practical Exercise

Create an AI for Social Good review brief for one public health use case. The brief should define the public health problem, intended public benefit, affected communities, equity considerations, non-AI alternatives, proposed AI role, safeguards, community engagement plan, evaluation measures, and sunset or rollback criteria.

The exercise should result in a decision-support document that helps leaders and governance reviewers determine whether the proposed AI use advances public health values.

Expected Artifact or Evidence

AI for Social Good review brief

Public value and equity assessment

Community engagement and accessibility plan

Impact and unintended-consequence measurement plan

References and Resources for Additional Information

WHO Ethics and Governance of Artificial Intelligence for Health: <https://www.who.int/publications/i/item/9789240029200>

WHO Guidance on Large Multi-Modal Models: <https://www.who.int/publications/i/item/9789240084759>

NIST Artificial Intelligence Risk Management Framework: <https://www.nist.gov/itl/ai-risk-management-framework>

NIST AI RMF Generative AI Profile:
<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>

CDC Public Health Data Strategy: <https://www.cdc.gov/public-health-data-strategy/>

Gap-Filling Modules: Core Track Gaps

Source file: AI_Gap_Filling_Training_Modules_for_Public_Health(1).docx

Gap-Filling AI Training Modules for Public Health

Recommended modules to complete shared, technical, analytics, operations, governance, and role-based learning tracks

Purpose. This document contains ten additional training modules recommended to fill gaps across the learning tracks. The modules are designed to complement the existing foundational, technical, analytics, operations, governance, communications, epidemiology, policy, executive leadership, and program management content. Each module follows the established training format with paragraph-based instruction, applied public health examples, practical exercises, expected evidence, and knowledge checks with varied answer positions.

Module List and Track Alignment

Knowledge Check

1. What is the best definition of AI for Social Good in public health?

- A. Using AI because it is new.
- B. Using AI intentionally to create measurable public benefit, reduce inequities, and support well-being.
- C. Replacing public health staff with automation.
- D. Using any free AI tool for public programs.

Correct answer: B

2. Why is efficiency alone not enough to show social good?

- A. Efficiency may not improve equity, access, trust, or outcomes and may shift burden onto communities.
- B. Efficiency always proves public value.
- C. Efficiency removes the need for community engagement.
- D. Efficiency prevents unintended consequences.

Correct answer: A

3. What does equity by design require?

- A. Reviewing equity only after launch.
- B. Avoiding community input.
- C. Identifying inequities, affected communities, barriers, and safeguards early in design.
- D. Using the same design for every population.

Correct answer: C

4. Which use best reflects AI for Social Good?

- A. A tool that speeds decisions by removing appeals.
- B. A tool that labels communities without providing support.
- C. A tool that replaces language access services.

D. A tool that improves access to services with community input, human fallback, privacy safeguards, and outcome monitoring.

Correct answer: D

5. What is strong completion evidence for this module?

A. An AI for Social Good review brief with public benefit, equity, community engagement, safeguards, and measurement.

B. A vendor brochure.

C. A generic statement that AI is helpful.

D. A list of trendy use cases.

Correct answer: A

References and Resources

WHO Ethics and Governance of Artificial Intelligence for Health

<https://www.who.int/publications/i/item/9789240029200>

WHO Guidance on Large Multi-Modal Models

<https://www.who.int/publications/i/item/9789240084759>

NIST Artificial Intelligence Risk Management Framework

<https://www.nist.gov/itl/ai-risk-management-framework>

CDC Public Health Data Strategy

<https://www.cdc.gov/public-health-data-strategy/php/about/index.html>

Agency equity, community engagement, accessibility, and language access policies

INT 200

AI Use Case Intake and Triage Basics

Course ID	INT 200
Track	Shared Foundational / Introductory Course
Level	applied/foundational
Estimated time	30
Audience	All public health staff, Supervisors, Program staff, Governance participants

Learning Objectives

- Describe an AI use case in plain language using the public health workflow, users, data, intended output, and decision context.
- Distinguish between low-risk, moderate-risk, and high-risk AI use cases based on data sensitivity, workflow impact, public-facing use, automation level, and potential harm.
- Identify when a use case should be routed to privacy, legal, IT, cybersecurity, equity, procurement, communications, or executive review.
- Recognize use cases that should not proceed because AI is unnecessary, inappropriate, unsafe, or unsupported by available data.
- Produce a use case intake and triage worksheet that can support governance review and portfolio prioritization.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Public Health Example

A communicable disease program wants to use generative AI to summarize case investigation notes and identify missing information. Intake would document the workflow, the staff who would use the output, the source data, whether protected health information is involved, whether the output would be saved in the system of record, and whether the tool would only summarize or also recommend action. Triage would likely route the use case to privacy, legal, IT, epidemiology, equity, and governance review before any pilot begins.

Technical and Operational Deep Dive

An effective intake process begins by defining the problem before selecting the technology. Staff should describe the current workflow, the pain point, the people affected, the desired improvement, and the reason AI is being considered. This prevents technology-first proposals that begin with a tool and then search for a public health problem to justify it. A clear problem statement also helps reviewers decide whether AI is needed or whether a simpler change, such as a template, dashboard, workflow redesign, or rules-based automation, would be safer and more appropriate.

The next step is to identify the data and output. Reviewers need to know whether the use case involves protected health information, personally identifiable information, confidential public health records, public documents, internal policies, vendor data, community feedback, or synthetic data. They also need to know what the AI output will be: a draft, summary, classification, prediction, recommendation, score, translation, alert, or automated action. Different outputs create different risks and require different levels of review.

Triage should evaluate both impact and readiness. Impact includes the potential effect on people, programs, public trust, legal obligations, and resource allocation. Readiness includes data quality, system integration, staff capacity, policy alignment, validation feasibility, and operational ownership. A high-impact use case with low readiness may not be rejected permanently, but it should not move into implementation until the readiness gaps are addressed.

Use case intake also supports portfolio management. When agencies collect use cases in a consistent format, leaders can compare proposals across programs, identify common infrastructure needs, avoid duplicative procurements, and align AI investments with strategic priorities. The intake record becomes the first governance artifact in the AI lifecycle.

Risks, Failure Modes, and Guardrails

Risks and failure modes include the following:

Informal AI experimentation with sensitive or confidential data before review.

Tool-first proposals that do not solve a real workflow problem.

Use cases moving forward without ownership, validation, or monitoring plans.

High-risk public-facing or decision-support uses being treated as low-risk productivity tools.

Equity, accessibility, and community impacts being considered too late.

Recommended guardrails include the following:

Require a standard intake worksheet before piloting or procuring AI tools.

Classify proposed uses by data sensitivity, automation level, public-facing impact, and decision consequence.

Route use cases to required reviewers before implementation begins.

Document use cases that are rejected, deferred, or approved with conditions.

Maintain a portfolio view of approved, active, paused, and retired AI uses.

Reflection Questions

Where does this topic appear in your agency, program, or workflow?

Which staff roles would need to understand or apply this concept before AI use is approved?

What could go wrong if this topic is not addressed before implementation?

What evidence would governance, leadership, or operations staff need before moving forward?

Which policy, data, equity, privacy, security, or workflow questions remain unresolved?

Practical Exercise

Complete a use case intake and triage worksheet for one real or realistic public health workflow. Include the problem statement, workflow, intended users, data sources, AI-supported task, output, potential benefit, risk level, required reviewers, readiness gaps, and recommendation to proceed, defer, redesign, or reject.

Expected Artifact or Evidence

Completed AI use case intake worksheet

Risk and readiness triage decision

List of required reviewers and next steps

Documented rationale for proceeding, deferring, redesigning, or rejecting the use case

References and Resources for Additional Information

NIST Artificial Intelligence Risk Management Framework

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP guidance for LLM application security

Agency policies for privacy, cybersecurity, records, procurement, accessibility, and public communications

Relevant public health program SOPs, data governance documentation, and implementation plans

Knowledge Check

1. What is the main purpose of AI use case intake?

- A. To collect standard information before an AI idea is reviewed or approved
- B. To let vendors define the agency workflow
- C. To replace privacy and IT review
- D. To approve all AI ideas quickly

Correct answer: A

2. Which proposal should usually receive higher-risk triage?

- A. A public-domain glossary rewrite
- B. A staff brainstorming exercise using no data

- C. A draft meeting agenda
- D. An AI tool that recommends which cases receive priority field investigation

Correct answer: D

3. What should happen when a proposed workflow problem can be solved without AI?

- A. The agency should consider a simpler, safer solution before proceeding with AI
- B. The agency should deploy AI anyway
- C. The agency should skip documentation
- D. The agency should treat the use as high-value automatically

Correct answer: A

4. Which item belongs in an intake worksheet?

- A. The staff cafeteria menu
- B. The intended users, data sources, output, risk level, and review pathway
- C. A general statement that AI is innovative
- D. A vendor logo

Correct answer: B

5. What is strong evidence of completion?

- A. Opening an AI tool once
- B. A list of popular AI products
- C. A completed use case intake and triage worksheet with rationale and required next steps
- D. A generic AI definition

Correct answer: C

References and Resources

NIST Artificial Intelligence Risk Management Framework

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP guidance for LLM application security

Agency policies for privacy, cybersecurity, records, procurement, accessibility, and public communications

Relevant public health program SOPs, data governance documentation, and implementation plans

OPS 200

Data Quality Engineering for AI

Course ID	OPS 200
Track	Operations and Data Quality Track
Level	applied/foundational
Estimated time	45
Audience	Operations leads, Data quality staff, Program managers, Implementation teams, Supervisors

Learning Objectives

- Identify major data quality dimensions relevant to public health AI.
- Explain the difference between general data quality and fitness for a specific use.
- Assess how missingness, timeliness, duplication, coding variation, and representativeness affect AI outputs.
- Define quality thresholds, review rules, and escalation criteria for a workflow.
- Produce a data quality review plan for an AI-supported use case.

Why This Topic Matters for Public Health AI

AI systems can create a false sense of precision when they produce polished outputs from weak data. A summary may sound authoritative even when key fields are missing. A risk score may look objective even when the data underrepresent communities with limited healthcare access. A classification model may appear efficient while encoding historical reporting patterns. Data quality engineering helps agencies identify these problems before they become operational decisions.

Public health data quality should always be evaluated in relation to the intended use. A dataset may be adequate for annual descriptive reporting but inadequate for real-time triage. A field may be useful for aggregate trends but too incomplete for individual-level decision support. Fitness for use is therefore a public health judgment as well as a technical assessment.

Data quality review also supports trust. Staff are more likely to use AI responsibly when they understand the limitations of the data behind the output. Leaders and communities are more likely to trust AI-supported workflows when limitations, uncertainty, and safeguards are documented.

Core Concepts

Completeness. Completeness refers to whether required data fields are present. Missing fields may affect case classification, risk prediction, subgroup analysis, and outreach planning.

Timeliness. Timeliness refers to whether data are available quickly enough for the intended use. Delayed data may be acceptable for retrospective analysis but unsafe for urgent response.

Validity. Validity refers to whether values follow expected formats, ranges, and rules. Invalid dates, codes, units, or categories can undermine AI performance.

Representativeness. Representativeness refers to whether data reflect the population, events, or settings relevant to the use case. Underrepresentation can produce biased or incomplete AI outputs.

Data drift. Data drift is a change in data patterns over time. Drift can occur when reporting partners, testing practices, disease patterns, coding rules, or community behaviors change.

Public Health Example

A health department may use AI to identify cases needing rapid follow-up. If demographic fields are missing more often for some populations, the model may be less reliable for those groups. If the data feed from certain clinics is delayed, the prioritization queue may reflect reporting timeliness rather than true need. A data quality review can reveal these issues before the tool is deployed.

A community health assessment team may use AI to summarize survey comments and administrative data. If survey responses are not representative of residents with limited English proficiency, unstable housing, disability, or limited internet access, the AI-generated themes may miss important community needs. Data quality engineering should include equity and representativeness review.

Technical Deep Dive

Data quality engineering begins by defining required fields and acceptable thresholds for the intended use. For an AI-supported case triage workflow, required fields may include condition, lab evidence, onset date, specimen date, jurisdiction, age, pregnancy status, hospitalization, and contact information. The agency should decide which missing fields stop automated support, which trigger a warning, and which can be handled through human review.

Quality measurement should include completeness, timeliness, validity, consistency, uniqueness, accuracy, and representativeness. Some dimensions can be measured automatically. Others require program knowledge, source review, or comparison to external benchmarks. Data quality should not be treated as a one-time predeployment activity. It should be monitored throughout the AI lifecycle.

Documentation is essential. A data quality plan should describe known limitations, review thresholds, monitoring frequency, responsible owners, escalation criteria, and decisions about acceptable use. When data quality is insufficient, the correct response may be to restrict the use case, add human review, improve the data pipeline, delay deployment, or avoid AI use for that workflow.

Risks, Failure Modes, and Guardrails

Confident output from weak data. AI may produce polished outputs from incomplete or biased data. Guardrails include uncertainty labeling, source review, and limits on use.

Fitness-for-use mismatch. Data suitable for one purpose may be unsafe for another. Guardrails include use-case-specific quality criteria and approval conditions.

Unmonitored drift. Data patterns may change after deployment. Guardrails include drift indicators, periodic review, and retraining or recalibration triggers.

Equity blind spots. Missingness and underrepresentation may affect some groups more than others. Guardrails include subgroup review, community context, and documentation of limitations.

Application to AI-Supported Workflows

Generative AI may summarize datasets or records, but the output must reflect missingness and uncertainty. Predictive analytics depend heavily on feature quality, label quality, and representative training data. NLP may be affected by language variation, abbreviations, and documentation practices. RAG depends on document quality, metadata, and currency. Agentic AI should be blocked from acting automatically when required data are missing or quality thresholds are not met.

Reflection Questions

Which data quality problems most often affect your program area?

Which missing fields would make an AI output unsafe or unreliable?

Which populations or settings may be underrepresented in the data?

How would staff detect data drift?

Who has authority to pause or limit AI use when data quality is insufficient?

Practical Exercise

Create a data quality review plan for one AI-supported workflow. Include the use case, required fields, quality dimensions, thresholds or review rules, known limitations, equity concerns, drift monitoring, documentation expectations, and escalation criteria.

Expected Artifact or Evidence

Data quality review plan

Data quality checklist

Thresholds or review criteria

Known limitations statement

References and Resources for Additional Information

CDC surveillance and data quality resources

CDC Public Health Data Strategy

NIST AI Risk Management Framework

Agency data governance policies

Knowledge Check

1. What does fitness for use mean?

- A. Data are perfect for every possible use.
- B. Data can be used without documentation.
- C. Data quality only matters to IT.
- D. Data are appropriate for a specific purpose.

Correct answer: D

2. Why can poor data quality be dangerous in AI-supported workflows?

- A. AI refuses to run with poor data.
- B. AI may produce confident outputs from incomplete, biased, or inaccurate data.
- C. Poor data quality always appears clearly on screen.
- D. Human review is unnecessary.

Correct answer: B

3. What is data drift?

- A. A change in data patterns over time that may affect AI performance.
- B. A type of meeting agenda.
- C. A procurement method.
- D. A privacy notice.

Correct answer: A

4. Which item belongs in a data quality review plan?

- A. Staff birthdays.
- B. Vendor slogans.
- C. Missingness, timeliness, duplicates, coding consistency, and known limitations.
- D. General AI definitions only.

Correct answer: C

5. When should an AI use case be paused or limited?

- A. Only when a vendor agrees.
- B. When data quality problems may materially affect safety, accuracy, equity, or public health action.
- C. Never.
- D. Only after deployment is complete.

Correct answer: B

References and Resources

CDC surveillance and data quality resources

CDC Public Health Data Strategy

NIST AI Risk Management Framework

Agency data governance policies

OPS 210

Record Linkage, Entity Resolution, and Deduplication

Course ID	OPS 210
Track	Operations and Data Quality Track
Level	applied/foundational
Estimated time	45
Audience	Operations leads, Data quality staff, Program managers, Implementation teams, Supervisors

Learning Objectives

- Explain why record linkage matters for public health AI.
- Distinguish between deterministic matching, probabilistic matching, entity resolution, and deduplication.
- Identify risks from false matches and missed matches.
- Describe how thresholds, clerical review, and data quality affect matching.
- Produce a linkage risk review for an AI-supported workflow.

Why This Topic Matters for Public Health AI

Record linkage matters because public health work often depends on combining information across systems. A surveillance system may need to associate a lab report with an existing case. An IIS may need to match vaccination records from multiple providers. A vital records linkage may support maternal and child health analysis. An AI tool using these linked records may summarize history, predict risk, classify cases, or identify duplicates. If linkage is wrong, the AI output may be wrong.

False matches and missed matches have different consequences. A false match can combine information from different people or events, creating privacy, safety, and accuracy risks. A missed match can split one person or event into multiple records, leading to duplicate work, incomplete histories, or undercounting. Both errors can affect public health decisions.

Linkage can also create equity risks. Names, addresses, phone numbers, language, housing status, migration patterns, and documentation practices vary across communities. Matching rules that work well for one population may perform poorly for another. AI workflows that rely on linked data should therefore include subgroup and equity review.

Core Concepts

Record linkage. Record linkage is the process of determining whether records from one or more sources refer to the same entity. The entity may be a person, provider, facility, address, organization, event, or outbreak.

Entity resolution. Entity resolution is the broader process of identifying, matching, merging, separating, and managing records about the same entity. It includes matching rules, thresholds, review processes, and stewardship.

Deterministic matching. Deterministic matching uses exact or rule-based agreement across selected fields. It is easier to explain but may miss matches when fields are incomplete or inconsistent.

Probabilistic matching. Probabilistic matching estimates the likelihood that records refer to the same entity even when fields differ. It can improve recall but requires thresholds, validation, and review.

Clerical review. Clerical review is human review of uncertain matches. It is essential when automatic matching is not reliable enough for the downstream use.

Public Health Example

An AI-supported case summary may pull together lab reports, case notes, hospitalization information, and prior reports for the same person. If the person matching process incorrectly combines two people, the summary may include information that belongs to someone else. If the process misses a prior record, the summary may omit important history. The investigator needs visibility into linkage confidence and source records.

An outbreak detection workflow may link complaints, lab results, facility inspections, and geographic data. Entity resolution may be needed to identify whether different facility names refer to the same establishment. AI-generated cluster summaries can be misleading if the facility resolution process is weak.

Technical Deep Dive

Matching design begins with the entity being matched and the decision that depends on the match. Person matching may use name, date of birth, address, phone, medical record number, email, sex, or other identifiers. Facility matching may use name, address, license number, geocode, and ownership information. Event matching may use dates, locations, condition, exposure, and related records.

Thresholds determine when records are automatically linked, automatically separated, or sent for review. Low thresholds may increase false matches. High thresholds may increase missed matches. The right balance depends on the use case. A workflow that affects individual services may need a different tolerance for error than an aggregate trend analysis.

AI readiness requires documentation of matching methods, fields, thresholds, review processes, error rates, and known limitations. If an AI tool uses linked records, staff should know whether linkage is deterministic, probabilistic, manual, or hybrid. They should also know whether uncertain matches are visible to users and whether AI outputs indicate linkage uncertainty.

Risks, Failure Modes, and Guardrails

False matches. False matches can combine data from different people or entities. Guardrails include conservative thresholds for consequential uses, clerical review, and source display.

Missed matches. Missed matches can fragment histories and create duplicate work. Guardrails include matching quality monitoring, feedback loops, and review queues.

Equity variation. Matching may perform differently across populations. Guardrails include subgroup evaluation, alternative identifiers, and review of naming and address patterns.

Hidden linkage uncertainty. AI outputs may present linked data as certain. Guardrails include confidence indicators, source records, and user review of uncertain matches.

Application to AI-Supported Workflows

Generative AI may summarize linked records, so source attribution and linkage confidence should be visible. Predictive analytics may use features created from linked records, making match quality part of model validity. NLP may extract identifiers from text to support linkage, which requires privacy and accuracy review. Agentic AI should not automatically merge or update records without strong approval and rollback procedures.

Reflection Questions

Which workflows in your agency depend on matching records across systems?

What fields are used for matching, and how complete are they?

Which populations may be more likely to experience missed or incorrect matches?

What happens when records are incorrectly merged or not merged?

Where should human review be required?

Practical Exercise

Create a linkage risk review for one AI-supported public health workflow. Include the matching purpose, data sources, matching fields, known field quality issues, match method, thresholds, clerical review process, false match risks, missed match risks, equity considerations, and AI output affected.

Expected Artifact or Evidence

Linkage risk review

Deduplication workflow map

Match review questions

Equity and escalation notes

References and Resources for Additional Information

Public health IIS and surveillance data quality resources

CDC data modernization resources

Health information exchange data quality materials

Agency matching and deduplication procedures

Knowledge Check

1. What is a false match?

- A. Two records are correctly separated.
- B. Two records are incorrectly linked.
- C. A record is updated on time.
- D. A model is retrained.

Correct answer: B

2. What is a missed match?

- A. A duplicate record is deleted correctly.
- B. A public report is published.
- C. Records that belong together are not linked.
- D. A dashboard refreshes.

Correct answer: C

3. Why does record linkage matter for AI?

- A. AI never uses linked records.
- B. Linkage eliminates all bias.
- C. Linkage only matters for finance.
- D. AI outputs may rely on linked records to summarize, classify, predict, or route information.

Correct answer: D

4. Who may need to review uncertain matches?

- A. Trained staff using a clerical review process.
- B. No one.
- C. The public.
- D. The AI model alone.

Correct answer: A

5. What is strong completion evidence for this module?

- A. A list of unrelated acronyms.
- B. A generic vendor brochure.
- C. A linkage risk review with false match, missed match, threshold, and equity considerations.
- D. A meeting title.

Correct answer: C

References and Resources

Public health IIS and surveillance data quality resources

CDC data modernization resources

Health information exchange data quality materials

Agency matching and deduplication procedures

OPS 220

Human Factors, Workflow Integration, and Usability Testing

Course ID	OPS 220
Track	Operations and Data Quality Track
Level	applied/foundational
Estimated time	45
Audience	Operations leads, Data quality staff, Program managers, Implementation teams, Supervisors

Learning Objectives

- Explain the purpose of human factors, workflow integration, and usability testing in public health AI work.
- Identify the technical, operational, privacy, security, equity, and governance issues associated with the topic.
- Apply the topic to a real or realistic public health workflow, system, or implementation scenario.
- Recognize common failure modes that can affect safety, accuracy, trust, workflow fit, or sustainability.
- Identify practical guardrails that should be documented before implementation.
- Produce a workflow usability test plan that can support readiness assessment, governance review, procurement, implementation, or monitoring.

Why This Topic Matters for Public Health AI

Human Factors, Workflow Integration, and Usability Testing matters because AI systems are embedded in public health programs, data systems, and decision processes rather than operating in isolation. The topic affects whether AI outputs are accurate, explainable, safe, equitable, secure, and sustainable. Agencies that ignore this area may create tools that work in a demonstration but fail in actual practice.

The topic also affects accountability. Public health agencies must be able to explain how data are used, why a tool was approved, what evidence supports the workflow, who reviews outputs, and what happens when the system fails. These responsibilities become more important as AI tools move from experimentation into operational use.

Core Concepts

Human factors is the discipline of designing systems around human abilities, limits, environments, and tasks. In AI-supported public health work, human factors help agencies ensure that AI outputs are understandable, actionable, and reviewed appropriately.

Usability testing is structured observation of how users interact with a tool or workflow. It can reveal confusion, workarounds, over-trust, alert fatigue, documentation burden, or steps that are inconsistent with actual practice. Usability testing is a safety activity, not simply a design preference.

Cognitive load refers to the mental effort required to complete a task. AI can reduce burden by summarizing information, but it can also increase burden if staff must verify unclear outputs, interpret unexplained scores, or resolve conflicting recommendations.

Public Health Example

A disease investigator may receive an AI-generated case summary in the surveillance system. If the summary hides uncertainty or makes unsupported claims look official, the investigator may over-trust it. If the summary is too long, poorly organized, or disconnected from source fields, it may not reduce burden. Usability testing would observe investigators using the summary and identify changes needed before launch.

Technical Deep Dive

Workflow integration begins with current-state mapping. Agencies should document how work happens today, including informal workarounds, duplicate entry, handoffs, delays, and decision points. AI should be introduced only after the team understands which step is being improved and how the change affects users.

Future-state design should define how staff will see, review, correct, accept, reject, or escalate AI outputs. The user interface should make source information available, show uncertainty when appropriate, and avoid presenting AI outputs as final decisions. Good design supports human judgment rather than bypassing it.

Usability testing should include realistic users, realistic tasks, and realistic data. Testing should evaluate not only whether users can complete a task, but also whether they understand the AI output, recognize limitations, know when to escalate, and can document their final decision.

Risks, Failure Modes, and Guardrails

A common failure mode is treating the topic as a technical detail rather than a public health control. When staff do not connect technical decisions to authority, workflow, equity, privacy, security, and accountability, the agency may adopt a tool that appears useful but cannot be sustained or defended.

Another risk is relying on vendor claims without agency-defined evidence. Vendors may provide demonstrations, dashboards, or summary metrics that do not match the agency workflow or data environment. A guardrail is to require documentation, test results, acceptance criteria, and agency-controlled review.

A third risk is failing to plan for change. Public health data, policies, systems, and emergencies change. Any technical approach must include monitoring, review, versioning, escalation, and retirement pathways so the agency can respond when assumptions no longer hold.

Application to AI-Supported Workflows

Generative AI, predictive analytics, RAG, automation, and agentic workflows all depend on the controls described in this module. The controls should be scaled to the risk of the use case. A tool that drafts internal training text may require lighter review than a tool that updates case records, prioritizes outreach, or supports public communication during an emergency.

The module should be applied during readiness assessment, procurement, implementation planning, pilot testing, and post-deployment monitoring. Learners should document how the topic affects data, systems, users, safeguards, and decision-making for the selected workflow.

Reflection Questions

Where does this topic appear in one current or proposed AI workflow in your agency?

What decisions, data sources, users, or partners would be affected if this area is poorly designed?

What evidence would governance reviewers need before approving the workflow?

Which safeguards should be required before pilot testing?

Who owns ongoing monitoring and review?

Practical Exercise

Create a workflow usability test plan for one real or realistic public health AI workflow. The artifact should describe the use case, affected users, data sources, system dependencies, risks, guardrails, review points, and unresolved questions.

The exercise should produce a work product that could be used in a governance meeting, procurement review, implementation planning session, pilot evaluation, or monitoring review. Learners should document assumptions and identify which staff or partners need to be consulted.

Expected Artifact or Evidence

The expected artifact is a completed workflow usability test plan. It should be specific to a public health workflow and should include technical dependencies, governance questions, human review expectations, monitoring indicators, and unresolved risks.

References and Resources for Additional Information

NIST AI Risk Management Framework

AHRQ human factors and usability resources

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP Top 10 for LLM Applications, when security or AI integrations are relevant

Agency privacy, cybersecurity, data governance, procurement, and records management policies

Knowledge Check

1. What is the best reason to address human factors, workflow integration, and usability testing before deployment?

- A. It makes governance unnecessary.
- B. It helps ensure the AI workflow is safe, useful, reviewable, and fit for public health operations.
- C. It allows the agency to skip testing.

D. It guarantees that the vendor product will work.

Correct answer: B

2. Which practice best supports responsible implementation?

A. Defining intended use, evidence requirements, safeguards, monitoring, and escalation before launch.

B. Waiting until after launch to decide whether the tool is risky.

C. Allowing users to decide individually how to use the tool.

D. Treating vendor documentation as the only required evidence.

Correct answer: A

3. Which statement best describes a common failure mode?

A. Human review is always unnecessary when AI is accurate.

B. Documentation makes AI less useful.

C. The agency assumes the tool is ready because the technology works in a demonstration, even though workflow, data, and governance have not been tested.

D. Monitoring should stop after the pilot.

Correct answer: C

4. What should be included in the practical artifact for this module?

A. Only a general definition.

B. A marketing description of the AI tool.

C. A list of unrelated data systems.

D. Workflow-specific risks, safeguards, roles, evidence requirements, and unresolved questions.

Correct answer: D

5. When should the issue be escalated for governance or leadership review?

A. Only if the tool is public-facing.

B. When risks affect privacy, safety, equity, legal authority, public trust, system integrity, or operational decisions.

C. Never, because technical staff can resolve all issues alone.

D. Only after an incident has already occurred.

Correct answer: B

References and Resources

NIST AI Risk Management Framework

AHRQ human factors and usability resources

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP Top 10 for LLM Applications, when security or AI integrations are relevant

Agency privacy, cybersecurity, data governance, procurement, and records management policies

OPS 230

AI Workflow Standard Operating Procedures

Course ID	OPS 230
Track	Operations and Data Quality Track
Level	applied/foundational
Estimated time	45
Audience	Operations leads, Data quality staff, Program managers, Implementation teams, Supervisors

Learning Objectives

- Explain why approved AI tools still require workflow-specific SOPs.
- Identify the core elements of an AI workflow SOP.
- Define human review, documentation, override, escalation, and fallback expectations.
- Distinguish between acceptable and prohibited uses within a specific AI workflow.
- Produce a draft AI workflow SOP for a public health use case.

Jurisdiction and Agency Policy Note

This national-level module is intended for training and planning purposes. It does not replace legal, privacy, procurement, civil rights, records, or agency policy review. Learners should confirm applicable requirements in their own jurisdiction and organization before implementing AI-supported workflows.

Public Health Example

A health department approves an AI tool to draft summaries of environmental health complaints. The SOP explains which complaint types are eligible, which data fields may be used, how staff must compare the draft to the original complaint, where the final summary is stored, how unsupported statements are removed, and when errors must be reported. The SOP also explains that the tool may not issue inspection decisions, enforcement recommendations, or public notices without human review.

Technical and Operational Deep Dive

An AI SOP should begin with scope. The scope explains which workflow the SOP covers, which staff roles may use the tool, which systems or data are involved, and which AI-supported tasks are allowed. Scope is important because many AI tools can technically perform more tasks than the agency has approved. A clear scope prevents an approved tool from drifting into unapproved uses.

The SOP should define the standard workflow step by step. This includes preparing data, using the approved tool, reviewing outputs, correcting errors, documenting the final decision, and escalating concerns. For generative AI, the SOP should require staff to verify outputs against source material. For predictive tools, it should require staff to interpret scores in context and avoid treating the output as an automatic decision.

A strong SOP also includes exception handling. Staff need to know what to do when the AI tool is unavailable, the output is inconsistent with source data, the output appears biased or unsafe, the workflow involves a sensitive or unusual case,

or the user is uncertain whether the use is allowed. The SOP should identify a fallback process and escalation pathway so staff are not forced to improvise.

SOPs should be maintained as living documents. AI workflows may change when policies, data feeds, models, prompts, vendors, or public health conditions change. SOP ownership, review frequency, and change control should be documented so the procedure remains aligned with approved use.

Risks, Failure Modes, and Guardrails

Risks and failure modes include the following:

Staff applying an AI tool beyond the approved workflow.

Outputs being copied into official records without required review.

Errors, bias concerns, or privacy issues not being escalated.

No fallback process when the tool fails or is unavailable.

Outdated SOPs that no longer match the deployed workflow.

Recommended guardrails include the following:

Create a workflow-specific SOP for each approved AI use.

Define allowed uses, prohibited uses, human review, and documentation requirements.

Require staff training before independent use.

Include override, escalation, and fallback procedures.

Review SOPs after incidents, major updates, or policy changes.

Reflection Questions

Where does this topic appear in your agency, program, or workflow?

Which staff roles would need to understand or apply this concept before AI use is approved?

What could go wrong if this topic is not addressed before implementation?

What evidence would governance, leadership, or operations staff need before moving forward?

Which policy, data, equity, privacy, security, or workflow questions remain unresolved?

Practical Exercise

Draft an AI workflow SOP for one approved or realistic public health use case. The SOP should include purpose, scope, roles, allowed uses, prohibited uses, step-by-step workflow, human review requirements, documentation requirements,

escalation triggers, fallback process, and review schedule.

Expected Artifact or Evidence

Draft AI workflow SOP

Role and responsibility matrix

Allowed/prohibited use section

Escalation and fallback process

Review and maintenance schedule

References and Resources for Additional Information

NIST Artificial Intelligence Risk Management Framework

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP guidance for LLM application security

Agency policies for privacy, cybersecurity, records, procurement, accessibility, and public communications

Relevant public health program SOPs, data governance documentation, and implementation plans

Knowledge Check

1. Why does an approved AI tool still need a workflow-specific SOP?

- A. Because approval automatically explains every operational step
- B. Because staff need clear instructions for allowed use, review, documentation, escalation, and fallback
- C. Because SOPs replace governance
- D. Because SOPs are only needed for vendors

Correct answer: B

2. What is an override?

- A. A vendor demonstration
- B. A privacy policy
- C. A user decision to reject, change, or not follow an AI output
- D. A model training dataset

Correct answer: C

3. Which item should an AI SOP include?

- A. Only the model name

- B. Only the vendor contact
- C. A general statement that AI is helpful
- D. Allowed uses, prohibited uses, human review, documentation, and escalation expectations

Correct answer: D

4. When should an SOP be reviewed?

- A. After incidents, major updates, policy changes, or workflow changes
- B. Never after approval
- C. Only when a vendor requests it
- D. Only after media attention

Correct answer: A

5. What is strong completion evidence?

- A. A draft workflow SOP with scope, roles, steps, review, escalation, fallback, and maintenance expectations
- B. A screenshot of the tool
- C. A list of AI benefits
- D. An informal email approving use

Correct answer: A

References and Resources

NIST Artificial Intelligence Risk Management Framework

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP guidance for LLM application security

Agency policies for privacy, cybersecurity, records, procurement, accessibility, and public communications

Relevant public health program SOPs, data governance documentation, and implementation plans

OPS 240

User Acceptance Testing for AI-Supported Workflows

Course ID	OPS 240
Track	Operations and Data Quality Track
Level	applied/foundational
Estimated time	45
Audience	Operations leads, Data quality staff, Program managers, Implementation teams, Supervisors

Learning Objectives

- Explain the purpose of UAT for AI-supported workflows.
- Develop role-based UAT scenarios and test scripts.
- Include routine cases, edge cases, failure cases, and equity-sensitive scenarios in UAT.
- Define acceptance criteria for usability, safety, documentation, human review, and workflow fit.
- Produce a UAT plan and sign-off recommendation for an AI-supported workflow.

Public Health Example

Before launching an AI tool that summarizes case investigation notes, investigators complete UAT scenarios using routine cases, incomplete records, conflicting notes, language access needs, and cases with sensitive information. Testers verify whether the summary is understandable, source information is visible, errors can be corrected, outputs can be documented properly, and escalation steps are clear. The tool does not launch until high-priority issues are resolved.

Technical and Operational Deep Dive

UAT should be designed around real users and real work. Testers should include the roles that will actually use or supervise the AI-supported workflow, such as epidemiologists, investigators, program staff, supervisors, data quality staff, communications staff, or support staff. If only technical staff test the tool, usability and workflow problems may be missed.

Test scenarios should include more than ideal cases. AI-supported workflows should be tested with incomplete data, conflicting information, duplicate records, unusual cases, sensitive data, ambiguous outputs, language access issues, accessibility needs, high-volume periods, system downtime, and outputs that should be overridden. Edge cases reveal whether the workflow is safe when conditions are imperfect.

Acceptance criteria should be defined before testing begins. Criteria may include whether users can complete the workflow, understand the AI output, inspect source information, identify unsupported statements, correct errors, document decisions, override outputs, and escalate concerns. Criteria should also address workflow burden, accessibility, performance, and user confidence.

UAT results should lead to a launch decision. Issues should be categorized by severity, assigned to owners, and resolved or accepted with documented rationale. A launch should not proceed simply because testing occurred. It should proceed only if acceptance criteria are met and unresolved issues do not create unacceptable risk.

Risks, Failure Modes, and Guardrails

Risks and failure modes include the following:

Testing only ideal scenarios and missing high-risk workflow failures.

Excluding real end users from testing.

Launching despite unresolved issues that affect safety, equity, or documentation.

Treating model validation as a substitute for workflow acceptance testing.

Failing to test override, escalation, fallback, and accessibility requirements.

Recommended guardrails include the following:

Develop role-based UAT scripts before launch.

Include routine, edge, failure, and equity-sensitive cases.

Define acceptance criteria before testing.

Track issues by severity, owner, and resolution status.

Require documented sign-off or remediation before deployment.

Reflection Questions

Where does this topic appear in your agency, program, or workflow?

Which staff roles would need to understand or apply this concept before AI use is approved?

What could go wrong if this topic is not addressed before implementation?

What evidence would governance, leadership, or operations staff need before moving forward?

Which policy, data, equity, privacy, security, or workflow questions remain unresolved?

Practical Exercise

Create a UAT plan for one AI-supported workflow. Include user roles, test scenarios, edge cases, acceptance criteria, tester instructions, issue log fields, severity categories, sign-off criteria, and launch recommendation.

Expected Artifact or Evidence

UAT plan

Role-based test scripts

Acceptance criteria

Issue log template

Sign-off recommendation

References and Resources for Additional Information

NIST Artificial Intelligence Risk Management Framework

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP guidance for LLM application security

Agency policies for privacy, cybersecurity, records, procurement, accessibility, and public communications

Relevant public health program SOPs, data governance documentation, and implementation plans

Gap-Filling Modules: Remaining Role and Governance Gaps

Source file: AI_Remaining_Gap_Filling_Training_Modules_for_Public_Health(1).docx

Additional Gap-Filling AI Training Modules for Public Health

Modules to complete the recommended learning tracks

This document completes the remaining gaps identified across the shared foundational course, technical architecture track, analytics/modeling track, operations/data quality track, governance/security track, and role-based tracks for communications, epidemiology, policy, executive leadership, and program management. These modules are designed to complement the previously created foundational, technical, leadership/governance, and gap-filling modules.

Each module follows the same structure used across the curriculum: training overview, learning objectives, definitions, public health example, deep-dive content, risks and guardrails, application to AI-supported workflows, reflection questions, practical exercise, expected evidence, knowledge check, and lesson resources.

Module List

1. Public Health System Integration Patterns (Technical architecture; operations/data quality; program management)
2. Causal Inference and AI in Public Health Decision-Making (Analytics/modeling; epidemiology; policy)
3. Communicating Model Results and Uncertainty (Analytics/modeling; communications; executive leadership; epidemiology)
4. Operational Monitoring and Continuous Improvement for AI (Operations/data quality; governance/security; program management)

5. AI Policy Development and Approval Workflow (Governance/security; policy; executive leadership)
6. Third-Party AI Risk Management (Governance/security; technical architecture; program management)
7. AI-Assisted Translation, Language Access, and Readability Review (Communications; policy; shared foundation)
8. Public Comment, Social Listening, and Community Feedback Analysis (Communications; policy; AI for social good)
9. AI for Surveillance Signal Review and Epidemiologic Judgment (Epidemiology; analytics/modeling; operations/data quality)
10. Ethical Use of AI in Resource Prioritization and Field Response (Epidemiology; policy; executive leadership; AI for social good)
11. AI Authority, Liability, and Public Health Decision-Making (Policy; governance/security; executive leadership)
12. AI Equity Policy and Civil Rights Considerations (Policy; governance/security; communications; executive leadership)
13. Executive AI Briefing for Public Health Leaders (Executive leadership; governance/security)
14. AI Governance Board Chair Training (Executive leadership; governance/security; program management)
15. AI Crisis Communications and Executive Response (Executive leadership; communications; governance/security)
16. AI Vendor Implementation Management (Program management; governance/security; technical architecture)
17. Grant and Funding Management for AI Projects (Program management; executive leadership; policy)

Knowledge Check

1. What is the purpose of UAT for AI-supported workflows?

- A. To replace privacy review
- B. To prove the model is perfect
- C. To eliminate user training
- D. To determine whether representative users can safely and effectively use the workflow

Correct answer: D

2. Who should participate in UAT?

- A. Representative users who will use or supervise the workflow
- B. Only the vendor sales team
- C. Only executive leadership
- D. Only people unfamiliar with the workflow

Correct answer: A

3. Which scenario should be included in AI UAT?

- A. Only the easiest case
- B. Routine, edge, failure, and equity-sensitive scenarios
- C. Only a vendor demonstration
- D. Only a screenshot review

Correct answer: B

4. What should happen to UAT issues?

- A. They should be ignored after testing
- B. They should be deleted before launch
- C. They should be tracked by severity, owner, status, and resolution
- D. They should be sent only to the public

Correct answer: C

5. What is strong completion evidence?

- A. A generic statement that users liked the tool
- B. A model accuracy score only
- C. A vendor invoice
- D. A UAT plan with scripts, acceptance criteria, issue log, and sign-off recommendation

Correct answer: D

References and Resources

CDC Evaluation Framework

<https://www.cdc.gov/evaluation/>

NIST Artificial Intelligence Risk Management Framework

<https://www.nist.gov/itl/ai-risk-management-framework>

NIST AI RMF Generative AI Profile

<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>

CDC Public Health Data Strategy

<https://www.cdc.gov/public-health-data-strategy/php/about/index.html>

Agency user acceptance testing, quality assurance, accessibility, and implementation policies

OPS 300

MLOps and LLMOps for Public Health

Course ID	OPS 300
Track	Operations and Data Quality Track
Level	intermediate/practitioner
Estimated time	60
Audience	Operations leads, Data quality staff, Program managers, Implementation teams, Supervisors

Learning Objectives

- Explain the purpose of mlops and llmops for public health in public health AI work.
- Identify the technical, operational, privacy, security, equity, and governance issues associated with the topic.
- Apply the topic to a real or realistic public health workflow, system, or implementation scenario.
- Recognize common failure modes that can affect safety, accuracy, trust, workflow fit, or sustainability.
- Identify practical guardrails that should be documented before implementation.
- Produce a AI operations plan that can support readiness assessment, governance review, procurement, implementation, or monitoring.

Why This Topic Matters for Public Health AI

MLOps and LLMOps for Public Health matters because AI systems are embedded in public health programs, data systems, and decision processes rather than operating in isolation. The topic affects whether AI outputs are accurate, explainable, safe, equitable, secure, and sustainable. Agencies that ignore this area may create tools that work in a demonstration but fail in actual practice.

The topic also affects accountability. Public health agencies must be able to explain how data are used, why a tool was approved, what evidence supports the workflow, who reviews outputs, and what happens when the system fails. These responsibilities become more important as AI tools move from experimentation into operational use.

Core Concepts

MLOps refers to operational practices for managing machine learning models through their lifecycle. LLMOps applies similar practices to large language model applications, including prompts, retrieval sources, evaluation datasets, guardrails, and model versions. These practices make AI systems more reliable, auditable, and maintainable.

Model versioning is the practice of tracking which version of a model, prompt, retrieval corpus, evaluation dataset, or configuration was used. Versioning matters because public health staff may need to reconstruct why an output changed or why an incident occurred.

Drift occurs when data patterns, user behavior, or model performance change over time. In public health, drift may occur because a new lab begins reporting, a case definition changes, a provider changes coding practices, or community conditions shift. Monitoring drift helps agencies decide when retesting or retraining is needed.

Public Health Example

A syndromic surveillance team may use an anomaly detection model to flag unusual emergency department visit patterns. The model may perform well during testing, but performance can change when chief complaint text changes, new facilities join the feed, or seasonal patterns shift. An MLOps plan would define monitoring indicators, evaluation frequency, responsible staff, retraining rules, and rollback procedures.

Technical Deep Dive

Operational AI management starts before deployment. Teams should define an intended use, baseline performance, evaluation datasets, release criteria, monitoring indicators, and an owner responsible for review. Without these elements, agencies may not know whether the deployed system is still behaving as approved.

LLMOps adds additional elements because generative systems depend on prompts, retrieval documents, context windows, model settings, safety filters, and human review workflows. A change in any of these components can change outputs. Agencies should track prompt templates, source document updates, model upgrades, and evaluation results.

Rollback is a critical safeguard. If a new model version, prompt, data pipeline, or retrieval index creates errors, staff need a documented way to return to the prior approved version or disable the AI function. Rollback authority should be defined before deployment.

Risks, Failure Modes, and Guardrails

A common failure mode is treating the topic as a technical detail rather than a public health control. When staff do not connect technical decisions to authority, workflow, equity, privacy, security, and accountability, the agency may adopt a tool that appears useful but cannot be sustained or defended.

Another risk is relying on vendor claims without agency-defined evidence. Vendors may provide demonstrations, dashboards, or summary metrics that do not match the agency workflow or data environment. A guardrail is to require documentation, test results, acceptance criteria, and agency-controlled review.

A third risk is failing to plan for change. Public health data, policies, systems, and emergencies change. Any technical approach must include monitoring, review, versioning, escalation, and retirement pathways so the agency can respond when assumptions no longer hold.

Application to AI-Supported Workflows

Generative AI, predictive analytics, RAG, automation, and agentic workflows all depend on the controls described in this module. The controls should be scaled to the risk of the use case. A tool that drafts internal training text may require lighter review than a tool that updates case records, prioritizes outreach, or supports public communication during an emergency.

The module should be applied during readiness assessment, procurement, implementation planning, pilot testing, and post-deployment monitoring. Learners should document how the topic affects data, systems, users, safeguards, and decision-making for the selected workflow.

Reflection Questions

Where does this topic appear in one current or proposed AI workflow in your agency?

What decisions, data sources, users, or partners would be affected if this area is poorly designed?

What evidence would governance reviewers need before approving the workflow?

Which safeguards should be required before pilot testing?

Who owns ongoing monitoring and review?

Practical Exercise

Create a AI operations plan for one real or realistic public health AI workflow. The artifact should describe the use case, affected users, data sources, system dependencies, risks, guardrails, review points, and unresolved questions.

The exercise should produce a work product that could be used in a governance meeting, procurement review, implementation planning session, pilot evaluation, or monitoring review. Learners should document assumptions and identify which staff or partners need to be consulted.

Expected Artifact or Evidence

The expected artifact is a completed AI operations plan. It should be specific to a public health workflow and should include technical dependencies, governance questions, human review expectations, monitoring indicators, and unresolved risks.

References and Resources for Additional Information

NIST AI Risk Management Framework

Platform-specific MLOps and LLMOps documentation for approved agency tools

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP Top 10 for LLM Applications, when security or AI integrations are relevant

Agency privacy, cybersecurity, data governance, procurement, and records management policies

Knowledge Check

1. What is the best reason to address mlops and llmops for public health before deployment?

- A. It guarantees that the vendor product will work.
- B. It allows the agency to skip testing.
- C. It makes governance unnecessary.
- D. It helps ensure the AI workflow is safe, useful, reviewable, and fit for public health operations.

Correct answer: D

2. Which practice best supports responsible implementation?

- A. Allowing users to decide individually how to use the tool.
- B. Treating vendor documentation as the only required evidence.
- C. Defining intended use, evidence requirements, safeguards, monitoring, and escalation before launch.
- D. Waiting until after launch to decide whether the tool is risky.

Correct answer: C

3. Which statement best describes a common failure mode?

- A. Human review is always unnecessary when AI is accurate.
- B. Documentation makes AI less useful.
- C. The agency assumes the tool is ready because the technology works in a demonstration, even though workflow, data, and governance have not been tested.
- D. Monitoring should stop after the pilot.

Correct answer: C

4. What should be included in the practical artifact for this module?

- A. Only a general definition.
- B. A marketing description of the AI tool.
- C. A list of unrelated data systems.
- D. Workflow-specific risks, safeguards, roles, evidence requirements, and unresolved questions.

Correct answer: D

5. When should the issue be escalated for governance or leadership review?

- A. Only if the tool is public-facing.
- B. When risks affect privacy, safety, equity, legal authority, public trust, system integrity, or operational decisions.
- C. Never, because technical staff can resolve all issues alone.
- D. Only after an incident has already occurred.

Correct answer: B

References and Resources

NIST AI Risk Management Framework

Platform-specific MLOps and LLMOps documentation for approved agency tools

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP Top 10 for LLM Applications, when security or AI integrations are relevant

Agency privacy, cybersecurity, data governance, procurement, and records management policies

OPS 310

Measuring AI Value, Performance, and Operational Impact

Course ID	OPS 310
Track	Operations and Data Quality Track
Level	intermediate/practitioner
Estimated time	60
Audience	Operations leads, Data quality staff, Program managers, Implementation teams, Supervisors

Learning Objectives

- Explain the purpose of measuring ai value, performance, and operational impact in public health AI work.
- Identify the technical, operational, privacy, security, equity, and governance issues associated with the topic.
- Apply the topic to a real or realistic public health workflow, system, or implementation scenario.
- Recognize common failure modes that can affect safety, accuracy, trust, workflow fit, or sustainability.
- Identify practical guardrails that should be documented before implementation.
- Produce a AI value and impact measurement plan that can support readiness assessment, governance review, procurement, implementation, or monitoring.

Why This Topic Matters for Public Health AI

Measuring AI Value, Performance, and Operational Impact matters because AI systems are embedded in public health programs, data systems, and decision processes rather than operating in isolation. The topic affects whether AI outputs are accurate, explainable, safe, equitable, secure, and sustainable. Agencies that ignore this area may create tools that work in a demonstration but fail in actual practice.

The topic also affects accountability. Public health agencies must be able to explain how data are used, why a tool was approved, what evidence supports the workflow, who reviews outputs, and what happens when the system fails. These responsibilities become more important as AI tools move from experimentation into operational use.

Core Concepts

AI value is the benefit created by an AI-supported workflow relative to its costs, risks, and alternatives. Value may include faster response, improved quality, reduced burden, better access, stronger equity, improved documentation, or better use of limited resources.

Model performance measures how well the AI system performs a technical task. Operational impact measures whether the workflow actually improves public health practice. A model can have strong technical metrics and still fail to reduce burden, improve timeliness, or support equitable outcomes.

Evaluation should be planned before implementation. Agencies should define baseline measures, target outcomes, data sources, review frequency, and decision rules for continuation, modification, scaling, or retirement.

Public Health Example

A health department may deploy AI to summarize case reports with the goal of reducing investigator burden. Accuracy alone is not enough to prove value. The agency should also measure review time, editing burden, error types, investigator satisfaction, case completion timeliness, escalation frequency, subgroup performance, and whether staff continue to trust and use the tool appropriately.

Technical Deep Dive

Value measurement begins with a theory of change. The team should explain how the AI capability is expected to improve the workflow and which outcomes should change. For example, a summarization tool may reduce review time, improve completeness, and allow staff to focus on complex cases. These assumptions should be tested rather than assumed.

Measures should include technical, operational, equity, and governance indicators. Technical indicators may include accuracy, error rates, latency, and uptime. Operational indicators may include timeliness, backlog, staff burden, user adoption, and quality. Equity indicators may include subgroup performance, access, language, disability, and community impact. Governance indicators may include documentation completeness, incident reports, and monitoring review completion.

Agencies should define action thresholds. If the tool produces too many errors, increases burden, performs poorly for a subgroup, or creates trust concerns, the agency should know whether to retrain, revise, restrict, pause, or retire the tool. Value measurement should support decisions, not merely reporting.

Risks, Failure Modes, and Guardrails

A common failure mode is treating the topic as a technical detail rather than a public health control. When staff do not connect technical decisions to authority, workflow, equity, privacy, security, and accountability, the agency may adopt a tool that appears useful but cannot be sustained or defended.

Another risk is relying on vendor claims without agency-defined evidence. Vendors may provide demonstrations, dashboards, or summary metrics that do not match the agency workflow or data environment. A guardrail is to require documentation, test results, acceptance criteria, and agency-controlled review.

A third risk is failing to plan for change. Public health data, policies, systems, and emergencies change. Any technical approach must include monitoring, review, versioning, escalation, and retirement pathways so the agency can respond when assumptions no longer hold.

Application to AI-Supported Workflows

Generative AI, predictive analytics, RAG, automation, and agentic workflows all depend on the controls described in this module. The controls should be scaled to the risk of the use case. A tool that drafts internal training text may require lighter review than a tool that updates case records, prioritizes outreach, or supports public communication during an emergency.

The module should be applied during readiness assessment, procurement, implementation planning, pilot testing, and post-deployment monitoring. Learners should document how the topic affects data, systems, users, safeguards, and decision-making for the selected workflow.

Reflection Questions

Where does this topic appear in one current or proposed AI workflow in your agency?

What decisions, data sources, users, or partners would be affected if this area is poorly designed?

What evidence would governance reviewers need before approving the workflow?

Which safeguards should be required before pilot testing?

Who owns ongoing monitoring and review?

Practical Exercise

Create a AI value and impact measurement plan for one real or realistic public health AI workflow. The artifact should describe the use case, affected users, data sources, system dependencies, risks, guardrails, review points, and unresolved questions.

The exercise should produce a work product that could be used in a governance meeting, procurement review, implementation planning session, pilot evaluation, or monitoring review. Learners should document assumptions and identify which staff or partners need to be consulted.

Expected Artifact or Evidence

The expected artifact is a completed AI value and impact measurement plan. It should be specific to a public health workflow and should include technical dependencies, governance questions, human review expectations, monitoring indicators, and unresolved risks.

References and Resources for Additional Information

NIST AI Risk Management Framework

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP Top 10 for LLM Applications, when security or AI integrations are relevant

Agency privacy, cybersecurity, data governance, procurement, and records management policies

Leadership and Governance Modules

Source file: Leadership_and_Governance_AI_Training_Modules(1).docx

Leadership and Governance for Public Health AI

Training Modules 1-10

This document provides a leadership and governance training track for public health agencies implementing responsible AI. The modules are designed to complement foundational responsible AI lessons and the technical readiness modules by focusing on executive decisions, governance operating models, investment strategy, risk appetite, public

accountability, community oversight, and sustainability.

Module List

1. AI Governance Charter and Operating Model
2. Executive AI Decision Rights and Accountability
3. AI Portfolio Strategy and Use Case Prioritization
4. AI Risk Appetite and Risk Tolerance for Public Health
5. AI Investment Strategy and Business Case Review
6. Organizational Change Management for AI Adoption
7. Workforce Strategy and AI Role Design
8. Public Accountability and AI Transparency Reporting
9. Community Oversight and Participatory Governance
10. AI Sustainability and Lifecycle Funding

Common Source Base

NIST Artificial Intelligence Risk Management Framework (AI RMF 1.0)

NIST Generative AI Profile

NIST Privacy Framework

CDC Public Health Data Strategy and Data Modernization resources

WHO Ethics and Governance of Artificial Intelligence for Health

OMB Memorandum M-24-10 on federal AI governance, risk management, and innovation, where relevant to public-sector governance models

Agency strategic plans, procurement policies, privacy policies, cybersecurity policies, records policies, and equity plans

Knowledge Check

1. What is the best reason to address measuring ai value, performance, and operational impact before deployment?

- A. It guarantees that the vendor product will work.
- B. It allows the agency to skip testing.
- C. It makes governance unnecessary.

D. It helps ensure the AI workflow is safe, useful, reviewable, and fit for public health operations.

Correct answer: D

2. Which practice best supports responsible implementation?

A. Defining intended use, evidence requirements, safeguards, monitoring, and escalation before launch.

B. Waiting until after launch to decide whether the tool is risky.

C. Allowing users to decide individually how to use the tool.

D. Treating vendor documentation as the only required evidence.

Correct answer: A

3. Which statement best describes a common failure mode?

A. Human review is always unnecessary when AI is accurate.

B. Documentation makes AI less useful.

C. The agency assumes the tool is ready because the technology works in a demonstration, even though workflow, data, and governance have not been tested.

D. Monitoring should stop after the pilot.

Correct answer: C

4. What should be included in the practical artifact for this module?

A. Only a general definition.

B. A marketing description of the AI tool.

C. A list of unrelated data systems.

D. Workflow-specific risks, safeguards, roles, evidence requirements, and unresolved questions.

Correct answer: D

5. When should the issue be escalated for governance or leadership review?

A. Only if the tool is public-facing.

B. When risks affect privacy, safety, equity, legal authority, public trust, system integrity, or operational decisions.

C. Never, because technical staff can resolve all issues alone.

D. Only after an incident has already occurred.

Correct answer: B

References and Resources

CDC Evaluation Framework

<https://www.cdc.gov/evaluation/>

CDC Public Health Data Strategy

<https://www.cdc.gov/public-health-data-strategy/php/about/index.html>

NIST Artificial Intelligence Risk Management Framework

<https://www.nist.gov/itl/ai-risk-management-framework>

NIST AI RMF Generative AI Profile

<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>

Agency performance management, quality improvement, equity, and evaluation policies

OPS 320

Operational Monitoring and Continuous Improvement for AI

Course ID	OPS 320
Track	Operations and Data Quality Track
Level	intermediate/practitioner
Estimated time	60
Audience	Operations leads, Data quality staff, Program managers, Implementation teams, Supervisors

Learning Objectives

- Explain the purpose of operational monitoring and continuous improvement for ai in public health AI.
- Identify the roles, decisions, risks, and documentation needed for the topic.
- Apply the topic to a real or realistic public health workflow or decision.
- Recognize equity, privacy, security, legal, operational, and public trust implications.
- Produce a practical artifact that supports readiness, governance, implementation, monitoring, or accountability.

Definitions

Decision right	The authority to approve, deny, modify, pause, or escalate an AI use case or operational decision.
Evidence artifact	A document, log, checklist, test result, approval record, or review note that demonstrates responsible AI practice.
Escalation pathway	A defined route for raising concerns to the appropriate reviewer or decision-maker.
Accountability owner	The person or role responsible for ensuring that an AI use case is governed, monitored, and corrected when needed.
Operational safeguard	A procedure, control, or review step that reduces risk during day-to-day use.

Jurisdiction and Agency Policy Note

This national-level module is intended for training and planning purposes. It does not replace legal, privacy, procurement, civil rights, records, or agency policy review. Learners should confirm applicable requirements in their own jurisdiction and organization before implementing AI-supported workflows.

Why This Topic Matters for Public Health AI

Operational Monitoring and Continuous Improvement for AI matters because AI adoption changes how public health agencies make, document, communicate, and review decisions. Without a structured approach, staff may rely on informal practices, vendor assurances, or one-time approvals that do not match the operational consequences of

AI-supported work.

Public health agencies operate under legal authorities, public accountability, resource constraints, and community trust obligations. A module on this topic helps agencies define repeatable expectations before urgent situations arise. It also helps staff distinguish low-risk administrative support from uses that affect surveillance, public communications, field response, resource allocation, civil rights, privacy, or official records.

Public Health Example

A public health agency is preparing to use AI in a workflow related to post-launch monitoring of AI-supported workflows. The project team initially treats the tool as a helpful efficiency aid, but governance reviewers recognize that the workflow may affect public trust, equity, official records, or public health action.

The agency responds by creating a clear artifact that defines intended use, prohibited use, responsible roles, review points, documentation expectations, escalation triggers, and monitoring measures. This converts a broad AI idea into a manageable public health practice.

Technical and Operational Deep Dive

The first step in this topic is to define the decision or workflow clearly. Staff should identify who is affected, what data are used, what AI output is produced, who reviews the output, and what action may follow. This prevents the module from becoming abstract and anchors the work in public health practice.

The second step is to assign responsibilities. AI-supported workflows often cross program, IT, legal, privacy, communications, procurement, and leadership boundaries. Responsibilities should be named by role rather than assumed. The artifact should specify who prepares materials, who reviews risks, who approves use, who monitors performance, and who can pause or revise the workflow.

The third step is to define evidence. Responsible AI depends on more than policy statements. Agencies need records that show what was reviewed, what decisions were made, what safeguards were applied, what incidents occurred, and how the workflow changed over time. Evidence also supports audits, public records review, quality improvement, and leadership accountability.

Risks, Failure Modes, and Guardrails

A common failure mode is treating the topic as a one-time checklist. AI risks can change when data sources, models, vendors, policies, users, or public expectations change. A guardrail is to define review intervals and change triggers.

Another failure mode is unclear ownership. When everyone is responsible in general, no one may be responsible in practice. A guardrail is to assign an accountability owner and name escalation pathways.

A third failure mode is relying on vendor claims or informal staff confidence without documentation. A guardrail is to require evidence artifacts that can be reviewed by governance, leadership, privacy, legal, equity, security, or program teams.

Application to AI-Supported Workflows

This topic applies across generative AI, predictive analytics, NLP, RAG, automation, and agentic workflows. The more sensitive the data, the more consequential the output, or the more automated the action, the stronger the

documentation, review, and monitoring expectations should be.

Reflection Questions

What public health decision or workflow does this topic affect?

Who has authority to approve or pause the work?

What evidence would show that the workflow is responsible and effective?

What equity, privacy, legal, security, or public trust risks require review?

How will the agency know when the workflow needs to be revised?

Practical Exercise

Create a practical operational monitoring and continuous improvement for ai artifact for one AI-supported public health use case. The artifact should define the use case, affected roles, review criteria, safeguards, documentation expectations, escalation triggers, monitoring indicators, and next review date.

Expected Artifact or Evidence

Operational Monitoring and Continuous Improvement for AI artifact

Role and responsibility table

Risk and safeguard summary

Escalation and monitoring plan

Approval or review note

References and Resources for Additional Information

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

Knowledge Check

1. What is the strongest evidence that operational monitoring and continuous improvement for ai has been applied responsibly?

- A. A general statement that AI is useful
- B. A vendor demonstration with no agency review
- C. A documented artifact that defines use, roles, safeguards, escalation, monitoring, and evidence
- D. A decision to skip review because the use case seems efficient

Correct answer: C

2. What is the best reason to assign an accountability owner?

- A. To reduce the need for documentation
- B. To ensure someone is responsible for monitoring, escalation, and correction
- C. To make the vendor solely responsible
- D. To prevent staff from raising concerns

Correct answer: B

3. When should an AI-supported workflow be escalated for review?

- A. Only after media coverage
- B. Never if the tool is efficient
- C. When outputs may affect safety, equity, privacy, legal authority, public trust, or official action
- D. Only at the end of the fiscal year

Correct answer: C

4. Which practice best supports public accountability?

- A. Keeping AI use informal
- B. Avoiding written records
- C. Letting each user decide independently
- D. Maintaining evidence artifacts, decision records, review notes, and monitoring results

Correct answer: D

5. What should happen when data sources, policy, vendor configuration, or workflow use changes?

- A. The use case should be reviewed to determine whether safeguards, validation, documentation, or approvals need updating
- B. The agency should ignore the change if staff like the tool
- C. The model should be trusted automatically
- D. The glossary should be deleted

Correct answer: A

References and Resources

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

PMG 200

Business Analysis for AI-Enabled Public Health Systems

Course ID	PMG 200
Track	Program Management Role-Based Track
Level	applied/foundational
Estimated time	45
Audience	Project managers, Program managers, Product owners, Implementation leads, Procurement leads

Learning Objectives

- Explain the purpose of business analysis for ai-enabled public health systems in public health AI work.
- Identify the technical, operational, privacy, security, equity, and governance issues associated with the topic.
- Apply the topic to a real or realistic public health workflow, system, or implementation scenario.
- Recognize common failure modes that can affect safety, accuracy, trust, workflow fit, or sustainability.
- Identify practical guardrails that should be documented before implementation.
- Produce a AI business analysis package that can support readiness assessment, governance review, procurement, implementation, or monitoring.

Jurisdiction and Agency Policy Note

This national-level module is intended for training and planning purposes. It does not replace legal, privacy, procurement, civil rights, records, or agency policy review. Learners should confirm applicable requirements in their own jurisdiction and organization before implementing AI-supported workflows.

Why This Topic Matters for Public Health AI

Business Analysis for AI-Enabled Public Health Systems matters because AI systems are embedded in public health programs, data systems, and decision processes rather than operating in isolation. The topic affects whether AI outputs are accurate, explainable, safe, equitable, secure, and sustainable. Agencies that ignore this area may create tools that work in a demonstration but fail in actual practice.

The topic also affects accountability. Public health agencies must be able to explain how data are used, why a tool was approved, what evidence supports the workflow, who reviews outputs, and what happens when the system fails. These responsibilities become more important as AI tools move from experimentation into operational use.

Core Concepts

Business analysis clarifies the problem, workflow, users, requirements, constraints, and success criteria before technology decisions are made. AI projects often fail when teams start with a tool rather than a well-defined public health need.

A user story describes what a user needs to do and why. In AI work, user stories should include human review, data requirements, documentation, and limitations. A user story should not simply say that the user wants AI; it should

describe the task the AI is supporting.

Acceptance criteria are testable conditions that determine whether a requirement has been met. For AI systems, acceptance criteria should include accuracy, workflow fit, usability, security, privacy, accessibility, and human oversight.

Public Health Example

A local health department may want AI to help classify community complaints. A business analyst would clarify whether the real problem is high complaint volume, inconsistent routing, delayed response, lack of category definitions, or insufficient staffing. The resulting requirements may call for classification support with human review rather than fully automated complaint resolution.

Technical Deep Dive

Business analysis begins with problem definition. Teams should document the public health objective, the operational pain point, the affected users, the current workflow, the decision being supported, and the consequences of error. A vague goal such as “use AI to improve efficiency” is not sufficient for implementation.

Requirements should be grounded in workflow. Analysts should define current-state and future-state processes, data inputs, outputs, decision points, exception handling, reporting needs, and approval steps. This helps prevent tools from being procured without a clear implementation pathway.

AI business analysis also requires boundary setting. Teams should define intended uses, prohibited uses, users, data permissions, human review points, and escalation triggers. These boundaries should become requirements, training expectations, and governance criteria.

Risks, Failure Modes, and Guardrails

A common failure mode is treating the topic as a technical detail rather than a public health control. When staff do not connect technical decisions to authority, workflow, equity, privacy, security, and accountability, the agency may adopt a tool that appears useful but cannot be sustained or defended.

Another risk is relying on vendor claims without agency-defined evidence. Vendors may provide demonstrations, dashboards, or summary metrics that do not match the agency workflow or data environment. A guardrail is to require documentation, test results, acceptance criteria, and agency-controlled review.

A third risk is failing to plan for change. Public health data, policies, systems, and emergencies change. Any technical approach must include monitoring, review, versioning, escalation, and retirement pathways so the agency can respond when assumptions no longer hold.

Application to AI-Supported Workflows

Generative AI, predictive analytics, RAG, automation, and agentic workflows all depend on the controls described in this module. The controls should be scaled to the risk of the use case. A tool that drafts internal training text may require lighter review than a tool that updates case records, prioritizes outreach, or supports public communication during an emergency.

The module should be applied during readiness assessment, procurement, implementation planning, pilot testing, and post-deployment monitoring. Learners should document how the topic affects data, systems, users, safeguards, and

decision-making for the selected workflow.

Reflection Questions

Where does this topic appear in one current or proposed AI workflow in your agency?

What decisions, data sources, users, or partners would be affected if this area is poorly designed?

What evidence would governance reviewers need before approving the workflow?

Which safeguards should be required before pilot testing?

Who owns ongoing monitoring and review?

Practical Exercise

Create a AI business analysis package for one real or realistic public health AI workflow. The artifact should describe the use case, affected users, data sources, system dependencies, risks, guardrails, review points, and unresolved questions.

The exercise should produce a work product that could be used in a governance meeting, procurement review, implementation planning session, pilot evaluation, or monitoring review. Learners should document assumptions and identify which staff or partners need to be consulted.

Expected Artifact or Evidence

The expected artifact is a completed AI business analysis package. It should be specific to a public health workflow and should include technical dependencies, governance questions, human review expectations, monitoring indicators, and unresolved risks.

References and Resources for Additional Information

NIST AI Risk Management Framework

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP Top 10 for LLM Applications, when security or AI integrations are relevant

Agency privacy, cybersecurity, data governance, procurement, and records management policies

Knowledge Check

1. What is the best reason to address business analysis for ai-enabled public health systems before deployment?

- A. It makes governance unnecessary.
- B. It helps ensure the AI workflow is safe, useful, reviewable, and fit for public health operations.

- C. It allows the agency to skip testing.
- D. It guarantees that the vendor product will work.

Correct answer: B

2. Which practice best supports responsible implementation?

- A. Defining intended use, evidence requirements, safeguards, monitoring, and escalation before launch.
- B. Waiting until after launch to decide whether the tool is risky.
- C. Allowing users to decide individually how to use the tool.
- D. Treating vendor documentation as the only required evidence.

Correct answer: A

3. Which statement best describes a common failure mode?

- A. Human review is always unnecessary when AI is accurate.
- B. Documentation makes AI less useful.
- C. The agency assumes the tool is ready because the technology works in a demonstration, even though workflow, data, and governance have not been tested.
- D. Monitoring should stop after the pilot.

Correct answer: C

4. What should be included in the practical artifact for this module?

- A. Only a general definition.
- B. A marketing description of the AI tool.
- C. A list of unrelated data systems.
- D. Workflow-specific risks, safeguards, roles, evidence requirements, and unresolved questions.

Correct answer: D

5. When should the issue be escalated for governance or leadership review?

- A. Only if the tool is public-facing.
- B. When risks affect privacy, safety, equity, legal authority, public trust, system integrity, or operational decisions.
- C. Never, because technical staff can resolve all issues alone.
- D. Only after an incident has already occurred.

Correct answer: B

References and Resources

NIST AI Risk Management Framework

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP Top 10 for LLM Applications, when security or AI integrations are relevant

Agency privacy, cybersecurity, data governance, procurement, and records management policies

PMG 210

Requirements Writing for AI Procurement and Implementation

Course ID	PMG 210
Track	Program Management Role-Based Track
Level	applied/foundational
Estimated time	45
Audience	Project managers, Program managers, Product owners, Implementation leads, Procurement leads

Learning Objectives

- Explain the purpose of requirements writing for ai procurement and implementation in public health AI work.
- Identify the technical, operational, privacy, security, equity, and governance issues associated with the topic.
- Apply the topic to a real or realistic public health workflow, system, or implementation scenario.
- Recognize common failure modes that can affect safety, accuracy, trust, workflow fit, or sustainability.
- Identify practical guardrails that should be documented before implementation.
- Produce a AI requirements matrix that can support readiness assessment, governance review, procurement, implementation, or monitoring.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Why This Topic Matters for Public Health AI

Requirements Writing for AI Procurement and Implementation matters because AI systems are embedded in public health programs, data systems, and decision processes rather than operating in isolation. The topic affects whether AI outputs are accurate, explainable, safe, equitable, secure, and sustainable. Agencies that ignore this area may create tools that work in a demonstration but fail in actual practice.

The topic also affects accountability. Public health agencies must be able to explain how data are used, why a tool was approved, what evidence supports the workflow, who reviews outputs, and what happens when the system fails. These responsibilities become more important as AI tools move from experimentation into operational use.

Core Concepts

A requirement is a clear statement of what a system must do or how it must perform. AI requirements must cover functional capabilities, nonfunctional expectations, data handling, validation, transparency, monitoring, security,

accessibility, interoperability, and support.

Functional requirements describe what the tool must do, such as summarize a case report, classify a complaint, or generate a draft message. Nonfunctional requirements describe how the tool must perform, such as uptime, response time, security, auditability, explainability, and accessibility.

Procurement requirements must be testable. If a requirement cannot be evaluated during proposal review, implementation, or acceptance testing, it is unlikely to protect the agency. Requirements should avoid vague claims such as “uses responsible AI” unless they are translated into evidence, controls, and deliverables.

Public Health Example

A state health agency procuring an AI-supported case triage tool should require interoperability with the surveillance system, role-based access, source traceability, validation results, subgroup performance reporting, audit logs, human review controls, incident reporting, documentation, and data deletion at contract end. These requirements make vendor claims testable and enforceable.

Technical Deep Dive

Strong AI requirements begin with intended use. The agency should state the specific workflow, users, data sources, outputs, and decisions the system may support. It should also state prohibited uses. This prevents scope creep and reduces the chance that a tool will be used for decisions it was not designed or approved to support.

Data requirements should specify allowed data sources, minimum necessary fields, terminology standards, data quality expectations, retention, ownership, reuse limits, and deletion requirements. If vendors will use agency data to improve models, that use should be explicitly approved or prohibited.

Acceptance and monitoring requirements should specify what evidence the vendor must provide before launch and throughout operations. This may include validation reports, test cases, model cards, change logs, uptime reports, incident reports, bias assessments, security documentation, and user training materials.

Risks, Failure Modes, and Guardrails

A common failure mode is treating the topic as a technical detail rather than a public health control. When staff do not connect technical decisions to authority, workflow, equity, privacy, security, and accountability, the agency may adopt a tool that appears useful but cannot be sustained or defended.

Another risk is relying on vendor claims without agency-defined evidence. Vendors may provide demonstrations, dashboards, or summary metrics that do not match the agency workflow or data environment. A guardrail is to require documentation, test results, acceptance criteria, and agency-controlled review.

A third risk is failing to plan for change. Public health data, policies, systems, and emergencies change. Any technical approach must include monitoring, review, versioning, escalation, and retirement pathways so the agency can respond when assumptions no longer hold.

Application to AI-Supported Workflows

Generative AI, predictive analytics, RAG, automation, and agentic workflows all depend on the controls described in this module. The controls should be scaled to the risk of the use case. A tool that drafts internal training text may require

lighter review than a tool that updates case records, prioritizes outreach, or supports public communication during an emergency.

The module should be applied during readiness assessment, procurement, implementation planning, pilot testing, and post-deployment monitoring. Learners should document how the topic affects data, systems, users, safeguards, and decision-making for the selected workflow.

Reflection Questions

Where does this topic appear in one current or proposed AI workflow in your agency?

What decisions, data sources, users, or partners would be affected if this area is poorly designed?

What evidence would governance reviewers need before approving the workflow?

Which safeguards should be required before pilot testing?

Who owns ongoing monitoring and review?

Practical Exercise

Create a AI requirements matrix for one real or realistic public health AI workflow. The artifact should describe the use case, affected users, data sources, system dependencies, risks, guardrails, review points, and unresolved questions.

The exercise should produce a work product that could be used in a governance meeting, procurement review, implementation planning session, pilot evaluation, or monitoring review. Learners should document assumptions and identify which staff or partners need to be consulted.

Expected Artifact or Evidence

The expected artifact is a completed AI requirements matrix. It should be specific to a public health workflow and should include technical dependencies, governance questions, human review expectations, monitoring indicators, and unresolved risks.

References and Resources for Additional Information

NIST AI Risk Management Framework

ONC interoperability resources and agency procurement templates

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP Top 10 for LLM Applications, when security or AI integrations are relevant

Agency privacy, cybersecurity, data governance, procurement, and records management policies

Knowledge Check

1. What is the best reason to address requirements writing for ai procurement and implementation before deployment?

- A. It guarantees that the vendor product will work.
- B. It allows the agency to skip testing.
- C. It makes governance unnecessary.
- D. It helps ensure the AI workflow is safe, useful, reviewable, and fit for public health operations.

Correct answer: D

2. Which practice best supports responsible implementation?

- A. Defining intended use, evidence requirements, safeguards, monitoring, and escalation before launch.
- B. Waiting until after launch to decide whether the tool is risky.
- C. Allowing users to decide individually how to use the tool.
- D. Treating vendor documentation as the only required evidence.

Correct answer: A

3. Which statement best describes a common failure mode?

- A. Human review is always unnecessary when AI is accurate.
- B. Documentation makes AI less useful.
- C. The agency assumes the tool is ready because the technology works in a demonstration, even though workflow, data, and governance have not been tested.
- D. Monitoring should stop after the pilot.

Correct answer: C

4. What should be included in the practical artifact for this module?

- A. Only a general definition.
- B. A marketing description of the AI tool.
- C. A list of unrelated data systems.
- D. Workflow-specific risks, safeguards, roles, evidence requirements, and unresolved questions.

Correct answer: D

5. When should the issue be escalated for governance or leadership review?

- A. Only if the tool is public-facing.
- B. When risks affect privacy, safety, equity, legal authority, public trust, system integrity, or operational decisions.
- C. Never, because technical staff can resolve all issues alone.
- D. Only after an incident has already occurred.

Correct answer: B

References and Resources

NIST AI Risk Management Framework

ONC interoperability resources and agency procurement templates

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP Top 10 for LLM Applications, when security or AI integrations are relevant

Agency privacy, cybersecurity, data governance, procurement, and records management policies

PMG 220

AI Project Management and Implementation Planning

Course ID	PMG 220
Track	Program Management Role-Based Track
Level	applied/foundational
Estimated time	45
Audience	Project managers, Program managers, Product owners, Implementation leads, Procurement leads

Learning Objectives

- Define the major phases of an AI implementation project.
- Develop an AI project charter that includes scope, objectives, roles, risks, dependencies, and governance gates.
- Identify implementation dependencies across data, systems, privacy, security, procurement, workflow, training, and evaluation.
- Create a launch readiness checklist for an AI-supported workflow.
- Plan transition from pilot to operations, including ownership, monitoring, maintenance, and sustainability.

Public Health Example

A public health agency plans to pilot an AI tool that drafts summaries of incoming case reports. The project plan identifies governance approval, data access, system integration, validation, user training, SOP development, incident response, and monitoring as required workstreams. The project cannot launch simply because the tool works in a demonstration. It must pass readiness checks showing that staff know how to review the output, errors can be reported, and the tool is monitored after launch.

Technical and Operational Deep Dive

AI project planning should begin with a charter that defines the use case and sets boundaries. The charter should explain the public health problem, intended users, expected benefits, data involved, systems affected, approved use, prohibited use, and success measures. It should also identify the sponsor, product owner, implementation lead, technical lead, privacy or legal contacts, security contact, evaluation lead, and operational owner.

Implementation should be organized into phases. Early phases usually include discovery, intake, feasibility review, governance approval, and business analysis. Middle phases include design, procurement or configuration, data access setup, integration, validation, workflow design, SOP development, and training. Later phases include user acceptance testing, launch readiness review, deployment, monitoring, and transition to operations. Governance gates should be placed between phases to prevent premature launch.

AI project risk management should be more detailed than a general project risk log. Risks may include poor data quality, unavailable interfaces, unclear authority, model performance problems, vendor delays, privacy concerns, security vulnerabilities, staff resistance, workflow burden, community trust concerns, and sustainability gaps. Each risk should have an owner, mitigation plan, status, and escalation threshold.

Transition to operations is often neglected. A pilot may show promise, but routine use requires support, monitoring, maintenance, budget, staff training, user feedback, incident response, and periodic review. The project plan should define when the implementation team ends and who becomes responsible for ongoing operations.

Risks, Failure Modes, and Guardrails

Risks and failure modes include the following:

Pilots launching without governance approval or operational ownership.

Project plans overlooking data, privacy, security, workflow, or training dependencies.

Vendors controlling implementation timelines without agency readiness gates.

No transition plan from pilot to routine operations.

Success measured only by deployment rather than public health value and safe use.

Recommended guardrails include the following:

Use an AI-specific project charter and risk log.

Define governance gates before implementation begins.

Create separate workstreams for data, technology, governance, workflow, training, validation, and evaluation.

Require launch readiness review before deployment.

Document transition to operations and sustainability responsibilities.

Reflection Questions

Where does this topic appear in your agency, program, or workflow?

Which staff roles would need to understand or apply this concept before AI use is approved?

What could go wrong if this topic is not addressed before implementation?

What evidence would governance, leadership, or operations staff need before moving forward?

Which policy, data, equity, privacy, security, or workflow questions remain unresolved?

Practical Exercise

Create an AI project charter and implementation plan for one proposed AI-supported public health workflow. Include project purpose, scope, objectives, roles, governance gates, milestones, workstreams, dependencies, risks, launch readiness criteria, and transition-to-operations plan.

Expected Artifact or Evidence

AI project charter

Implementation workplan

AI-specific risk log

Launch readiness checklist

Transition-to-operations plan

References and Resources for Additional Information

NIST Artificial Intelligence Risk Management Framework

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP guidance for LLM application security

Agency policies for privacy, cybersecurity, records, procurement, accessibility, and public communications

Relevant public health program SOPs, data governance documentation, and implementation plans

Knowledge Check

1. What distinguishes AI project management from ordinary software implementation?

- A. AI projects never require privacy review
- B. AI projects are always shorter
- C. AI projects should be run only by vendors
- D. AI projects require attention to data, model behavior, governance, workflow, validation, monitoring, and human review

Correct answer: D

2. What is a governance gate?

- A. A decision point where approval is required before moving to the next phase
- B. A user password
- C. A model parameter
- D. A public dashboard

Correct answer: A

3. Which item belongs in an AI project charter?

- A. Only the tool name
- B. Only the vendor price
- C. Scope, intended use, roles, risks, dependencies, and success measures

D. Only a launch date

Correct answer: C

4. Why is transition to operations important?

A. It eliminates monitoring

B. It defines who owns support, monitoring, maintenance, training, and improvement after implementation

C. It ends governance review permanently

D. It prevents staff feedback

Correct answer: B

5. What is strong completion evidence?

A. A vendor brochure

B. A list of AI terms

C. A meeting agenda

D. A charter, workplan, risk log, launch readiness checklist, and transition plan

Correct answer: D

References and Resources

NIST Artificial Intelligence Risk Management Framework

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP guidance for LLM application security

Agency policies for privacy, cybersecurity, records, procurement, accessibility, and public communications

Relevant public health program SOPs, data governance documentation, and implementation plans

PMG 300

AI Vendor Implementation Management

Course ID	PMG 300
Track	Program Management Role-Based Track
Level	intermediate/practitioner
Estimated time	50
Audience	Project managers, Program managers, Product owners, Implementation leads, Procurement leads

Learning Objectives

- Explain the purpose of ai vendor implementation management in public health AI.
- Identify the roles, decisions, risks, and documentation needed for the topic.
- Apply the topic to a real or realistic public health workflow or decision.
- Recognize equity, privacy, security, legal, operational, and public trust implications.
- Produce a practical artifact that supports readiness, governance, implementation, monitoring, or accountability.

Definitions

Decision right	The authority to approve, deny, modify, pause, or escalate an AI use case or operational decision.
Evidence artifact	A document, log, checklist, test result, approval record, or review note that demonstrates responsible AI practice.
Escalation pathway	A defined route for raising concerns to the appropriate reviewer or decision-maker.
Accountability owner	The person or role responsible for ensuring that an AI use case is governed, monitored, and corrected when needed.
Operational safeguard	A procedure, control, or review step that reduces risk during day-to-day use.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Why This Topic Matters for Public Health AI

AI Vendor Implementation Management matters because AI adoption changes how public health agencies make, document, communicate, and review decisions. Without a structured approach, staff may rely on informal practices,

vendor assurances, or one-time approvals that do not match the operational consequences of AI-supported work.

Public health agencies operate under legal authorities, public accountability, resource constraints, and community trust obligations. A module on this topic helps agencies define repeatable expectations before urgent situations arise. It also helps staff distinguish low-risk administrative support from uses that affect surveillance, public communications, field response, resource allocation, civil rights, privacy, or official records.

Public Health Example

A public health agency is preparing to use AI in a workflow related to managing vendor delivery, configuration, testing, issue resolution, acceptance, and transition to operations. The project team initially treats the tool as a helpful efficiency aid, but governance reviewers recognize that the workflow may affect public trust, equity, official records, or public health action.

The agency responds by creating a clear artifact that defines intended use, prohibited use, responsible roles, review points, documentation expectations, escalation triggers, and monitoring measures. This converts a broad AI idea into a manageable public health practice.

Technical and Operational Deep Dive

The first step in this topic is to define the decision or workflow clearly. Staff should identify who is affected, what data are used, what AI output is produced, who reviews the output, and what action may follow. This prevents the module from becoming abstract and anchors the work in public health practice.

The second step is to assign responsibilities. AI-supported workflows often cross program, IT, legal, privacy, communications, procurement, and leadership boundaries. Responsibilities should be named by role rather than assumed. The artifact should specify who prepares materials, who reviews risks, who approves use, who monitors performance, and who can pause or revise the workflow.

The third step is to define evidence. Responsible AI depends on more than policy statements. Agencies need records that show what was reviewed, what decisions were made, what safeguards were applied, what incidents occurred, and how the workflow changed over time. Evidence also supports audits, public records review, quality improvement, and leadership accountability.

Risks, Failure Modes, and Guardrails

A common failure mode is treating the topic as a one-time checklist. AI risks can change when data sources, models, vendors, policies, users, or public expectations change. A guardrail is to define review intervals and change triggers.

Another failure mode is unclear ownership. When everyone is responsible in general, no one may be responsible in practice. A guardrail is to assign an accountability owner and name escalation pathways.

A third failure mode is relying on vendor claims or informal staff confidence without documentation. A guardrail is to require evidence artifacts that can be reviewed by governance, leadership, privacy, legal, equity, security, or program teams.

Application to AI-Supported Workflows

This topic applies across generative AI, predictive analytics, NLP, RAG, automation, and agentic workflows. The more sensitive the data, the more consequential the output, or the more automated the action, the stronger the documentation, review, and monitoring expectations should be.

Reflection Questions

What public health decision or workflow does this topic affect?

Who has authority to approve or pause the work?

What evidence would show that the workflow is responsible and effective?

What equity, privacy, legal, security, or public trust risks require review?

How will the agency know when the workflow needs to be revised?

Practical Exercise

Create a practical ai vendor implementation management artifact for one AI-supported public health use case. The artifact should define the use case, affected roles, review criteria, safeguards, documentation expectations, escalation triggers, monitoring indicators, and next review date.

Expected Artifact or Evidence

AI Vendor Implementation Management artifact

Role and responsibility table

Risk and safeguard summary

Escalation and monitoring plan

Approval or review note

References and Resources for Additional Information

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

Knowledge Check

1. What is the strongest evidence that ai vendor implementation management has been applied responsibly?

- A. A general statement that AI is useful
- B. A vendor demonstration with no agency review
- C. A documented artifact that defines use, roles, safeguards, escalation, monitoring, and evidence
- D. A decision to skip review because the use case seems efficient

Correct answer: C

2. What is the best reason to assign an accountability owner?

- A. To reduce the need for documentation
- B. To ensure someone is responsible for monitoring, escalation, and correction
- C. To make the vendor solely responsible
- D. To prevent staff from raising concerns

Correct answer: B

3. When should an AI-supported workflow be escalated for review?

- A. Only after media coverage
- B. Never if the tool is efficient
- C. When outputs may affect safety, equity, privacy, legal authority, public trust, or official action
- D. Only at the end of the fiscal year

Correct answer: C

4. Which practice best supports public accountability?

- A. Keeping AI use informal
- B. Avoiding written records
- C. Letting each user decide independently
- D. Maintaining evidence artifacts, decision records, review notes, and monitoring results

Correct answer: D

5. What should happen when data sources, policy, vendor configuration, or workflow use changes?

- A. The use case should be reviewed to determine whether safeguards, validation, documentation, or approvals need updating
- B. The agency should ignore the change if staff like the tool
- C. The model should be trusted automatically
- D. The glossary should be deleted

Correct answer: A

References and Resources

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

PMG 310

Grant and Funding Management for AI Projects

Course ID	PMG 310
Track	Program Management Role-Based Track
Level	intermediate/practitioner
Estimated time	50
Audience	Project managers, Program managers, Product owners, Implementation leads, Procurement leads

Learning Objectives

- Explain the purpose of grant and funding management for ai projects in public health AI.
- Identify the roles, decisions, risks, and documentation needed for the topic.
- Apply the topic to a real or realistic public health workflow or decision.
- Recognize equity, privacy, security, legal, operational, and public trust implications.
- Produce a practical artifact that supports readiness, governance, implementation, monitoring, or accountability.

Definitions

Decision right	The authority to approve, deny, modify, pause, or escalate an AI use case or operational decision.
Evidence artifact	A document, log, checklist, test result, approval record, or review note that demonstrates responsible AI practice.
Escalation pathway	A defined route for raising concerns to the appropriate reviewer or decision-maker.
Accountability owner	The person or role responsible for ensuring that an AI use case is governed, monitored, and corrected when needed.
Operational safeguard	A procedure, control, or review step that reduces risk during day-to-day use.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Why This Topic Matters for Public Health AI

Grant and Funding Management for AI Projects matters because AI adoption changes how public health agencies make, document, communicate, and review decisions. Without a structured approach, staff may rely on informal

practices, vendor assurances, or one-time approvals that do not match the operational consequences of AI-supported work.

Public health agencies operate under legal authorities, public accountability, resource constraints, and community trust obligations. A module on this topic helps agencies define repeatable expectations before urgent situations arise. It also helps staff distinguish low-risk administrative support from uses that affect surveillance, public communications, field response, resource allocation, civil rights, privacy, or official records.

Public Health Example

A public health agency is preparing to use AI in a workflow related to connecting AI projects to allowable costs, deliverables, reporting, performance measures, and sustainability. The project team initially treats the tool as a helpful efficiency aid, but governance reviewers recognize that the workflow may affect public trust, equity, official records, or public health action.

The agency responds by creating a clear artifact that defines intended use, prohibited use, responsible roles, review points, documentation expectations, escalation triggers, and monitoring measures. This converts a broad AI idea into a manageable public health practice.

Technical and Operational Deep Dive

The first step in this topic is to define the decision or workflow clearly. Staff should identify who is affected, what data are used, what AI output is produced, who reviews the output, and what action may follow. This prevents the module from becoming abstract and anchors the work in public health practice.

The second step is to assign responsibilities. AI-supported workflows often cross program, IT, legal, privacy, communications, procurement, and leadership boundaries. Responsibilities should be named by role rather than assumed. The artifact should specify who prepares materials, who reviews risks, who approves use, who monitors performance, and who can pause or revise the workflow.

The third step is to define evidence. Responsible AI depends on more than policy statements. Agencies need records that show what was reviewed, what decisions were made, what safeguards were applied, what incidents occurred, and how the workflow changed over time. Evidence also supports audits, public records review, quality improvement, and leadership accountability.

Risks, Failure Modes, and Guardrails

A common failure mode is treating the topic as a one-time checklist. AI risks can change when data sources, models, vendors, policies, users, or public expectations change. A guardrail is to define review intervals and change triggers.

Another failure mode is unclear ownership. When everyone is responsible in general, no one may be responsible in practice. A guardrail is to assign an accountability owner and name escalation pathways.

A third failure mode is relying on vendor claims or informal staff confidence without documentation. A guardrail is to require evidence artifacts that can be reviewed by governance, leadership, privacy, legal, equity, security, or program teams.

Application to AI-Supported Workflows

This topic applies across generative AI, predictive analytics, NLP, RAG, automation, and agentic workflows. The more sensitive the data, the more consequential the output, or the more automated the action, the stronger the documentation, review, and monitoring expectations should be.

Reflection Questions

What public health decision or workflow does this topic affect?

Who has authority to approve or pause the work?

What evidence would show that the workflow is responsible and effective?

What equity, privacy, legal, security, or public trust risks require review?

How will the agency know when the workflow needs to be revised?

Practical Exercise

Create a practical grant and funding management for AI projects artifact for one AI-supported public health use case. The artifact should define the use case, affected roles, review criteria, safeguards, documentation expectations, escalation triggers, monitoring indicators, and next review date.

Expected Artifact or Evidence

Grant and Funding Management for AI Projects artifact

Role and responsibility table

Risk and safeguard summary

Escalation and monitoring plan

Approval or review note

References and Resources for Additional Information

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

Appendix. Common Source Base Used Across Modules

CDC Public Health Data Strategy and CDC Data Modernization resources.

CDC electronic case reporting, electronic laboratory reporting, surveillance, and public health data quality resources.

HL7 FHIR, HL7 US Core, public health implementation guides, CDA, and HL7 v2 implementation materials.

ONC USCDI, health IT, interoperability, and information blocking resources.

NIST Artificial Intelligence Risk Management Framework, NIST Generative AI Profile, NIST Privacy Framework, and related NIST cybersecurity guidance.

OWASP Top 10 for Large Language Model Applications and related secure AI application guidance.

WHO ethics and governance guidance for artificial intelligence in health.

Agency policies and procedures for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, records retention, incident response, and emergency communications.

Knowledge Check

1. What is the strongest evidence that grant and funding management for ai projects has been applied responsibly?

- A. A general statement that AI is useful
- B. A vendor demonstration with no agency review
- C. A documented artifact that defines use, roles, safeguards, escalation, monitoring, and evidence
- D. A decision to skip review because the use case seems efficient

Correct answer: C

2. What is the best reason to assign an accountability owner?

- A. To reduce the need for documentation
- B. To ensure someone is responsible for monitoring, escalation, and correction
- C. To make the vendor solely responsible
- D. To prevent staff from raising concerns

Correct answer: B

3. When should an AI-supported workflow be escalated for review?

- A. Only after media coverage
- B. Never if the tool is efficient
- C. When outputs may affect safety, equity, privacy, legal authority, public trust, or official action
- D. Only at the end of the fiscal year

Correct answer: C

4. Which practice best supports public accountability?

- A. Keeping AI use informal
- B. Avoiding written records
- C. Letting each user decide independently
- D. Maintaining evidence artifacts, decision records, review notes, and monitoring results

Correct answer: D

5. What should happen when data sources, policy, vendor configuration, or workflow use changes?

- A. The use case should be reviewed to determine whether safeguards, validation, documentation, or approvals need updating
- B. The agency should ignore the change if staff like the tool
- C. The model should be trusted automatically
- D. The glossary should be deleted

Correct answer: A

References and Resources

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

Appendix. Common Source Base Used Across Modules

CDC Public Health Data Strategy and CDC Data Modernization resources.

CDC electronic case reporting, electronic laboratory reporting, surveillance, and public health data quality resources.

HL7 FHIR, HL7 US Core, public health implementation guides, CDA, and HL7 v2 implementation materials.

ONC USCDI, health IT, interoperability, and information blocking resources.

NIST Artificial Intelligence Risk Management Framework, NIST Generative AI Profile, NIST Privacy Framework, and related NIST cybersecurity guidance.

OWASP Top 10 for Large Language Model Applications and related secure AI application guidance.

WHO ethics and governance guidance for artificial intelligence in health.

Agency policies and procedures for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, records retention, incident response, and emergency communications.

POL 300

Public Health AI Policy Landscape and Regulatory Readiness

Course ID	POL 300
Track	Policy Role-Based Track
Level	intermediate/practitioner
Estimated time	50
Audience	Policy staff, Legal staff, Executive leaders, Government affairs staff, Program leadership

Learning Objectives

- Identify categories of law, policy, standards, and guidance that may affect public health AI.
- Explain why regulatory readiness requires monitoring, interpretation, documentation, and policy maintenance.
- Distinguish between binding requirements, agency policy, professional guidance, standards, and best practices.
- Create a policy scan process for AI-related developments.
- Produce an AI regulatory readiness memo or policy update plan.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Public Health Example

A local health department is considering AI tools for translation, complaint triage, and case note summarization. A policy scan identifies privacy, public records, procurement, accessibility, cybersecurity, and civil rights considerations. The agency creates a readiness memo that distinguishes existing requirements from emerging guidance, recommends interim safeguards, and establishes a quarterly process to update AI policies as requirements evolve.

Technical and Operational Deep Dive

Policy readiness begins by identifying relevant domains rather than looking for a single AI law. Public health AI may be affected by health privacy rules, state confidentiality laws, public records requirements, procurement standards, accessibility laws, cybersecurity policies, civil rights obligations, data sharing agreements, records retention rules, and professional ethics. AI-specific policy should be layered onto this existing policy environment.

Agencies should distinguish between different levels of authority. A statute, regulation, contract clause, grant requirement, executive order, agency policy, professional guideline, and technical standard do not all carry the same legal weight. However, nonbinding guidance can still be important because it may reflect emerging expectations for

responsible practice, risk management, transparency, documentation, or auditability.

A policy scan should be assigned to a responsible role or group. The scan should identify sources to monitor, frequency of review, criteria for relevance, and a process for escalating items that may affect agency policy or active AI use cases. The process should avoid overreacting to every news item while still ensuring that meaningful changes are not missed.

Regulatory readiness also requires change management. When policy requirements change, the agency may need to update SOPs, procurement language, risk tiers, training, documentation, vendor oversight, public transparency materials, or incident response procedures. A policy update is not complete until it has been translated into operational practice.

Risks, Failure Modes, and Guardrails

Risks and failure modes include the following:

Assuming there is no policy risk because there is no single AI statute.

Treating nonbinding guidance as irrelevant even when it reflects expected practice.

Failing to update internal policies after legal, procurement, privacy, or cybersecurity changes.

Allowing vendors to define compliance obligations without agency review.

Maintaining policies that are too vague to guide real decisions.

Recommended guardrails include the following:

Assign responsibility for AI policy scanning and interpretation.

Distinguish binding requirements from guidance and best practices.

Maintain a crosswalk from external requirements to internal policies and procedures.

Update training, SOPs, and procurement language when policies change.

Document assumptions, decisions, and unresolved legal or policy questions.

Reflection Questions

Where does this topic appear in your agency, program, or workflow?

Which staff roles would need to understand or apply this concept before AI use is approved?

What could go wrong if this topic is not addressed before implementation?

What evidence would governance, leadership, or operations staff need before moving forward?

Which policy, data, equity, privacy, security, or workflow questions remain unresolved?

Practical Exercise

Create a public health AI policy scan and regulatory readiness memo. Identify relevant policy domains, sources to monitor, current agency policy gaps, likely operational implications, recommended interim safeguards, and a process for ongoing updates.

Expected Artifact or Evidence

AI policy scan plan

Regulatory readiness memo

Policy-domain crosswalk

Recommended policy updates

Ongoing monitoring process

References and Resources for Additional Information

NIST Artificial Intelligence Risk Management Framework

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP guidance for LLM application security

Agency policies for privacy, cybersecurity, records, procurement, accessibility, and public communications

Relevant public health program SOPs, data governance documentation, and implementation plans

Knowledge Check

1. What does regulatory readiness mean in this module?

- A. Waiting until every AI law is final
- B. The ability to identify, interpret, document, and respond to requirements that affect AI use
- C. Relying only on vendor assurances
- D. Avoiding all AI use permanently

Correct answer: B

2. Which statement is most accurate?

- A. Only AI-specific laws matter
- B. Guidance is always binding
- C. Public health AI is unrelated to privacy or public records
- D. AI may be affected by privacy, procurement, records, cybersecurity, accessibility, civil rights, and agency policies

Correct answer: D

3. What is a policy scan?

- A. A structured review of developments that may require policy or practice changes
- B. A model evaluation score
- C. A source code review only
- D. A public dashboard

Correct answer: A

4. What should happen after a policy change affects AI use?

- A. Nothing until a vendor acts
- B. The agency should delete all documentation
- C. Relevant SOPs, training, procurement language, and governance practices should be reviewed and updated
- D. Staff should improvise case by case

Correct answer: C

5. What is strong completion evidence?

- A. A policy scan and readiness memo with sources, gaps, implications, safeguards, and update process
- B. A single news headline
- C. A generic statement that policy is changing
- D. A vendor presentation

Correct answer: A

References and Resources

NIST Artificial Intelligence Risk Management Framework

NIST Generative AI Profile

CDC Public Health Data Strategy and Data Modernization resources

OWASP guidance for LLM application security

Agency policies for privacy, cybersecurity, records, procurement, accessibility, and public communications

Relevant public health program SOPs, data governance documentation, and implementation plans

POL 310

AI Policy Development and Approval Workflow

Course ID	POL 310
Track	Policy Role-Based Track
Level	intermediate/practitioner
Estimated time	50
Audience	Policy staff, Legal staff, Executive leaders, Government affairs staff, Program leadership

Learning Objectives

- Explain the purpose of ai policy development and approval workflow in public health AI.
- Identify the roles, decisions, risks, and documentation needed for the topic.
- Apply the topic to a real or realistic public health workflow or decision.
- Recognize equity, privacy, security, legal, operational, and public trust implications.
- Produce a practical artifact that supports readiness, governance, implementation, monitoring, or accountability.

Definitions

Decision right	The authority to approve, deny, modify, pause, or escalate an AI use case or operational decision.
Evidence artifact	A document, log, checklist, test result, approval record, or review note that demonstrates responsible AI practice.
Escalation pathway	A defined route for raising concerns to the appropriate reviewer or decision-maker.
Accountability owner	The person or role responsible for ensuring that an AI use case is governed, monitored, and corrected when needed.
Operational safeguard	A procedure, control, or review step that reduces risk during day-to-day use.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Why This Topic Matters for Public Health AI

AI Policy Development and Approval Workflow matters because AI adoption changes how public health agencies make, document, communicate, and review decisions. Without a structured approach, staff may rely on informal practices,

vendor assurances, or one-time approvals that do not match the operational consequences of AI-supported work.

Public health agencies operate under legal authorities, public accountability, resource constraints, and community trust obligations. A module on this topic helps agencies define repeatable expectations before urgent situations arise. It also helps staff distinguish low-risk administrative support from uses that affect surveillance, public communications, field response, resource allocation, civil rights, privacy, or official records.

Public Health Example

A public health agency is preparing to use AI in a workflow related to creating, approving, implementing, and updating internal AI policies. The project team initially treats the tool as a helpful efficiency aid, but governance reviewers recognize that the workflow may affect public trust, equity, official records, or public health action.

The agency responds by creating a clear artifact that defines intended use, prohibited use, responsible roles, review points, documentation expectations, escalation triggers, and monitoring measures. This converts a broad AI idea into a manageable public health practice.

Technical and Operational Deep Dive

The first step in this topic is to define the decision or workflow clearly. Staff should identify who is affected, what data are used, what AI output is produced, who reviews the output, and what action may follow. This prevents the module from becoming abstract and anchors the work in public health practice.

The second step is to assign responsibilities. AI-supported workflows often cross program, IT, legal, privacy, communications, procurement, and leadership boundaries. Responsibilities should be named by role rather than assumed. The artifact should specify who prepares materials, who reviews risks, who approves use, who monitors performance, and who can pause or revise the workflow.

The third step is to define evidence. Responsible AI depends on more than policy statements. Agencies need records that show what was reviewed, what decisions were made, what safeguards were applied, what incidents occurred, and how the workflow changed over time. Evidence also supports audits, public records review, quality improvement, and leadership accountability.

Risks, Failure Modes, and Guardrails

A common failure mode is treating the topic as a one-time checklist. AI risks can change when data sources, models, vendors, policies, users, or public expectations change. A guardrail is to define review intervals and change triggers.

Another failure mode is unclear ownership. When everyone is responsible in general, no one may be responsible in practice. A guardrail is to assign an accountability owner and name escalation pathways.

A third failure mode is relying on vendor claims or informal staff confidence without documentation. A guardrail is to require evidence artifacts that can be reviewed by governance, leadership, privacy, legal, equity, security, or program teams.

Application to AI-Supported Workflows

This topic applies across generative AI, predictive analytics, NLP, RAG, automation, and agentic workflows. The more sensitive the data, the more consequential the output, or the more automated the action, the stronger the

documentation, review, and monitoring expectations should be.

Reflection Questions

What public health decision or workflow does this topic affect?

Who has authority to approve or pause the work?

What evidence would show that the workflow is responsible and effective?

What equity, privacy, legal, security, or public trust risks require review?

How will the agency know when the workflow needs to be revised?

Practical Exercise

Create a practical ai policy development and approval workflow artifact for one AI-supported public health use case. The artifact should define the use case, affected roles, review criteria, safeguards, documentation expectations, escalation triggers, monitoring indicators, and next review date.

Expected Artifact or Evidence

AI Policy Development and Approval Workflow artifact

Role and responsibility table

Risk and safeguard summary

Escalation and monitoring plan

Approval or review note

References and Resources for Additional Information

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

Knowledge Check

1. What is the strongest evidence that ai policy development and approval workflow has been applied responsibly?

- A. A general statement that AI is useful
- B. A vendor demonstration with no agency review
- C. A documented artifact that defines use, roles, safeguards, escalation, monitoring, and evidence
- D. A decision to skip review because the use case seems efficient

Correct answer: C

2. What is the best reason to assign an accountability owner?

- A. To reduce the need for documentation
- B. To ensure someone is responsible for monitoring, escalation, and correction
- C. To make the vendor solely responsible
- D. To prevent staff from raising concerns

Correct answer: B

3. When should an AI-supported workflow be escalated for review?

- A. Only after media coverage
- B. Never if the tool is efficient
- C. When outputs may affect safety, equity, privacy, legal authority, public trust, or official action
- D. Only at the end of the fiscal year

Correct answer: C

4. Which practice best supports public accountability?

- A. Keeping AI use informal
- B. Avoiding written records
- C. Letting each user decide independently
- D. Maintaining evidence artifacts, decision records, review notes, and monitoring results

Correct answer: D

5. What should happen when data sources, policy, vendor configuration, or workflow use changes?

- A. The use case should be reviewed to determine whether safeguards, validation, documentation, or approvals need updating
- B. The agency should ignore the change if staff like the tool
- C. The model should be trusted automatically
- D. The glossary should be deleted

Correct answer: A

References and Resources

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

POL 320

AI Authority, Liability, and Public Health Decision-Making

Course ID	POL 320
Track	Policy Role-Based Track
Level	intermediate/practitioner
Estimated time	50
Audience	Policy staff, Legal staff, Executive leaders, Government affairs staff, Program leadership

Learning Objectives

- Explain the purpose of ai authority, liability, and public health decision-making in public health AI.
- Identify the roles, decisions, risks, and documentation needed for the topic.
- Apply the topic to a real or realistic public health workflow or decision.
- Recognize equity, privacy, security, legal, operational, and public trust implications.
- Produce a practical artifact that supports readiness, governance, implementation, monitoring, or accountability.

Definitions

Decision right	The authority to approve, deny, modify, pause, or escalate an AI use case or operational decision.
Evidence artifact	A document, log, checklist, test result, approval record, or review note that demonstrates responsible AI practice.
Escalation pathway	A defined route for raising concerns to the appropriate reviewer or decision-maker.
Accountability owner	The person or role responsible for ensuring that an AI use case is governed, monitored, and corrected when needed.
Operational safeguard	A procedure, control, or review step that reduces risk during day-to-day use.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Why This Topic Matters for Public Health AI

AI Authority, Liability, and Public Health Decision-Making matters because AI adoption changes how public health agencies make, document, communicate, and review decisions. Without a structured approach, staff may rely on

informal practices, vendor assurances, or one-time approvals that do not match the operational consequences of AI-supported work.

Public health agencies operate under legal authorities, public accountability, resource constraints, and community trust obligations. A module on this topic helps agencies define repeatable expectations before urgent situations arise. It also helps staff distinguish low-risk administrative support from uses that affect surveillance, public communications, field response, resource allocation, civil rights, privacy, or official records.

Public Health Example

A public health agency is preparing to use AI in a workflow related to legal authority, delegated responsibility, accountability, and liability in AI-supported decisions. The project team initially treats the tool as a helpful efficiency aid, but governance reviewers recognize that the workflow may affect public trust, equity, official records, or public health action.

The agency responds by creating a clear artifact that defines intended use, prohibited use, responsible roles, review points, documentation expectations, escalation triggers, and monitoring measures. This converts a broad AI idea into a manageable public health practice.

Technical and Operational Deep Dive

The first step in this topic is to define the decision or workflow clearly. Staff should identify who is affected, what data are used, what AI output is produced, who reviews the output, and what action may follow. This prevents the module from becoming abstract and anchors the work in public health practice.

The second step is to assign responsibilities. AI-supported workflows often cross program, IT, legal, privacy, communications, procurement, and leadership boundaries. Responsibilities should be named by role rather than assumed. The artifact should specify who prepares materials, who reviews risks, who approves use, who monitors performance, and who can pause or revise the workflow.

The third step is to define evidence. Responsible AI depends on more than policy statements. Agencies need records that show what was reviewed, what decisions were made, what safeguards were applied, what incidents occurred, and how the workflow changed over time. Evidence also supports audits, public records review, quality improvement, and leadership accountability.

Risks, Failure Modes, and Guardrails

A common failure mode is treating the topic as a one-time checklist. AI risks can change when data sources, models, vendors, policies, users, or public expectations change. A guardrail is to define review intervals and change triggers.

Another failure mode is unclear ownership. When everyone is responsible in general, no one may be responsible in practice. A guardrail is to assign an accountability owner and name escalation pathways.

A third failure mode is relying on vendor claims or informal staff confidence without documentation. A guardrail is to require evidence artifacts that can be reviewed by governance, leadership, privacy, legal, equity, security, or program teams.

Application to AI-Supported Workflows

This topic applies across generative AI, predictive analytics, NLP, RAG, automation, and agentic workflows. The more sensitive the data, the more consequential the output, or the more automated the action, the stronger the documentation, review, and monitoring expectations should be.

Reflection Questions

What public health decision or workflow does this topic affect?

Who has authority to approve or pause the work?

What evidence would show that the workflow is responsible and effective?

What equity, privacy, legal, security, or public trust risks require review?

How will the agency know when the workflow needs to be revised?

Practical Exercise

Create a practical ai authority, liability, and public health decision-making artifact for one AI-supported public health use case. The artifact should define the use case, affected roles, review criteria, safeguards, documentation expectations, escalation triggers, monitoring indicators, and next review date.

Expected Artifact or Evidence

AI Authority, Liability, and Public Health Decision-Making artifact

Role and responsibility table

Risk and safeguard summary

Escalation and monitoring plan

Approval or review note

References and Resources for Additional Information

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

Knowledge Check

1. What is the strongest evidence that ai authority, liability, and public health decision-making has been applied responsibly?

- A. A general statement that AI is useful
- B. A vendor demonstration with no agency review
- C. A documented artifact that defines use, roles, safeguards, escalation, monitoring, and evidence
- D. A decision to skip review because the use case seems efficient

Correct answer: C

2. What is the best reason to assign an accountability owner?

- A. To reduce the need for documentation
- B. To ensure someone is responsible for monitoring, escalation, and correction
- C. To make the vendor solely responsible
- D. To prevent staff from raising concerns

Correct answer: B

3. When should an AI-supported workflow be escalated for review?

- A. Only after media coverage
- B. Never if the tool is efficient
- C. When outputs may affect safety, equity, privacy, legal authority, public trust, or official action
- D. Only at the end of the fiscal year

Correct answer: C

4. Which practice best supports public accountability?

- A. Keeping AI use informal
- B. Avoiding written records
- C. Letting each user decide independently
- D. Maintaining evidence artifacts, decision records, review notes, and monitoring results

Correct answer: D

5. What should happen when data sources, policy, vendor configuration, or workflow use changes?

- A. The use case should be reviewed to determine whether safeguards, validation, documentation, or approvals need updating
- B. The agency should ignore the change if staff like the tool
- C. The model should be trusted automatically
- D. The glossary should be deleted

Correct answer: A

References and Resources

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

POL 330

AI Equity Policy and Civil Rights Considerations

Course ID	POL 330
Track	Policy Role-Based Track
Level	intermediate/practitioner
Estimated time	50
Audience	Policy staff, Legal staff, Executive leaders, Government affairs staff, Program leadership

Learning Objectives

- Explain the purpose of ai equity policy and civil rights considerations in public health AI.
- Identify the roles, decisions, risks, and documentation needed for the topic.
- Apply the topic to a real or realistic public health workflow or decision.
- Recognize equity, privacy, security, legal, operational, and public trust implications.
- Produce a practical artifact that supports readiness, governance, implementation, monitoring, or accountability.

Definitions

Decision right	The authority to approve, deny, modify, pause, or escalate an AI use case or operational decision.
Evidence artifact	A document, log, checklist, test result, approval record, or review note that demonstrates responsible AI practice.
Escalation pathway	A defined route for raising concerns to the appropriate reviewer or decision-maker.
Accountability owner	The person or role responsible for ensuring that an AI use case is governed, monitored, and corrected when needed.
Operational safeguard	A procedure, control, or review step that reduces risk during day-to-day use.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Why This Topic Matters for Public Health AI

AI Equity Policy and Civil Rights Considerations matters because AI adoption changes how public health agencies make, document, communicate, and review decisions. Without a structured approach, staff may rely on informal

practices, vendor assurances, or one-time approvals that do not match the operational consequences of AI-supported work.

Public health agencies operate under legal authorities, public accountability, resource constraints, and community trust obligations. A module on this topic helps agencies define repeatable expectations before urgent situations arise. It also helps staff distinguish low-risk administrative support from uses that affect surveillance, public communications, field response, resource allocation, civil rights, privacy, or official records.

Public Health Example

A public health agency is preparing to use AI in a workflow related to civil rights, nondiscrimination, accessibility, language access, disparate impact, and equity safeguards. The project team initially treats the tool as a helpful efficiency aid, but governance reviewers recognize that the workflow may affect public trust, equity, official records, or public health action.

The agency responds by creating a clear artifact that defines intended use, prohibited use, responsible roles, review points, documentation expectations, escalation triggers, and monitoring measures. This converts a broad AI idea into a manageable public health practice.

Technical and Operational Deep Dive

The first step in this topic is to define the decision or workflow clearly. Staff should identify who is affected, what data are used, what AI output is produced, who reviews the output, and what action may follow. This prevents the module from becoming abstract and anchors the work in public health practice.

The second step is to assign responsibilities. AI-supported workflows often cross program, IT, legal, privacy, communications, procurement, and leadership boundaries. Responsibilities should be named by role rather than assumed. The artifact should specify who prepares materials, who reviews risks, who approves use, who monitors performance, and who can pause or revise the workflow.

The third step is to define evidence. Responsible AI depends on more than policy statements. Agencies need records that show what was reviewed, what decisions were made, what safeguards were applied, what incidents occurred, and how the workflow changed over time. Evidence also supports audits, public records review, quality improvement, and leadership accountability.

Risks, Failure Modes, and Guardrails

A common failure mode is treating the topic as a one-time checklist. AI risks can change when data sources, models, vendors, policies, users, or public expectations change. A guardrail is to define review intervals and change triggers.

Another failure mode is unclear ownership. When everyone is responsible in general, no one may be responsible in practice. A guardrail is to assign an accountability owner and name escalation pathways.

A third failure mode is relying on vendor claims or informal staff confidence without documentation. A guardrail is to require evidence artifacts that can be reviewed by governance, leadership, privacy, legal, equity, security, or program teams.

Application to AI-Supported Workflows

This topic applies across generative AI, predictive analytics, NLP, RAG, automation, and agentic workflows. The more sensitive the data, the more consequential the output, or the more automated the action, the stronger the documentation, review, and monitoring expectations should be.

Reflection Questions

What public health decision or workflow does this topic affect?

Who has authority to approve or pause the work?

What evidence would show that the workflow is responsible and effective?

What equity, privacy, legal, security, or public trust risks require review?

How will the agency know when the workflow needs to be revised?

Practical Exercise

Create a practical AI equity policy and civil rights considerations artifact for one AI-supported public health use case. The artifact should define the use case, affected roles, review criteria, safeguards, documentation expectations, escalation triggers, monitoring indicators, and next review date.

Expected Artifact or Evidence

AI Equity Policy and Civil Rights Considerations artifact

Role and responsibility table

Risk and safeguard summary

Escalation and monitoring plan

Approval or review note

References and Resources for Additional Information

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications

Knowledge Check

1. What is the strongest evidence that ai equity policy and civil rights considerations has been applied responsibly?

- A. A general statement that AI is useful
- B. A vendor demonstration with no agency review
- C. A documented artifact that defines use, roles, safeguards, escalation, monitoring, and evidence
- D. A decision to skip review because the use case seems efficient

Correct answer: C

2. What is the best reason to assign an accountability owner?

- A. To reduce the need for documentation
- B. To ensure someone is responsible for monitoring, escalation, and correction
- C. To make the vendor solely responsible
- D. To prevent staff from raising concerns

Correct answer: B

3. When should an AI-supported workflow be escalated for review?

- A. Only after media coverage
- B. Never if the tool is efficient
- C. When outputs may affect safety, equity, privacy, legal authority, public trust, or official action
- D. Only at the end of the fiscal year

Correct answer: C

4. Which practice best supports public accountability?

- A. Keeping AI use informal
- B. Avoiding written records
- C. Letting each user decide independently
- D. Maintaining evidence artifacts, decision records, review notes, and monitoring results

Correct answer: D

5. What should happen when data sources, policy, vendor configuration, or workflow use changes?

- A. The use case should be reviewed to determine whether safeguards, validation, documentation, or approvals need updating
- B. The agency should ignore the change if staff like the tool
- C. The model should be trusted automatically
- D. The glossary should be deleted

Correct answer: A

References and Resources

CDC Public Health Data Strategy and Data Modernization resources

NIST Artificial Intelligence Risk Management Framework and NIST Generative AI Profile

WHO Ethics and Governance of Artificial Intelligence for Health guidance

OWASP Top 10 for Large Language Model Applications

ONC health IT, interoperability, and information blocking resources

Agency policies for privacy, public records, cybersecurity, procurement, civil rights, accessibility, language access, and emergency communications