AI Playbook and Toolkit for Public Health Departments

AI Audit Readiness and Evidence Management

GOV 340

AI audit readiness is the agency’s ability to show what AI systems are used, why they were approved, what data they rely on, how they were evaluated, who is accountable, what safeguards are in place, and how issues are monitored and corrected. Public health agencies may need this evidence for internal oversight, public records requests, legal review, grant reporting, procurement monitoring, incident response, or public accountability. This module teaches governance, compliance, program, and leadership staff how to manage evidence across the AI lifecycle. The emphasis is practical: if a concern arises six months after deployment, the agency should be able to reconstruct the decision pathway, validation evidence, user guidance, vendor commitments, incidents, and changes over time. Audit readiness does not mean creating paperwork for its own sake. It means preserving enough evidence to demonstrate responsible decision-making, support improvement, and protect public trust.

Level	intermediate/practitioner
Primary track	Governance and Security Track
Appears in tracks	governance-and-security, policy, program-management, public-health-executive-leadership

Learning Objectives

- Explain why AI systems require lifecycle evidence management.
- Identify the core records that support AI audit readiness.
- Distinguish between approval evidence, validation evidence, operational evidence, and incident evidence.
- Create an AI evidence inventory for an approved or proposed use case.
- Develop retention, ownership, and review expectations for AI governance artifacts.

Module Overview

AI audit readiness is the agency's ability to show what AI systems are used, why they were approved, what data they rely on, how they were evaluated, who is accountable, what safeguards are in place, and how issues are monitored and corrected. Public health agencies may need this evidence for internal oversight, public records requests, legal review, grant reporting, procurement monitoring, incident response, or public accountability.

This module teaches governance, compliance, program, and leadership staff how to manage evidence across the AI lifecycle. The emphasis is practical: if a concern arises six months after deployment, the agency should be able to reconstruct the decision pathway, validation evidence, user guidance, vendor commitments, incidents, and changes over time.

Audit readiness does not mean creating paperwork for its own sake. It means preserving enough evidence to demonstrate responsible decision-making, support improvement, and protect public trust.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Public Health Example

A health department uses an AI tool to help classify public inquiries. Six months later, a concern is raised that the tool may be misclassifying language access requests. Audit-ready evidence would include the approved use case, validation plan, subgroup review, SOP, training records, user override logs, incident reports, vendor documentation, and monitoring results. This evidence allows the agency to investigate the concern and decide whether to modify, pause, or retire the workflow.

Technical and Operational Deep Dive

AI evidence begins at intake. The agency should preserve the use case description, risk triage, required reviewers, approval conditions, and rationale for the decision. If a use case is rejected or deferred, that decision should also be documented. These early records help explain why a tool exists and what limits were placed on its use.

Validation evidence should show how the agency determined whether the AI-supported workflow was fit for purpose. This may include test datasets, evaluation rubrics, accuracy measures, subgroup performance, user acceptance testing, security review, privacy review, accessibility review, and pilot findings. Validation evidence should be understandable to decision-makers, not only technical staff.

Operational evidence should show what happened after deployment. This may include monitoring dashboards, user override rates, error reports, support tickets, incident records, model or prompt changes, vendor updates, SOP revisions, training completion, and governance meeting minutes. Operational evidence is especially important because AI performance and workflow use can change over time.

Evidence management requires ownership and retention rules. Governance bodies should define who stores each artifact, where it is stored, how long it is retained, who can access it, and how evidence is updated when tools change. Audit readiness depends on consistent records, not scattered files and informal emails.

Risks, Failure Modes, and Guardrails

Risks and failure modes include the following:

Approvals without records of conditions or rationale.

Validation results stored informally or not retained.

Vendor documentation accepted without agency review or local evidence.

Incidents that cannot be reconstructed because logs or decisions were not preserved.

Public records or audit responses delayed because artifacts are scattered.

Recommended guardrails include the following:

Maintain an AI evidence inventory for each approved AI use.

Use decision logs for approvals, deferrals, conditions, changes, and retirement.

Preserve validation, testing, monitoring, training, SOP, and incident records.

Assign artifact ownership and retention expectations.

Review evidence completeness during periodic governance review.

Reflection Questions

Where does this topic appear in your agency, program, or workflow?

Which staff roles would need to understand or apply this concept before AI use is approved?

What could go wrong if this topic is not addressed before implementation?

What evidence would governance, leadership, or operations staff need before moving forward?

Which policy, data, equity, privacy, security, or workflow questions remain unresolved?

Practical Exercise

Create an AI audit readiness evidence inventory for one AI use case. Identify required artifacts from intake through retirement, the owner of each artifact, storage location, retention expectation, update trigger, and current status.

Expected Artifact or Evidence

AI evidence inventory

Decision log template

Artifact ownership and retention table

Gap list for missing evidence

Review schedule

Expected Assignment

- AI evidence inventory
- Decision log template
- Artifact ownership and retention table
- Gap list for missing evidence
- Review schedule

Knowledge Check

- 1. What is AI audit readiness?
- 2. Which record should be preserved from AI governance review?
- 3. What is validation evidence?
- 4. Why is operational evidence important?
- 5. What is strong completion evidence?

References and Resources

- NIST Artificial Intelligence Risk Management Framework: https://www.nist.gov/itl/ai-risk-management-framework
- NIST Generative AI Profile: https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence
- CDC Public Health Data Strategy and Data Modernization resources: https://www.cdc.gov/public-health-data-strategy/php/about/index.html
- OWASP guidance for LLM application security
- Agency policies for privacy, cybersecurity, records, procurement, accessibility, and public communications
- Relevant public health program SOPs, data governance documentation, and implementation plans