



AI Playbook and Toolkit for Public Health Departments

AI Governance Charter and Operating Model

GOV 400

AI governance needs a formal home before an agency begins approving tools, pilots, vendor products, or workflow automation. A governance charter defines the purpose, scope, authority, membership, cadence, decision rules, and documentation expectations for the group that will oversee AI. Without a charter, AI decisions often become scattered across IT, program areas, legal, privacy, procurement, and individual project teams, which can lead to inconsistent standards and unclear accountability. This module helps leaders design an AI governance operating model that is practical for a public health agency. The operating model should connect strategy, risk management, technical review, privacy review, equity review, procurement, implementation oversight, and incident response. The goal is not to create another committee for its own sake. The goal is to make sure AI decisions are reviewed consistently, documented clearly, and connected to public health authority, community trust, and operational realities.

Level	advanced/leadership
Primary track	Governance and Security Track
Appears in tracks	governance-security

Learning Objectives

- Define the purpose and minimum contents of an AI governance charter.
- Distinguish between a governance committee, technical review group, executive sponsor, and implementation team.
- Identify which agency functions should participate in AI governance and when their review is required.
- Develop operating procedures for intake, review, approval, monitoring, exceptions, escalation, and retirement of AI use cases.
- Produce a draft AI governance charter or operating model that can be adapted by a public health agency.

Module Overview

AI governance needs a formal home before an agency begins approving tools, pilots, vendor products, or workflow automation. A governance charter defines the purpose, scope, authority, membership, cadence, decision rules, and documentation expectations for the group that will oversee AI. Without a charter, AI decisions often become scattered across IT, program areas, legal, privacy, procurement, and individual project teams, which can lead to inconsistent standards and unclear accountability.

This module helps leaders design an AI governance operating model that is practical for a public health agency. The operating model should connect strategy, risk management, technical review, privacy review, equity review, procurement, implementation oversight, and incident response. The goal is not to create another committee for its own sake. The goal is to make sure AI decisions are reviewed consistently, documented clearly, and connected to public health authority, community trust, and operational realities.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Why This Topic Matters for Public Health AI

Public health agencies operate under legal authorities, public accountability obligations, data-sharing agreements, privacy requirements, procurement rules, and equity commitments. AI governance must therefore be more than a technology review process. It must define how the agency decides which AI uses are acceptable, which safeguards are required, who has authority to approve use, and how concerns are escalated.

A governance operating model also helps avoid the problem of uneven adoption. One program may use a public AI tool informally, another may procure a vendor product, and a third may build an internal workflow using sensitive data. If these activities are not reviewed through a common process, the agency may create inconsistent privacy practices, duplicative investments, and uneven risk tolerance. Governance makes AI activity visible and manageable.

Public Health Example

A state health department wants to allow staff to use AI tools for program support, document review, code assistance, case summarization, and public-facing communications. Without a governance charter, each program

may develop its own rules. The communications office may require human review, the surveillance program may prohibit sensitive data entry, IT may limit tool access, and procurement may approve a vendor product without a clear validation plan.

A governance operating model brings these decisions into one structured process. The agency can define low-risk uses that require staff training and supervisor approval, moderate-risk uses that require privacy and program review, and high-risk uses that require executive approval, validation evidence, monitoring, and incident response planning. The charter does not eliminate flexibility. It creates a consistent path for responsible flexibility.

Technical and Governance Deep Dive

A strong governance charter begins with purpose and scope. Purpose explains why the governance body exists, such as protecting public trust, supporting responsible innovation, aligning AI investments with agency priorities, and managing risk across the AI lifecycle. Scope explains what falls under the governance process. For public health, scope should include internally developed AI tools, vendor AI products, generative AI tools, predictive models, agentic workflows, AI features embedded in existing software, and AI-supported uses of sensitive or operational data.

The operating model should define roles clearly. An executive sponsor provides authority and connects governance to agency strategy. A governance chair manages the review process. Privacy, legal, cybersecurity, equity, procurement, records, program, data, and informatics representatives bring specialized review perspectives. Technical reviewers assess architecture, validation, interoperability, security, and monitoring. Program owners remain responsible for the public health workflow and for ensuring that AI supports, rather than replaces, appropriate judgment.

The model should also define the life of an AI use case. Intake captures the proposed use, data involved, intended users, expected benefits, risks, and requested decision. Triage determines the review level. Governance review evaluates authority, data use, privacy, equity, security, accessibility, technical readiness, validation, procurement, and monitoring. Approval should include conditions, documentation requirements, and an owner. Post-approval oversight should review performance, incidents, changes, and whether the use case should continue.

Risks, Failure Modes, and Guardrails

A common failure mode is governance that exists only on paper. A charter may sound strong, but if staff do not know when to submit use cases, who decides, or what evidence is required, AI decisions will still happen informally. A practical guardrail is to create a simple intake form, review rubric, meeting cadence, decision log, and escalation pathway.

Another failure mode is over-centralization. If every minor use requires extensive review, staff may avoid the process or use unapproved tools. Governance should be tiered so low-risk uses move quickly while high-impact uses receive deeper review. This balance helps the agency support innovation while protecting public health values and obligations.

Application to AI-Supported Workflows

Generative AI governance should define approved tools, prohibited data, review expectations, and when outputs may be used externally. Predictive model governance should require validation, subgroup performance review, monitoring, and clear limits on decision use. Agentic AI governance should define allowed actions, approval checkpoints, logging, rollback procedures, and incident triggers. In each case, the charter gives the agency a repeatable way to make and document decisions.

Reflection Questions

Which current or proposed AI use case in your agency raises this leadership or governance issue most clearly?

Who currently has authority to make decisions about this issue, and is that authority documented?

What evidence would governance or leadership need before approving the use case?

Which staff, partners, or communities would be affected if this issue were handled poorly?

What policy, process, or artifact should be created or updated after this module?

Practical Exercise

Draft a one-page AI governance charter and a one-page operating model. Include scope, membership, review levels, decision rights, documentation, escalation, monitoring, and meeting cadence.

Before You Begin

Choose a real or realistic public health AI use case. Document assumptions, unresolved questions, and which agency roles would need to review the artifact before it is adopted. The exercise should produce a work product that could support governance review, leadership discussion, implementation planning, or policy development.

Expected Artifact or Evidence

Draft AI governance charter; governance operating model; intake and decision-flow diagram; list of unresolved policy questions.

Expected Assignment

- Draft AI governance charter; governance operating model; intake and decision-flow diagram; list of unresolved policy questions.

Knowledge Check

- 1. What is the strongest reason this module topic belongs in AI leadership and governance training?
- 2. Which practice best supports accountable public health AI governance?
- 3. When should an AI use case be escalated for leadership or governance review?
- 4. What is weak evidence of completion for a leadership and governance module?
- 5. What should leaders do when an AI use case has meaningful benefits but unresolved risks?

References and Resources

- NIST Artificial Intelligence Risk Management Framework (AI RMF 1.0): <https://www.nist.gov/itl/ai-risk-management-framework>
- NIST Generative AI Profile: <https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>
- NIST Privacy Framework
- CDC Public Health Data Strategy and Data Modernization resources: <https://www.cdc.gov/public-health-data-strategy/php/about/index.html>
- WHO Ethics and Governance of Artificial Intelligence for Health: <https://www.who.int/publications/i/item/9789240029200>

-
- OMB Memorandum M-24-10 on federal AI governance, risk management, and innovation, where relevant to public-sector governance models
 - Agency strategic plans, procurement policies, privacy policies, cybersecurity policies, records policies, and equity plans