



AI Playbook and Toolkit for Public Health Departments

Public Accountability and AI Transparency Reporting

GOV 420

Public health agencies have a responsibility to explain how AI is used in ways that affect public services, public communication, surveillance, resource allocation, or community trust. Transparency reporting helps agencies communicate what AI tools are in use, why they are used, what safeguards are in place, and how the public or partners can raise concerns. This module helps leaders design AI transparency practices that are meaningful, not performative. Transparency should be understandable, accessible, and proportionate to risk. It should support public accountability while protecting security, privacy, confidential data, and sensitive operational details.

Level	advanced/leadership
Primary track	Governance and Security Track
Appears in tracks	governance-security

Learning Objectives

- Explain why transparency reporting matters for public health AI.
- Identify which AI uses should be disclosed internally, publicly, or to specific partners.
- Develop contents for an AI transparency notice or AI use registry.
- Balance transparency with privacy, security, confidentiality, and operational sensitivity.
- Produce a draft AI transparency reporting template.

Module Overview

Public health agencies have a responsibility to explain how AI is used in ways that affect public services, public communication, surveillance, resource allocation, or community trust. Transparency reporting helps agencies communicate what AI tools are in use, why they are used, what safeguards are in place, and how the public or partners can raise concerns.

This module helps leaders design AI transparency practices that are meaningful, not performative. Transparency should be understandable, accessible, and proportionate to risk. It should support public accountability while protecting security, privacy, confidential data, and sensitive operational details.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Why This Topic Matters for Public Health AI

Public health depends on trust. If communities discover that AI is being used in surveillance, communication, service routing, or resource allocation without explanation, they may question whether decisions are fair, accurate, or accountable. Transparency does not require agencies to disclose every technical detail, but it does require clear communication about the purpose and safeguards of consequential AI uses.

Transparency also improves internal accountability. A registry of AI use cases helps leaders know what tools are active, who owns them, what data they use, and when they were last reviewed. It reduces the chance that tools continue operating after their purpose, vendor terms, data sources, or risk profile changes.

Public Health Example

A health department uses AI to draft public health messages, summarize internal reports, and prioritize inspection follow-up. Some uses may not require broad public disclosure, while others may affect public services and should be listed in an AI use registry. The registry does not need to reveal sensitive security details, but it should identify the purpose, owner, data categories, review process, and human oversight.

For a public-facing tool, the agency may also create a plain-language notice explaining that AI assists with initial responses, that staff review high-risk content, that personal information should not be entered unless requested through approved channels, and that users can contact the agency for help or correction. This type of notice supports trust and sets realistic expectations.

Technical and Governance Deep Dive

Transparency reporting should be risk-based. Low-risk internal tools may be documented in an internal registry. Tools that influence public communication, prioritization, access, enforcement, or public-facing interactions may require public notice. High-impact tools may require more detailed documentation, such as intended use, data sources, validation approach, human review, limitations, contact information, and review schedule.

The reporting process should be accessible. Public-facing explanations should use plain language, avoid technical jargon, and be available in appropriate languages and accessible formats. Transparency is not achieved by publishing a dense technical document that affected communities cannot understand. Agencies should consider how communities, partners, and staff will actually interpret the information.

Transparency should also include mechanisms for accountability. A notice or registry should identify how concerns can be raised, how corrections can be requested when appropriate, and how the agency reviews incidents or complaints. Public reporting should connect to internal governance processes so that feedback can lead to action.

Risks, Failure Modes, and Guardrails

One risk is transparency theater. Agencies may publish broad statements about responsible AI without disclosing actual uses, owners, or safeguards. A guardrail is to maintain an AI use registry with concrete information about approved and operational use cases.

Another risk is over-disclosure of sensitive details. Publishing system architecture, vulnerabilities, or confidential data flows could create security or privacy risks. A guardrail is to review transparency materials through privacy, legal, security, and communications before release.

Application to AI-Supported Workflows

Generative AI transparency may explain when AI assists drafting and how human review occurs. Predictive analytics transparency may explain what the model supports, what decisions remain human, and how fairness is monitored. Agentic AI transparency may explain what actions are automated, what approvals are required, and how errors can be reported.

Reflection Questions

Which current or proposed AI use case in your agency raises this leadership or governance issue most clearly?

Who currently has authority to make decisions about this issue, and is that authority documented?

What evidence would governance or leadership need before approving the use case?

Which staff, partners, or communities would be affected if this issue were handled poorly?

What policy, process, or artifact should be created or updated after this module?

Practical Exercise

Create a transparency reporting template and apply it to one AI use case. Include purpose, owner, users, data categories, risk level, safeguards, human review, limitations, review date, and contact pathway.

Before You Begin

Choose a real or realistic public health AI use case. Document assumptions, unresolved questions, and which agency roles would need to review the artifact before it is adopted. The exercise should produce a work product that could support governance review, leadership discussion, implementation planning, or policy development.

Expected Artifact or Evidence

AI transparency template; sample AI use registry entry; public notice draft; internal review checklist.

Expected Assignment

- AI transparency template; sample AI use registry entry; public notice draft; internal review checklist.

Knowledge Check

- 1. What is the strongest reason this module topic belongs in AI leadership and governance training?
- 2. Which practice best supports accountable public health AI governance?
- 3. When should an AI use case be escalated for leadership or governance review?
- 4. What is weak evidence of completion for a leadership and governance module?
- 5. What should leaders do when an AI use case has meaningful benefits but unresolved risks?

References and Resources

- NIST Artificial Intelligence Risk Management Framework (AI RMF 1.0):
<https://www.nist.gov/itl/ai-risk-management-framework>
- NIST Generative AI Profile: <https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>
- NIST Privacy Framework
- CDC Public Health Data Strategy and Data Modernization resources:
<https://www.cdc.gov/public-health-data-strategy/php/about/index.html>
- WHO Ethics and Governance of Artificial Intelligence for Health:
<https://www.who.int/publications/i/item/9789240029200>
- OMB Memorandum M-24-10 on federal AI governance, risk management, and innovation, where relevant to public-sector governance models
- Agency strategic plans, procurement policies, privacy policies, cybersecurity policies, records policies, and equity plans