



AI Playbook and Toolkit for Public Health Departments

Community Oversight and Participatory Governance

GOV 430

Public health AI governance should not be limited to internal agency review. When AI affects communities, services, surveillance, communication, prioritization, or resource allocation, agencies should consider how community perspectives will inform design, review, implementation, and monitoring. Participatory governance creates structured ways for affected people and community partners to shape AI use. This module helps leaders distinguish between informing the public, consulting communities, involving partners, and sharing oversight. It emphasizes that community oversight must be meaningful and appropriately supported. Agencies should not ask communities to approve technical systems without plain-language information, time, context, and a clear explanation of how input will influence decisions.

Level	advanced/leadership
Primary track	Governance and Security Track
Appears in tracks	governance-security

Learning Objectives

- Explain why community oversight matters for public health AI.
- Identify AI use cases that warrant community engagement or participatory governance.
- Design a community engagement approach matched to risk, impact, and community context.
- Recognize risks of tokenism, extractive engagement, and inaccessible technical language.
- Produce a participatory governance plan for one AI use case.

Module Overview

Public health AI governance should not be limited to internal agency review. When AI affects communities, services, surveillance, communication, prioritization, or resource allocation, agencies should consider how community perspectives will inform design, review, implementation, and monitoring. Participatory governance creates structured ways for affected people and community partners to shape AI use.

This module helps leaders distinguish between informing the public, consulting communities, involving partners, and sharing oversight. It emphasizes that community oversight must be meaningful and appropriately supported. Agencies should not ask communities to approve technical systems without plain-language information, time, context, and a clear explanation of how input will influence decisions.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Why This Topic Matters for Public Health AI

Public health agencies make decisions that affect communities, often using data collected through clinical care, surveillance, service systems, public programs, and emergency response. AI can intensify concerns about surveillance, bias, exclusion, language access, disability access, and government trust. Community oversight helps agencies identify concerns that may not be visible through technical review alone.

Participatory governance can also improve AI quality and usefulness. Community partners may identify missing context, harmful assumptions, access barriers, communication needs, and unintended consequences. Their input can help agencies design safer workflows, clearer notices, better escalation pathways, and more meaningful impact measures.

Public Health Example

A city health department considers using an AI tool to help prioritize outreach for maternal and child health services. The model would use program, demographic, and service utilization data to identify neighborhoods where outreach may be needed. Although the goal is supportive, communities may worry that the tool could stigmatize neighborhoods, miss families not represented in data, or direct resources away from groups with less complete records.

A participatory governance plan would involve community-based organizations, maternal health advocates, service users, and equity staff before the tool is finalized. The agency would explain the intended use, data limitations,

safeguards, and decision points in plain language. Community input could shape the variables used, how results are communicated, what services are offered, and how harm or misclassification concerns are handled.

Technical and Governance Deep Dive

Participatory governance should be matched to the impact of the use case. Internal productivity tools may not require community oversight, although they still require policy and privacy review. Tools that influence public communication, service access, enforcement, prioritization, outreach, or resource allocation may require deeper engagement. The higher the impact on communities, the stronger the case for involving community perspectives early.

Engagement should be designed with care. Agencies should provide plain-language descriptions, realistic examples, data-use explanations, and clear questions for feedback. Meetings should be accessible, compensated when appropriate, linguistically inclusive, and scheduled in ways that support participation. Agencies should also explain what decisions are open for influence and what constraints exist because of law, funding, or public health authority.

Follow-through is essential. Community members should know how their input was used, what changed because of engagement, and how concerns will be monitored after launch. Participatory governance is not a single listening session. It is an ongoing relationship between the agency and affected communities, especially when the AI use is consequential or long-term.

Risks, Failure Modes, and Guardrails

One risk is tokenistic engagement. Agencies may seek community input after key decisions have already been made. A guardrail is to engage before requirements, safeguards, and success measures are finalized.

Another risk is technical opacity. Community participants may be asked to respond to vague or overly technical descriptions. A guardrail is to provide plain-language examples, visual workflows, and concrete descriptions of what the AI will and will not do.

Application to AI-Supported Workflows

Generative AI tools that produce public messages may benefit from community review for tone, trust, accessibility, and language. Predictive models used for outreach or allocation should include review of data representativeness and potential stigma. Agentic tools that trigger actions affecting residents should have clear community-informed escalation and correction pathways.

Reflection Questions

Which current or proposed AI use case in your agency raises this leadership or governance issue most clearly?

Who currently has authority to make decisions about this issue, and is that authority documented?

What evidence would governance or leadership need before approving the use case?

Which staff, partners, or communities would be affected if this issue were handled poorly?

What policy, process, or artifact should be created or updated after this module?

Practical Exercise

Develop a participatory governance plan for one AI use case. Include affected groups, engagement purpose, timing, materials, questions, accessibility supports, decision influence, feedback loop, and monitoring plan.

Before You Begin

Choose a real or realistic public health AI use case. Document assumptions, unresolved questions, and which agency roles would need to review the artifact before it is adopted. The exercise should produce a work product that could support governance review, leadership discussion, implementation planning, or policy development.

Expected Artifact or Evidence

Participatory governance plan; community briefing outline; engagement questions; feedback response plan.

Expected Assignment

- Participatory governance plan; community briefing outline; engagement questions; feedback response plan.

Knowledge Check

- 1. What is the strongest reason this module topic belongs in AI leadership and governance training?
- 2. Which practice best supports accountable public health AI governance?
- 3. When should an AI use case be escalated for leadership or governance review?
- 4. What is weak evidence of completion for a leadership and governance module?
- 5. What should leaders do when an AI use case has meaningful benefits but unresolved risks?

References and Resources

- NIST Artificial Intelligence Risk Management Framework (AI RMF 1.0):
<https://www.nist.gov/itl/ai-risk-management-framework>
- NIST Generative AI Profile: <https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>
- NIST Privacy Framework
- CDC Public Health Data Strategy and Data Modernization resources:
<https://www.cdc.gov/public-health-data-strategy/php/about/index.html>
- WHO Ethics and Governance of Artificial Intelligence for Health:
<https://www.who.int/publications/i/item/9789240029200>
- OMB Memorandum M-24-10 on federal AI governance, risk management, and innovation, where relevant to public-sector governance models
- Agency strategic plans, procurement policies, privacy policies, cybersecurity policies, records policies, and equity plans