



AI Playbook and Toolkit for Public Health Departments

Executive AI Decision Rights and Accountability

EXE 400

AI governance can only work when decision rights are clear. Leaders need to know who can approve a pilot, who can accept residual risk, who can authorize use of sensitive data, who can pause a tool, and who is accountable when an AI-supported workflow affects public health action. Ambiguous accountability is especially risky in public health because AI outputs may influence investigations, communications, referrals, resource allocation, and community trust. This module focuses on executive decision rights and accountability structures. Learners examine how to assign ownership across strategy, risk, technical implementation, program operations, privacy, legal, cybersecurity, procurement, communications, and community accountability. The module emphasizes that accountability cannot be delegated to a vendor or to the AI system itself. Public health agencies remain accountable for how AI is selected, used, monitored, and explained.

Level	advanced/leadership
Primary track	Public Health Executive Leadership Track
Appears in tracks	governance-security

Learning Objectives

- Explain why AI accountability must be assigned to people and roles rather than to tools or vendors.
- Identify executive, program, technical, privacy, legal, security, and governance responsibilities for AI decisions.
- Create a decision-rights matrix for AI use case approval, risk acceptance, monitoring, escalation, and retirement.
- Distinguish between consultative input, approval authority, and operational ownership.
- Define accountability expectations for vendor-supported AI tools and internally developed tools.

Module Overview

AI governance can only work when decision rights are clear. Leaders need to know who can approve a pilot, who can accept residual risk, who can authorize use of sensitive data, who can pause a tool, and who is accountable when an AI-supported workflow affects public health action. Ambiguous accountability is especially risky in public health because AI outputs may influence investigations, communications, referrals, resource allocation, and community trust.

This module focuses on executive decision rights and accountability structures. Learners examine how to assign ownership across strategy, risk, technical implementation, program operations, privacy, legal, cybersecurity, procurement, communications, and community accountability. The module emphasizes that accountability cannot be delegated to a vendor or to the AI system itself. Public health agencies remain accountable for how AI is selected, used, monitored, and explained.

Jurisdiction and Agency Policy Note

This module provides national-level training guidance for public health agencies and partners. It is not legal advice and does not replace review by agency legal counsel, privacy officers, procurement officials, civil rights staff, records officers, or other authorized decision-makers. Requirements may vary by state, locality, territory, Tribe, agency, funding source, data-sharing agreement, emergency authority, and organizational policy. Learners should use this module to identify the issues that require review and should confirm applicable requirements before approving, procuring, deploying, or publicly communicating about AI-supported work.

Why This Topic Matters for Public Health AI

AI systems can distribute responsibility across many actors. A vendor may provide the model, IT may manage access, a program may define the workflow, an epidemiologist may review the output, and leadership may approve implementation. Without a formal accountability model, each actor may assume that another party is responsible for validation, monitoring, public explanation, or incident response.

Clear decision rights support trust and speed. When leaders know which decisions they own, they can act quickly when a tool performs poorly, when a privacy concern emerges, or when a public-facing communication needs review. Decision rights also prevent informal escalation patterns in which difficult decisions are delayed because no one is clearly authorized to accept, modify, or reject a proposed use.

Public Health Example

A local health department implements an AI-assisted tool to classify incoming environmental health complaints by topic and urgency. The vendor configured the model, IT manages access, environmental health staff review the queue, and the program director owns the workflow. A complaint is misclassified and a serious housing-related concern is not routed for timely follow-up.

A decision-rights matrix helps the agency respond. The program owner is accountable for workflow performance, the vendor must provide technical support and documentation, IT reviews logging and access, governance reviews whether use should continue, and an executive sponsor decides whether the risk requires pausing the tool. Because responsibilities are defined in advance, the agency can respond without confusion or blame shifting.

Technical and Governance Deep Dive

Decision rights should be defined across the full AI lifecycle. During ideation, leadership may decide whether a use case aligns with agency priorities. During review, governance may decide whether the use case is low, moderate, or high risk. During procurement, authorized officials decide whether vendor claims, contract language, and technical requirements are sufficient. During implementation, program and technical owners decide whether acceptance criteria have been met. During operations, defined owners decide whether monitoring results require adjustment, escalation, or retirement.

A public health accountability model should separate approval authority from operational ownership. An executive may approve a use case, but the program owner remains accountable for how the workflow functions. A privacy officer may advise on data safeguards, but the program cannot treat that consultation as a substitute for ongoing compliance. A vendor may provide model documentation, but the agency must decide whether the documentation is sufficient for its context and authorities.

Decision rights should also include stop authority. Agencies often define who can approve a tool but fail to define who can pause it. Stop authority is essential for AI because model performance, data flows, user behavior, and public conditions can change over time. A clear pause process helps staff report concerns early rather than waiting for a serious incident.

Risks, Failure Modes, and Guardrails

One risk is accountability diffusion. When everyone is involved but no one is accountable, important concerns may go unresolved. A guardrail is to assign a named program owner and governance owner for every AI use case.

Another risk is inappropriate risk acceptance by staff who lack authority. A frontline supervisor may feel pressure to continue a tool despite known concerns. A guardrail is to define which risks can be accepted at the program level and which require executive or governance approval.

Application to AI-Supported Workflows

For generative AI, decision rights should clarify who approves public release of AI-assisted content. For predictive analytics, leaders should decide whether the score may inform resource allocation or only support review. For agentic AI, decision rights should define who authorizes system actions, rollback procedures, and changes to permissions. Accountability should remain tied to public health roles and legal authority.

Reflection Questions

Which current or proposed AI use case in your agency raises this leadership or governance issue most clearly?

Who currently has authority to make decisions about this issue, and is that authority documented?

What evidence would governance or leadership need before approving the use case?

Which staff, partners, or communities would be affected if this issue were handled poorly?

What policy, process, or artifact should be created or updated after this module?

Practical Exercise

Create a decision-rights matrix for one AI use case. Include approval, data access, validation, launch, monitoring, incident escalation, risk acceptance, public communication, and retirement.

Before You Begin

Choose a real or realistic public health AI use case. Document assumptions, unresolved questions, and which agency roles would need to review the artifact before it is adopted. The exercise should produce a work product that could support governance review, leadership discussion, implementation planning, or policy development.

Expected Artifact or Evidence

Decision-rights matrix; RACI table; risk acceptance thresholds; escalation authority list.

Expected Assignment

- Decision-rights matrix; RACI table; risk acceptance thresholds; escalation authority list.

Knowledge Check

- 1. What is the strongest reason this module topic belongs in AI leadership and governance training?
- 2. Which practice best supports accountable public health AI governance?
- 3. When should an AI use case be escalated for leadership or governance review?
- 4. What is weak evidence of completion for a leadership and governance module?
- 5. What should leaders do when an AI use case has meaningful benefits but unresolved risks?

References and Resources

- NIST Artificial Intelligence Risk Management Framework (AI RMF 1.0):
<https://www.nist.gov/itl/ai-risk-management-framework>
- NIST Generative AI Profile: <https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>
- NIST Privacy Framework
- CDC Public Health Data Strategy and Data Modernization resources:
<https://www.cdc.gov/public-health-data-strategy/php/about/index.html>
- WHO Ethics and Governance of Artificial Intelligence for Health:
<https://www.who.int/publications/i/item/9789240029200>
- OMB Memorandum M-24-10 on federal AI governance, risk management, and innovation, where relevant to public-sector governance models
- Agency strategic plans, procurement policies, privacy policies, cybersecurity policies, records policies, and equity plans