



AI Playbook and Toolkit for Public Health Departments

AI Pilot Planning and Governance Coordination

PGM 120

This module prepares public health program managers to plan AI pilots that are limited, governed, measurable, and safe. It focuses on pilot scope, governance conditions, data boundaries, user training, evaluation criteria, risk controls, and go/no-go decisions. The module emphasizes that pilots are structured learning activities, not automatic paths to scale.

Level	applied
Primary track	Program Management Track
Appears in tracks	program-management

Learning Objectives

- Define the purpose, scope, and boundaries of an AI pilot.
- Coordinate governance, privacy, security, equity, technical, and program review before pilot launch.
- Develop pilot success criteria, monitoring measures, and pause criteria.
- Create a pilot plan that supports responsible go/no-go decisions.

Module Overview

This module prepares public health program managers to plan AI pilots that are limited, governed, measurable, and safe. It focuses on pilot scope, governance conditions, data boundaries, user training, evaluation criteria, risk controls, and go/no-go decisions. The module emphasizes that pilots are structured learning activities, not automatic paths to scale.

Jurisdiction and Agency Policy Note

This national-level training module is intended for public health education, planning, and workforce development. It does not replace agency policy, legal review, privacy review, security review, procurement review, communications clearance, civil rights review, Tribal consultation, or governance approval. Learners should confirm applicable requirements, approved tools, data rules, and review pathways within their own jurisdiction and organization before implementing AI-supported workflows.

Why This Topic Matters for Public Health AI

AI pilots can be valuable because they allow agencies to test feasibility, user experience, workflow fit, and risk controls before broader implementation. However, pilots can also create risk if they use real data without approval, expand beyond scope, or become permanent without evaluation. A pilot should be time-limited, documented, reviewed, and connected to a clear decision. Program managers play a key role in making sure the pilot remains a learning activity rather than an uncontrolled deployment.

Governance coordination is essential before a pilot begins. The team should know what has been approved, what conditions apply, which data can be used, who may participate, what outputs may be used for, and what activities are prohibited. The pilot plan should also identify what evidence will be collected and who will decide whether to revise, scale, pause, or stop. Without those elements, a pilot can bypass the very safeguards it was supposed to test.

Public Health Example

A disease surveillance team wants to pilot an AI tool that summarizes notifiable condition case notes for internal epidemiologist review. The pilot uses fictional and de-identified test data at first, limits users to a small trained group, requires human verification of every summary, and prohibits using AI output for official case classification. The program manager documents success measures, including time saved, summary accuracy, reviewer burden, missed information, and user concerns. At the end of the pilot, governance reviews the evidence before deciding whether any broader use should be considered.

Technical, Programmatic, and Operational Deep Dive

A pilot plan should begin with a specific question. The team may be testing whether the AI tool improves workflow efficiency, whether staff can review outputs safely, whether data can be processed in an approved environment, or whether the use case creates unacceptable risk. A vague pilot question leads to vague evaluation and makes it easier for the pilot to continue without a decision. The plan should also state what the pilot is not testing and what

actions are out of scope.

Governance coordination should be visible in the pilot plan. The plan should list required reviewers, approval conditions, data restrictions, security requirements, equity considerations, human review rules, documentation requirements, and user training. It should also identify who can authorize changes to scope, such as adding data, users, vendors, or operational uses. Scope changes should not occur informally during the pilot.

Evaluation should be designed before launch. Metrics may include accuracy, completeness, time, user burden, error types, equity concerns, accessibility, adoption, incidents, and public health value. The plan should also define pause criteria, such as data exposure, repeated unsafe outputs, unmanageable reviewer burden, user overreliance, or evidence that the tool worsens inequity. Pilot findings should be documented in a way that supports governance and leadership decisions.

Risks, Failure Modes, and Guardrails

A common failure mode is pilot creep. A small test may gradually add users, data, workflows, or real decisions without formal approval. Guardrails include written scope, change-control rules, named owners, and governance review of any expansion. Program managers should treat changes in data, users, automation, or decision consequence as governance events.

Another risk is evaluating only whether staff like the tool. User satisfaction matters, but it is not enough to determine safety, equity, or effectiveness. The pilot should evaluate output quality, workflow impact, human review burden, errors, and unintended consequences. Feedback from users should be paired with evidence from monitoring and review.

A third risk is failing to define what happens after the pilot. Without a go/no-go process, a pilot may become an unreviewed operational system. The pilot plan should identify the decision-makers, evidence package, timeline, and possible decisions. Possible outcomes should include revise, scale, defer, pause, or retire.

Application to AI-Supported Workflows

This module applies when an AI use case has passed initial intake and is being considered for controlled testing. The pilot should not begin until governance has reviewed scope, data, risks, and safeguards. Program managers should coordinate across teams so that the pilot produces useful evidence and does not create unapproved operational dependence. The pilot plan should become part of the project record.

Reflection Questions

What would your team need to learn from a pilot before deciding whether to scale?

What conditions should require an AI pilot to pause?

Who should approve changes to pilot scope?

Practical Exercise

Create an AI pilot planning brief. Include the pilot question, scope, data boundaries, users, approved-tool requirements, governance conditions, training plan, metrics, pause criteria, go/no-go decision process, and documentation requirements.

Example: For an AI-assisted document summarization pilot, the brief might limit the pilot to de-identified test documents, require reviewer scoring of accuracy and completeness, prohibit external sharing of summaries, define pause criteria for missing critical information, and require governance review before any real operational use.

Expected Artifact or Evidence

AI pilot planning brief.

Governance coordination checklist.

Pilot evaluation and pause criteria.

Knowledge Check

- 1. What is the best description of an AI pilot?
- 2. What is pilot creep?
- 3. Why should metrics be defined before launch?
- 4. What should happen if a pilot repeatedly produces unsafe outputs?
- 5. Who should approve scope changes?

References and Resources

- [object Object]
- [object Object]
- [object Object]
- [object Object]