



AI Playbook and Toolkit for Public Health Departments

AI Procurement Contract Clauses and Vendor Accountability

POL 210

This module teaches public health staff how to translate responsible AI expectations into procurement and contract requirements. AI vendor products may affect data privacy, security, model performance, public communication, decision support, system integration, documentation, monitoring, and lifecycle management. Procurement is therefore a governance control point, not only a purchasing process. The module does not provide legal advice or a universal contract template. Contract language must be reviewed by agency procurement officials, counsel, privacy officers, IT security, records staff, and program leadership. The module provides the concepts and categories that public health teams should raise before entering vendor discussions, pilots, or agreements.

Level	Intermediate
Primary track	Policy
Appears in tracks	policy

Learning Objectives

- Identify AI-specific procurement and contract risks for public health agencies.
- Describe key contract topics for AI vendor accountability, data rights, model changes, monitoring, and incident response.
- Develop AI procurement requirements tied to a public health use case and risk classification.
- Explain why procurement is part of AI governance and lifecycle oversight.
- Identify when legal, privacy, security, procurement, accessibility, and governance review are needed.

Module Overview

This module teaches public health staff how to translate responsible AI expectations into procurement and contract requirements. AI vendor products may affect data privacy, security, model performance, public communication, decision support, system integration, documentation, monitoring, and lifecycle management. Procurement is therefore a governance control point, not only a purchasing process.

The module does not provide legal advice or a universal contract template. Contract language must be reviewed by agency procurement officials, counsel, privacy officers, IT security, records staff, and program leadership. The module provides the concepts and categories that public health teams should raise before entering vendor discussions, pilots, or agreements.

Jurisdiction and Agency Policy Note

This national-level module is intended for training and planning purposes. It does not replace legal, privacy, procurement, civil rights, records, cybersecurity, Tribal consultation, labor relations, human resources, or agency policy review. Learners should confirm applicable requirements in their own jurisdiction and organization before implementing AI-supported workflows.

Why This Topic Matters for Public Health AI

Many AI risks are shaped by vendor terms, data rights, system architecture, support practices, model updates, subcontractors, and monitoring commitments. If procurement documents do not address these issues, the agency may discover too late that data can be reused, model changes are not disclosed, audit logs are limited, outputs cannot be explained, or the vendor cannot support required monitoring. Strong procurement language helps align AI products with public health accountability.

Public health agencies also need to avoid being locked into tools that cannot meet privacy, security, equity, accessibility, interoperability, records, or transparency requirements. Procurement should require evidence, documentation, and safeguards before a product becomes operational. This is especially important when AI tools are embedded inside larger platforms and may not be obvious to end users.

Public Health Example

A health department is considering a vendor tool that summarizes call center notes and suggests referral categories. The vendor states that the tool is AI-enabled and can improve efficiency, but the initial proposal does not explain whether notes are stored, whether data are used to improve the model, whether subcontractors have access, or how output accuracy is monitored. The program is interested because staff are overwhelmed.

Before procurement moves forward, the agency develops AI-specific requirements. The requirements prohibit vendor reuse of sensitive data without explicit approval, require security documentation, define human review expectations, require audit logs, require notice of model changes, require accessibility and language support, and

require incident reporting. The agency also asks for evidence of performance and limitations in similar public health or health service contexts.

Technical, Programmatic, and Operational Deep Dive

AI procurement should begin with the use case and risk classification. The agency should define the problem, workflow, data involved, users, affected populations, decision consequences, and required safeguards before selecting a vendor. Tool-first procurement can lead to contracts that do not match actual operational needs or risk controls.

Contract clauses should address data rights, privacy, cybersecurity, access controls, logging, retention, deletion, subcontractors, model training, model updates, performance evidence, documentation, human oversight, accessibility, interoperability, and termination. Requirements should be specific enough to evaluate proposals and hold vendors accountable. Agencies should also define what evidence must be provided before pilot, go-live, scale, renewal, or retirement decisions.

Vendor accountability continues after contract signing. The agency should monitor whether the vendor meets service commitments, reports incidents, discloses material changes, provides documentation, supports audits, and responds to corrective action. Governance bodies should receive vendor monitoring information when the tool affects consequential workflows.

Risks, Failure Modes, and Guardrails

A major failure mode is accepting vendor claims without independent review. Vendors may describe AI capabilities in broad terms while providing limited evidence about performance, limitations, bias, data handling, or model updates. Guardrails include requiring documentation, demonstrations using appropriate test scenarios, security review, privacy review, accessibility review, and evidence-based acceptance criteria.

Another risk is weak data rights language. If contracts allow vendor reuse of prompts, uploaded files, outputs, logs, or derived data, public health agencies may lose control over sensitive information. Guardrails include explicit limits on data use, retention, training, subcontractor access, deletion, and breach notification.

A third risk is inadequate exit planning. Agencies may become dependent on a vendor tool without a way to retrieve data, preserve records, transfer workflows, or continue operations if the contract ends. Guardrails include data return provisions, transition support, documentation, open standards where possible, and a plan for pausing or retiring the tool.

Application to AI-Supported Workflows

Procurement requirements should be connected to the workflow the AI tool will support. If the tool drafts public communication, the contract should support review, audit, source traceability, and correction. If the tool supports analytic prioritization, the contract should support validation, subgroup evaluation, monitoring, and access to performance documentation.

Program, IT, privacy, security, legal, procurement, and governance reviewers should collaborate before vendor engagement becomes too advanced. Once a preferred vendor has been selected informally, it becomes harder to change requirements or reject a tool. Early collaboration helps ensure that safeguards shape the procurement instead of being added after the decision is effectively made.

Reflection Questions

What AI-related vendor claims would require evidence before procurement?

What data rights should be explicitly limited in an AI contract?

How should model updates or feature changes be communicated and approved?

Practical Exercise

Draft a procurement requirements checklist for a fictional AI vendor product. Choose a use case such as call center summarization, public message drafting, surveillance signal review, inspection prioritization, or program reporting. Do not name a real vendor unless using publicly available information for illustration only.

Your checklist should include use case description, data involved, prohibited vendor data uses, security requirements, documentation requirements, model change notice, human review, performance evidence, incident reporting, accessibility, interoperability, termination, and monitoring requirements.

Example: For an AI tool that summarizes call center notes, the checklist prohibits vendor training on agency data, requires encryption and role-based access, requires audit logs, requires human review before referral categorization is used, requires notice of model changes, and requires a process to export records if the contract ends.

Expected Artifact or Evidence

AI Procurement Requirements Checklist

Vendor Accountability Review Questions

Knowledge Check

- 1. What is the most important first step after completing the AI Procurement Contract Clauses and Vendor Accountability module?
- 2. Which practice best supports responsible AI use in public health?
- 3. Which statement best reflects the risk or guardrail for this module: A major failure mode is accepting vendor claims without independent review. Vendors may describe AI capabilities in broad terms while providing limited evidence about performance, limitations, bias, data handling, or model updates. Guardrails include requiring documentation, demonstrations using appropriate test scenarios, security review, privacy review, accessibility review, and evidence-based acceptance criteria.
- Another risk is weak data rights language. If contracts allow vendor reuse of prompts, uploaded files, outputs, logs, or derived data, public health agencies may lose control over sensitive information. Guardrails include explicit limits on data use, retention, training, subcontractor access, deletion, and breach notification.
- A third risk is inadequate exit planning. Agencies may become dependent on a vendor tool without a way to retrieve data, preserve records, transfer workflows, or continue operations if the contract ends. Guardrails include data return provisions, transition support, documentation, open standards where possible, and a plan for pausing or retiring the tool.
- 4. When should a staff member escalate an AI-related concern?
- 5. What counts as stronger completion evidence for a training module?

References and Resources

- [object Object]
- [object Object]
- [object Object]
- [object Object]

-
- [object Object]